

*IBM Tape Device Drivers  
Programming Reference*



---

# Contents

<b>IBM Tape Device Drivers Programming Reference.....</b>	<b>1</b>
<b>IBM Tape Device Drivers Programming Reference.....</b>	<b>1</b>
Introduction.....	2
Common extended features.....	3
Tape drive functions and device driver ioctl.....	3
Media partitioning.....	4
Data safe (append-only) mode.....	5
Read Position long/extended form and Locate(16) commands.....	6
Logical Block Protection.....	6
Programmable Early Warning (PEW).....	7
Log Sense page and subpage.....	7
Mode Sense page and subpage.....	7
Verify Tape.....	7
RAO - Recommended Access Order.....	7
AIX tape and medium changer device driver.....	8
Software interface for tape devices.....	8
Software interface for medium changer devices.....	9
Special files.....	9
Device and volume information logging.....	14
Persistent reservation support and IOCTL operations.....	16
General IOCTL operations.....	23
Tape IOCTL operations.....	37
Medium changer IOCTL operations.....	75
Return codes.....	84
Linux tape and medium changer device driver.....	87
Software interface.....	87
General IOCTL operations.....	89
Tape drive IOCTL operations.....	99
Tape drive compatibility IOCTL operations.....	129
Medium changer IOCTL operations.....	129
Return codes.....	137
Windows tape and medium changer device drivers.....	139
Windows programming interface.....	139
Event log.....	167
Notices.....	168
How to send your comments.....	170
Terms and conditions for IBM Documentation.....	170
<b>Index.....</b>	<b>172</b>

# IBM Tape Device Drivers Programming Reference

---

Welcome to the IBM Tape Device Drivers Programming Reference documentation, where you can find information about writing for applications to interfaces with IBM Tape Device Drivers on multiple platforms.

## Programming Reference

[“Introduction” on page 2](#)

These publications and URLs provide user information about writing for applications to interfaces for IBM tape drives, medium changers, and library device drivers.

[“Common extended features” on page 3](#)

[“AIX tape and medium changer device driver” on page 8](#)

[“Linux tape and medium changer device driver” on page 87](#)

[“Windows tape and medium changer device drivers” on page 139](#)

## Troubleshooting and support

[IBM tape storage support](#)

[IBM Support home page](#)

## More information

[IBM Tape Device Drivers Programming Reference \(Legacy\)](#)

[IBM Diamondback tape library](#)

[IBM TS4500 tape library](#)

[IBM TS4300 tape library](#)

[IBM TS2900 tape autoloader](#)

[IBM TS2270 tape drive](#)

[IBM TS2280 tape drive](#)

[IBM TS2290 tape drive](#)

[IBM Community \(Community platform\)](#)

[IBM Support content \(Product support\)](#)

[Redbooks home page](#)

# IBM Tape Device Drivers Programming Reference

---

Welcome to the IBM Tape Device Drivers Programming Reference documentation, where you can find information about writing for applications to interfaces with IBM Tape Device Drivers on multiple platforms.

## Programming Reference

[“Introduction” on page 2](#)

These publications and URLs provide user information about writing for applications to interfaces for IBM tape drives, medium changers, and library device drivers.

[“Common extended features” on page 3](#)

[“AIX tape and medium changer device driver” on page 8](#)

[“Linux tape and medium changer device driver” on page 87](#)

[“Windows tape and medium changer device drivers” on page 139](#)

## Troubleshooting and support

[IBM tape storage support](#)

[IBM Support home page](#)

## More information

[IBM Tape Device Drivers Programming Reference \(Legacy\)](#)

[IBM Diamondback tape library](#)

[IBM TS4500 tape library](#)

[IBM TS4300 tape library](#)

[IBM TS2900 tape autoloader](#)

[IBM TS2270 tape drive](#)

[IBM TS2280 tape drive](#)

[IBM TS2290 tape drive](#)

[IBM Community \(Community platform\)](#)

[IBM Support content \(Product support\)](#)

[Redbooks home page](#)

## Introduction

---

These publications and URLs provide user information about writing for applications to interfaces for IBM tape drives, medium changers, and library device drivers.

### Purpose

The IBM tape and medium changer device drivers are designed specifically to take advantage of the features that are provided by the IBM tape drives and medium changer devices. The goal is to give applications access to the functions required for basic tape functions (such as backup and restore) and medium changer operations (such as cartridge mount and unmount), and to the advanced functions needed by full tape management systems. Whenever possible, the driver is designed to take advantage of the device features transparent to the application.

The most current information on hardware and software requirements for IBM tape and medium changer device drivers can be found in the individual platform readme files or with the subsequent links.

### Hardware requirements

The tape drivers and the IBM Tape Diagnostic Tool (ITDT) are developed to support various versions of different platforms. For the latest support, refer to the Interoperation Center website - [IBM System Storage Interoperation Center \(SSIC\)](#).

**Note:** A single Fibre Channel host bus adapter (HBA) for concurrent tape and disk operations is not recommended. Tape and disk devices require incompatible HBA settings for reliable operation and optimal performance characteristics. Under stress conditions (high I/O rates for either tape, disk, or both) where disk and tape subsystems share a common HBA, stability problems are observed. These issues are resolved by separating disk and tape I/O streams onto separate HBAs and by using SAN zoning to minimize contention. IBM is focused on assuring server/storage configuration interoperability. It is strongly recommended that your overall implementation plan includes provisions for separating disk and tape workloads.

For information about this issue, see the following Redbook, section 4.1.3 in <http://www.redbooks.ibm.com/abstracts/sg246502.html?Open>.

### Software requirements

If you use a third-party application, consult with your application provider about the compatibility with IBM tape device drivers.

For detailed driver requirements for each operating system, refer to the appropriate chapter.

## **IBM tape products**

The IBM tape product family provides an excellent solution for customers with small to large storage and performance requirements.

- IBM TS2250/TS2260/TS2270/TS2280/TS2290 tape drive
- IBM TS2350/TS2360 tape drive
- IBM TS1160/TS1170 tape drive (Enterprise)
- IBM TS4500/Diamondback Tape Libraries (also known as IBM Tape Library 3584)
- IBM TS4300 Tape Library

## **Related information**

Reference material, including the Adobe pdf version of this publication, is available at <http://www-01.ibm.com/support/docview.wss?uid=ssg1S7003032>.

A companion publication that covers installation and user aspects for the device drivers is *IBM Tape Device Drivers: Installation and Users Guide*, GC27-2130-00, at <http://www-01.ibm.com/support/docview.wss?uid=ssg1S7002972>.

## **AIX**

The following URL points to information about IBM System p (also known as pSeries) servers: <https://www.ibm.com/power>.

## **Linux**

The following URLs relate to Linux® distributions: <http://www.redhat.com> and <http://www.suse.com>.

## **Microsoft Windows**

The following URL relates to Microsoft Windows systems: <http://www.microsoft.com>.

## **Additional information**

The following publication contains information that is related to the IBM tape drive, medium changer, and library device drivers: *American National Standards Institute Small Computer System Interface X3T9.2/86-109 X3.180, X3B5/91-173C, X3B5/91-305, X3.131-199X Revision 10H, and X3T9.9/91-11 Revision 1*.

## **Common extended features**

---

### **Tape drive functions and device driver ioctls**

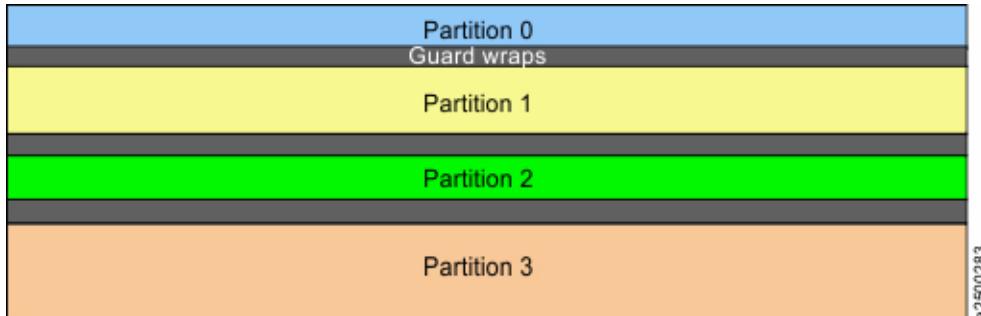
Beginning with the TS1140, TS2250, and TS2350 (LTO 5) generation of tape drives, functions are supported that previous generations of LTO and TS11xx tape drives do not support. The device drivers provide *ioctls* that applications can use for these functions. Refer to the appropriate platform section for the specific *ioctls* and data structures that are not included in this section.

- Media Partitioning
  - Supported tape drives: LTO 5 and TS1140 and later models
- Data Safe (Append-Only) Mode
  - Supported tape drives: LTO 5 and TS1140 and later models
- Read Position SCSI Command for Long and Extended forms
  - Supported tape drives: LTO 5 and TS1140 and later models

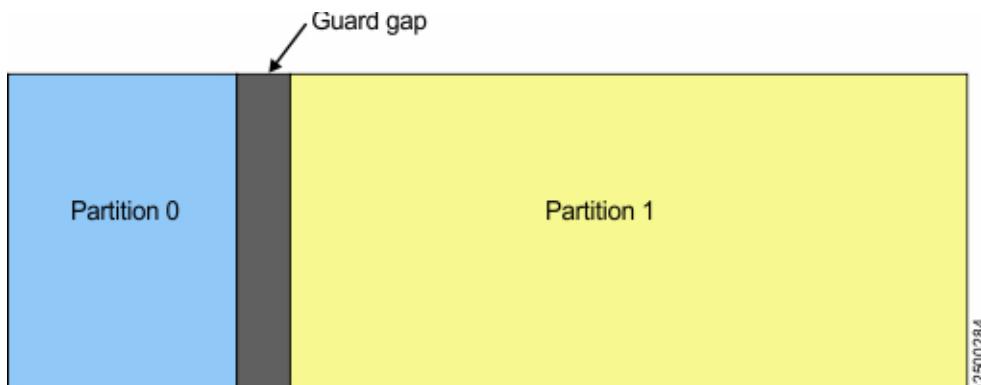
- Locate(16) SCSI Command  
Supported tape drives: LTO 5 and TS1140 and later models
- Logical Block Protection  
Supported tape drives: LTO 5 and TS1120 and later models
- Programmable Early Warning (PEW)  
Supported tape drives: LTO 5 and TS1120 and later models
- Log Sense Page and Subpage  
Supported tape drives: LTO 5 and TS1130 and later models
- Mode Sense Page and Subpage  
Supported tape drives: LTO 4 and TS1120 and later models
- Verify Tape  
Supported tape drives: LTO 5 and TS1120 and later models

## Media partitioning

There are two types of partitioning: Wrap-wise partitioning and Longitudinal partitioning (maximum 2 partitions).



*Figure 1. Wrap-wise partitioning*



*Figure 2. Longitudinal partitioning*

In Wrap-wise partitioning, media can be partitioned into 1 or 2 partitions for LTO 5 and TS1140. For later generations, see drive documentation for the number of partitions supported. The data partition (the default) for a single partition always exists as partition 0. WORM media cannot be partitioned.

The ioctl's the device drivers provide for tape partitioning are

- **Query Partition**

The Query Partition ioctl returns the partition information for the current media in the tape drive. It also returns the current active partition the tape drive is using for the media.

**Note:** If the **Create Partition** ioctl fails, then the **Query Partition** ioctl does not return the correct partition information. To get the correct information, the application must unload and reload the tape again.

- **Create Partition**

The **Create Partition** ioctl is used to format the current media in the tape driver to either 1 or 2 partitions. When two partitions are created, the FDP, SDP, or IDP partition type is specified by the application. The tape must be positioned at the beginning of tape (partition 0 logical block id 0) before this ioctl is used or the ioctl fails.

If the **number\_of\_partitions** field to create in the ioctl structure is one partition, all other fields are ignored and not used. The tape drive formats the media by using its default partitioning type and size for a single partition.

When the type field in the ioctl structure is set to either FDP or SDP, the **size\_unit** and **size** fields in the ioctl structure are not used. When the type field in the ioctl structure is set to IDP, the **size\_unit** and **size** fields are used to specify the size for each partition. One of the two partition sizes for either partition 0 or 1 must be specified as 0xFFFF to use the remaining capacity. The other partition is created by using the **size\_unit** and **size** field for the partition.

- **Set Active Partition**

The **Set Active Partition** ioctl is used to position the tape drive to a specific partition. It becomes the current active partition for subsequent commands and a specific logical block id in the partition. To position to the beginning of the partition, the **logical\_block\_id** field in the ioctl structure must be set to 0.

## Data safe (append-only) mode

Data safe (append-only) mode sets the drive into a logical WORM mode so any non-WORM tape when loaded is handled similarly to a WORM tape. After data or filemarks are written to the tape, it cannot normally be overwritten. New data or filemarks can be appended only at the end of previously written data. Data safe mode applies only to drive operation. When a non-WORM tape is unloaded, it does not change and is still a non-WORM tape.

Conditions exist when the drive is in data safe mode an application might want to explicitly overwrite previously written data by issuing a **write**, **write filemark**, or **erase** command. These commands are referred to as write type commands. An application might also want to explicitly partition the tape with the **Create Partition** ioctl that issues a **format** command. The drive supports a new **Allow Data Overwrite SCSI** command for this purpose.

The ioctls that the device drivers provide for data safe mode are

- **Querying and setting data safe mode**

All platform device drivers except Windows added a data safe mode parameter to existing ioctls that are used to query or set tape drive parameters. The Windows device driver added two new ioctls to query or set data safe mode.

A query ioctl returns the current drive mode, either data safe mode off (normal mode) or data safe mode on. A set ioctl sets the drive to either data safe mode off (normal mode) or data safe mode on. Data safe mode can be set whether a tape is loaded in the drive or not. Data safe mode can be set back to normal mode only when a tape is not currently loaded in the drive.

- **Allow Data Overwrite**

The **Allow Data Overwrite** ioctl is used to allow previously written data on the tape to be overwritten when data safe mode is enabled on the drive, for a subsequent write type command, or to allow a format command with the **Create Partition** ioctl.

To allow a subsequent write type command, the tape position must be set to the correct partition and logical block address within the partition before the ioctl is used. The **partition\_number** and

**logical\_block\_id** fields in the ioctl structure must be set to that partition and logical block ID. The **allow\_format\_overwrite** field in the ioctl structure must be set to 0.

To allow a subsequent **Create Partition** ioctl to format the tape, the **allow\_format\_overwrite** field in the ioctl structure must be set to 1. The **partition\_number** and **logical\_block\_id** fields are not used. But, the tape must be at the beginning of tape (partition 0 logical block id 0) before the **Create Partition** ioctl is issued.

## Read Position long/extended form and Locate(16) commands

Because of the increased tape media capacity and depending on the block sizes and number of files an application can write on tape, the 4-byte fields such as the logical block id the current **Read Position** command (referred to as the short form) that returns 20 bytes might overflow. The same applies to the **Locate(10)** command for the logical block id.

LTO 5 and later supports new forms of the existing **Read Position** command in addition to the current short form. The short form continues to return 4-byte fields in 20 bytes of return data. The long form returns 8-byte fields in 32 bytes of return data with the current position information for the logical block id and logical filemark. The extended form returns 8-byte fields in 32 bytes of return data with the current position information for the logical block id and buffer status. The format of return data in the **Read Position** command is specified by using a service action field in the Read Position SCSI CDB.

LTO 5 and later also supports the **Locate(16)** command that uses 8-byte fields. This command can either position the tape to a logical block id or a logical filemark by setting the **dest\_type** field in the Locate(16) SCSI CDB. After the locate command completes, the tape is positioned at the BOP side of the tape.

The ioctls the device drivers provide are

- **Read Tape Position**

The **Read Tape Position** ioctl returns the **Read Position** command data in either the short, long, or extended form. The form to be returned is specified by setting the **data\_format** field in the ioctl structure.

- **Set Tape Position**

The **Set Tape Position** ioctl issues a **Locate(16)** command to position the tape in the current active partition to either a logical block id or logical filemark. The **logical\_id\_type** field in the ioctl structure specifies either a logical block or logical filemark.

## Logical Block Protection

The ioctls the device drivers provide are

- **Query Logical Block Protection**

This ioctl queries whether the drive can support this feature, what lbp method is used, and where the protection information is included.

The **lbp\_capable** field indicates that the drive has the logical block protection (LBP) capability or not. The **lbp\_method** field is shown if LBP is enabled and what the protection method is. The LBP information length is shown in the **lbp\_info\_length** field. The fields of **lbp\_w**, **lbp\_r**, and **rbdp** present that the protection information is included in write, read, or recover buffer data. The **rbdp** field is not supported for the LTO drive.

- **Set Logical Block Protection**

This ioctl enables or disables Logical Block Protection, sets up what method is used, and where the protection information is included.

The **lbp\_capable** field is ignored in this ioctl by the tape driver. If the **lbp\_method** field is 0 (LBP\_DISABLE), all other fields are ignored and not used. When the **lbp\_method** field is set to a valid non-zero method, all other fields are used to specify the setup for LBP.

## Programmable Early Warning (PEW)

With the tape parameter, the application is allowed to request the tape drive to create a zone that is called the programmable early warning zone (PEWZ) in front of Early Warning (EW).



Figure 3. Programmable Early Warning Zone (PEWZ)

This parameter establishes the programmable early warning zone size. It is a 2-byte numerical value that specifies how many MB before the standard end-of-medium early warning zone to place the programmable early warning indicator. If the value is set to a positive integer, a user application is warned that the tape is running out of space when the tape head reaches the PEW location. If **pew** is set to 0, then there is no early warning zone and the user is notified only at the standard early warning location.

## Log Sense page and subpage

This ioctl of the **SIOC\_LOG\_SENSE10\_PAGE** issues a **Log Sense(10)** command and returns log sense data for a specific page and subpage. This ioctl command is enhanced to add a subpage variable from the log sense page. It returns a log sense page or subpage from the device. The wanted page is selected by specifying the **page\_code** or **subpage\_code** in the structure. Optionally, a specific parm pointer, also known as a parm code, and the number of parameter bytes can be specified with the command.

## Mode Sense page and subpage

This ioctl of the **SIOC\_MODE\_SENSE** issues a **Mode Sense(10)** or **(6)** command and returns the whole mode sense data. The data includes the header, block descriptor, and page code for a specific page or subpage from the device.

## Verify Tape

The ioctl of **VERIFY\_DATA\_TAPE** issues the **VERIFY** command to cause data to be read from the tape and passed through the drive's error detection and correction hardware. This action determines whether data can be recovered from the tape. Also, whether the protection information is present and validates correctly on logical block on the medium. The driver returns a failure or success signal if the **VERIFY SCSI** command is completed in a **Good SCSI** status. The **Verify** command is supported on all LTO libraries. Verify to EOD (ETD) or verify by filemark (VBF) is supported on drives that support Logical Block Protection (LBP).

## RAO - Recommended Access Order

The 3592 E07 implements a function that is called Recommended Access Order. This function provides the capability to improve multiple block recall and retrieval times. It provides an application with the optimized order in which a list of blocks must be recalled to minimize the required total time period.

An application uses the **GRAO** command to request that the drive generate a recommended access order for the User Data Segments that are sent in this command. After a **GRAO** command completes, use the **RRAO** command to receive the results.

## QUERY\_RAO\_INFO IOCTL

The IOCTL queries the maximum number and size of User Data Segments (UDS) that are supported from tape drive and driver for the wanted **uds\_type**: with or without geometry, the geometry can be used to build a representation of the physical layout of the UDS on tape.. The application calls this IOCTL before the **GENERATE\_RAO** and **RECEIVE\_RAO** IOCTLs are issued. The return in this IOCTL is to be used by the application to limit the number of UDS requested in calls to the **GENERATE\_RAO** IOCTL.

## **GENERATE\_RAO**

The IOCTL is called to send a **GRAO** list (UDS's descriptors list) to request the drive to generate a Recommending Access Order list. The process method to create a RAO list is either 1 or 2. 1 does not reorder the UDS's list, but does calculate the estimated locate time for each UDS in the list, and 2 reorders the UDS's list and calculates the estimated locate time for each UDS in its resultant position. The type of UDS is either with or without the geometry. The **uds\_number** must be not larger than **max\_host\_uds\_number** returned in the **QUERY\_RAO\_INFO** IOCTL. A UDS descriptor has a name, partition number and beginning and ending logical object identifiers.

## **RECEIVE\_RAO**

After a **GENERATE\_RAO** IOCTL is completed, the application calls the **RECEIVE\_RAO** IOCTL to receive a recommended access order of UDS from the drive. The application must allocate to the buffer an accurate size to receive the list.

The **gra0** list, for the **generate\_rao** and **receive\_rao** structures, is in the following format when lists are sent or received. The structures for the headers and UDS segments are not provided by the device driver but is the responsibility of the calling application. It is defined in the *IBM Enterprise Tape System 3592 SCSI Reference*.

```
-- List Header  
-- UDS Segment Descriptor (first)  
.....  
-- UDS Segment Descriptor (last)
```

# **AIX tape and medium changer device driver**

---

This chapter provides an introduction to the IBM AIX® Enhanced Tape and Medium Changer Device Driver (Atape) programming interface to IBM Tape Storage Family tape and medium changer devices.

## **Software interface for tape devices**

The AIX tape and medium changer device driver provides the following entry points for tape devices.

### **Open**

This entry point is driven by **open**, **openx**, and **creat** subroutines.

### **Write**

This entry point is driven by **write**, **writev**, **writex**, and **writevx** subroutines.

### **Read**

This entry point is driven by **read**, **readv**, **readx**, and **readvx** subroutines.

### **Close**

This entry point is driven explicitly by the **close** subroutine and implicitly by the operating system at program termination.

### **ioctl**

This entry point provides a set of tape and SCSI-specific functions. It allows AIX applications to access and control the features and attributes of the tape device programmatically. For the medium changer devices, it also provides a set of medium changer functions that is accessed through the tape device special files or independently through an extra special file for the medium changer only.

### **Dump**

This entry point allows the use of the AIX dump facility with the driver.

The standard set of AIX device management commands is available. The **chdev**, **xmdev**, **mkdev**, and **lsdev** commands are used to bring the device online or change the attributes that determine the status of the tape device.

## Software interface for medium changer devices

The AIX tape and medium changer device driver provides the following AIX entry points for the medium changer devices.

### Open

This entry point is driven by **open** and **openx** subroutines.

### Close

This entry point is driven explicitly by the **close** subroutine and implicitly by the operating system at program termination.

### IOCTL

This entry point provides a set of medium changer and SCSI-specific functions. It allows AIX applications to access and control the features and attributes of the tape system robotic device programmatically.

The standard set of AIX device management commands is available. The **chdev**, **xmdev**, **mkdev**, and **lsdev** commands are used to bring the device online or change the attributes that determine the status of the tape system robotic device.

## Special files

After the driver is installed and a tape device is configured and made available for use, access is provided through the special files. These special files, which consist of the standard AIX special files for tape devices (with other files unique to the Atape driver), are in the **/dev** directory.

### Special files for tape devices

Each tape device has a set of special files that provides access to the same physical drive but to different types of functions. In addition to the tape special files, a special file is provided to tape devices that allow access to the medium changer as a separate device. See [Table 1 on page 9](#). The asterisk (\*) represents a number that is assigned to a particular device (such as **rmt0**).

Table 1. Special files for tape devices					
Special File Name	Rewind on Close <sup>1</sup>	Retention on Open <sup>2</sup>	Bytes per Inch <sup>3</sup>	Trailer Label	Unload on Close
/dev/rmt*	Yes	No	N/A	No	No
/dev/rmt*.1	No	No	N/A	No	No
/dev/rmt*.2	Yes	Yes	N/A	No	No
/dev/rmt*.3	No	Yes	N/A	No	No
/dev/rmt*.4	Yes	No	N/A	No	No
/dev/rmt*.5	No	No	N/A	No	No
/dev/rmt*.6	Yes	Yes	N/A	No	No
/dev/rmt*.7	No	Yes	N/A	No	No
/dev/rmt*.10 <sup>4</sup>	No	No	N/A	No	No
/dev/rmt*.20	Yes	No	N/A	No	Yes
/dev/rmt*.40	Yes	No	N/A	Yes	No
/dev/rmt*.41	No	No	N/A	Yes	No
/dev/rmt*.60	Yes	No	N/A	Yes	Yes

*Table 1. Special files for tape devices (continued)*

Special File Name	Rewind on Close <sup>1</sup>	Retention on Open <sup>2</sup>	Bytes per Inch <sup>3</sup>	Trailer Label	Unload on Close
/dev/rmt*.null <sup>5</sup>	Yes	No	N/A	No	No
/dev/rmt*.smc <sup>6</sup>	N/A	N/A	N/A	N/A	N/A

**Note:**

1. The **Rewind on Close** special files for the Ultrium tape drives write filemarks under certain conditions before rewinding. See [“Opening the special file for I/O” on page 11](#).
2. The **Retention on Open** special files rewind the tape on open only. Retensioning is not completed because these tape products run the retension operation automatically when needed.
3. The **Bytes per Inch** options are ignored for the tape devices that this driver supports. The density selection is automatic.
4. The **rmt\*.10** file bypasses normal close processing, and the tape is left at the current position.
5. The **rmt\*.null** file is a pseudo device similar to the **/dev/null** AIX special file. The IOCTL calls can be issued to this file without a real device that is attached to it, and the device driver returns a successful completion. Read and write system calls return the requested number of bytes. This file can be used for application development or debugging problems.
6. The **rmt\*.smc** file can be opened independently of the other tape special files.

For tape drives with attached SCSI medium changer devices, the **rmt\*.smc** special file provides a separate path for issuing commands to the medium changer. When this special file is opened, the application can view the medium changer as a separate SCSI device.

This special file and the **rmt\*** special file can be opened at the same time. The file descriptor that results from opening the **rmt\*.smc** special file does not support the following operations.

- Read
- Write
- Open in diagnostic mode
- Commands that are designed for a tape device

If a tape drive has an attached SCSI medium changer device, all operations (including the medium changer operations) are supported through the interface to the **rmt\*** special file.

## Special files for medium changer devices

After the driver is installed and a medium changer device is configured and made available for use, access to the robotic device is provided through the **smc\*** special file in the **/dev** directory.

Table 2 on page 10 shows the attributes of the special file. The asterisk (\*) represents a number that is assigned to a particular device (such as **smc0**). The term **smc** is used for a SCSI medium changer device. The **smc\*** special file provides a path for issuing commands to control the medium changer robotic device.

Table 2. Special files	
Special file name	Description
/dev/smc*	Access to the medium changer robotic device
/dev/smc*.null	Pseudo medium changer device

**Note:** The **smc\*.null** file is a pseudo device similar to the **/dev/null** AIX special file. The commands can be issued to this file without a real device that is attached to it, and the device driver returns a successful completion. This file can be used for application development or debugging problems.

The file descriptor that results from opening the **smc** special file does not support the following operations.

- Read
- Write
- Commands that are designed for a tape device

## Opening the special file for I/O

Several options are available when a file is opened for access. These options, which are known as **O\_FLAGS**, affect the characteristics of the opened tape device or the result of the **open** operation. The **Open** command is

```
tapefd=open("/dev/rmt0",O_FLAGS);  
smcfd=open("/dev/smc0",O_FLAGS);
```

The **O\_FLAGS** parameter has the following flags.

- **O\_RDONLY**

This flag allows only operations that do not change the content of the tape. The flag is ignored if it is used to open the **smc** special files.

- **O\_RDWR**

This flag allows complete access to the tape. The flag is ignored if it is used to open the **smc** special files.

- **O\_WRONLY**

This flag does not allow the tape to be read. All other operations are allowed. The flag is ignored if it is used to open the **smc** special files.

- **O\_NDELAY** or **O\_NONBLOCK**

These two flags complete the same function. The driver does not wait until the device is ready before it opens and allows commands to be sent. If the device is not ready, subsequent commands (which require that the device is ready or a physical tape is loaded) fail with ENOTREADY. Other commands, such as gathering the inquiry data, complete successfully.

- **O\_APPEND**

When the tape drive is opened with this flag, the driver rewinds the tape. Then, it seeks to the first two consecutive filemarks, and places the initial tape position between them. This status is the same if the tape was previously opened with a **No Rewind on Close** special file. This process can take several minutes for a full tape. The flag is ignored if it is used to open the **smc** special files.

This flag must be used with the **O\_WRONLY** flag to append data to the end of the current data on the tape. The **O\_RDONLY** or **O\_RDWR** flag is illegal in combination with the **O\_APPEND** flag.

**Note:** This flag cannot be used with the **Retention on Open** special files, such as **rmx.2**.

If the **open** system call fails, the *errno* value contains the error code. See “Return codes” on page 84 for a description of the *errno* values.

## The extended open operation

An extended **open** operation is also supported on the device. This operation allows special types of processing during the opening and subsequent closing of the tape device. The **Extended Open** command is

```
tapefd=openx("/dev/rmt0",O_FLAGS,NULL,E_FLAGS);  
smcfd=openx("/dev/smc0",O_FLAGS,NULL,E_FLAGS);
```

The **O\_FLAGS** parameter provides the same options that are described in “[Opening the special file for I/O](#)” on page 11. The third parameter is always **NULL**. The **E\_FLAGS** parameter provides the extended options. The **E\_FLAGS** values can be combined during an **open** operation or they can be used with an OR operation.

The **E\_FLAGS** parameter has the following flags.

- **SC\_RETAIN\_RESERVATION**

This flag prevents the **SCSI Release** command from being sent during a **close** operation.

- **SC\_FORCED\_OPEN**

The flag forces the release of any current reservation on the device by an initiator. The reservation can either be a **SCSI Reserve** or **SCSI Persistent Reserve**.

- **SC\_KILL\_OPEN**

This flag kills all currently open processes and then exits the open with *errno* EINPROGRESS returned.

- **SC\_PR\_SHARED\_REGISTER**

This flag overrides the configuration reservation type attribute whether it was set to **reserve\_6** or persistent. It sets the device driver to use **Persistent Reserve** while the device is open until closed. The configuration reservation type attribute is not changed and the next open without using this flag uses the configuration reservation type. The device driver also registers the host reservation key on the device. This flag can be used with the other extended flags.

- **SC\_DIAGNOSTIC**

The device is opened in diagnostic mode, and no SCSI commands are sent to the device during an **open** operation or a **close** operation. All operations (such as reserve and mode select) must be processed by the application.

- **SC\_NO\_RESERVE**

This flag prevents the **SCSI Reserve** command from being sent during an **open** operation.

- **SC\_PASSTHRU**

No SCSI commands are sent to the device during an **open** operation or a **close** operation. All operations (such as reserve the device, release the device, and set the tape parameters) must be processed explicitly by the application. This flag is the same as the **SC\_DIAGNOSTIC** flag. The exception is that a **SCSI Test Unit Unit Ready** command is issued to the device during an open operation to clear any unit attentions.

- **SC\_FEL**

This flag turns on the forced error logging in the tape device for read and write operations.

- **SC\_NO\_ERRORLOG**

This flag turns off the AIX error logging for all read, write, or IOCTL operations.

- **SC\_TMCP**

This flag allows up to eight processes to concurrently open a device when the device is already open by another process. There is no restriction for medium changer IOCTL commands that can be issued when this flag is used. However, for tape devices only a limited set of IOCTL commands can be issued. If an IOCTL command cannot be used with this flag, then *errno* EINVAL is returned.

If another process already has the device open with this flag, the **open** fails, and the *errno* is set to EAGAIN.

If the **open** system call fails, the *errno* value contains the error code. See “[Return codes](#)” on page 84 for a description of the *errno* values.

## Writing to the special file

Several subroutines allow writing data to a tape. The basic **write** command is

```
count=write(tapefd, buffer, numbytes);
```

The **write** operation returns the number of bytes written during the operation. It can be less than the value in **numbytes**. If the block size is fixed (block\_size≠0), the **numbytes** value must be a multiple of the block size. If the block size is variable, the value that is specified in **numbytes** is written. If the **count** is less than zero, the *errno* value contains the error code that is returned from the driver.

See “Return codes” on page 84 for a description of the *errno* values.

The **writev**, **writex**, and **writevx** subroutines are also supported. Any values that are passed in the **ext** field with the extended write operation are ignored.

## Reading from the special file

Several subroutines allow reading data from a tape. The basic **read** command is

```
count=read(tapefd, buffer, numbytes);
```

The **read** operation returns the number of bytes read during the operation. It can be less than the value in **numbytes**. If the block size is fixed (block\_size≠0), the **numbytes** value must be a multiple of the block size. If the **count** is less than zero, the *errno* value contains the error code that is returned from the driver.

See “Return codes” on page 84 for a description of the *errno* values.

If the block size is variable, then the value that is specified in **numbytes** is read. If the blocks read are smaller than requested, the block is returned up to the maximum size of one block. If the blocks read are greater than requested, an error occurs with the error set to ENOMEM.

Reading a filemark returns a value of zero and positions the tape after the filemark. Continuous reading (after EOM is reached) results in a value of zero and no further change in the tape position.

The **readv** subroutine is also supported.

## Reading with the TAPE\_SHORT\_READ extended parameter

For normal read operations, if the block size is set to variable (0) and the amount of data in a block on the tape is more than the number of bytes requested in the call, an ENOMEM error is returned. An application can read fewer bytes without an error by using the **readx** or **readvx** subroutine and specifying the **TAPE\_SHORT\_READ** extended parameter.

```
count=readx(tapefd, buffer, numbytes, TAPE_SHORT_READ);
```

The **TAPE\_SHORT\_READ** parameter is defined in the **/usr/include/sys/tape.h** header file.

## Reading with the TAPE\_READ\_REVERSE extended parameter

The **TAPE\_READ\_REVERSE** extended read parameter reads data from the tape in the reverse direction. The order of the data that is returned in the buffer for each block that is read from the tape is the same as if it were read in the forward direction. However, the last block that is written is the first block in the buffer. This parameter can be used with both fixed and variable block sizes. The **TAPE\_SHORT\_READ** extended parameter can be used with this parameter, if the block size is set to variable (0).

Use this parameter with the **readx** or **readvx** subroutine that specifies the **TAPE\_READ\_REVERSE** extended parameter.

```
count=readx(tapefd, buffer, numbytes, TAPE_READ_REVERSE);
```

The **TAPE\_READ\_REVERSE** parameter is defined in the **/usr/include/sys/Atape.h** header file.

## Closing the special file

Closing a special file is a simple process. The file descriptor that is returned by the **Open** command is used to close the command.

```
rc=close(tapefd);
rc=close(smcfd);
```

The return code from the **close** operation must be checked by the application. If the return code is not zero, the *errno* value is set during a **close** operation to indicate that a problem occurred while the special file was closing. The **close** subroutine tries to run as many operations as possible even if there are failures during portions of the **close** operation. If the device driver cannot terminate the file correctly with filemarks, it tries to close the connection. If the **close** operation fails, consider the device closed and try another **open** operation to continue processing the tape. After a **close** failure, assume that either the data or the tape is inconsistent.

For tape drives, the result of a **close** operation depends on the special file that was used during the **open** operation and the tape operation that was run while it was opened. The SCSI commands are issued according to the following logic.

```
If the last tape operation was a WRITE command
    Write 2 filemarks on tape
    If special file is Rewind on Close (Example: /dev/rmt0)
        Rewind tape
    If special file is a No-Rewind on Close (Example: /dev/rmt0.1)
        Backward space 1 filemark (tape is positioned to append next file)

If the last tape operation was a WRITE FILEMARK command
    Write 1 filemark on tape
    If special file is Rewind on Close (Example: /dev/rmt0)
        Rewind tape
    If special file is a No-Rewind on Close (Example: /dev/rmt0.1)
        Backward space 1 filemark (tape is positioned to append next file)

If the last tape operation was a READ command
    If special file is Rewind on Close (Example: /dev/rmt0)
        Rewind tape
    If special file is a No-Rewind on Close (Example: /dev/rmt0.1)
        Forward space to next filemark (tape is positioned to read or append next file)

If the last tape operation was NOT a READ, WRITE, or WRITE FILEMARK command
    If special file is Rewind on Close (Example: /dev/rmt0)
        Rewind tape
    If special file is a No-Rewind on Close (Example: /dev/rmt0.1)
        No commands are issued, tape remains at the current position
```

## Device and volume information logging

The device driver provides a logging facility that saves information about the device and the media. The information is extensive for some devices and limited for other devices. If this feature is set to On, either by configuration or the **STIOCSETP** IOCTL, the device driver logging facility gathers all available information through the **SCSI Log Sense** command.

This process is separate from error logging. Error logging is routed to the system error log. Device information logging is sent to a separate file.

The following parameters control this utility.

- Logging
- Maximum size of the log file
- Volume ID for logging

See the *IBM Tape Storage Tape Device Drivers: Installation and User's Guide*: <https://www.ibm.com/docs/en/ts4300-tape-library?topic=tape-device-drivers-diagnostic-tool-users-guide>

Each time an **Unload** command or the **STIOC\_LOG\_SENSE** IOCTL command is issued, the log sense data is collected, and an entry is added to the log. Each time a new cartridge is loaded, the log sense data in the tape device is reset so that the log data is gathered on a per-volume basis.

## Log file

The data is logged in the **/usr/adm/ras** directory. The file name is dependent on each device so each device has a separate log. An example of the **rmt1** device file is

```
/usr/adm/ras/Atape.rmt1.log
```

The files are in binary format. Each entry has a header followed by the raw **Log Sense** pages as defined for a particular device.

The first log page is always page 0x00. This page, as defined in the SCSI-2 ANSI specification, contains all the pages that are supported by the device. Page 0x00 is followed by all the pages that are specified in page 0x00. The format of each following page is defined in the SCSI specification and the device manual.

The format of the file is defined by the data structure. The **logfile\_header** is followed by **max\_log\_size** (or a fewer number of entries for each file). The **log\_record\_header** is followed by a log entry.

The data structure for log recording is

```
struct logfile_header
{
    char owner[16];           /* module that created the file */
    time_t when;              /* time when file created */
    unsigned long count;       /* number of entries in file */
    unsigned long first;       /* first entry number in wrap queue */
    unsigned long max;         /* maximum entries allowed before wrap */
    unsigned long size;        /* size of entry (bytes), entry size is fixed */
};

struct log_record_header
{
    time_t when;              /* time when log entry made */
    ushort type;               /* log entry type */
#define LOGDEMOUNT 1          /* demount log entry */
#define LOGSENSE 2             /* log sense ioctl entry */
#define LOGOVERFLOW 3          /* log overflow entry */
    char device_type[8];        /* device type that made entry */
    char volid[16];             /* volume ID of entry */
    char serial[12];            /* serial number of device */
    reserved[12];
};
```

The format of the log file is

logfile_header
log_record_header
log_record_entry
•
•
•
•
log_record_header
log_record_entry

Each **log\_record\_entry** contains multiple log sense pages. The log pages are placed in order one after another. Each log page contains a header followed by the page contents.

The data structure for the header of the log page is

```

struct log_page_header
{
    char code;           /* page code */
    char res;            /* reserved */
    unsigned short len; /* length of data in page after header */
};

```

## Persistent reservation support and IOCTL operations

### ODM attributes and configuring persistent reserve support

Two new ODM attributes are added for PR (Persistent Reservation) support:

- `reserve_type`
- `reserve_key`

The `reserve_type` attribute determines the type of reservation that the device driver uses for the device. The values can be `reserve_6`, which is the default for the device driver or persistent. This attribute can be set by either using the **AIX SMIT** menu to **Change/Show Characteristics of a Tape Drive** or from a command line with the AIX command

```
chdev -l rmtx -a reserve_type=persistent or -a reserve_type=reserve_6
```

The `reserve_key` attribute is used to optionally set a user-defined host reservation key for the device when the `reserve_type` is set to persistent. The default for this attribute is blank (NULL). The default uses a device driver unique host reservation key that is generated for the device. This attribute can be set by either using the **AIX SMIT** menu to **Change/Show Characteristics of a Tape Drive** or from a command line with the AIX command

```
chdev -l rmtx -a reserve_key=key
```

The key value can be specified as a 1-8 character ASCII alphanumeric key or a 1-16 hexadecimal key that has the format 0xkey. If fewer than 8 characters are used for an ASCII key such as hostA, the remaining characters are set to 0x00 (NULL).

**Note:** If the Data Path Failover (DPF) feature is enabled for a logical device by setting the `alternate_pathing` attribute to yes, the configuration `reserve_type` attribute is not used and the device driver uses persistent reservation. Either the user-defined `reserve_key` value or if not defined the default device driver host reservation key is used.

### Default device driver host reservation key

If a user-defined host reservation key is not specified, then the device driver uses a unique static host reservation key for the device. This key is generated when the first device is configured and the device driver is initially loaded into kernel memory. The key is 16 hexadecimal digits in the format 0xApppppppsssssss, where pppppp is the configuration process id that loaded the device driver. Also, ssssssss is the 32-bit value of the TOD clock when the device driver was loaded. When any device is configured and the `reserve_key` value is NULL, then the device driver sets the `reserve_key` value to this default internally for the device.

### Preempting and clearing another host reservation

When another host initiator is no longer using the device but has left either an **SCSI-2 Reserve 6** or a **Persistent Reserve** active preventing by using the device, either type of reservation can be cleared by using the `openx()` extended parameter **SC\_FORCED\_OPEN**.

**Note:** This parameter must be used only when the application or user is sure that the reservation must be cleared.

## Openx() extended parameters

The following openx() extended parameters are provided for managing device driver reserve during open processing and release during close processing. These parameters apply to either **SCSI-2 Reserve 6** or **Persistent Reserve**. The **SC\_PASSTHRU** parameter applies only to the Atape device driver and is defined in **/usr/include/sys/Atape.h**. All other parameters are AIX system parameters that are defined in **/usr/include/sys/scsi.h**. AIX base tape device drivers might not support all of these parameters.

- **SC\_PASSTHRU**
- **SC\_DIAGNOSTIC**
- **SC\_NO\_RESERVE**
- **SC\_RETAIN\_RESERVATION**
- **SC\_PR\_SHARED\_REGISTER**
- **SC\_FORCED\_OPEN**

The **SC\_PASSTHRU** parameter bypasses all commands that are normally issued on open and close by the device driver. In addition to bypassing the device driver that reserves on open and releases the device on close, all other open commands except test unit ready such as mode selects and rewind on close (if applicable) are also bypassed. A test unit ready is still issued on open to clear any pending unit attentions from the device. This is the only difference in using the **SC\_DIAGNOSTIC** parameter.

The **SC\_DIAGNOSTIC** parameter bypasses all commands that are normally issued on open and close by the device driver. In addition to bypassing the device driver that reserves on open and releases the device on close, all other open commands such test unit ready, mode selects, and rewind on close (if applicable) are also bypassed.

The **SC\_NO\_RESERVE** parameter bypasses the device driver that issues a reserve on open only. All other normal open device driver commands are still issued such as test unit ready and mode selects.

The **SC\_RETAIN\_RESERVATION** parameter bypasses the device driver that issues a release on close only. All other normal close device driver commands are still issued such as rewind (if applicable).

The **SC\_PR\_SHARED\_REGISTER** parameter sets the device driver **reserve\_type** to persistent and overrides the configuration **reserve\_type** attribute whether it was set to **reserve\_6** or **persistent**. A subsequent reserve on the current open by the device driver (if applicable) uses **Persistent Reserve**. The **reserve\_type** is only changed for the current open. The next open without using this parameter uses the configuration **reserve\_type**. In addition to setting the **reserve\_type** to persistent, the device driver registers the host reservation key on the device. This parameter can also be used with the extended parameters.

The **SC\_FORCED\_OPEN** parameter first clears either a **SCSI-2 Reserve 6** or a **Persistent Reservation** if one currently exists on the device from another host. The device driver open processing then continues according to the type of open. This parameter can also be used with the extended parameters.

## AIX tape persistent reserve IOCTLs

The Atape device driver supports the AIX common tape **Persistent Reserve** IOCTLs for application programs to manage their own **Persistent Reserve** support. The IOCTLs are defined in the header file **/usr/include/sys/tape.h**.

The following two IOCTLs return **Persistent Reserve** information by using the **SCSI Persistent Reserve In** command.

- **STPRES\_READKEYS**
- **STPRES\_READRES**

The following four IOCTLs complete **Persistent Reserve** functions by using the **SCSI Persistent Reserve Out** command.

- **STPRES\_CLEAR**

- **STPRES\_PREEMPT**
- **STPRES\_PREEMPT\_ABORT**
- **STPRES\_REGISTER**

Except for the **STPRES\_REGISTER** IOCTL, the other three IOCTLS require that the host reservation key is registered on the device first. This action can be done by either issuing the **STPRES\_REGISTER** IOCTL before the IOCTLS are issued or by opening the device with the **SC\_PR\_SHARED\_REGISTER** parameter.

The **STPRES\_READKEYS** IOCTL issues the **persistent reserve in** command with the read keys service action. The following structure is the argument for this IOCTL.

```
struct st_pres_in {
    ushort      version;
    ushort      allocation_length;
    uint        generation;
    ushort      returned_length;
    uchar       scsi_status;
    uchar       sense_key;
    uchar       scsi_asc;
    uchar       scsi_ascq;
    uchar       *reservation_info;
}
```

The `allocation_length` is the maximum number of bytes of key values that are returned in the `reservation_info` buffer. The `returned_length` value indicates how many bytes of key values that device reported in the parameter data. Also, it shows the list of key values that are returned by the device up to `allocation_length` bytes. If the `returned_length` is greater than the `allocation_length`, then the application did not provide an `allocation_length` large enough for all of the keys the device registered. The device driver does not consider it an error.

The **SYPRES\_READRES** IOCTL issues the **persistent reserve in** command with the read reservations service action. The **STPRES\_READRES** IOCTL uses the same following IOCTL structure as the **STPRES\_READKEYS** IOCTL.

```
struct st_pres_in {
    ushort      version;
    ushort      allocation_length;
    uint        generation;
    ushort      returned_length;
    uchar       scsi_status;
    uchar       sense_key;
    uchar       scsi_asc;
    uchar       scsi_ascq;
    uchar       *reservation_info;
}
```

The `allocation_length` is the maximum number of bytes of reservation descriptors that are returned in the `reservation_info` buffer. The `returned_length` value indicates how many bytes of reservation descriptor values that device reported in the parameter data. Also, it shows the list of reservation descriptor values that are returned by the device up to `allocation_length` bytes. If the `returned_length` is greater than the `allocation_length`, then the application did not provide an `allocation_length` large enough for all of the reservation descriptors the device registered. The device driver does not consider it an error.

The **STPRES\_CLEAR** IOCTL issues the **persistent reserve out** command with the clear service action. The following structure is the argument for this IOCTL.

```
struct st_pres_clear {
    ushort      version;
    uchar       scsi_status;
    uchar       sense_key;
    uchar       scsi_asc;
    uchar       scsi_ascq;
}
```

The **STPRES\_CLEAR** IOCTL clears a persistent reservation and all persistent reservation registrations on the device.

The **STPRES\_PREEMPT** IOCTL issues the **persistent reserve out** command with the preempt service action. The following structure is the argument for this IOCTL.

```
struct st_pres_preempt {  
    ushort           version;  
    unsigned long long preempt_key;  
    uchar            scsi_status;  
    uchar            sense_key;  
    uchar            scsi_asc;  
    uchar            scsi_ascq;  
};
```

The **STPRES\_PREEMPT** IOCTL preempts a persistent reservation or registration. The `preempt_key` contains the value of the registration key of the initiator that is to be preempted. The determination of whether it is the persistent reservation or registration that is preempted is made by the device. If the initiator corresponding to the `preempt_key` is associated with the reservation that is preempted, then the reservation is preempted and any matching registrations are removed. If the initiator corresponding to the `preempt_key` is not associated with the reservation that is preempted, then any matching registrations are removed. The SPC2 standard states that if a valid request for a preempt service action fails, it can be because of the condition in which another initiator has left the device. The suggested recourse in this case is for the preempting initiator to issue a logical unit reset and retry the preempting service action.

The **STPRES\_PREEMPT\_ABORT** IOCTL issues the **persistent reserve out** command with the preempt and abort service action. The **STPRES\_PREEMPT\_ABORT** IOCTL uses the same argument structure as the **STPRES\_PREEMPT** IOCTL.

```
struct st_pres_preempt {  
    ushort           version;  
    unsigned long long preempt_key;  
    uchar            scsi_status;  
    uchar            sense_key;  
    uchar            scsi_asc;  
    uchar            scsi_ascq;  
};
```

The **STPRES\_PREEMPT\_ABORT** IOCTL preempts a persistent reservation or registration and abort all outstanding commands from the initiators corresponding to the `preempt_key` registration key value. The `preempt_key` contains the value of the registration key of the initiator for which the preempt and abort is to apply. The determination of whether it is the persistent reservation or registration that is to be preempted is made by the device. If the initiator corresponding to the `preempt_key` is associated with the reservation that is preempted, then the reservation is preempted and any matching registrations are removed. If the initiator corresponding to the `preempt_key` is not associated with the reservation that is preempted, then any matching registrations are removed. Regardless of whether the preempted initiator holds the reservation, all outstanding commands from all initiators corresponding to the `preempt_key` are aborted.

The **STPRES\_REGISTER** IOCTL issues the **persistent reserve out** command with the register service action. The following structure is the argument for this IOCTL.

```
struct st_pres_register {  
    ushort           version;  
    uchar            scsi_status;  
    uchar            sense_key;  
    uchar            scsi_asc;  
    uchar            scsi_ascq;  
};
```

The **STPRES\_REGISTER** IOCTL registers the current host persistent reserve registration key value with the device. The **STPRES\_REGISTER** IOCTL is only supported if the device is opened with a `reserve_type` set to persistent, otherwise an error of EACCESS is returned. The intended use of this IOCTL is to allow a preempted host to regain access to a shared device without requiring that the device is closed and reopened.

If a persistent reserve IOCTL fails, the return code is set to **-1** and the `errno` value is set to one of the following.

- **ENOMEM** Device driver cannot obtain memory to run the command.
- **EFAULT** An error occurred while the caller's data buffer was manipulated
- **EACCES** The device is opened with a **reserve\_type** set to **reserve\_6**
- **EINVAL** The requested IOCTL is not supported by this version of the device driver or invalid parameter that is provided in the argument structure
- **ENXIO** The device indicated that the persistent reserve command is not supported
- **EBUSY** The device returned a SCSI status byte of **RESERVATION CONFLICT** or **BUSY**. Or, the reservation for the device was preempted by another host and the device driver does not issue further commands.
- **EIO** Unknown I/O failure occurred on the command

## Atape persistent reserve IOCTLs

The Atape device driver provides **Persistent Reserve** IOCTLs for application programs to manage their own Persistent Reserve support. These IOCTLs are defined in the header file **/usr/include/sys/Atape\_pr.h**.

The following IOCTLs return Persistent Reserve information by using the **SCSI Persistent Reserve In** command.

- **STIOC\_READ\_RESERVEDKEYS**
- **STIOC\_READ\_RESERVATIONS**
- **STIOC\_READ\_RESERVED\_FULL\_STATUS**

The following IOCTLs complete Persistent Reserve functions by using the **SCSI Persistent Reserve Out** command.

- **STIOC\_REGISTER\_KEY**
- **STIOC\_REMOVE\_REGISTRATION**
- **STIOC\_CLEAR\_ALL\_REGISTRATIONS**
- **STIOC\_PREEEMPT\_RESERVATION**
- **STIOC\_PREEEMPT\_ABORT**
- **STIOC\_CREATE\_PERSISTENT\_RESERVE**

The following IOCTLs are modified to handle both **SCSI-2 Reserve 6** and **Persistent Reserve** based on the current **reserve\_type** setting.

- **SIOC\_RESERVE**
- **SIOC\_RELEASE**

The **STIOC\_READ\_RESERVEDKEYS** IOCTL returns the reservation keys from the device. The argument for this IOCTL is the address of a **read\_keys** structure. If the **reserve\_key\_list** pointer is NULL, then only the generation and length fields are returned. This action allows an application to first obtain the length of the **reserve\_key\_list** and malloc a return buffer before the IOCTL is issued with a **reserve\_key\_list** pointer to that buffer. If the return length is 0, then no reservation keys are registered with the device.

The following structure is used for this IOCTL.

```
struct read_keys
{
    uint generation;           /* counter for PERSISTENT RESERVE OUT requests */
    uint length;               /* number of bytes in the Reservation Key list */
    ullong *reserve_key_list;  /* list of reservation keys */
};
```

The **STIOC\_READ\_RESERVATIONS** IOCTL returns the current reservations from the device if any exist. The argument for this IOCTL is the address of a **read\_reserves** structure. If the **reserve\_list** pointer is NULL, then only the generation and length fields are returned. This action allows an application to

first obtain the length of the **reserve\_list** and malloc a return buffer before the IOCTL is issued with a **reserve\_list** pointer to that buffer. If the return length is 0, then no reservations currently exist on the device.

The following structures are used for this IOCTL.

```
struct reserve_descriptor
{
    ullong      key;          /* reservation key */ */
    uint        scope_spec_addr; /* scope-specific address */ */
    uchar       reserved; */
    uint        scope:4,      /* persistent reservation scope */ */
                type:4;      /* reservation type */ */
    ushort      ext_length;   /* extent length */ */
};

struct read_reserves
{
    uint         generation; /* counter for PERSISTENT RESERVE OUT requests */
    uint         length;     /* number of bytes in the Reservation list */
    struct reserve_descriptor* reserve_list; /* list of reservation key descriptors */
};
```

The **STIOC\_READ\_RESERVE\_FULL\_STATUS** IOCTL returns extended information for all reservation keys and reservations from the device if any exist. The argument for this IOCTL is the address of a **read\_full\_status** structure. If the **status\_list** pointer is NULL, then only the generation and length fields are returned. This action allows an application to first obtain the length of the **status\_list** and malloc a return buffer before the IOCTL is issued with a **status\_list** pointer to that buffer. If the return length is 0, then no reservation keys or reservations currently exist on the device.

The following structures are used for this IOCTL.

```
struct transport_id
{
    uint format_code:2,
        rsvd:2,
        protocol_id:4;
};

struct fcp_transport_id
{
    uint format_code:2,
        rsvd:2,
        protocol_id:4;
    char   reserved1[7];
    ullong n_port_name;
    char   reserved2[8];
};

struct scsi_transport_id
{
    uint format_code:2,
        rsvd:2,
        protocol_id:4;
    char   reserved1[1];
    ushort scsi_address;
    ushort obsolete;
    ushort target_port_id;
    char   reserved2[16];
};

struct sas_transport_id
{
    uint format_code:2,
        rsvd:2,
        protocol_id:4;
    char   reserved1[3];
    ullong sas_address;
    char   reserved2[12];
};

struct status_descriptor
{
    ullong      key;          /* reservation key */ */
    char        reserved1[4];
    uint        rsvd:5,
                spc2_r:1; /* future use for SCSI-2 reserve */
```

```

        all_tg_pt:1;          /* all target ports          */
        r_holder:1;           /* reservation holder        */
        uint scope:4;          /* persistent reservation scope */
        type:4;               /* reservation type          */
        char reserved2[4];
        ushort target_port_id; /* relative target port id */
        uint descriptor_length; /* additional descriptor length */
    union {
        struct transport_id transport_id; /* transport ID             */
        struct fcp_transport_id fcp_id;   /* FCP transport ID         */
        struct sas_transport_id sas_id;   /* SAS transport ID         */
        struct scsi_transport_id scsi_id; /* SCSI transport ID        */
    };
};

struct read_full_status
{
    uint generation;      /* counter for PERSISTENT RESERVE OUT requests */
    uint length;          /* number of bytes for total status descriptors */
    struct status_descriptor *status_list; /* list of reserve status descriptors */
};

```

The **STIOC\_REGISTER\_KEY** IOCTL registers a host reservation key on the device. The argument for this IOCTL is the address of an unsigned long key that can be 1 - 16 hexadecimal digits. If the key value is 0, then the device driver registers the configuration reserve key on the device. This key is either a user-specified host key or the device driver default host key.

If the host has a current persistent reservation on the device and the key is different from the current reservation key, the reservation is retained and the host reservation key is changed to the new key.

The **STIOC\_REMOVE\_REGISTRATION** IOCTL removes the host reservation key and reservation if one exists from the device. There is no argument for this IOCTL. The **SIOC\_RELEASE** IOCTL can also be used to complete the same function.

The **STIOC\_CLEAR\_ALL\_REGISTRATIONS** IOCTL clears all reservation keys and reservations on the device (if any exist) for the same host and any other host. There is no argument for this IOCTL.

The **STIOC\_PREEMPT\_RESERVATION** IOCTL registers a host reservation key on the device and then preempts the reservation that is held by another host if one exists. Or, it creates a new persistent reservation by using the host reservation key. The argument for this IOCTL is the address of an unsigned long key that can be 1 - 16 hexadecimal digits. If the key value is 0, then the device driver registers the configuration reserve key on the device. This key is either a user-specified host key or the device driver default host key.

The **STIOC\_PREEMPT\_ABORT** IOCTL registers a host reservation key on the device, preempts the reservation that is held by another host, and clears the task that is set for the preempted initiator if one exists. Or, it creates a new persistent reservation by using the host reservation key. The argument for this IOCTL is the address of an unsigned long key that can be 1 - 16 hexadecimal digits. If the key value is 0, then the device driver registers the configuration reserve key on the device. This key is either a user-specified host key or the device driver default host key.

The **STIOC\_CREATE\_PERSISTENT\_RESERVE** IOCTL creates a persistent reservation on the device by using the host reservation key that was registered with the **STIOC\_REGISTER\_KEY** IOCTL. There is no argument for this IOCTL. The **SIOC\_RESERVE** IOCTL can also be used to complete the same function.

The **SIOC\_RESERVE** IOCTL reserves the device. If the **reserve\_type** is set to **reserve\_6**, the device driver issues a **SCSI Reserve 6** command. If the **reserve\_type** is set to persistent, the device driver first registers the current host reservation key and then creates a persistent reservation. The current host reservation key can be either the configuration key for the device or a key that was registered previously with the **STIOC\_REGISTER\_KEY** IOCTL.

The **SIOC\_RELEASE** IOCTL releases the device. If the **reserve\_type** is set to **reserve\_6**, the device driver issues a **SCSI Release 6** command. If the **reserve\_type** is set to persistent, the device driver removes the host reservation key and reservation if one exists from the device.

If a persistent reserve IOCTL fails, the return code is set to **-1** and the **errno** value is set to one of the following.

- **ENOMEM** Device driver cannot obtain memory to complete the command.

- **EFAULT** An error occurred while the caller's data buffer was manipulated
- **EACCES** The current open is using a **reserve\_type** set to **reserve\_6**
- **EINVAL** Device does not support either the **SCSI Persistent Reserve In/Out** command, the service action for the command, or the sequence of the command such as issuing the **STIOC\_REMOVE\_REGISTRATION** IOCTL when no reservation key was registered for the host.
- **EBUSY** Device failed the command with reservation conflict. Either a **SCSI-2 Reserve 6** reservation is active, the sequence of the command such as issuing the **STIOC\_CREATE\_PERSISTENT\_RESERVE** IOCTL when no reservation key was registered for the host, or the reservation for the device was preempted by another host and the device driver does not issue further commands.
- **EIO** Unknown I/O failure occurred on the command.

## General IOCTL operations

This chapter describes the IOCTL commands that provide control and access to the tape and medium changer devices. These commands are available for all tape and medium changer devices. They can be issued to any **rmt\***, **rmt\*.smc**, or **smc\*** special file.

### Overview

The following IOCTL commands are supported.

#### **IOCINFO**

Return device information.

#### **STIOCMD**

Issue the **AIX Pass-through** command.

#### **STPASSTHRU**

Issue the **AIX Pass-through** command.

#### **SIOC\_PASSTHRU\_COMMAND**

Issue the **Atape Pass-through** command.

#### **SIOC\_INQUIRY**

Return inquiry data.

#### **SIOC\_REQSENSE**

Return sense data.

#### **SIOC\_RESERVE**

Reserve the device.

#### **SIOC\_RELEASE**

Release the device.

#### **SIOC\_TEST\_UNIT\_READY**

Issue a **SCSI Test Unit Ready** command.

#### **SIOC\_LOG\_SENSE\_PAGE**

Return log sense data for a specific page.

#### **SIOC\_LOG\_SENSE10\_PAGE**

Return log sense data for a specific page and Subpage.

#### **SIOC\_MODE\_SENSE\_PAGE**

Return mode sense data for a specific page.

#### **SIOC\_MODE\_SENSE\_SUBPAGE**

Return mode sense data for a specific page and subpage.

#### **SIOC\_MODE\_SENSE**

Return whole mode sense data include header, block descriptor, and page for a specific page.

#### **SIOC\_MODE\_SELECT\_PAGE**

Set mode sense data for a specific page.

**SIOC\_MODE\_SELECT\_SUBPAGE**

Set mode sense data for a specific page and subpage.

**SIOC\_INQUIRY\_PAGE**

Return inquiry data for a specific page.

**SIOC\_DISABLE\_PATH**

Manually disable (fence) a SCSI path for a device.

**SIOC\_ENABLE\_PATH**

Enable a manually disabled (fenced) SCSI path for a device.

**SIOC\_SET\_PATH**

Explicitly set the current path that is used by the device driver.

**SIOC\_QUERY\_PATH**

Query device and path information for the primary and first alternate SCSI path for a device. This IOCTL is obsolete but still supported. The **SIOC\_DEVICE\_PATHS** IOCTL can be used instead of this IOCTL.

**SIOC\_DEVICE\_PATHS**

Query device and path information for the primary and all alternate SCSI paths for the device.

**SIOC\_RESET\_PATH**

Issue an **Inquiry** command on each SCSI path that is not manually disabled (fenced) and enable the path if the **Inquiry** command succeeds.

**SIOC\_CHECK\_PATH**

Completes the same function as the **SIOC\_RESET\_PATH** IOCTL.

**SIOC\_QUERY\_OPEN**

Returns the process ID that currently has the device opened.

**SIOC\_RESET\_DEVICE**

Issues a SCSI target reset or SCSI lun reset (for FCP or SAS attached) to the device.

**SIOC\_DRIVER\_INFO**

Query the device driver information.

These IOCTL commands and their associated structures are defined by including the **/usr/include/sys/Atape.h** header file in the C program by using the functions.

**IOCINFO**

This IOCTL command provides access to information about the tape or medium changer device. It is a standard AIX IOCTL function.

An example of the **IOCINFO** command is

```
#include <sys/devinfo.h>
#include <sys/Atape.h>
struct devinfo info;

if (!ioctl (fd, IOCINFO, &info))
{
    printf ("The IOCINFO ioctl succeeded\n");
}
else
{
    perror ("The IOCINFO ioctl failed");
}
```

An example of the output data structure for a tape drive **rmt\*** special file is

```
info.devtype=DD_SCTAPE
info.devsubtype=ATAPE_3590
info.un.scmt.type=DT_STREAM
info.un.scmt.blksize=tape block size (0=variable)
```

An example of the output data structure for an integrated medium changer **rmt\*.smc** special file is

```
info.devtype=DD_MEDIUM_CHANGER;  
info.devs subtype=ATAPE_3590;
```

An example of the output data structure for an independent medium changer **smc\*** special file is

```
info.devtype=DD_MEDIUM_CHANGER;  
info.devs subtype=ATAPE_7337;
```

See the **Atape.h** header file for the defined **devs subtype** values.

## **STIOCMD**

This IOCTL command issues the **SCSI Pass-through** command. It is used by the diagnostic and service aid routines. The structure for this command is in the **/usr/include/sys/scsi.h** file.

This IOCTL is supported on both SCSI adapter attached devices and FCP adapter attached devices. For FCP adapter devices, the returned **adapter status** field is converted from the FCP codes that are defined in **/usr/include/sys/scsi\_buf.h** to the SCSI codes defined in **/usr/include/sys/scsi.h**, if possible. This action is to provide downward compatibility with existing applications that use the **STIOCMD** IOCTL for SCSI attached devices.

**Note:** There is no interaction by the device driver with this command. The error handling and logging functions are disabled. If the command results in a check condition, the application must issue a **Request Sense** command to clear any contingent allegiance with the device.

An example of the **STIOCMD** command is

```
struct sc_iocmd sciocmd;  
struct inquiry_data inqdata;  
  
bzero(&sciocmd, sizeof(struct sc_iocmd));  
bzero(&inqdata, sizeof(struct inquiry_data));  
  
/* issue inquiry */  
sciocmd.scsi_cdb[0]=0x12;  
sciocmd.timeout_value=200; /* SECONDS */  
sciocmd.command_length=6;  
sciocmd.buffer=(char *)&inqdata;  
sciocmd.data_length=sizeof(struct inquiry_data);  
sciocmd.scsi_cdb[4]=sizeof(struct inquiry_data);  
sciocmd.flags=B_READ;  
  
if (!ioctl (sffd, STIOCMD, &sciocmd))  
{  
    printf ("The STIOCMD ioctl for Inquiry Data succeeded\n");  
    printf ("\nThe inquiry data is:\n");  
    dump_bytes (&inqdata, sizeof(struct inquiry_data), "Inquiry Data");  
}  
else  
{  
    perror ("The STIOCMD ioctl for Inquiry Data failed");  
}
```

## **STPASSTHRU**

This IOCTL command issues the **AIX Pass-through** command that is supported by base AIX tape device drivers. The IOCTL command and structure are defined in the header files **/usr/include/sys/scsi.h** and **/usr/include/sys/tape.h**. Refer to AIX documentation for information about using the command.

## **SIOC\_PASSTHRU\_COMMAND**

This IOCTL command issues the Atape device driver **Pass-through** command. The data structure that is used on this IOCTL is

```
struct scsi_passthru_cmd {  
    uchar command_length; /* Length of SCSI command 6, 10, 12 or 16 */  
    uchar scsi_cdb[16]; /* SCSI command descriptor block */  
    uint timeout_value; /* Timeout in seconds or 0 for command default */
```

```

    uint  buffer_length;      /* Length of data buffer or 0          */
    char *buffer;            /* Pointer to data buffer or NULL      */
    uint  number_bytes;      /* Number of bytes transferred to/from buffer */
    uchar sense_length;     /* Number of valid sense bytes        */
    uchar sense[MAXSENSE]; /* Sense data when sense length > 0   */
    uint  trace_length;     /* Number bytes in buffer to trace, 0 for none */
    char  read_data_command; /* Input flag, set it to 1 for read type cmds */
    char  reserved[27];
};


```

The **arg** parameter for the IOCTL is the address of a **scsi\_passthru\_cmd** structure.

The device driver issues the SCSI command by using the **command\_length** and **scsi\_cdb** fields. If the command receives data from the device (such as SCSI Inquiry), then the application must also set the **buffer\_length** and **buffer pointer** for the return data along with the **read\_data\_command** set to **1**. For commands that send data to the device (such as **SCSI Mode Select**), the **buffer\_length** and pointer is set for the send data and the **read\_data\_command** set to **0**. If the command has no data transfer, the buffer length is set to 0 and buffer pointer that is set to NULL.

The specified **timeout\_value** field is used if not 0. If 0, then the device driver assigns its internal timeout value that is based on the SCSI command.

The **trace\_length** field is normally used only for debug. It specifies the number of bytes on a data transfer type command that is traced when the AIX Atape device driver trace is running.

If the SCSI command fails, then the IOCTL returns **-1** and **errno** value is set for the failing command. If the device returned sense data for the failure, then the **sense\_length** is set to the number of sense bytes returned in the sense field. If there was no sense data for the failure, the **sense\_length** is 0.

If the SCSI command transfers data either to or from the device, then the **number\_bytes** fields indicate how many bytes were transferred.

## **SIOC\_INQUIRY**

This IOCTL command collects the inquiry data from the device.

The data structure is

```

struct inquiry_data
{
    uint  qual:3,           /* peripheral qualifier */
    type:5;                /* device type */
    uint  rm:1,              /* removable medium */
    mod:7;                 /* device type modifier */
    uint  iso:2,             /* ISO version */
    ecma:3;                /* ECMA version */
    ansi:3;                /* ANSI version */
    uint  aenc:1,             /* asynchronous event notification */
    trmiop:1,               /* terminate I/O process message */
    :2,                     /* reserved */
    rdf:4;                 /* response data format */
    uchar len;               /* additional length */
    uchar resvd1;            /* reserved */
    uint  :4,                /* reserved */
    mchngr:1;               /* Medium Changer mode (SCSI-3 only) */
    :3;                     /* reserved */
    uint  reladr:1,           /* relative addressing */
    wbus32:1,                /* 32-bit wide data transfers */
    wbus16:1,                /* 16-bit wide data transfers */
    sync:1,                  /* synchronous data transfers */
    linked:1,                /* linked commands */
    :1,                     /* reserved */
    cmdque:1,                /* command queuing */
    sftre:1;                 /* soft reset */
    uchar vid[8];            /* vendor ID */
    uchar pid[16];            /* product ID */
    uchar revision[4];        /* product revision level */
    uchar vendor1[20];         /* vendor specific */
    uchar resvd2[40];          /* reserved */
    uchar vendor2[31];         /* vendor specific (padded to 127) */
};


```

An example of the **SIOC\_INQUIRY** command is

```
#include <sys/Atape.h>

struct inquiry_data inquiry_data;

if (!ioctl (fd, SIOC_INQUIRY, &inquiry_data))
{
    printf ("The SIOC_INQUIRY ioctl succeeded\n");
    printf ("\nThe inquiry data is:\n");
    dump_bytes ((uchar *)&inquiry_data, sizeof (struct inquiry_data));
}
else
{
    perror ("The SIOC_INQUIRY ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_REQSENSE**

This IOCTL command returns the device sense data. If the last command resulted in an input/output error (EIO), the sense data is returned for the error. Otherwise, a new sense command is issued to the device.

The data structure is

```
struct request_sense
{
    uint      valid:1,           /* sense data is valid */
              err_code:7;        /* error code */
    uchar     segnum;           /* segment number */
    uint      fm:1,             /* filemark detected */
              eom:1,             /* end of medium */
              ili:1,             /* incorrect length indicator */
              resvd1:1,          /* reserved */
              key:4;             /* sense key */
    signed int info;            /* information bytes */
    uchar     addrlen;          /* additional sense length */
    uint      cmdinfo;          /* command specific information */
    uchar     asc;               /* additional sense code */
    uchar     ascq;              /* additional sense code qualifier */
    uchar     fru;               /* field replaceable unit code */
    uint      sksv:1,            /* sense key specific valid */
              cd:1,              /* control/data */
              resvd2:2,           /* reserved */
              bpv:1,              /* bit pointer valid */
              sim:3;              /* system information message */
    uchar     field[2];          /* field pointer */
    uchar     vendor[109];        /* vendor specific (padded to 127) */
};
```

An example of the **SIOC\_REQSENSE** command is

```
#include <sys/Atape.h>

struct request_sense sense_data;

if (!ioctl (smcfd, SIOC_REQSENSE, &sense_data))
{
    printf ("The SIOC_REQSENSE ioctl succeeded\n");
    printf ("\nThe request sense data is:\n");
    dump_bytes ((uchar *)&sense_data, sizeof (struct request_sense));
}
else
{
    perror ("The SIOC_REQSENSE ioctl failed");
}
```

## **SIOC\_RESERVE**

This IOCTL command reserves the device to the device driver. The specific SCSI command that is issued to the device depends on the current reservation type that is used by the device driver, either a **SCSI Reserve** or **Persistent Reserve**.

There are no arguments for this IOCTL command.

An example of the **SIOC\_RESERVED** command is

```
#include <sys/Atape.h>

if (!ioctl (fd, SIOC_RESERVED, NULL))
{
    printf ("The SIOC_RESERVED ioctl succeeded\n");
}
else
{
    perror ("The SIOC_RESERVED ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_RELEASE**

This IOCTL command releases the current device driver reservation on the device. The specific SCSI command that is issued to the device depends on the current reservation type that is used by the device driver, either a **SCSI Reserve** or **Persistent Reserve**.

There are no arguments for this IOCTL command.

An example of the **SIOC\_RELEASE** command is

```
#include <sys/Atape.h>

if (!ioctl (fd, SIOC_RELEASE, NULL))
{
    printf ("The SIOC_RELEASE ioctl succeeded\n");
}
else
{
    perror ("The SIOC_RELEASE ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_TEST\_UNIT\_READY**

This IOCTL command issues the **SCSI Test Unit Ready** command to the device.

There are no arguments for this IOCTL command.

An example of the **SIOC\_TEST\_UNIT\_READY** command is

```
#include <sys/Atape.h>

if (!ioctl (fd, SIOC_TEST_UNIT_READY, NULL))
{
    printf ("The SIOC_TEST_UNIT_READY ioctl succeeded\n");
}
else
{
    perror ("The SIOC_TEST_UNIT_READY ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_LOG\_SENSE\_PAGE**

This IOCTL command returns a log sense page from the device. The page is selected by specifying the **page\_code** in the **log\_sense\_page** structure. Optionally, a specific **parm** pointer, also known as a **parm code**, and the number of parameter bytes can be specified with the command.

To obtain the entire log page, the **len** and **parm\_pointer** fields are set to zero. To obtain the entire log page that starts at a specific parameter code, set the **parm\_pointer** field to the wanted code and the **len** field to zero. To obtain a specific number of parameter bytes, set the **parm\_pointer** field to the wanted code. Then, set the **len** field to the number of parameter bytes plus the size of the log page header (4 bytes). The first 4 bytes of returned data are always the log page header.

See the appropriate device manual to determine the supported log pages and content.

The data structure is

```

struct log_sense_page
{
    char page_code;
    unsigned short len;
    unsigned short parm_pointer;
    char data[LOGSENSEPAGE];
};

```

An example of the **SIOC\_LOG\_SENSE\_PAGE** command is

```

#include <sys/Atape.h>

struct log_sense_page log_page;
int temp;

/* get log page 0, list of log pages */
log_page.page_code = 0x00;
log_page.len = 0;
log_page.parm_pointer = 0;

if (!ioctl (fd, SIOC_LOG_SENSE_PAGE, &log_page))
{
    printf ("The SIOC_LOG_SENSE_PAGE ioctl succeeded\n");
    dump_bytes(log_page.data, LOGSENSEPAGE);
}
else
{
    perror ("The SIOC_LOG_SENSE_PAGE ioctl failed");
    sioc_request_sense();
}

/* get 3590 fraction of volume traversed */
log_page.page_code = 0x38;
log_page.len = 0;
log_page.parm_pointer = 0x000F;

if (!ioctl (fd, SIOC_LOG_SENSE_PAGE, &log_page))
{
    temp = log_page.data[(sizeof(log_page_header) + 4)];
    printf ("The SIOC_LOG_SENSE_PAGE ioctl succeeded\n");
    printf ("Fractional Part of Volume Traversed %x\n",temp);
}
else
{
    perror ("The SIOC_LOG_SENSE_PAGE ioctl failed");
    sioc_request_sense();
}

```

## **SIOC\_LOG\_SENSE10\_PAGE**

This IOCTL command is enhanced to add a subpage variable from **SIOC\_LOG\_SENSE\_PAGE**. It returns a log sense page or subpage from the device. The page is selected by specifying the **page\_code** or **subpage\_code** in the **log\_sense10\_page** structure. Optionally, a specific **parm** pointer, also known as a **parm** code, and the number of parameter bytes can be specified with the command.

To obtain the entire log page, the **len** and **parm\_pointer** fields are set to zero. To obtain the entire log page that starts at a specific parameter code, set the **parm\_pointer** field to the wanted code and the **len** field to zero. To obtain a specific number of parameter bytes, set the **parm\_pointer** field to the wanted code. Then, set the **len** field to the number of parameter bytes plus the size of the log page header (4 bytes). The first 4 bytes of returned data are always the log page header. See the appropriate device manual to determine the supported log pages and content.

The data structure is

```

/* log sense page and subpage structure */
struct log_sense10_page
{
    uchar page_code;           /* [IN] log sense page code          */
    uchar subpage_code;        /* [IN] log sense Subpage code       */
    uchar reserved[2];
    unsigned short len;        /* [IN] specific allocation length for the data */
                                /* [OUT] number of valid bytes in   */
                                /*      data(log_page_header_size+page_length) */

```

```

    unsigned short parm_pointer;
        /* [IN] specific parameter number at which
           the data begins */
    char data[LOGSENSEPAGE]; /* [OUT] log sense page and Subpage data */
    */
};


```

An example of the **SIOC\_LOG\_SENSE10\_PAGE** command is

```

#include <sys/Atape.h>

struct log_sense10_page logdata10;
struct log_page_header *page_header;
char text[80];

logdata10.page_code = page;
logdata10.subpage_code = subpage;
logdata10.len = len;
logdata10.parm_pointer = parm;
page_header = (struct log_page_header *)logdata10.data;

printf("Issuing log sense for page 0x%02X and subpage 0x%02X...\n",page,subpage);

if (!ioctl (fd, SIOC_LOG_SENSE10_PAGE, &logdata10))
{
    sprintf(text,"Log Sense Page 0x%02X, Subpage 0x%02X, Page Length %d
Data",page,subpage,logdata10.len);
    dump_bytes(logdata10.data,logdata10.len,text);
}
else
{
    perror ("The SIOC_LOG_SENSE10_PAGE ioctl failed");
    sioc_request_sense();
}


```

### ***SIOC\_MODE\_SENSE\_PAGE***

This IOCTL command returns a mode sense page from the device. The page is selected by specifying the **page\_code** in the **mode\_sense\_page** structure.

See the appropriate device manual to determine the supported mode pages and content.

The data structure is

```

struct mode_sense_page
{
    char page_code;
    char data[MODESENSEPAGE];
};


```

An example of the **SIOC\_MODE\_SENSE\_PAGE** command is

```

#include <sys/Atape.h>

struct mode_sense_page mode_page;

/* get Medium Changer mode */
mode_page.page_code = 0x20;
if (!ioctl (fd, SIOC_MODE_SENSE_PAGE, &mode_page))
{
    printf ("The SIOC_MODE_SENSE_PAGE ioctl succeeded\n");
    if (mode_page.data[2] == 0x02)
        printf ("The library is in Random mode.\n");
    else
        if (mode_page.data[2] == 0x05)
            printf ("The library is in Automatic (Sequential) mode.\n");
}
else
{
    perror ("The SIOC_MODE_SENSE_PAGE ioctl failed");
    sioc_request_sense();
}


```

## **SIOC\_MODE\_SENSE\_SUBPAGE**

This IOCTL command returns the whole mode sense data, including header, block descriptor, and page code for a specific page or subpage from the device. The wanted page or subpage is inputted by specifying the **page\_code** and **subpage\_code** in the **mode\_sense** structure.

The data structure is

```
struct mode_sense
{
    uchar page_code; /* [IN] mode sense page code */
    uchar subpage_code; /* [IN] mode sense subpage code */
    uchar reserved[6];
    uchar cmd_code; /* [OUT] SCSI Command Code: this field is set with */
                    /* SCSI command code which the device responded. */
                    /* x'5A' = Mode Sense (10) */
                    /* x'1A' = Mode Sense (6) */
    char data[MODESENSEPAGE]; /* [OUT] whole mode sense data include header,
                                block descriptor and page */
};
```

An example of the **SIOC\_MODE\_SENSE** command is

```
#include <sys/Atape.h>

struct mode_sense modedata;
char text[80];
bzero(&modedata, sizeof(struct mode_sense));
modedata.page_code = page;
modedata.subpage_code = subpage;

printf("Issuing mode sense subpage for page 0x%02X subpage 0x%02X...\n",
       page, subpage);

if (!ioctl (fd, SIOC_MODE_SENSE, &modedata))
{
    sprintf(text,"Mode Sense 0x%02X Subpage 0x%02X cmd_code 0x%02X",
           modedata.page_code,modedata.subpage_code,modedata.cmd_code);
    dump_bytes((char *)&modedata, sizeof(struct mode_sense), text);
}
else
{
    perror ("The SIOC_MODE_SENSE ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_MODE\_SELECT\_PAGE**

This IOCTL command sets device parameters in a specific mode page. The wanted page is selected by specifying the **page\_code** in the **mode\_sense\_page** structure. See the appropriate device manual to determine the supported mode pages and parameters that can be modified. The **arg** parameter for the IOCTL is the address of a **mode\_sense\_page** structure.

The data structure is

```
struct mode_sense_page
{
    uchar page_code; /* mode sense page code */
    char data[MODESENSEPAGE];
};
```

This data structure is also used for the **SIOC\_MODE\_SENSE\_PAGE** IOCTL. The application must issue the **SIOC\_MODE\_SENSE\_PAGE** IOCTL, and modify the wanted bytes in the returned **mode\_sense\_page** structure data field. Then, it issues this IOCTL with the modified fields in the structure.

## **SIOC\_MODE\_SELECT\_SUBPAGE**

This IOCTL command sets device parameters in a specific mode page and subpage. The wanted page and subpage are selected by specifying the **page\_code** and **subpage\_code** in the **mode\_sense\_subpage** structure. See the appropriate device manual to determine the supported mode pages, subpages,

and parameters that can be modified. The **arg** parameter for the IOCTL is the address of a **mode\_sense\_subpage** structure.

The data structure is

```
struct mode_sense_subpage
{
    uchar page_code;           /* mode sense page code */
    uchar subpage_code;        /* mode sense subpage code */
    uint reserved:7;
    sp_bit:1;                 /* mode select save page bit */
    char data[MODESENSEPAGE];
};
```

This data structure is also used for the **SIOC\_MODE\_SENSE\_SUBPAGE** IOCTL. The application must issue the **SIOC\_MODE\_SENSE\_SUBPAGE** IOCTL, and modify the wanted bytes in the returned **mode\_sense\_subpage** structure data field. Then, it issues this IOCTL with the modified fields in the structure. If the device supports setting the **sp** bit for the mode page to **1**, then the **sp\_bit** field can be set to **0** or **1**. If the device does not support the **sp** bit, then the **sp\_bit** field must be set to **0**.

## **SIOC\_QUERY\_OPEN**

This IOCTL command returns the ID of the process that currently has a device open. There is no associated data structure. The **arg** parameter specifies the address of an **int** for the return process ID.

If the application opened the device by using the extended **open** parameter **SC\_TMCP**, the process ID is returned for any other process that has the device open currently. Or, zero is returned if the device is not currently open. If the application opened the device without the extended **open** parameter **SC\_TMCP**, the process ID of the current application is returned.

An example of the **SIOC\_QUERY\_OPEN** command is

```
#include <sys/Atape.h>

int sioc_query_open (void)
{
    int pid = 0;

    if (ioctl(fd, SIOC_QUERY_OPEN, &pid) == 0)
    {
        if (pid)
            printf("Device is currently open by process id %d\n",pid)
        else
            printf("Device is not open\n");
    }
    else
        printf("Error querying device open...\n");

    return errno;
}
```

## **SIOC\_INQUIRY\_PAGE**

This IOCTL command returns an inquiry page from the device. The page is selected by specifying the **page\_code** in the **inquiry\_page** structure.

See the appropriate device manual to determine the supported inquiry pages and content.

The data structure is

```
struct inquiry_page
{
    char page_code;
    char data[INQUIRYPAGE];
};
```

An example of the **SIOC\_INQUIRY\_PAGE** command is

```
#include <sys/Atape.h>
```

```

struct inquiry_page inq_page;

/* get inquiry page x83 */
inq_page.page_code = 0x83;
if (!ioctl(fd, SIOC_INQUIRY_PAGE, &inq_page))
{
    printf ("The SIOC_INQUIRY_PAGE ioctl succeeded\n");
}
else
{
    perror ("The SIOC_INQUIRY_PAGE ioctl failed");
    sioc_request_sense();
}

```

## **SIOC\_DISABLE\_PATH**

This IOCTL command manually disables (fences) the device driver from using either the primary or an alternate SCSI path to a device until the **SIOC\_ENABLE\_PATH** command is issued for the same path that is manually disabled. The **arg** parameter on the IOCTL command specifies the path to be disabled. The primary path is path 1, the first alternate path 2, the second alternate path 3, and so on. This command can be used concurrently when the device is already open by another process by using the openx() extended parameter **SC\_TMCP**.

This IOCTL command is valid only if the device has one or more alternate paths configured. Otherwise, the IOCTL command fails with *errno* set to EINVAL. The **SIOC\_DEVICE\_PATHS** IOCTL command can be used to determine the paths that are enabled or manually disabled.

An example of the **SIOC\_DISABLE\_PATH** command is

```

#include <sys/Atape.h>

/* Disable primary SCSI path */
ioctl(fd, SIOC_DISABLE_PATH, PRIMARY_SCSI_PATH);

/* Disable alternate SCSI path */
ioctl(fd, SIOC_DISABLE_PATH, ALTERNATE_SCSI_PATH);

```

## **SIOC\_ENABLE\_PATH**

This IOCTL command enables a manually disabled (fenced) path to a device that is disabled by **SIOC\_DISABLE\_PATH** IOCTL. The **arg** parameter on the IOCTL command specifies the path to be enabled. The primary path is path 1, the first alternate path 2, the second alternate path 3, and so on. This command can be used concurrently when the device is already open by another process by using the openx() extended parameter **SC\_TMCP**.

The **SIOC\_DEVICE\_PATHS** IOCTL command can be used to determine the paths that are enabled or manually disabled.

## **SIOC\_SET\_PATH**

This IOCTL command explicitly sets the current path to a device that the device driver uses. The **arg** parameter on the IOCTL command specifies the path to be set to the current path. The primary path is path 1, the first alternate path 2, the second alternate path 3, and so on. This command can be used concurrently when the device is already open by another process by using the openx() extended parameter **SC\_TMCP**.

The **SIOC\_DEVICE\_PATHS** IOCTL command can be used to determine the current path the device driver is using for the device.

## **SIOC\_DEVICE\_PATHS**

This IOCTL command returns a **device\_paths** structure. The number of paths are configured to a device and a **device\_path\_t** path structure for each configured path. The device, HBA, and path information for the primary path are configured along with all alternate SCSI paths. This IOCTL command must be used instead of the **SIOC\_QUERY\_PATH** IOCTL that is obsolete. This command can be used concurrently when the device is already open by another process by using the openx() extended parameter **SC\_TMCP**.

The data structures are

```
struct device_path_t {
    char name[15];                                /* logical device name */
    char parent[15];                               /* logical parent name */
    uchar id_valid;                                /* obsolete and not set */
    uchar id;                                     /* SCSI target address of device */
    uchar lun;                                    /* SCSI logical unit of device */
    uchar bus;                                    /* SCSI bus for device */
    uchar fcp_id_valid;                            /* FCP scsi/lun id fields valid */
    unsigned long long fcp_scsi_id;                /* FCP SCSI id of device */
    unsigned long long fcp_lun_id;                 /* FCP logical unit of device */
    unsigned long long fcp_ww_name;                /* FCP world wide name */
    uchar enabled;                                 /* path enabled */
    uchar drive_port_valid;                        /* drive port field valid */
    uchar drive_port;                             /* drive port number */
    uchar fenced;                                 /* path fenced by disable ioctl */
    uchar current_path;                           /* Current path assignment */
    uchar dynamic_tracking;                      /* FCP Dynamic tracking enabled */
    unsigned long long fcp_node_name;              /* FCP node name */
    char type[16];                                /* Device type and model */
    char serial[16];                              /* Device serial number */
    uchar sas_id_valid;                            /* FCP scsi/lun id fields valid */
    char cpname[15];                             /* logical name of control path drive */
    uchar last_path;                             /* Last failure path */
    char reserved[4];
};

struct device_paths {
    int number_paths;                            /* number of paths configured */
    struct device_path_t path[MAX_SCSI_PATH];
};
```

The **arg** parameter for the IOCTL is the address of a **device\_paths** structure.

The **current\_path** in the return structures is set to the current path the device uses for the device. If this IOCTL is issued to a medium changer smc logical driver, the **cpname** has the logical rmt name that is the control path drive for each smc logical path.

## **SIOC\_QUERY\_PATH**

This IOCTL command returns information about the device and SCSI paths, such as logical parent, SCSI IDs, and status of the SCSI paths.

**Note:** This IOCTL is obsolete but still supported. The **SIOC\_DEVICE\_PATHS** IOCTL must be used instead.

The data structure is

```
struct scsi_path {
    char primary_name[15];                         /* Primary logical device name */
    char primary_parent[15];                        /* Primary SCSI parent name */
    uchar primary_id;                             /* Primary target address of device */
    uchar primary_lun;                            /* Primary logical unit of device */
    uchar primary_bus;                            /* Primary SCSI bus for device */
    unsigned long long primary_fcp_scsi_id;        /* Primary FCP SCSI id of device */
    unsigned long long primary_fcp_lun_id;          /* Primary FCP logical unit of device */
    unsigned long long primary_fcp_ww_name;         /* Primary FCP world wide name */
    uchar primary_enabled;                         /* Primary path enabled */
    uchar primary_id_valid;                        /* Primary id/lun/bus fields valid */
    uchar primary_fcp_id_valid;                    /* Primary FCP scsi/lun id fields valid */
    uchar alternate_configured;                   /* Alternate path configured */
    char alternate_name[15];                        /* Alternate logical device name */
    char alternate_parent[15];                      /* Alternate SCSI parent name */
    uchar alternate_id;                            /* Alternate target address of device */
    uchar alternate_lun;                           /* Alternate logical unit of device */
    uchar alternate_bus;                           /* Alternate SCSI bus for device */
    unsigned long long alternate_fcp_scsi_id;       /* Alternate FCP SCSI id of device */
    unsigned long long alternate_fcp_lun_id;         /* Alternate FCP logical unit of device */
    unsigned long long alternate_fcp_ww_name;        /* Alternate FCP world wide name */
    uchar alternate_enabled;                        /* Alternate path enabled */
    uchar alternate_id_valid;                      /* Alternate id/lun/bus fields valid */
};
```

```

uchar alternate_fcp_id_valid;           /* Alternate FCP scsi/lun id fields
                                         valid */
uchar primary_drive_port_valid;        /* Primary drive port field valid */
uchar primary_drive_port;              /* Primary drive port number */
uchar alternate_drive_port_valid;      /* Alternate drive port field valid */
uchar alternate_drive_port;            /* Alternate drive port number */
uchar primary_fenced;                 /* Primary fenced by disable ioctl */
uchar alternate_fenced;               /* Alternate fenced by disable ioctl */
uchar current_path;                  /* Current path assignment */
uchar primary_sas_id_valid;           /* Primary FCP scsi/lun id fields
                                         valid */
uchar alternate_sas_id_valid;          /* Alternate FCP scsi/lun id fields
                                         valid */
char reserved[55];  };

```

An example of the **SIOC\_QUERY\_PATH** command is

```

#include <sys/Atape.h>

int sioc_query_path(void)
{
    struct scsi_path path;

    printf("Querying SCSI paths...\n");

    if (ioctl(fd, SIOC_QUERY_PATH, &path) == 0)
        show_path(&path);

    return errno;
}

void show_path(struct scsi_path *path)
{
    printf("\n");
    if (path->alternate_configured)
        printf("Primary Path Information:\n");
    printf("  Logical Device..... %s\n",path->primary_name);
    printf("  SCSI Parent..... %s\n",path->primary_parent);
    if (path->primary_fcp_id_valid)
    {
        if (path->primary_id_valid)
        {
            printf("    Target ID..... %d\n",path->primary_id);
            printf("    Logical Unit..... %d\n",path->primary_lun);
            printf("    SCSI Bus..... %d\n",path->primary_bus);
        }
        printf("    FCP SCSI ID..... 0x%llx\n",path->primary_fcp_scsi_id);
        printf("    FCP Logical Unit..... 0x%llx\n",path->primary_fcp_lun_id);
        printf("    FCP World Wide Name..... 0x%llx\n",path->primary_fcp_ww_name);
    }
    else
    {
        printf("    Target ID..... %d\n",path->primary_id);
        printf("    Logical Unit..... %d\n",path->primary_lun);
    }
    if (path->primary_drive_port_valid)
        printf("    Drive Port Number..... %d\n",path->primary_drive_port);
    if (path->primary_enabled)
        printf("    Path Enabled..... Yes\n");
    else
        printf("    Path Enabled..... No \n");
    if (path->primary_fenced)
        printf("    Path Manually Disabled..... Yes\n");
    else
        printf("    Path Manually Disabled..... No \n");

    if (!path->alternate_configured)
        printf("  Alternate Path Configured..... No\n");
    else
    {
        printf("  Alternate Path Configured..... Yes\n");
        printf("\nAlternate Path Information:\n");
        printf("  Logical Device..... %s\n",path->alternate_name);
        printf("  SCSI Parent..... %s\n",path->alternate_parent);
        if (path->alternate_fcp_id_valid)
        {
            if (path->alternate_id_valid)
            {
                printf("    Target ID..... %d\n",path->alternate_id);

```

```

        printf(" Logical Unit..... %d\n",path->alternate_lun);
        printf(" SCSI Bus..... %d\n",path->alternate_bus);
    }
    printf(" FCP SCSI ID..... 0x%llx\n",path->alternate_fcp_scsi_id);
    printf(" FCP Logical Unit..... 0x%llx\n",path->alternate_fcp_lun_id);
    printf(" FCP World Wide Name..... 0x%llx\n",path->alternate_fcp_ww_name);
}
else
{
    printf(" Target ID..... %d\n",path->alternate_id);
    printf(" Logical Unit..... %d\n",path->alternate_lun);
}
if (path->alternate_drive_port_valid)
    printf(" Drive Port Number..... %d\n",path->alternate_drive_port);
if (path->alternate_enabled)
    printf(" Path Enabled..... Yes\n");
else
    printf(" Path Enabled..... No \n");
if (path->alternate_fenced)
    printf(" Path Manually Disabled..... Yes\n");
else
    printf(" Path Manually Disabled..... No \n");

}
}

```

### ***SIOC\_RESET\_PATH and SIOC\_CHECK\_PATH***

Both of these IOCTL commands check all SCSI paths to a device that are not manually disabled by the **SIOC\_DISABLE\_PATH** IOCTL. It is done by issuing a **SCSI Inquiry** command on each path to verify communication. If the command succeeds, then the path is enabled. If it fails, the path is disabled and is not used by the device driver. This command can be used concurrently when the device is already open by another process by using the openx() extended parameter **SC\_TMCP**.

This IOCTL command returns the same data structure as the **SIOC\_QUERY\_PATH** IOCTL command with the updated path information for the primary and first alternate path. See the **SIOC\_QUERY\_PATH** IOCTL command for a description of the data structure and output information. If more than one alternate path is configured for the device, then the **SIOC\_DEVICE\_PATHS** IOCTL must be used to determine the paths that are enabled.

An example of the **SIOC\_RESET\_PATH** command is

```

#include <sys/Atape.h>

int sioc_reset_path(void)
{
    struct scsi_path path;

    printf("Resetting SCSI paths...\n");
    if (ioctl(fd, SIOC_RESET_PATH, &path) == 0)
        show_path(&path);
    return errno;
}

```

### ***SIOC\_RESET\_DEVICE***

This IOCTL command issues a SCSI target reset to the device if parallel SCSI is attached or a SCSI lun reset if FCP/SAS is attached to the device. This IOCTL command can be used to clear a SCSI Reservation that is active on the device. This command can be used concurrently when the device is already open by another process by using the openx() extended parameter **SC\_TMCP**.

There is no argument for this IOCTL and the **arg** parameter is ignored.

### ***SIOC\_DRIVER\_INFO***

This command returns the information about the currently installed Atape driver.

The following data structure is filled out and returned by the driver.

```

struct driver_info {
    uchar dd_name[16];           /* Atape driver name (Atape) */
    uchar dd_version[16];        /* Atape driver version e.g. 12.0.8.0 */
    uchar os[16];                /* Operating System (AIX) */
    uchar os_version[32];        /* Running OS Version e.g. 6.1 */
    uchar sys_arch[16];          /* Sys Architecture (POWER or others) */
    uchar reserved[32];          /* Reserved for IBM Development Use */
};

```

An example of the **SIOC\_DRIVER\_INFO** command is

```

#include <sys/Atape.h>

int sioc_driver_info()
{
    struct driver_info dd_info;

    printf("Issuing driver info...\n");

    if (!ioctl (fd, SIOC_DRIVER_INFO, &dd_info))
    {
        printf("Driver Name:      %s\n",dd_info.dd_name);
        printf("Driver Version:   %s\n",dd_info.dd_version);
        printf("Operating System: %s\n",dd_info.os);
        printf("OS Version:       %s\n",dd_info.os_version);
        printf("System Arch:      %s\n",dd_info.sys_arch);
    }
    return errno;
}

```

## Tape IOCTL operations

The device driver supports the tape IOCTL commands available with the base AIX operating system. In addition, it supports a set of expanded tape IOCTL commands that give applications access to extra features and functions of the tape drives.

### Overview

The following IOCTL commands are supported.

#### **STIOCHGP**

Set the block size.

#### **STIOCTOP**

Complete the IOCTL tape operation.

#### **STIOCQRYP**

Query the tape device, device driver, and media parameters.

#### **STIOCSETP**

Change the tape device, device driver, and media parameters.

#### **STIOCSYNC**

Synchronize the tape buffers with the tape.

#### **STIOCDM**

Display the message on the display panel.

#### **STIOCQRYPOS**

Query the tape position and the buffered data.

#### **STIOCSETPOS**

Set the tape position.

#### **STIOCQRYSENSE**

Query the sense data from the tape device.

#### **STIOCQRYINQUIRY**

Return the inquiry data.

**STIOC\_LOG\_SENSE**

Return the log sense data.

**STIOC\_RECOVER\_BUFFER**

Recover the buffered data from the tape device.

**STIOC\_LOCATE**

Locate to the tape position.

**STIOC\_READ\_POSITION**

Read the current tape position.

**STIOC\_SET\_VOLID**

Set the volume name for the current mounted tape. The name is used for tape volume logging only.

**STIOC\_DUMP**

Force and read a dump from the device.

**STIOC\_FORCE\_DUMP**

Force a dump on the device.

**STIOC\_READ\_DUMP**

Read a dump from the device.

**STIOC\_LOAD\_UCODE**

Download the microcode to the device.

**STIOC\_RESET\_DRIVE**

Issue a **SCSI Send Diagnostic** command to reset the tape drive.

**STIOC\_FMR\_TAPE**

Create an FMR tape.

**MTDEVICE**

Obtain the device number of a drive in an IBM Enterprise Tape Library 3494.

**STIOC\_PREVENT\_MEDIUM\_REMOVAL**

Prevent medium removal by an operator.

**STIOC\_ALLOW\_MEDIUM\_REMOVAL**

Allow medium removal by an operator.

**STIOC\_REPORT\_DENSITY\_SUPPORT**

Return supported densities from the tape device.

**STIOC\_GET\_DENSITY**

Get the current write density settings from the tape device.

**STIOC\_SET\_DENSITY**

Set the write density settings on the tape device.

**STIOC\_CANCEL\_ERASE**

Cancel an **erase immediate** command that is in progress.

**GET\_ENCRYPTION\_STATE**

This IOCTL can be used for application, system, and library-managed encryption. It allows a query only of the encryption status.

**SET\_ENCRYPTION\_STATE**

This IOCTL can be used only for application-managed encryption. It sets encryption state for application-managed encryption.

**SET\_DATA\_KEY**

This IOCTL can be used only for application-managed encryption. It sets the data key for application-managed encryption.

**READ\_TAPE\_POSITION**

Read current tape position in either short, long, or extended form.

**SET\_TAPE\_POSITION**

Set the current tape position to either a logical object or logical file position.

**CREATE\_PARTITION**

Create one or more tape partitions and format the media.

**QUERY\_PARTITION**

Query tape partitioning information and current active partition.

**SET\_ACTIVE\_PARTITION**

Set the current active tape partition.

**ALLOW\_DATA\_OVERWRITE**

Set the drive to allow a subsequent data overwrite type command at the current position or allow a **CREATE\_PARTITION** IOCTL when data safe (append-only) mode is enabled.

**QUERY\_LOGICAL\_BLOCK\_PROTECTION**

Query Logical Block Protection (LBP) support and its setup.

**SET\_LOGICAL\_BLOCK\_PROTECTION**

Enable or disable Logical Block Protection (LBP), set the protection method, and how the protection information is transferred.

**STIOC\_READ\_ATTRIBUTE**

Read attribute values from medium auxiliary memory.

**STIOC\_WRITE\_ATTRIBUTE**

Write attribute values to medium auxiliary memory.

**VERIFY\_TAPE\_DATA**

Read the data from tape and verify its correction.

**QUERY\_RAO\_INFO**

Query the maximum number and size of User Data Segments (UDS).

**GENERATE\_RAO**

Send a GRAO list to request the drive to generate a **Recommended Access Order** list.

**RECEIVE\_RAO**

Receive a **Recommended Access Order** list of UDS from the drive.

**QUERY\_ARCHIVE\_MODE\_UNTHREAD**

Query for Archive Mode Unthread support in tape cartridge loading.

**SET\_ARCHIVE\_MODE\_UNTHREAD**

Set Archive Mode Unthread enable or disable for tape cartridge loading.

**TAPE\_LOAD\_UNLOAD**

Load or unload tape cartridge with various behaviors.

**GET\_VHF\_DEVICE\_STATUS**

Report very high frequency data for the tape drive and medium statuses.

These IOCTL commands and their associated structures are defined in the **/usr/include/sys/Atape.h** header file, which is included in the corresponding C program that uses the functions.

**STIOCHGP**

This IOCTL command sets the current block size. A block size of zero is a variable block. Any other value is a fixed block.

An example of the **STIOCHGP** command is

```
#include <sys/Atape.h>

struct stchgp stchgp;
stchgp.st_blksize = 512;
if (ioctl(tapefd,STIOCHGP,&stchgp)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## **STIOCTOP**

This IOCTL command runs basic tape operations. The *st\_count* variable is used for many of its operations. Normal error recovery applies to these operations. The device driver can issue several tries to complete them.

For all **space** operations, the tape position finishes on the end-of-tape side of the record or filemark for forward movement and on the beginning-of-tape side of the record or filemark for backward movement. The only exception occurs for forward and backward **space record** operations over a filemark if the device is configured for the AIX **record space** mode.

The input data structure is

```
struct stop
{
    short st_op;           /* operations defined below */
    daddr_t st_count;      /* how many of them to do (if applicable) */
};
```

The *st\_op* variable is set to one of the following operations.

### **STOFFL**

Unload the tape. The **st\_count** parameter does not apply.

### **STREW**

Rewind the tape. The **st\_count** parameter does not apply.

### **STERASE**

Erase the entire tape. The **st\_count** parameter does not apply.

### **STERASE\_IMM**

Erase the entire tape with the immediate bit set. The **st\_count** parameter does not apply.

This action issues the **erase** command to the device with the immediate bit set in the SCSI CDB. When this command is used, another process can cancel the erase operation by issuing the **STIOC\_CANCEL\_ERASE** IOCTL. The application that issued the **STERASE\_IMM** still waits for the **erase** command to complete like the **STERASE st\_op** if the **STIOC\_CANCEL\_ERASE** IOCTL is not issued. Refer to for a description of the **STIOC\_CANCEL\_ERASE** IOCTL.

### **STERASEGAP**

Erase the gap that was written to the tape. The **st\_count** parameter does not apply.

### **STRETEN**

Start the rewind operation. The tape devices run the retension operation automatically when needed.

### **STWEOF**

Write the **st\_count** number of filemarks.

### **STWEOF\_IMM**

Write the **st\_count** number of filemarks with the immediate bit set.

This action issues a **write filemark** command to the device with the immediate bit set in the SCSI CDB. The device returns immediate status and the IOCTL also returns immediately. Unlike the **STWEOF st\_op**, any buffered write data are not flushed to tape before the filemarks are written. This action can improve the time that it takes for a **write filemark** command to complete.

### **STFSF**

Space forward the **st\_count** number of filemarks.

### **STRSF**

Space backward the **st\_count** number of filemarks.

### **STFSR**

Space forward the **st\_count** number of records.

### **STRSR**

Space backward the **st\_count** number of records.

### **STTUR**

Issue the **Test Unit Ready** command. The **st\_count** parameter does not apply.

### **STLOAD**

Issue the **SCSI Load** command. The **st\_count** parameter does not apply. The operation of the **SCSI Load** command varies depending on the type of device. See the appropriate hardware reference manual.

### **STSEOD**

Space forward to the end of the data. The **st\_count** parameter does not apply. This operation is supported except on the IBM 3490E tape devices.

### **STFSSF**

Space forward to the first **st\_count** number of contiguous filemarks.

### **STRSSF**

Space backward to the first **st\_count** number of contiguous filemarks.

### **STEJECT**

Unload the tape. The **st\_count** parameter does not apply.

### **STINSLT**

Issue the **SCSI Load** command. The **st\_count** parameter does not apply.

**Note:** If zero is used for operations that require the **count** parameter, the command is not issued to the device, and the device driver returns a successful completion.

An example of the **STIOCTOP** command is

```
#include <sys/Atape.h>

struct stop stop;
stop.st_op=STWEOF;
stop.st_count=3;
if (ioctl(tapefd,STIOCTOP,&stop)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

### **STIOCQRYP or STIOCSETP**

The **STIOCQRYP** IOCTL command allows the program to query the tape device, device driver, and media parameters. The **STIOCSETP** IOCTL command allows the program to change the tape device, device driver, and media parameters. Before the **STIOCSETP** IOCTL command is issued, use the **STIOCQRYP** IOCTL command to query and fill the fields of the data structure that you do not want to change. Then, issue the **STIOCSETP** command to change the selected fields.

Changing certain fields (such as **buffered\_mode**) impacts performance. If the **buffered\_mode** field is false, then each record that is written to the tape is transferred to the tape immediately. This operation guarantees that each record is on the tape, but it impacts performance.

### **STIOCQRYP parameters that cannot be changed with the STIOCSETP IOCTL command**

The following parameters that are returned by the **STIOCQRYP** IOCTL command cannot be changed by the **STIOCSETP** IOCTL command.

- trace**

This parameter is the current setting of the AIX system tracing for channel 0. All Atape device driver events are traced in channel 0 with other kernel events. If set to On, device driver tracing is active.

- hwrd**

This parameter is the trace hookword used for Atape events.

- write\_protect**

If the currently mounted tape is write-protected, this field is set to TRUE. Otherwise, it is set to FALSE.

- **min\_blksize**

This parameter is the minimum block size for the device. The driver sets this field by issuing the **SCSI Read Block Limits** command.

- **max\_blksize**

This parameter is the maximum block size for the device. The driver sets this field by issuing the **SCSI Read Block Limits** command.

- **max\_scsi\_xfer**

This parameter is the maximum transfer size of the parent SCSI adapter for the device.

- **acf\_mode**

If the tape device has the ACF installed, this parameter returns the current mode of the ACF. Otherwise, the value of ACF\_NONE is returned. The ACF mode can be set from the operator panel on the tape device.

- **alt\_pathing**

This parameter is the configuration setting for path failover support. If the path failover support is enabled, this parameter is set to TRUE.

- **medium\_type**

This parameter is the media type of the current loaded tape. Some tape devices support multiple media types and report different values in this field. See the documentation for the specific tape device to determine the possible values.

- **density\_code**

This parameter is the density setting for the current loaded tape. Some tape devices support multiple densities and report the current setting in this field. See the documentation for the specific tape device to determine the possible values.

- **reserve\_type**

This parameter is the configuration setting for the reservation type that the device driver uses when the device is reserved, either a **SCSI Reserve 6** command or a **SCSI Persistent Reserve** command.

- **reserve\_key**

This parameter is the reservation key the device driver uses with SCSI Persistent Reserve. If a configuration reservation key was specified, then this key can be either a 1-8 ASCII character key or a 1-16 hexadecimal key. If a configuration key was not specified, then the reservation key is a 16 hexadecimal key that the device driver generates.

## Parameters that can be changed with **STIOCSETP IOCTL** command

The following parameters can be changed with the **STIOCSETP** IOCTL command.

- **blksize**

This parameter specifies the effective block size for the tape device.

- **autoload**

This parameter turns the autoload feature On and Off in the device driver. If set to On, the cartridge loader is treated as a large virtual tape.

- **buffered\_mode**

This parameter turns the buffered mode write On and Off.

- **compression**

This parameter turns the hardware compression On and Off.

- **trailer\_labels**

If this parameter is set to On, writing a record past the early warning mark on the tape is allowed. The first **write** operation to detect EOM returns the ENOSPC error code. This **write** operation does not complete successfully. All subsequent **write** operations are allowed to continue despite the check conditions that result from EOM. When the end of the physical volume is reached, EIO is returned. This parameter can be used before EOM or after EOM is reached.

- **rewind\_immediate**

This parameter turns the immediate bit On and Off in rewind commands. If set to On, the STREW tape operation runs faster. However, the next command takes a long time to finish unless the rewind operation is physically complete.

- **logging**

This parameter turns the volume logging On and Off. If set to On, the volume log data is collected and saved in the tape log file when the **Rewind** and **Unload** command is issued to the tape drive.

- **volid**

This parameter is the volume ID of the current loaded tape. If it is not set, the device driver initializes the **volid** to UNKNOWN. If logging is active, the parameter is used to identify the volume in the tape log file entry. It is reset to UNKNOWN when the tape is unloaded.

- **emulate\_autoloader**

This parameter turns the emulate autoloader feature On and Off.

- **record\_space\_mode**

This parameter specifies how the device driver operates when a forward or backward **space record** operation encounters a filemark. The two modes of operation are SCSI and AIX.

- **logical\_write\_protect**

This parameter sets or resets the logical write protect of the current tape.

**Note:** The tape position must be at the beginning of the tape to change this parameter from its current value.

- **capacity\_scaling and capacity\_scaling\_value**

The **capacity\_scaling** parameter queries the capacity or logical length of the current tape or on a set operation changes the current tape capacity. On a query operation, this parameter returns the current capacity for the tape. It is one of the defined values such as SCALE\_100, SCALE\_75, SCALE\_VALUE. If the query returns SCALE\_VALUE, then the **capacity\_scaling\_value** parameter is the current capacity. Otherwise, the **capacity\_scaling** parameter is the current capacity.

On a set operation, if the **capacity\_scaling** parameter is set to SCALE\_VALUE then the **capacity\_scaling\_value** parameter is used to set the tape capacity. Otherwise, one of the other defined values for the **capacity\_scaling** parameter is used.

**Note:**

1. The tape position must be at the beginning of the tape to change this parameter from its current value.
2. Changing this parameter destroys any existing data on the tape.

- **retain\_reservation**

When this parameter is set to 1, the device driver does not release the device reservation when the device is closed for the current open and any subsequent opens and closes until the **STIOCSETP** IOCTL is issued with **retain\_reservation** parameter set to 0. The device driver still reserves the device on open to make sure that the previous reservation is still valid.

- **data\_safe\_mode**

This parameter queries the current drive setting for data safe (append-only) mode. Also, on a set operation it changes the current data safe mode setting on the drive. On a set operation, a parameter value of zero sets the drive to normal (non-data safe) mode and a value of 1 sets the drive to data safe mode.

- **disable\_sim\_logging**

This parameter turns the automatic logging of tape SIM/MIM data On and Off. By default, the device driver reads **Log Sense Page X'31'** automatically when device sense data indicates that data is available. The data is saved in the AIX error log. Reading **Log Sense Page X'31'** clears the current SIM/MIM data.

Setting this bit disables the device driver from reading the Log Sense Page so an application can read and manage its own SIM/MIM data. The SIM/MIM data is saved in the AIX error log if an application reads the data with the **SIOC\_LOG\_SENSE\_PAGE** or **STIOC\_LOG\_SENSE** IOCTLs.

- **read\_sili\_bit**

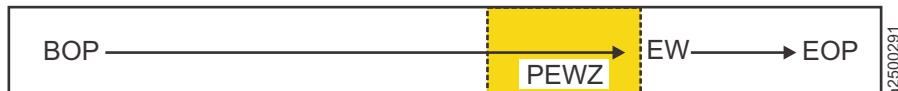
This parameter turns the **Suppress Incorrect Length Indication** (SILI) bit On and Off for variable length read commands. The device driver sets this bit when the device is configured, if it detects that the adapter can support this setting. When this bit is Off, variable length read commands results in a SCSI check condition if less data is read than the read system call requested. This action can have a significant impact on read performance.

The input or output data structure is

```
struct stchgp_s
{
    int blksize;                      /* new block size */          */
    boolean trace;                    /* TRUE=trace on */          */
    uint hkwrd;                      /* trace hook word */        */
    int sync_count;                  /* obsolete - not used */   */
    boolean autoload;                /* on/off autoload feature */ */
    boolean buffered_mode;           /* on/off buffered mode */  */
    boolean compression;             /* on/off compression */    */
    boolean trailer_labels;          /* on/off allow writing after EOM */
    boolean rewind_immediate;       /* on/off immediate rewinds */
    boolean bus_domination;         /* obsolete - not used */   */
    boolean logging;                 /* volume logging */        */
    boolean write_protect;           /* write_protected media */ */
    uint min_blksize;                /* minimum block size */    */
    uint max_blksize;                /* maximum block size */   */
    uint max_scsi_xfer;              /* maximum scsi tranfer len */
    char valid[16];                 /* volume id */            */
    uchar acf_mode;                 /* automatic cartridge facility mode */
    #define ACF_NONE      0
    #define ACF_MANUAL    1
    #define ACF_SYSTEM    2
    #define ACF_AUTOMATIC 3
    #define ACF_ACCUMULATE 4
    #define ACF_RANDOM    5
    uchar record_space_mode;        /* fsr/bsr space mode */   */
    #define SCSI_SPACE_MODE 1
    #define AIX_SPACE_MODE 2
    uchar logical_write_protect;     /* logical write protect */ */
    #define NO_PROTECT    0
    #define ASSOCIATED_PROTECT 1
    #define PERSISTENT_PROTECT 2
    #define WORM_PROTECT   3
    uchar capacity_scaling;         /* capacity scaling */    */
    #define SCALE_100     0
    #define SCALE_75      1
    #define SCALE_50      2
    #define SCALE_25      3
    #define SCALE_VALUE   4 /* use capacity_scaling_value below */
    uchar retain_reservation;       /* retain reservation */   */
    uchar alt_pathing;              /* alternate pathing active */
    boolean emulate_autoloader;     /* emulate autoloader in random mode */
    uchar medium_type;              /* tape medium type */    */
    uchar density_code;             /* tape density code */   */
    boolean disable_sim_logging;    /* disable sim/mim error logging */
    boolean read_sili_bit;           /* SILI bit setting for read commands */
    uchar capacity_scaling_value;   /* capacity scaling provided value */
    uchar reserve_type;             /* reservation type */    */
    #define RESERVE_RESERVE 0 /* SCSI Reserve 6 type */
    #define PERSISTENT_RESERVE 1 /* persistent reservation type */
    uchar reserve_key[8];           /* persistent reservation key */
    uchar data_safe_mode;           /* data safe mode */    */
    ushort pew_size;                /* programmable early warning size */
    uchar reserved[9];
};
```

- **pew\_size**

With the tape parameter, the application is allowed to request the tape drive to create a zone that is called the programmable early warning zone (PEWZ) in the front of Early Warning (EW).



a2500291

*Figure 4. Programmable Early Warning Zone (PEWZ)*

When a WRITE or WRITE FILE MARK (WFM) command writes data or filemark upon first reaching the PEWZ, Atape driver sets ENOSPC for Write and WFM to indicate that the current position reaches the PEWZ. After PEWZ is reached and before Early Warning is reached, all further writes and WFs are allowed. The **TRAILER** parameter and the current design for LEOM (Logical End of Medium/Partition, or Early Warning Zone) and PEOM (Physical End of Medium/Partition) have no effect on the driver behavior in PEWZ.

For the application developers:

1. Two methods are used to determine PEWZ when the *errno* is set to ENOSPC for **Write** or **Write FileMark** command, since ENOSPC is returned for either EW or PEW.
  - Method 1: Issue a **Request Sense** IOCTL, check the sense key and ASC-ASCQ, and if it is 0x0/0x0007 (PROGRAMMABLE EARLY WARNING DETECTED), the tape is in PEW. If the sense key ASC-ASCQ is 0x0/0x0000 or 0x0/0x0002, the tape is in EW.
  - Method 2: Call **Read Position** IOCTL in long or extended form and check bpew and eop bits. If bpew = 1 and eop = 0, the tape is in PEW. If bpew = 1 and eop = 1, the tape is in EW.

Atape driver requests the tape drive to save the mode page indefinitely. The PEW size is modified in the drive until a new setup is requested from the driver or application. The application must be programmed to issue the **Set** IOCTL to zero when PEW support is no longer needed, as Atape drivers do not complete this function. PEW is a setting of the drive and not tape. Therefore, it is the same on each partition, should partitions exist.

2. Encountering the PEWZ does not cause the device server to run a synchronize operation or terminate the command. It means that the data or filemark is written in the cartridge when a check condition with PROGRAMMABLE EARLY WARNING DETECTED is returned. But, the Atape driver still returns the counter to less than zero (-1) for a **write** command or a failure for **Write FileMark** IOCTL call with ENOSPC error. In this way, it forces the application to use one of the methods to check PEW or EW. When the application determines ENOSPC comes from PEW, it reads the requested write data or filemark that are written into the cartridge and reach or pass the PEW point. The application can issue a **Read position** IOCTL to validate the tape position.

An example of the **STIOCQRYP** and **STIOCSETP** commands is

```
#include <sys/Atape.h>
struct stchgp_s stchgp;

/* get current parameters */
if (ioctl(tapefd,STIOCQRYP,&stchgp)<0)
{
    {
        printf("IOCTL failure. errno=%d",errno);
        exit(errno);
    }

/* set new parameters */
stchgp.rewind_immediate=1;
stchgp.trailer_labels=1;
if (ioctl(tapefd,STIOCSETP,&stchgp)<0)
{
    {
        printf("IOCTL failure. errno=%d",errno);
        exit(errno);
    }
}
```

## **STIOCSYNC**

This input/output control (IOCTL) command flushes the tape buffers to the tape immediately.

There are no arguments for this IOCTL command.

An example of the **STIOCSYNC** command is

```
if (ioctl(tapefd,STIOCSYNC,NULL)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## **STIOCDM**

This IOCTL command displays and manipulates one or two messages on the message display. The message that is sent with this call does not always remain on the display. It depends on the current state of the tape device.

The input data structure is

```
#define MAXMSGLEN      8
struct stdm_s
{
    char dm_func;           /* function code */
    /* function selection */
    #define DMSTATUSMSG 0x00 /* general status message */
    #define DMDVMSG     0x20 /* demount/verify message */
    #define DMMIMMED   0x40 /* mount with immediate action indicator*/
    #define DMDEMIMMED 0xE0 /* demount/mount with immediate action */
    /* message control */
    #define DMMSG0      0x00 /* display message 0 */
    #define DMMSG1      0x04 /* display message 1 */
    #define DMFLASHMSG0 0x08 /* flash message 0 */
    #define DMFLASHMSG1 0x0C /* flash message 1 */
    #define DMAALTERNATE 0x10 /* alternate message 0 and message 1 */
    char dm_msg0[MAXMSGLEN];
    char dm_msg1[MAXMSGLEN];
};
```

An example of the **STIOCDM** command is

```
#include <sys/Atape.h>
struct stdm_s stdm;
stdm.dm_func=DMSTATUSMSG|DMMSG0;
bcopy("SSD",stdm.dm_msg0,8);
if (ioctl(tapefd,STIOCDM,&stdm)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## **STIOCQRYPOS or STIOCSETPOS**

The **STIOCQRYPOS** IOCTL command queries the position on the tape. The **STIOCSETPOS** IOCTL command sets the position on the tape. Only the **block\_type** and **curpos** fields are used during a **set** operation. The tape position is defined as where the next **read** or **write** operation occurs. The **query** function can be used independently or with the **set** function. Also, the **set** function can be used independently or with the **query** function.

The **block\_type** field is set to **QP\_LOGICAL** when a SCSI logical **blockid** format is wanted. During a query operation, the **curpos** field is set to a simple **unsigned int**.

On IBM 3490 tape drives only, the **block\_type** field can be set to **QP\_PHYSICAL**. Setting this **block\_type** on any other device is ignored and defaults to **QP\_LOGICAL**. After a **set** operation, the position is at the logical block that is indicated by the **curpos** field. If the **block\_type** field is set to **QP\_PHYSICAL**, the **curpos** field that is returned is a vendor-specific **blockid** format from the tape device. When **QP\_PHYSICAL** is used for a **query** operation, the **curpos** field is used only in a subsequent set operation with **QP\_PHYSICAL**. This function completes a high speed **locate** operation. Whenever possible, use **QP\_PHYSICAL** because it is faster. This advantage is obtained only when the **set** operation uses the **curpos** field from the **QP\_PHYSICAL** query.

After a **query** operation, the **lbot** field indicates the last block of the data that was transferred physically to the tape. If the application writes 12 (0 - 11) blocks and **lbot** equals 8, then three blocks are in the tape buffer. This field is valid only if the last command was a **write** command. This field does not reflect the number of application **write** operations. A **write** operation can translate into multiple blocks. It reflects tape blocks as indicated by the block size. If an attempt is made to obtain this information and the last command is not a write command, the value of **LBOT\_UNKNOWN** is returned.

The driver sets the **bot** field to TRUE if the tape position is at the beginning of the tape. Otherwise, it is set to FALSE. The driver sets the **eot** field to TRUE if the tape is positioned between the early warning and the physical end of the tape. Otherwise, it is set to FALSE.

The number of blocks and number of bytes currently in the tape device buffers is returned in the **num\_blocks** and **num\_bytes** fields. The **bcu** and **bycu** settings indicate whether these fields contain valid data. The block ID of the next block of data that transferred to or from the physical tape is returned in the **tapepos** field.

The partition number field that is returned is the current partition of the loaded tape.

The input or output data structure is

```
typedef unsigned int blockid_t;
struct stpos_s
{
    char block_type; /* format of block ID information */
    #define QP_LOGICAL 0 /* SCSI logical block ID format */
    #define QP_PHYSICAL 1 /* 3490 only, vendor-specific block ID format */
    boolean eot; /* ignored for all other devices */
    /* position is after early warning,
     * before physical end of tape */
    blockid_t curpos; /* for query, current position,
                        * for set, position to go to */
    blockid_t lbot; /* last block written to tape */
    #define LBOT_NONE 0xFFFFFFFF /* no blocks were written to tape */
    #define LBOT_UNKNOWN 0xFFFFFFFF /* unable to determine information */
    uint num_blocks; /* number of blocks in buffer */
    uint num_bytes; /* number of bytes in buffer */
    boolean bot; /* position is at beginning of tape */
    uchar partition_number; /* current partition number on tape */
    boolean bcu; /* number of blocks in buffer is unknown */
    boolean bycu; /* number of bytes in buffer is unknown */
    blockid_t tapepos; /* next block transferred */
    uchar reserved2[48];
};
```

An example of the **STIOCQRYPOS** and **STIOCSETPOS** commands is

```
#include <sys/Atape.h>
struct stpos_s stpos;
stpos.block_type=QP_PHYSICAL;
if (ioctl(tapefd,STIOCQRYPOS,&stpos)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
oldposition=stpos.curpos;
.
.
stpos.curpos=oldposition;
stpos.block_type=QP_PHYSICAL;
if (ioctl(tapefd,STIOCSETPOS,&stpos)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## **STIOCQRYSENSE**

This IOCTL command returns the last sense data that is collected from the tape device, or it issues a new **Request Sense** command and returns the collected data. If **LASTERROR** is requested, the sense data is valid only if the last tape operation has an error that issued a sense command to the device. If the sense

data is valid, the IOCTL command completes successfully and the **len** field is set to a value greater than zero.

The **residual\_count** field contains the residual count from the last operation.

The input or output data structure is

```
#define MAXSENSE      255
struct stsense_s
{
    /* input */
    char sense_type;          /* fresh (new sense) or sense from last error */
    #define FRESH     1          /* initiate a new sense command */
    #define LASTERROR 2          /* return sense gathered from
                                    the last SCSI sense command */

    /* output */
    uchar sense[MAXSENSE];    /* actual sense data */
    int len;                  /* length of valid sense data returned */
    int residual_count;       /* residual count from last operation */
    uchar reserved[60];
};
```

An example of the **STIOCQRYSENSE** command is

```
#include <sys/Atape.h>
struct stsense_s stsense;
stsense.sense_type=LASTERROR;
#define MEDIUM_ERROR 0x03
if (ioctl(tapefd,STIOCQRYSENSE,&stsense)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
if ((stsense.sense[2]&0x0F)==MEDIUM_ERROR)
{
    printf("We're in trouble now!");
    exit(SENSE_KEY(&stsense.sense));
}
```

## **STIOCQRYINQUIRY**

This IOCTL command returns the inquiry data from the device. The data is divided into standard and vendor-specific portions.

The output data structure is

```
/* inquiry data info */
struct inq_data_s
{
    BYTE b0;
    /* macros for accessing fields of byte 1 */
    #define PERIPHERAL_QUALIFIER(x) ((x->b0 & 0xE0)>>5)
    #define PERIPHERAL_CONNECTED          0x00
    #define PERIPHERAL_NOT_CONNECTED      0x01
    #define LUN_NOT_SUPPORTED            0x03

    #define PERIPHERAL_DEVICE__TYPE(x) (x->b0 & 0x1F)
    #define DIRECT_ACCESS                0x00
    #define SEQUENTIAL_DEVICE            0x01
    #define PRINTER_DEVICE               0x02
    #define PROCESSOR_DEVICE              0x03
    #define CD_ROM_DEVICE                 0x05
    #define OPTICAL_MEMORY_DEVICE        0x07
    #define MEDIUM_CHANGER_DEVICE        0x08
    #define UNKNOWN                      0x1F

    BYTE b1;
    /* macros for accessing fields of byte 2 */
    #define RMB(x) ((x->b1 & 0x80)>>7)           /* removable media bit */
    #define FIXED          0
    #define REMOVABLE       1
    #define device_type_qualifier(x) (x->b1 & 0x7F) /* vendor specific */

    BYTE b2;
    /* macros for accessing fields of byte 3 */
    #define ISO_Version(x)  ((x->b2 & 0xC0)>>6)
    #define ECMA_Version(x) ((x->b2 & 0x38)>>3)
```

```

#define ANSI_Version(x) ((x->b2 & 0x07)
#define NONSTANDARD 0
#define SCSI1 1
#define SCSI2 2

BYTE b3;
/* macros for accessing fields of byte 4 */
#define AENC(x) ((x->b3 & 0x80)>>7) /* asynchronous event notification */
#ifndef TRUE
#define TRUE 1
#endif
#ifndef FALSE
#define FALSE 0
#endif
#define TrmIOP(x) ((x->b3 & 0x40)>>6) /* support terminate I/O process message? */
#define Response_Data_Format(x) (x->b3 & 0x0F)
#define SCSI1INQ 0 /* SCSI-1 standard inquiry data format */
#define CCSINQ 1 /* CCS standard inquiry data format */
#define SCSI2INQ 2 /* SCSI-2 standard inquiry data format */

BYTE additional_length; /* number of bytes following this field minus 4 */
BYTE res56[2];

BYTE b7;
/* macros for accessing fields of byte 7 */
#define RelAddr(x) ((x->b7 & 0x80)>>7) /* the following fields are true or false */
#define WBus32(x) ((x->b7 & 0x40)>>6)
#define WBus16(x) ((x->b7 & 0x20)>>5)
#define Sync(x) ((x->b7 & 0x10)>>4)
#define Linked(x) ((x->b7 & 0x08)>>3)
#define CmdQue(x) ((x->b7 & 0x02)>>1)
#define SftRe(x) ((x->b7 & 0x01)

char vendor_identification[8];
char product_identification[16];
char product_revision_level[4];
};

struct st_inquiry
{
    struct inq_data_s standard;
    BYTE vendor_specific[255-sizeof(struct inq_data_s)];
};

```

An example of the **STIOCQRYINQUIRY** command is

```

struct st_inquiry inqd;
if (ioctl(tapefd,STIOCQRYINQUIRY,&inqd)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
if (ANSI_Version(((struct inq_data_s *)&(inqd.standard)))==SCSI2)
    printf("Hey! We have a SCSI-2 device\n");

```

## **STIOC\_LOG\_SENSE**

This IOCTL command returns the log sense data from the device. If volume logging is set to On, the log sense data is saved in the tape log file.

The output data structure is

```

struct log_sense
{
    struct log_record_header header;
    char data[MAXLOGSENSE];
}

```

An example of the **STIOC\_LOG\_SENSE** command is

```

struct log_sense logdata;

if (ioctl(tapefd,STIOC_LOG_SENSE,&logdata)<0)
{
    printf("IOCTL failure. errno=%d",errno);
}

```

```
    exit(errno);
}
```

## **STIOC\_RECOVER\_BUFFER**

This IOCTL command recovers the buffer data from the tape device. It is typically used after an error occurs during a **write** operation that prevents the data in the tape device buffers from being written to tape. The **STIOCQRYPOS** command can be used before this IOCTL command to determine the number of blocks and the bytes of data that is in the device buffers.

Each **STIOC\_RECOVER\_BUFFER** IOCTL call returns one block of data from the device. This **ioctl** command can be issued multiple times to completely recover all the buffered data from the device.

After the IOCTL command is completed, the **ret\_len** field contains the number of bytes returned in the application buffer for the block. If no blocks are in the tape device buffer, then the **ret\_len** value is set to zero.

The output data structure is

```
struct buffer_data
{
    char *buffer;
    int bufsize;
    int ret_len;
};
```

An example of the **STIOC\_RECOVER\_BUFFER** command is

```
struct buffer_data bufdata;

bufdata.bufsize = 256 * 1024;
bufdata.buffer = malloc(256 * 1024);

if (ioctl(tapefd,STIOC_RECOVER_BUFFER,&bufdata)<0)
{
    printf("IOCTL failure. errno=%d",errno);
}
else
{
    printf("Returned bytes=%d",bufdata.ret_len);
```

## **STIOC\_LOCATE**

This IOCTL command causes the tape to be positioned at the specified block ID. The block ID used for the command must be obtained with the **STIOC\_READ\_POSITION** command.

An example of the **STIOC\_LOCATE** command is

```
#include <sys/Atape.h>

unsigned int current_blockid;

/* read current tape position */
if (ioctl(tapefd,STIOC_READ_POSITION,&current_blockid)<0)
{
    printf("IOCTL failure. errno=%d\n",errno);
    exit(1);
}

/* restore current tape position */
if (ioctl(tapefd,STIOC_LOCATE,current_blockid)<0)
{
    printf("IOCTL failure. errno=%d\n",errno);
    exit(1);
}
```

## **STIOC\_READ\_POSITION**

This IOCTL command returns the block ID of the current position of the tape. The block ID returned from this command can be used with the **STIOC\_LOCATE** command to set the position of the tape.

An example of the **STIOC\_READ\_POSITION** command is

```
#include <sys/Atape.h>

unsigned int current_blockid;

/* read current tape position */
if (ioctl(tapefd,STIOC_READ_POSITION,&current_blockid)<0)
{
    printf("IOCTL failure. errno=%d\n",errno);
    exit(1);
}

/* restore current tape position */
if (ioctl(tapefd,STIOC_LOCATE,current_blockid)<0)
{
    printf("IOCTL failure. errno=%d\n",errno);
    exit(1);
}
```

### ***STIOC\_SET\_VOLID***

This IOCTL command sets the volume name for the currently mounted tape. The volume name is used by the device driver for tape volume logging only and is not written or stored on the tape. The volume name is reset to unknown whenever an **unload** command is issued to unload the current tape. The volume name can be queried and set by using the **STIOCQRYP** and **STIOCSETP** IOCTLs.

The argument that is used for this command is a character pointer to a buffer that contains the name of the volume to be set.

An example of the **STIOC\_SET\_VOLID** command is

```
/* set the volume id for the current tape to VOL001 */
char *volid = "VOL001";
if (ioctl(tapefd,STIOC_SET_VOLID,volid)<0)
{
    printf("IOCTL failure. errno=%d\n",errno);
    exit(errno);
}
```

### ***STIOC\_DUMP***

This IOCTL command forces a dump on the tape device, then stores the dump to either a host-specified file or in the **/var/adm/ras** system directory. The device driver stores up to three dumps in this directory. The first dump that is created is named **Atape.rmtx.dump1**, where x is the device number, for example, **rmt0**. The second and third dumps are **dump2** and **dump3**. After a third dump file is created, the next dump starts at **dump1** again and overlays the previous **dump1** file.

The argument that is used for this command is NULL to dump to the system directory. Or, it is a character pointer to a buffer that contains the path and file name for the dump file. The dump can also be stored on a diskette by specifying **/dev/rfd0** for the name.

An example of the **STIOC\_DUMP** command is

```
/* generate drive dump and store in the system directory */
if (ioctl(tapefd,STIOC_DUMP,NULL)<0)
{
    printf("IOCTL failure. errno=%d\n",errno);
    exit(errno);
}

/* generate drive dump and store in file 3590.dump */
char *dump_name = "3590.dump";
if (ioctl(tapefd,STIOC_DUMP,dump_name)<0)
{
    printf("IOCTL failure. errno=%d\n",errno);
    exit(errno);
}
```

## ***STIOC\_FORCE\_DUMP***

This IOCTL command forces a dump on the tape device. The dump can be retrieved from the device by using the **STIOC\_READ\_DUMP** IOCTL.

There are no arguments for this IOCTL command.

An example of the **STIOC\_FORCE\_DUMP** command is

```
/* generate a drive dump */
if (ioctl(tapefd,STIOC_FORCE_DUMP,NULL)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## ***STIOC\_READ\_DUMP***

This IOCTL command reads a dump from the tape device. Then, it stores the dump to either a host specified file or in the **/var/adm/ras** system directory. The device driver stores up to three dumps in this directory. The first dump that is created is named **Atape.rmtx.dump1**, where x is the device number, for example **rmt0**. The second and third dumps are **dump2** and **dump3**. After a third dump file is created, the next dump starts at **dump1** again and overlays the previous **dump1** file.

Dumps are either generated internally by the tape drive or can be forced by using the **STIOC\_FORCE\_DUMP** IOCTL.

The argument that is used for this command is NULL to dump to the system directory. Or, it is a character pointer to a buffer that contains the path and file name for the dump file. The dump can also be stored on a diskette by specifying **/dev/rfd0** for the name.

An example of the **STIOC\_READ\_DUMP** command is

```
/* read drive dump and store in the system directory */
if (ioctl(tapefd,STIOC_READ_DUMP,NULL)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}

/* read drive dump and store in file 3590.dump */
char *dump_name = "3590.dump";
if (ioctl(tapefd,STIOC_READ_DUMP,dump_name)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## ***STIOC\_LOAD\_UCODE***

This IOCTL command downloads microcode to the device. The argument that is used for this command is a character pointer to a buffer that contains the path and file name of the microcode. Microcode can also be loaded from a diskette by specifying **/dev/rfd0** for the name.

An example of the **STIOC\_LOAD\_UCODE** command is

```
/* download microcode from file */
char *name = "/etc/microcode/D0I4_BB5.fmrz";
if (ioctl(tapefd,STIOC_LOAD_UCODE,name)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}

/* download microcode from diskette */
if (ioctl(tapefd,STIOC_LOAD_UCODE,"/dev/rfd0")<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## ***STIOC\_RESET\_DRIVE***

This IOCTL command issues a **SCSI Send Diagnostic** command to reset the tape drive. There are no arguments for this IOCTL command.

An example of the **STIOC\_RESET\_DRIVE** command is

```
/* reset the tape drive */
if (ioctl(tapefd,STIOC_RESET_DRIVE,NULL)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## ***STIOC\_FMR\_TAPE***

This IOCTL command creates an FMR tape. The tape is created with the current microcode loaded in the tape device.

There are no arguments for this IOCTL command.

An example of the **STIOC\_FMR\_TAPE** command is

```
/* create fmr tape */
if (ioctl(tapefd,STIOC_FMR_TAPE,NULL)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## ***MTDEVICE (Obtain device number)***

This IOCTL command obtains the device number that is used for communicating with the IBM TotalStorage Enterprise library 3494.

The structure of the IOCTL request is

```
int device;
if (ioctl(tapefd,MTDEVICE,&device)<0)
{
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## ***STIOC\_PREVENT\_MEDIUM\_REMOVAL***

This IOCTL command prevents an operator from removing medium from the device until the **STIOC\_ALLOW\_MEDIUM\_REMOVAL** command is issued or the device is reset.

There is no associated data structure.

An example of the **STIOC\_PREVENT\_MEDIUM\_REMOVAL** command is

```
#include <sys/Atape.h>

if (!ioctl (tapefd, STIOC_PREVENT_MEDIUM_REMOVAL, NULL))
    printf ("The STIOC_PREVENT_MEDIUM_REMOVAL ioctl succeeded\n");
else
{
    perror ("The STIOC_PREVENT_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}
```

## ***STIOC\_ALLOW\_MEDIUM\_REMOVAL***

This IOCTL command allows an operator to remove medium from the device. This command is used normally after an **STIOC\_PREVENT\_MEDIUM\_REMOVAL** command to restore the device to the default state.

There is no associated data structure.

An example of the **STIOC\_ALLOW\_MEDIUM\_REMOVAL** command is

```
#include <sys/Atape.h>

if (!ioctl (tapefd, STIOC_ALLOW_MEDIUM_REMOVAL, NULL))
    printf ("The STIOC_ALLOW_MEDIUM_REMOVAL ioctl succeeded\n");
else
{
    perror ("The STIOC_ALLOW_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}
```

## **STIOC\_REPORT\_DENSITY\_SUPPORT**

This IOCTL command issues the **SCSI Report Density Support** command to the tape device and returns either all supported densities or supported densities for the currently mounted media. The media field specifies which type of report is requested. The **number\_reports** field is returned by the device driver and indicates how many density reports in the **reports array** field were returned.

The data structures that are used with this IOCTL are

```
typedef struct density_report
{
    uchar primary_density_code;      /* primary density code */
    uchar secondary_density_code;    /* secondary density code */
    uint wrtok :1,                  /* write ok, device can write this format */
    dup :1,                         /* zero if density only reported once */
    deflt :1,                        /* current density is default format */
    res_1 :5;                       /* reserved */
    uchar reserved[2];              /* reserved */
    uchar bits_per_mm[3];           /* bits per mm */
    uint bits_per_mm:24;            /* bits per mm */
    ushort media_width;             /* media width in millimeters */
    ushort tracks;                 /* tracks */
    uint capacity;                  /* capacity in megabytes */
    char assigning_org[8];          /* assigning organization in ASCII */
    char density_name[8];           /* density name in ASCII */
    char description[20];            /* description in ASCII */
};

struct report_density_support
{
    uchar media;                   /* report all or current media as defined above */
    ushort number_reports;         /* number of density reports returned in array */
    struct density_report reports[MAX_DENSITY_REPORTS];
};
```

Examples of the **STIOC\_REPORT\_DENSITY\_SUPPORT** command are

```
#include <sys/Atape.h>

int stioc_report_density_support(void)
{
    int i;
    struct report_density_support density;

    printf("Issuing Report Density Support for ALL supported media...\n");

    density.media = ALL_MEDIA_DENSITY;

    if (ioctl(fd, STIOC_REPORT_DENSITY_SUPPORT, &density) != 0)
        return errno;
    printf("Total number of densities reported: %d\n", density.number_reports);
    for (i = 0; i < density.number_reports; i++)
    {
        printf("\n");
        printf("  Density Name.....%0.8s\n",
               density.reports[i].density_name);
        printf("  Assigning Organization.%0.8s\n",
               density.reports[i].assigning_org);
        printf("  Description.....%0.20s\n",
               density.reports[i].description);
        printf("  Primary Density Code....%02X\n",
               density.reports[i].primary_density_code);
```

```

    printf(" Secondary Density Code..%02X\n",
           density.reports[i].secondary_density_code);

    if (density.reports[i].wrtok)
        printf(" Write OK.....Yes\n");
    else
        printf(" Write OK.....No\n");

    if (density.reports[i].dup)
        printf(" Duplicate.....Yes\n");
    else
        printf(" Duplicate.....No\n");

    if (density.reports[i].deflt)
        printf(" Default.....Yes\n");
    else
        printf(" Default..... No\n");

    printf(" Bits per MM.....%d\n",
           density.reports[i].bits_per_mm);
    printf(" Media Width (millimeters)%d\n",
           density.reports[i].media_width);
    printf(" Tracks.....%d\n",
           density.reports[i].tracks);
    printf(" Capacity (megabytes)....%d\n",
           density.reports[i].capacity);
    if (opcode)
    {
        printf ("\nHit <enter> to continue...");
        getchar();
    }

printf("\nIssuing Report Density Support for CURRENT media...\n");
density.media = CURRENT_MEDIA_DENSITY;

if (ioctl(fd, STIOC_REPORT_DENSITY_SUPPORT, &density) != 0)
    return errno;

for (i = 0; i < density.number_reports; i++)
{
    printf("\n");
    printf(" Density Name.....%0.8s\n",
           density.reports[i].density_name);
    printf(" Assigning Organization..%0.8s\n",
           density.reports[i].assigning_org);
    printf(" Description.....%0.20s\n",
           density.reports[i].description);
    printf(" Primary Density Code....%02X\n",
           density.reports[i].primary_density_code);
    printf(" Secondary Density Code..%02X\n",
           density.reports[i].secondary_density_code);

    if (density.reports[i].wrtok)
        printf(" Write OK.....Yes\n");
    else
        printf(" Write OK.....No\n");

    if (density.reports[i].dup)
        printf(" Duplicate.....Yes\n");
    else
        printf(" Duplicate.....No\n");

    if (density.reports[i].deflt)
        printf(" Default.....Yes\n");
    else
        printf(" Default..... No\n");

    printf(" Bits per MM.....%d\n",density.reports[i].bits_per_mm);
    printf(" Media Width (millimeters)%d\n",density.reports[i].media_width);
    printf(" Tracks.....%d\n",density.reports[i].tracks);
    printf(" Capacity (megabytes)....%d\n",density.reports[i].capacity);
}

return errno;
}

```

## ***STIOC\_GET\_DENSITY and STIOC\_SET DENSITY***

The **STIOC\_GET\_DENSITY** IOCTL is used to query the current write density format settings on the tape drive. The current density code from the drive **Mode Sense** header, the **Read/Write Control Mode** page default density, and pending density are returned.

The **STIOC\_SET\_DENSITY** IOCTL is used to set a new write density format on the tape drive by using the default and pending density fields. The **density code** field is not used and ignored on this IOCTL. The application can specify a new write density for the current loaded tape only or as a default for all tapes. Refer to the examples.

The application must get the current density settings first before the current settings are modified. If the application specifies a new density for the current loaded tape only, then the application must issue another **Set Density** IOCTL after the current tape is unloaded and the next tape is loaded to either the default maximum density or a new density. This action ensures the tape drive uses the correct density. If the application specifies a new default density for all tapes, the setting remains in effect until changed by another set density IOCTL or the tape drive is closed by the application.

Following is the structure for the **STIOC\_GET\_DENSITY** and **STIOC\_SET\_DENSITY** IOCTLs.

```
struct density_data_t
{
    char density_code;          /* mode sense header density code */
    char default_density;       /* default write density */
    char pending_density;       /* pending write density */
    char reserved[9];           */
};
```

### **Note:**

1. The IOCTLs are only supported on tape drives that can write multiple density formats. Refer to the Hardware Reference for the specific tape drive to determine whether multiple write densities are supported. If the tape drive does not support the IOCTLs, *errno* EINVAL is returned.
2. The device driver always sets the default maximum write density for the tape drive on every open system call. Any previous **STIOC\_SET\_DENSITY** IOCTL values from the last open are not used.
3. If the tape drive detects an invalid density code or cannot complete the operation on the **STIOC\_SET\_DENSITY** IOCTL, the *errno* is returned. Then, the current drive density settings before the IOCTL are restored.
4. The `struct density_data_t` defined in the header file is used for both IOCTLs. The **density\_code** field is not used and ignored on the **STIOC\_SET\_DENSITY** IOCTL.

## **Examples**

```
struct density_data_t data;

/* open the tape drive      */
/* get current density settings */
rc = ioctl(fd, STIOC_GET_DENSITY, %data);

/* set 3592 J1A density format for current loaded tape only */
data.default_density = 0x7F;
data.pending_density = 0x51;
rc = ioctl(fd, STIOC_SET_DENSITY, %data);

/* unload tape   */
/* load next tape */
/* set 3592 E05 density format for current loaded tape only */
data.default_density = 0x7F;
data.pending_density = 0x52;
rc = ioctl(fd, STIOC_SET_DENSITY, %data);

/* unload tape   */
/* load next tape */
/* set default maximum density for current loaded tape */
data.default_density = 0;
```

```

data.pending_density = 0;
rc = ioctl(fd, STIOC_SET_DENSITY, %data);

/* close the tape drive          */
/* open the tape drive          */
/* set 3592 J1A density format for current loaded tape and all subsequent tapes */
data.default_density = 0x51;
data.pending_density = 0x51;

rc = ioctl(fd, STIOC_SET_DENSITY, %data);

```

### ***STIOC\_CANCEL\_ERASE***

The **STIOC\_CANCEL\_ERASE** IOCTL is used to cancel an erase operation currently in progress. This action happens when an application issued the **STIOCTOP** IOCTL with the **st\_op** field that specifies **STERASE\_IMM**. The application that issued the erase and is waiting for it to complete then returns immediately with *errno* ECANCELLED. This IOCTL always returns 0 whether an **erase immediate** operation is in progress or not.

This IOCTL can be issued only when the `openx()` extended parameter **SC\_TMCP** is used to open the device. It happens when the application that issued the erase still has the device currently open. There is no argument for this IOCTL and the **arg** parameter is ignored.

### ***GET\_ENCRYPTION\_STATE***

This IOCTL command queries the drive's encryption method and state. The data structure that is used for this IOCTL is for all of the supported operating systems.

```

struct encryption_status {
    uchar encryption_capable; /* (1)Set this field as a boolean based on the
                               capability of the drive */
    uchar encryption_method; /* (2)Set this field to one of the
                               #defines METHOD_* below */
#define METHOD_NONE 0           /* Only used in GET_ENCRYPTION_STATE */
#define METHOD_LIBRARY 1        /* Only used in GET_ENCRYPTION_STATE */
#define METHOD_SYSTEM 2         /* Only used in GET_ENCRYPTION_STATE */
#define METHOD_APPLICATION 3   /* Only used in GET_ENCRYPTION_STATE */
#define METHOD_CUSTOM 4         /* Only used in GET_ENCRYPTION_STATE */
#define METHOD_UNKNOWN 5        /* Only used in GET_ENCRYPTION_STATE */

    uchar encryption_state; /* (3) Set this field to one of the
                           #defines STATE_* below */
#define STATE_OFF 0             /* Used in GET/SET_ENCRYPTION_STATE */
#define STATE_ON 1               /* Used in GET/SET_ENCRYPTION_STATE */
#define STATE_NA 2               /* Only used in GET_ENCRYPTION_STATE */

    uchar[13] reserved;
};

```

An example of the **GET\_ENCRYPTION\_STATE** command is

```

int qry_encryption_state (void)
{
    int rc = 0;
    struct encryption_status encryption_status_t;

    printf("issuing query encryption status...\n");
    memset(&encryption_status_t, 0, sizeof(struct encryption_status));
    rc = ioctl(fd, GET_ENCRYPTION_STATE, &encryption_status_t);
    if(rc == 0)
    {
        if(encryption_status_t.encryption_capable)
            printf("encryption capable.....Yes\n");
        else
            printf("encryption capable.....No\n");
        switch(encryption_status_t.encryption_method)
        {
        case METHOD_NONE:
            printf("encryption method.....METHOD_NONE\n");
            break;
        case METHOD_LIBRARY:
            printf("encryption method.....METHOD_LIBRARY\n");
            break;
        case METHOD_SYSTEM:

```

```

        printf("encryption method.....METHOD_SYSTEM\n");
        break;
    case METHOD_APPLICATION:
        printf("encryption method.....METHOD_APPLICATION\n");
        break;
    case METHOD_CUSTOM:
        printf("encyrption method.....METHOD_CUSTOM\n");
        break;
    case METHOD_UNKNOWN:
        printf("encryption method.....METHOD_UNKNOWN\n");
        break;

    default:
        printf("encryption method.....Error\n");
    }

    switch(encryption_status_t.encryption_state)
    {
    case STATE_OFF:
        printf("encryption state.....OFF\n");
        break;
    case STATE_ON:
        printf("encryption state.....ON\n");
        break;
    case STATE_NA:
        printf("encryption state.....NA\n");
        break;

    default:
        printf("encryption state.....Error\n");
    }
}

return rc;
}

```

## **SET\_ENCRYPTION\_STATE**

This IOCTL command allows set encryption state only for application-managed encryption. On unload, some of the drive settings can be reset to default. To set the encryption state, the application must issue this IOCTL after a tape is loaded and at BOP.

The data structure used for this IOCTL is the same as the one for **GET\_ENCRYPTION\_STATE**. An example of the **SET\_ENCRYPTION\_STATE** command is

```

int set_encryption_state(int option)
{
    int rc = 0;
    struct encryption_status encryption_status_t;

    printf("issuing query encryption status...\n");
    memset(&encryption_status_t, 0, sizeof(struct encryption_status));
    rc = ioctl(fd, GET_ENCRYPTION_STATE, );&encryption_status_t
    if(rc < 0) return rc;

    if(option == 0)
        encryption_status_t.encryption_state = STATE_OFF;
    else if(option == 1)
        encryption_status_t.encryption_state = STATE_ON;
    else
    {
        printf("Invalid parameter.\n");
        return -EINVAL;
    }

    printf("Issuing set encryption state.....\n");
    rc = ioctl(fd, SET_ENCRYPTION_STATE, &encryption_status_t);

    return rc;
}

```

## **SET\_DATA\_KEY**

This IOCTL command allows set the data key only for application-managed encryption. The data structure that is used for this IOCTL is for all of the supported operating systems.

```

struct data_key
{
    uchar[12] data_key_index;
    uchar data_key_index_length;
    uchar[15] reserved1;
    uchar[32] data_key;
    uchar[48] reserved2;
};

```

An example of the **SET\_DATA\_KEY** command is

```

int set_datakey(void)
{
    int rc = 0;
    struct data_key encryption_data_key_t;

    printf("Issuing set encryption data key.....\n");
    memset(&encryption_data_key_t, 0, sizeof(struct data_key));
    /* fill in your data key here, then issue the following ioctl*/
    rc = ioctl(fd, SET_DATA_KEY, &encryption_data_key_t);
    return rc;
}

```

## ***READ\_TAPE\_POSITION***

The **READ\_TAPE\_POSITION** IOCTL is used to return **Read Position** command data in either the short, long, or extended form. The type of data to return is specified by setting the **data\_format** field to either **RP\_SHORT\_FORM**, **RP\_LONG\_FORM**, or **RP\_EXTENDED\_FORM**.

The data structures that are used with this IOCTL are

```

#define RP_SHORT_FORM      0x00
#define RP_LONG_FORM       0x06
#define RP_EXTENDED_FORM   0x08

struct short_data_format {
    uint bop:1,                      /* beginning of partition          */
    eop:1,                          /* end of partition                */
    locu:1,                         /* 1 means num_buffer_logical_obj field is
                                         unknown                         */
    bycu:1,                         /* 1 means the num_buffer_bytes field is
                                         unknown                         */
    rsvd :1,                         /* 1 means the first and last logical
                                         obj position fields are unknown */
    lolu:1,                         /* 1 means the position fields have
                                         overflowed and can not be reported */
    perr: 1,                         /* beyond programmable early warning */
    bpew :1,                         /* beyond programmable early warning */
    uchar active_partition;           /* current active partition         */
    char reserved[2];                /*                                     */
    uint first_logical_obj_position; /* current logical object position */
    uint last_logical_obj_position;  /* next logical object to be transferred
                                         to tape                         */
    uint num_buffer_logical_obj;     /* number of logical objects in buffer */
    uint num_buffer_bytes;           /* number of bytes in buffer        */
    char reserved1;                  /*                                     */
};

struct long_data_format {
    uint bop:1,                      /* beginning of partition          */
    eop:1,                          /* end of partition                */
    rsvd1:2,                         /* 1 means the logical file id field
                                         in unknown                      */
    mpu:1,                           /* 1 means either the partition number
                                         or logical obj number field
                                         are unknown                     */
    lonu:1,                          /* 1 means either the partition number
                                         or logical obj number field
                                         are unknown                     */
    rsvd2:1,                         /* beyond programmable early warning */
    bpew :1,                         /* beyond programmable early warning */
    char reserved[6];                /*                                     */
    uchar active_partition;           /* current active partition         */
    ullong logical_obj_number;       /* current logical object
                                         position                         */
    ullong logical_file_id;          /* number of filemarks from bop
                                         and current logical position */
    ullong obsolete;                 /*                                     */
};

```

```

};

struct extended_data_format {
    uint bop:1,           /* beginning of partition      */
    eop:1,           /* end of partition          */
    locu:1,           /* 1 means num_buffer_logical_obj field
                      is unknown               */
    bycu:1,           /* 1 means the num_buffer_bytes field
                      is unknown               */
    rsvd :1,           /* 1 means the first and last logical obj
                      position fields are unknown */
    lolu:1,           /* 1 means the position fields have
                      overflowed and can not be reported */
    perr: 1,           /* 1 means the position fields have
                      overflowed and can not be reported */
    bpew :1;           /* beyond programmable early warning */
uchar active_partition; /* current active partition */
ushort additional_length;
uint num_buffer_logical_obj; /* number of logical objects in buffer */
ullong first_logical_obj_position; /* current logical object position */
ullong last_logical_obj_position; /* next logical object to be transferred
                                    to tape */
ullong num_buffer_bytes; /* number of bytes in buffer */
char reserved;
};

struct read_tape_position{
    uchar data_format; /* Specifies the return data format either short,
long or extended as defined above */
union
{
    struct short_data_format rp_short;
    struct long_data_format rp_long;
    struct extended_data_format rp_extended;
    char reserved[64];
} rp_data;
};

```

#### Example of the **READ\_TAPE\_POSITION** IOCTL

```

#include <sys/Atape.h>

struct read_tape_position rpos;

printf("Reading tape position long form....\n");
rpos.data_format = RP_LONG_FORM;
if (ioctl (fd, READ_TAPE_POSITION, &rpos) <0)
    return errno;

    if (rpos.rp_data.rp_long.bop)
        printf("    Beginning of Partition ..... Yes\n");
    else
        printf("    Beginning of Partition ..... No\n");
    if (rpos.rp_data.rp_long.eop)
        printf("    End of Partition ..... Yes\n");
    else
        printf("    End of Partition ..... No\n");
    if (rpos.rp_data.rp_long.bpew)
        printf("    Beyond Early Warning ... .... Yes\n");
    else
        printf("    Beyond Early Warning ..... No\n");
    if (rpos.rp_data.rp_long.lonu)
    {
        printf("    Active Partition ..... UNKNOWN \n");
        printf("    Logical Object Number ..... UNKNOWN \n");
    }
    else
    {
        printf("    Active Partition ... ..... %u \n",
            rpos.rp_data.rp_long.active_partition);
        printf("    Logical Object Number ..... %llu \n",
            rpos.rp_data.rp_long.logical_obj_number);
    }

    if (rpos.rp_data.rp_long.mpu)
        printf("    Logical File ID ..... UNKNOWN \n");
    else
        printf("    Logical File ID ..... %llu \n",
            rpos.rp_data.rp_long.logical_file_id);

```

## ***SET\_TAPE\_POSITION***

The **SET\_TAPE\_POSITION** IOCTL is used to position the tape in the current active partition to either a logical block id or logical filemark. The **logical\_id\_type** field in the IOCTL structure specifies either a logical block or logical filemark.

The data structure that is used with this IOCTL is

```
#define LOGICAL_ID_BLOCK_TYPE 0x00
#define LOGICAL_ID_FILE_TYPE 0x01

struct set_tape_position{
    uchar logical_id_type; /* Block or file as defined above */
    ullong logical_id; /* logical object or logical file to position to */
    char reserved[32];
};
```

Examples of the **SET\_TAPE\_POSITION** IOCTL

```
#include <sys/Atape.h>

struct set_tape_position setpos;

/* position to logical block id 10 */
setpos.logical_id_type = LOGICAL_ID_BLOCK_TYPE
setpos.logical_id = 10;
ioctl(fd, SET_TAPE_POSITION, &setpos);

/* position to logical filemark 4 */
setpos.logical_id_type = LOGICAL_ID_FILE_TYPE
setpos.logical_id = 4;
ioctl(fd, SET_TAPE_POSITION, &setpos);
```

## ***SET\_ACTIVE\_PARTITION***

The **SET\_ACTIVE\_PARTITION** IOCTL is used to position the tape to a specific partition. Then, it becomes the current active partition for subsequent commands and a specific logical block id in the partition. To position to the beginning of the partition, the **logical\_block\_id** field is set to 0.

The data structure that is used with this IOCTL is

```
struct set_active_partition {
    uchar partition_number; /* Partition number 0-n to change to */
    ullong logical_block_id; /* Blockid to locate to within partition */
    char reserved[32];
};
```

Examples of the **SET\_ACTIVE\_PARTITION** IOCTL

```
#include <sys/Atape.h>

struct set_active_partition partition;

/* position the tape to partition 1 and logical block id 12 */
partition.partition_number = 1;
partition.logical_block_id = 12;
ioctl(fd, SET_ACTIVE_PARTITION, &partition);

/* position the tape to the beginning of partition 0 */
partition.partition_number = 0;
partition.logical_block_id = 0;
ioctl(fd, SET_ACTIVE_PARTITION, &partition);
```

## ***QUERY\_PARTITION***

The **QUERY\_PARTITION** IOCTL is used to return partition information for the tape drive and the current media in the tape drive. It includes the current active partition the tape drive is using for the media. The **number\_of\_partitions** field is the current number of partitions on the media and the **max\_partitions** is the maximum partitions that the tape drive supports. The **size\_unit** field can be either one of the defined values or another value such as 8. It is used with the **size\_array** field value for each partition to specify

the actual size partition sizes. The **partition\_method** field is either Wrap-wise Partitioning or Longitudinal Partitioning. Refer to “[CREATE\\_PARTITION](#)” on page 62 for details.

The data structure that is used with this IOCTL is

```
The define for "partition_method":
#define UNKNOWN_TYPE          0 /* vendor-specific or unknown */
#define WRAP_WISE_PARTITION    1 /* Wrap-wise Partitioning without RABF */
#define LONGITUDINAL_PARTITION 2 /* Longitudinal Partitioning */
#define WRAP_WISE_PARTITION_WITH_FASTSYNC 3 /* Wrap-wise Partitioning with RABF */

The define for "size_unit":
#define SIZE_UNIT_BYTES        0      /* Bytes */
#define SIZE_UNIT_KBYTES       3      /* Kilobytes */
#define SIZE_UNIT_MBYTES       6      /* Megabytes */
#define SIZE_UNIT_GBYTES       9      /* Gigabytes */
#define SIZE_UNIT_TBYTES       12     /* Terabytes */

struct query_partition {
    uchar max_partitions;           /* Max number of supported partitions */
    uchar active_partition;         /* current active partition on tape */
    uchar number_of_partitions;     /* Number of partitions from 1 to max */
    uchar size_unit;               /* Size unit of partition sizes below */
    ushort size[MAX_PARTITIONS];   /* Array of partition sizes in size units */
                                    /* for each partition, 0 to (number - 1) */
    uchar partition_method;        /* partitioning type */
    char reserved [31];
};
```

Examples of the **QUERY\_PARTITION** IOCTL

```
#include <sys/Atape.h>

struct query_partition partition;
int i;

if (ioctl(fd, QUERY_PARTITION, &partition) < 0)
    return errno;

printf("  Max supported partitions ... %d\n",partition.max_partitions);
printf("  Number of partitions ..... %d\n",partition.number_of_partitions);
printf("  Active partition ..... %d\n",partition.active_partition);
printf("  Partition Method ..... %d\n",partition.partition_method);
if (partition.size_unit == SIZE_UNIT_BYTES)
    printf("  Partition size unit ..... Bytes\n");
else if (partition.size_unit == SIZE_UNIT_KBYTES)
    printf("  Partition size unit ..... Kilobytes\n");
else if (partition.size_unit == SIZE_UNIT_MBYTES)
    printf("  Partition size unit ..... Megabytes\n");
else if (partition.size_unit == SIZE_UNIT_GBYTES)
    printf("  Partition size unit ..... Gigabytes\n");
else if (partition.size_unit == SIZE_UNIT_TBYTES)
    printf("  Partition size unit ..... Terabytes\n");
else
    printf("  Partition size unit ..... %d\n",partition.size_unit);

for (i=0; i < partition.number_of_partitions; i++)
    printf("  Partition %d size ..... %d\n",i,partition.size[i]);
```

## **CREATE\_PARTITION**

The **CREATE\_PARTITION** IOCTL is used to format the current media in the tape drive into 1 or more partitions. The number of partitions to create is specified in the **number\_of\_partitions** field. When more than one partition is created, the **type** field specifies the type of partitioning, either FDP, SDP, or IDP. The tape must be positioned at the beginning of tape (partition 0 logical block id 0) before this IOCTL is used.

If the **number\_of\_partitions** field to create in the IOCTL structure is one partition, all other fields are ignored and not used. The tape drive formats the media by using its default partitioning type and size for a single partition.

When the **type** field in the IOCTL structure is set to either FDP or SDP, the **size\_unit** and **size** fields in the IOCTL structure are not used. When the **type** field in the IOCTL structure is set to IDP, the **size\_unit** with the **size** fields are used to specify the size for each partition.

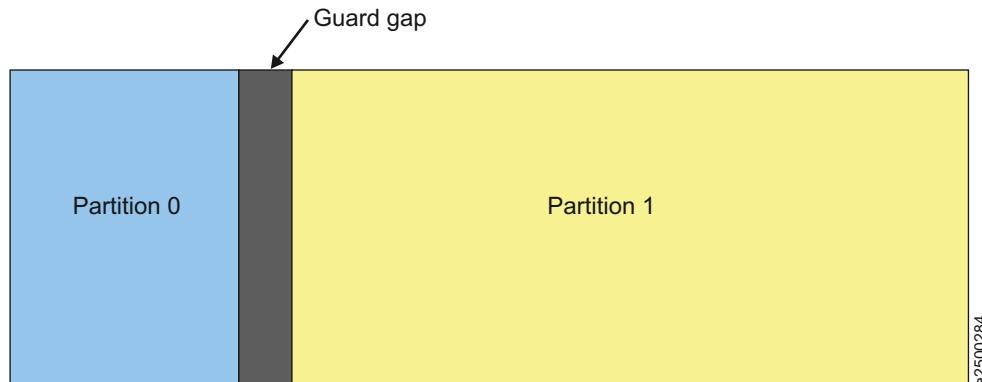
There are two partition types: Wrap-wise Partitioning (Figure 5 on page 63) optimized for streaming performance, and Longitudinal Partitioning (Figure 6 on page 63) optimized for random access performance. Media is always partitioned into 1 by default. Or, more than one partition where the data partition always exists as partition 0 and other extra index partition 1 to n can exist.

A WORM media cannot be partitioned and the **Format Medium** commands are rejected. Attempts to scale a partitioned media is accepted. However, only if you use the correct **FORMAT** field setting, as part of scaling the volume is set to a single data partition cartridge.



a2500283

Figure 5. Wrap-wise partitioning



a2500284

Figure 6. Longitudinal partitioning

The following chart lists the maximum number of partitions that the tape drive supports.

Table 3. Number of supported partitions	
Drive type	Maximum number of supported partitions
LTO 5 (TS2250 and TS2350) and later	2 in Wrap-wise Partitioning
3592 E07 (TS 1140)	4 in Wrap-wise Partitioning 2 in Longitudinal Partitioning

The data structure that is used with this IOCTL is

```

The define for "partition_method":
#define UNKNOWN_TYPE          0 /* vendor-specific or unknown */
#define WRAP_WISE_PARTITION    1 /* Wrap-wise Partitioning without RABF */
#define LONGITUDINAL_PARTITION 2 /* Longitudinal Partitioning */
#define WRAP_WISE_PARTITION_WITH_FASTSYNC 3 /* Wrap-wise Partitioning with RABF */

The define for "type":
#define IDP_PARTITION          1 /* Initiator Defined Partition type */
#define SDP_PARTITION           2 /* Select Data Partition type */
#define FDP_PARTITION           3 /* Fixed Data Partition type */

The define for "size_unit":
#define SIZE_UNIT_BYTES         0 /* Bytes */
#define SIZE_UNIT_KBYTES        3 /* Kilobytes */

```

```

#define SIZE_UNIT_MBYTES    6      /* Megabytes                      */
#define SIZE_UNIT_GBYTES    9      /* Gigabytes                       */
#define SIZE_UNIT_TBYTES   12      /* Terabytes                        */

struct tape_partition {
    uchar type;                  /* Type of tape partition to create */
    uchar number_of_partitions;  /* Number of partitions to create */
    uchar size_unit;             /* IDP size unit of partition sizes below */
    ushort size[MAX_PARTITIONS];/* Array of partition sizes in size units */
    uchar partition_method;     /* for each partition, 0 to (number - 1) */
    char reserved [31];         /* partitioning type */
};

```

### Examples of the **CREATE\_PARTITION** IOCTL

```

#include <sys/Atape.h>

struct tape_partition partition;

/* create 2 SDP partitions on LT0-5 */
partition.type = SDP_PARTITION;
partition.number_of_partitions = 2;
partition.partition_method = WRAP_WISE_PARTITION;
ioctl(fd, CREATE_PARTITION, &partition);

/* create 2 IDP partitions with partition 1 for 37 gigabytes and partition 0
for the remaining capacity on LT0-5 */
partition.type = IDP_PARTITION;
partition.number_of_partitions = 2;
partition.partition_method = WRAP_WISE_PARTITION;
partition.size_unit = SIZE_UNIT_GBYTES;
partition.size[0] = 0xFFFF;
partition.size[1] = 37;
ioctl(fd, CREATE_PARTITION, &partition);

/* format the tape into 1 partition */
partition.number_of_partitions = 1;
ioctl(fd, CREATE_PARTITION, &partition);

/* create 4 IDP partitions on 3592 JC volume in Wrap-wise partitioning
with partition 0 and 2 for 94.11 gigabytes (minimum size)and partition 1 and 3
to use the remaining capacity equally around 1.5 TB on 3592 E07 */
partition.type = IDP_PARTITION;
partition.number_of_partitions = 4;
partition.partition_method = WRAP_WISE_PARTITION;
partition.size_unit = 8;           /* 100 megabytes */
partition.size[0] = 0x03AD;
partition.size[1] = 0xFFFF;
partition.size[2] = 0x03AD;
partition.size[3] = 0x3AD2;

```

## **ALLOW\_DATA\_OVERWRITE**

The **ALLOW\_DATA\_OVERWRITE** IOCTL is used to set the drive to allow a subsequent data write type command at the current position. Or, it allows a **CREATE\_PARTITION** IOCTL when data safe (append-only) mode is enabled.

For a subsequent write type command, the **allow\_format\_overwrite** field must be set to 0. The **partition\_number** and **logical\_block\_id** fields must be set to the current partition and position within the partition where the overwrite occurs.

For a subsequent **CREATE\_PARTITION** IOCTL, the **allow\_format\_overwrite** field must be set to 1. The **partition\_number** and **logical\_block\_id** fields are not used. However, the tape must be at the beginning of tape (partition 0 logical block id 0) before the **CREATE\_PARTITION** IOCTL is issued.

The data structure that is used with this IOCTL is

```

struct allow_data_overwrite{
    uchar partition_number;        /* Partition number 0-n to overwrite */
    ullong logical_block_id;       /* Blockid to overwrite to within partition */
    uchar allow_format_overwrite; /* allow format if in data safe mode */
    char reserved[32];
};

```

## Examples of the **ALLOW\_DATA\_OVERWRITE** IOCTL

```
#include <sys/Atape.h>

struct read_tape_position rpos;
struct allow_data_overwrite data_overwrite;
struct set_active_partition partition;

/* get current tape position for a subsequent write type command and */
/* set the allow_data_overwrite fields with the current position for the next
write type command */
rpos.data_format = RP_LONG_FORMAT;
if (ioctl(fd, READ_TAPE_POSITION, &rpos) < 0)
    return errno;

data_overwrite.partition_number = rpos.rp_data.rp_long.active_partition;
data_overwrite.logical_block_id = rpos.rp_data.rp_long.logical_obj_number;
data_overwrite.allow_format_overwrite = 0;
ioctl (fd, ALLOW_DATA_OVERWRITE, &data_overwrite);

/* set the tape position to the beginning of tape and */
/* prepare a format overwrite for the CREATE_PARTITION ioctl */
partition.partition_number = 0;
partition.logical_block_id = 0;
if (ioctl(fd, SET_ACTIVE_PARTITION, &partition;) < 0)
    return errno;

data_overwrite.allow_format_overwrite = 1;
ioctl (fd, ALLOW_DATA_OVERWRITE, &data_overwrite);
```

## **QUERY\_LOGICAL\_BLOCK\_PROTECTION**

The IOCTL queries whether the drive can support this feature, what Logical Block Protection (LBP) method is used, and where the protection information is included.

The **lbp\_capable** field indicates whether the drive has logical block protection capability. The **lbp\_method** field displays if LBP is enabled and what the protection method is. The LBP information length is shown in the **lbp\_info\_length** field. The fields of **lbp\_w**, **lbp\_r**, and **rbdp** present that the protection information is included in write, read, or recover buffer data.

The data structure that is used with this IOCTL is

```
struct logical_block_protection
{
    uchar lbp_capable;      /* [OUTPUT] the capability of lbp for QUERY ioctl only */
    uchar lbp_method;       /* lbp method used for QUERY [OUTPUT]and SET [INPUT] ioctls */
#define LBP_DISABLE          0x00
#define REED_SOLOMON_CRC     0x01
    uchar lbp_info_length;  /* lbp info length for QUERY [OUTPUT] and SET [INPUT] ioctls */
    uchar lbp_w;             /* protection info included in write data */
    uchar lbp_r;             /* protection info included in read data */
    uchar rbdp;              /* protection info included in recover buffer data */
    uchar reserved[26];
};
```

## Examples of the **QUERY\_LOGICAL\_BLOCK\_PROTECTION** IOCTL

```
#include <sys/Atape.h>

struct logical_block_protection lbp_protect;

printf("Querying Logical Block Protection....\n");

if (ioctl(fd, QUERY_LOGICAL_BLOCK_PROTECTION, &lbp_protect) < 0)
    return errno;
printf(" Logical Block Protection capable..... %d\n",lbp_protect.lbp_capable);
printf(" Logical Block Protection method..... %d\n",lbp_protect.lbp_method);
printf(" Logical Block Protection Info Length.... %d\n",lbp_protect.lbp_info_length);
printf(" Logical Block Protection for Write..... %d\n",lbp_protect.lbp_w);
printf(" Logical Block Protection for Read..... %d\n",lbp_protect.lbp_r);
printf(" Logical Block Protection for RBDP..... %d\n",lbp_protect.rbdp);
```

## **SET\_LOGICAL\_BLOCK\_PROTECTION**

The IOCTL enables or disables Logical Block Protection, sets up what method is used, and where the protection information is included.

The **lbp\_capable** field is ignored in this IOCTL by the Atape driver. If the **lbp\_method** field is 0 (LBP\_DISABLE), all other fields are ignored and not used. When the **lbp\_method** field is set to a valid non-zero method, all other fields are used to specify the setup for LBP.

The data structure that is used with this IOCTL is

```
struct logical_block_protection
{
    uchar lbp_capable;      /* [OUTPUT] the capability of lbp for QUERY ioctl only */
    uchar lbp_method;       /* lbp method used for QUERY [OUTPUT] and SET [INPUT] ioctls */
    #define LBP_DISABLE      0x00
    #define REED_SOLOMON_CRC 0x01
    uchar lbp_info_length; /* lbp info length for QUERY [OUTPUT] and SET [INPUT] ioctls */
    uchar lbp_w;            /* protection info included in write data */
    /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    uchar lbp_r;            /* protection info included in read data */
    /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    uchar rbdp;             /* protection info included in recover buffer data */
    /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    uchar reserved[26];
};
```

Examples of the **SET\_LOGICAL\_BLOCK\_PROTECTION** IOCTL

```
#include <sys/Atape.h>

int rc;
struct logical_block_protection lbp_protect;

printf("Setting Logical Block Protection....\n\n");

printf ("Enter Logical Block Protection method:      ");
gets (buf);
lbp_protect.lbp_method= atoi(buf);
printf ("Enter Logical Block Protection Info Length: ");
gets (buf);
lbp_protect.lbp_info_length= atoi(buf);
printf ("Enter Logical Block Protection for Write:   ");
gets (buf);
lbp_protect.lbp_w= atoi(buf);
printf ("Enter Logical Block Protection for Read:     ");
gets (buf);
lbp_protect.lbp_r= atoi(buf);
printf ("Enter Logical Block Protection for RBDP:     ");
gets (buf);
lbp_protect.rbdp= atoi(buf);

rc = ioctl(fd, SET_LOGICAL_BLOCK_PROTECTION, &lbp_protect);

if (rc)
    printf ("Set Logical Block Protection Fails (rc %d)",rc);
else
    printf ("Set Logical Block Protection Succeeds");
```

### **Note:**

1. The drive always expects a CRC attached with a data block when LBP is enabled for **lbp\_r** and **lbp\_w**. Without the CRC bytes attachment, the drive fails the **Read and Write** command. To prevent the CRC block transfer between the drive and application, the maximum block size limit must be determined by application. Call the **STIOCQRYP** IOCTL and get the system maximum block size limit. Call the **Read Block Limits** command to get the drive maximum block size limit. Then, use the minimum of the two limits.
2. When a unit attention with a power-on and device reset (Sense key/Asc-Ascq x6/x2900) occurs, the LBP enable bits (**lbp\_w**, **lbp\_r**, and **rbdp**) are reset to OFF by default. Atape tape driver returns EIO for an IOCTL call in this situation. Once the application determines it is a reset unit attention in the sense data, it responses to query LBP setup and reissues this IOCTL to set up LBP properly.

3. The LBP setting is controlled by the application and not the device driver. If an application enables LBP, it must also disable LBP when it closes the drive, as this action is not done by the device driver.

### **STIOC\_READ\_ATTRIBUTE**

The IOCTL is issued to read attribute values that belongs to a specific partition from medium auxiliary memory.

The input or output data structure is

```
#define MAX_ATTR_LEN 1024
struct read_attribute
{
    uchar service_action; /* [IN] service action */
    uchar partition_number; /* [IN] the partition which the attributes belong to */
    ushort first_attr_id; /* [IN] first attribute id to be returned */
    uint attr_data_len; /* [OUT] length of attribute data returned */
    uchar reserved[8];
    char data[MAX_ATTR_LEN]; /* [OUT] read attributes data */
};
```

An example of the **STIOC\_READ\_ATTRIBUTE** command is

```
#include <sys/Atape.h>
int rc,attr_len;
struct read_attribute rd_attr;

memset(&rd_attr,0,sizeof(struct read_attribute));
rd_attr.service_action=0x00;
rd_attr.partition_number=1;
rd_attr.first_attr_id=0x800;

printf("Read attribute command ....\n");
rc=ioctl(fd, STIOC_READ_ATTRIBUTE, &rd_attr);

if (rc)
    printf ("Read Attribute failed (rc %d)",rc);
else
{
    printf ("Read Attribute Succeeds!");
    dump_bytes (rd_attr.data, min(MAX_ATTR_LEN, rd_attr.attr_data_len),
    "Attribute Data");
}
```

### **STIOC\_WRITE\_ATTRIBUTE**

The IOCTL sets the attributes in medium auxiliary memory at a specific partition.

Following is the structure for **STIOC\_WRITE\_ATTRIBUTE** IOCTL

```
struct write_attribute
{
    uchar write_cache; /* [IN] WTC - Write-through cache */
    uchar partition_number; /* [IN] the partition which the attribute is belonged to */
    uint parm_list_len; /* [IN] parameter list length */
    uchar reserved[10];
    char data[MAX_ATTR_LEN]; /* [IN] write attributes data */
};
```

An example of the **STIOC\_WRITE\_ATTRIBUTE** command is

```
#include <sys/Atape.h>

int rc;
struct write_attribute wr_attr;

memset(&wr_attr,0,sizeof(struct write_attribute));

wr_attr.write_cache=0;
wr_attr.parm_list_len=0x11;
wr_attr.data[3]=0x0D;
wr_attr.data[4]=0x08;
wr_attr.data[6]=0x01;
wr_attr.data[8]=0x08;
```

```

wr_attr.data[9]='I';
wr_attr.data[10]='B';
wr_attr.data[11]='M';
wr_attr.data[12]=' ';
wr_attr.data[13]='T';
wr_attr.data[14]='E';
wr_attr.data[15]='S';
wr_attr.data[16]='T';

printf("Issuing a sample Write Attribute command ....\n\n");
rc=ioctl(fd, STIOC_WRITE_ATTRIBUTE, &wr_attr);

if (rc)
    printf ("Write Attribute failed (rc %d)",rc);
else
    printf ("Write Attribute Succeeds");

```

## **VERIFY\_TAPE\_DATA**

The IOCTL issues a **VERIFY** command. This command causes data to be read from the tape and passed through the drive's error detection and correction hardware. This action determines whether it can be recovered from the tape, whether the protection information is present, and validates correctly on logical block on the medium. The driver returns the IOCTL a failure or a success if the **VERIFY SCSI** command is completed in a **Good SCSI** status.

### **Note:**

1. When an application sets the VBF method, it considers the driver's close operation in which the driver can write filemarks in its close, which the application did not explicitly request. For example, some drivers write two consecutive filemarks that mark the end of data on the tape in its close, if the last tape operation was a **WRITE** command.
2. Per the user's or application's request, Atape driver sets the block size in the field of **Block Length** in mode block descriptor for **Read and Write** commands. Then, it maintains this block size setting in a whole open. For instance, the tape driver sets a zero in the **Block Length** field for the variable block size. This act causes the missing of an overlength condition on a SILI Read. **Block Length** must be set to a non-zero value.

Before Fixed bit is set to ON with VTE or VBF ON in **Verify** IOCTL, the application is requested to set the block size in mode block descriptor. Then, the drive uses it to verify the length of each logical block. For example, a 256 KB length is set in **Block Length** field to verify the data. The setup overrides the early setting from IBM tape driver.

When the application completes the **Verify** IOCTL call, the original block size setting must be restored for **Read and Write** commands, the application either issues **set block size** IOCTL. Or, it closes the drive immediately and reopens the drive for the next tape operation. It is recommended to reopen the drive for the next tape operation. Otherwise, it causes next **Read and Write** command misbehavior.

3. To support DPF for **Verify** command with FIXED bit on, it is requested to issue an IBM tape driver to set **blksize** in **STIOCSETP** IOCTL. IBM tape driver sets the **block length** in mode block descriptor same as the block size and save the block size in kernel memory. The driver restores the **block length** before it tries the **Verify SCSI** command again. Otherwise, it causes the **Verify** command to fail.
4. The IOCTL can be returned longer than the timeout when DPF occurs.

The structure is defined for this IOCTL as

```

struct verify_data
{
    uint      : 2, /* reserved */
    vte: 1, /* [IN] verify to end-of-data */
    vlpbm: 1, /* [IN] verify logical block protection info */
    vbf: 1, /* [IN] verify by filemarks */
    immed: 1, /* [IN] return SCSI status immediately */
    bytcmp: 1, /* No use currently */
    fixed: 1; /* [IN] set Fixed bit to verify the length of each logical block */
uchar reseved[15];
uint verify_length; /* [IN] amount of data to be verified */
} ;

```

An example of the **VERIFY\_TAPE\_DATA** command is to verify all of logical block from the current position to end of data. It includes a verification that each logical block uses the logical block protection method that is specified in the **Control Data Protection** mode page, when vte is set to 1 with vlbpm on.

```
#include <sys/Atape.h>

int rc;
struct verify_data vrf_data;

memset(&vrf_data,0,sizeof(struct verify_data));
vrf_data.vte=1;
vrf_data.vlbpm=1;
vrf_data.vbf=0;
vrf_data.immed=0;
vrf_data.fixed=0;
vrf_data.verify_length=0;

printf("Verify Tape Data command ....\n");
rc=ioctl(fd,VERIFY_TAPE_DATA, &vrf_data);

if (rc)
    printf ("Verify Tape Data failed (rc %d)",rc);
else printf
    ("Verify Tape Data Succeeded!");
```

## **QUERY\_RAO\_INFO**

The IOCTL is used to query the maximum number and size of User Data Segments (UDS) that are supported from tape drive and driver for the wanted **uds\_type**. The application calls this IOCTL before the **GENERATE\_RAO** and **RECEIVE\_RAO** IOCTLs are issued. The application uses the return data to limit the number of UDS requested in the **GENERATE\_RAO** IOCTL.

The structure that is defined for this IOCTL as

```
struct query_rao_info {
    char uds_type;           /* [IN]  0: UDS_WITHOUT_GEOMETRY          */
                           /*           1: UDS_WITH_GEOMETRY           */
    char reserved[7];
    ushort max_uds_number;   /* [OUT] Max UDS number supported from drive */
                           /*           */
    ushort max_uds_size;     /* [OUT] Max single UDS size supported from drive in byte */
                           /*           */
    ushort max_host_uds_number; /* [OUT] Max UDS number supported from driver */
                           /*           */
}
```

An example of the **QUERY\_RAO\_INFO** command is

```
#include <sys/Atape.h>

int rc;
struct query_rao_info qrao;

bzero(&qrao,sizeof(struct query_rao_info));
qrao.udt=uds_type;

rc=ioctl(fd,QUERY_RAO_INFO,&qrao);

if (rc)
    printf("QUERY_RAO_INFO fails with rc %d\n",rc);
else
{
    max_host_uds_num=qrao.max_host_uds_number;
    max_uds_size=qrao.max_uds_size;
}

return rc;
```

## **GENERATE\_RAO**

The IOCTL is called to send a GRAO list to request the drive to generate a **Recommending Access Order** list.

The process method is either 1 or 2 to create a RAO list, and the type of UDS is either with or without the geometry. The **uds\_number** must be not larger than **max\_host\_uds\_number** in the **QUERY\_RAO\_INFO**

IOCTL. The application allocates a memory with **grao\_list\_leng** (`uds_number * sizeof(struct grao_uds_desc) + 8`) for the pointer of **grao\_list**. 8 bytes is the size that is needed for the header portion on of the return data.

The structures for the **GENERATE\_RAO** IOCTL are

```
struct generate_rao {
    char process; /* [IN] Requested process to generate RAO list */
    /* 0: no reorder UDS and no calculate locate time */
    /* (not currently supported by the drive) */
    /* 1: no reorder UDS but calculate locate time */
    /* 2: reorder UDS and calculate locate time */
    char uds_type; /* [IN] 0: UDS_WITHOUT_GEOMETRY */
    /* 1: UDS_WITH_GEOMETRY */
    char reserved1[2];
    uint grao_list_leng; /* [IN] The data length is allocated for GRAO list. */
    char *grao_list; /* [IN] the pointer is allocated to the size of grao_list_leng */
    /* (uds_number * sizeof(struct grao_uds_desc))
     * +sizeof(struct grao_list_header))
     * and contains the data of GRAO parameter list.
     * The uds number isn't larger than max_host_uds_number */
    /* in QUERY_RAO ioctl. */
    char reserved2[8];
}
```

The **grao** list is in the format and the parameter data can be generated by using the structures that are defined here.

```
-- List Header
-- UDS Segment Descriptor (first)
.....
-- UDS Segment Descriptor (last)

struct grao_list_header {
    uchar reserved[4];
    uint addl_data; /* additional data */
}

struct grao_uds_desc {
    ushort desc_leng; /* descriptor length */
    char reserved[3];
    char uds_name[10]; /* uds name given by application */
    char partition; /* Partition number 0-n to overwrite */
    ullong beginning_loi; /* Beginning logical object ID */
    ullong ending_loi; /* Ending logical object ID */
}
```

An example of the **GENERATE\_RAO** command is

```
#include<sys/Atape.h>

int rc;
struct generate_rao grao;

bzero(&grao,sizeof(struct generate_rao));

grao.process=2;
grao.udt_type=uds_type;
grao.grao_list_leng=max_host_uds_num * sizeof(struct grao_uds_desc)
+ sizeof(struct grao_list_header);

if (!(grao.grao_list=malloc(grao.grao_list_leng)))
{
    perror("Failure allocating memory");
    return (errno);
}
memset(grao.grao_list, 0, grao.grao_list_leng);
grao.grao_list[3]=36;

rc=ioctl(fd,GENERATE_RAO,&grao);
if (rc)
    printf("GENERATE_RAO fails with rc %d\n",rc);
else
    printf("GENERATE_RAO succeeds\n");

free(grao.grao_list);
```

```
    return rc;
```

## RECEIVE\_RAO

After a **GENERATE\_RAO** IOCTL is completed, the application calls the **RECEIVE\_RAO** IOCTL to receive a recommended access order of UDS from the drive. To avoid a system crash, it is important that the application allocates a large enough block of memory for the **\*rrao\_list** pointer and notifies the driver of the allocated size. It is done by indicating the size of the buffer in bytes to the *rrao\_list\_leng* variable as an input to the **receive\_rao\_list** structure.

The structure for the **RECEIVE\_RAO** IOCTL is

```
struct receive_rao_list {
    uint     rrao_list_offset; /* [IN] The offset of receive RAO list to */
    /* begin returning data */
    uint     rrao_list_leng;  /* [IN/OUT] number byte of data length */
    /* [IN] The data length is allocated for RRAO */
    /* list by application the length is */
    /* (max_uds_size * uds_number + */
    /* sizeof(struct rrao_list_header) */
    /* max_uds_size is reported in */
    /* sizeof(struct rrao_list_header) */
    /* uds_number is the total UDS number */
    /* requested from application in */
    /* GENERATE_RAO ioctl */
    /* [OUT] the data length is actual returned */
    /* in RRAO list from the driver */
    /* [IN/OUT] the data pointer of RRAO list */
    char    *rrao_list;
    char    reserved[8];
};
```

The **rrao** list is in this format.

```
List Header
UDS descriptor (first)
-- Basic UDS descriptor
-- Additional UDS info descriptor (first)
.....
-- Additional UDS info descriptor (last)

.....
UDS descriptor (last)
-- Basic UDS descriptor
-- Additional UDS info descriptor (first)
.....
-- Additional UDS info descriptor (last)
```

The sample code is

```
int rc;
struct receive_rao_list rrao;

bzero(&rrao,sizeof(struct receive_rao_list));

rrao.rrao_list_offset=0;
rrao.rrao_list_leng=max_host_uds_num * max_uds_size + 8;
/* 8 is the header of rrao list */

if (!(rrao.rrao_list=malloc(rrao.rrao_list_leng)))
{
    perror("Failure allocating memory");
    return (errno);
}
memset(rrao.rrao_list, 0, rrao.rrao_list_leng);

rc=ioctl(fd,RECEIVE_RAO,&rrao);
if (rc)
    printf("RECEIVE_RAO fails with rc %d\n",rc);
else
    printf("rrao_list_leng %d\n",rrao.rrao_list_leng);

free(rrao.rrao_list);
```

```
    return rc;
```

## **QUERY\_ARCHIVE\_MODE\_UNTHREAD**

This ioctl reports the feature of **Archive Mode Untread** enable or disable in the tape drive.

The structure for this IOCTL is

```
struct archive_mode_unthread
{
    uchar amu_enable;
    uchar reserve[7];
} archive_mode_unthread
```

An example of the **QUERY\_ARCHIVE\_MODE\_UNTHREAD** command is

```
#include <sys/Atape.h>
int rc;
struct archive_mode_unthread qamu;
bzero(&qamu,sizeof(struct archive_mode_unthread));
rc = ioctl(fd,QUERY_ARCHIVE_MODE_UNTHREAD,&qamu);
if (rc)
printf("QUERY_ARCHIVE_MODE_UNTHREAD fails with rc %d\n",rc);
else
printf("amu_enable %d \n", qamu.amu_enable);
return rc;
```

## **SET\_ARCHIVE\_MODE\_UNTHREAD**

This ioctl is used to enable or disable the feature of **Archive Mode Untread** in the tape drive.

The structure is defined for this IOCTL as

```
struct archive_mode_unthread
{
    uchar amu_enable;
    uchar reserve[7];
} archive_mode_unthread;
```

An example of the **SET\_ARCHIVE\_MODE\_UNTHREAD** command is

```
#include <sys/Atape.h>
int rc;
struct archive_mode_unthread set_amu;
set_amu.amu_enable=ENABLE;
rc=ioctl(fd,SET_ARCHIVE_MODE_UNTHREAD,&set_amu);
if (rc)
printf("SET_ARCHIVE_MODE_UNTHREAD fails with rc %d\n",rc);
else
printf("SET_ARCHIVE_MODE_UNTHREAD succeeds \n");
return rc;
```

## **TAPE\_LOAD\_UNLOAD**

This ioctl is called to complete the various behaviors of tape cartridge loading. When **Archive Mode Untread** is enabled, the variable of **RETEN** can be used to set the new unthread method in this new load command. See [Table 4 on page 72](#) for interaction with other variables.

*Table 4. Behavior for the combinations of the retention, load, hold variables*

Volume position	Hold	Load*	Retention	Description
U, M, I, or T	0	0	0	Unload the cartridge from the drive. Upon completion of the command, MAM is not accessible. If the cartridge is already unloaded, GOOD Status is returned.

Table 4. Behavior for the combinations of the retention, load, hold variables (continued)

Volume position	Hold	Load*	Retention	Description
U, M, or I	0	0	1	Unload the cartridge from the drive. Upon completion of the command, MAM is not accessible. If the cartridge is already unloaded, GOOD Status is returned.
T	0	0	1	Run an archive mode Unthread (see “SET_ARCHIVE_MODE_UNTHREAD” on page 72) and then unload the cartridge from the drive. Upon completion of the command, MAM is not accessible.
U, M, or I	0	1	-	Load the cartridge and become READY.
T	0	1	-	The logical position is set to logical block 0 of partition 0 (BOP 0) (Not equivalent to a <b>Rewind</b> command as the active partition is set to partition 0).
U, M, I, or T	1	-	1	The cartridge is moved to the seated position with MAM accessible and the tape not threaded.
U, M, or I	1	-	1	The cartridge is moved to the seated position with MAM accessible and the tape not threaded.
T	1	-	1	An archive mode Unthread (see “SET_ARCHIVE_MODE_UNTHREAD” on page 72) is completed and then the cartridge is moved to the seated position with MAM accessible and the tape not threaded.

**Key:**

U - Unloaded

M - MAM accessible not threaded

I = In the IDLE\_C power condition state (that is, power-saving mode) with a volume mounted

T - Threaded

- = Don't care

\*The **LOAD UNLOAD** command with the LOAD bit set to 0b is sometimes called an **unload** command. The **LOAD UNLOAD** command with the LOAD bit set to 1b is sometimes called a **reload** command.

The structure for the **TAPE\_LOAD\_UNLOAD** ioctl is defined as:

```
struct tape_load_unload
{
    uchar immediate; /* [IN] immediate return a SCSI status */
    uchar hold; /* [IN] request MAM accessible and volume not for access */
    uchar eot; /* [IN] end of tape, no use and always 0 now */
    uchar retention; /* [IN] archive mode unthread */
    uchar load; /* [IN] perform a load or unload */
    char reserved[11];
};
```

An example of the **TAPE\_LOAD\_UNLOAD** command is

```
#include <sys/Atape.h>
int rc;
struct tape_load_unload load_unload;
printf("Tape Load Unload Command....\n");
printf ("Enter immediate : ");
gets (buf);
load_unload.immediate = atoi(buf);
printf ("Enter hold : ");
```

```

gets (buf);
load_unload.hold = atoi(buf);
printf ("Enter retention : ");
gets (buf);
load_unload.retention = atoi(buf);
printf ("Enter load : ");
gets (buf);
load_unload.load = atoi(buf);
printf ("\n");
rc = ioctl(fd, TAPE_LOAD_UNLOAD, &load_unload);
if (rc)
printf ("Tape Load Unload Command Fails (rc %d)",rc);
else
printf ("Tape Load Unload Command Succeeds");
return rc;

```

## ***GET\_VHF\_DEVICE\_STATUS***

The device drivers provide application a friendly interface **GET\_VHF\_DEVICE\_STATUS** to report the tape drive and medium statuses. (VHF is short Very High Frequency).

The structure for the **GET\_VHF\_DEVICE\_STATUS** ioctl is

```

struct get_vhf_device_status
{
uchar pamr; /* [OUT] prevent/allow medium removal */
uchar hiu; /* [OUT] host initiated unload */
uchar macc; /* [OUT] medium auxiliary memory (MAM) accessible */
uchar cmpr; /* [OUT] data compress is enabled */
uchar wrtp; /* [OUT] volume is physically write protected when mprsnt is on */
uchar crqst; /* [OUT] cleaning requested */
uchar crqrd; /* [OUT] cleaning required */
uchar dinit; /* [OUT] this VHF data valid or invalid */
uchar inxtn; /* [OUT] device is transitioning to another state */
uchar rsvd1;
uchar raa; /* [OUT] robotic access allowed */
uchar mprsnt; /* [OUT] medium present */
uchar rsvd2;
uchar mstd; /* [OUT] medium seated */
uchar mthrd; /* [OUT] medium threaded */
uchar mounted; /* [OUT] medium is mounted */
uchar dev_act; /* [OUT] the current activity of the device below */
/* 00h No DT device activity
01h Cleaning operation in progress
02h Medium is being loaded
03h Medium is being unloaded
04h Other medium activity
05h Reading from medium
06h Writing to medium
07h Locating medium
08h Rewinding medium
09h Erasing medium
0Ah Formatting medium
0Bh Calibrating medium
0Ch Other DT device activity
0Dh Microcode update in progress
0Eh Reading encrypted from medium
0Fh Writing encrypted to medium */
uchar vs;
uchar rsvd3;
uchar tddc; /* [OUT] tape diagnostic data entry created */
uchar epp; /* [OUT] encryption parameters present */
uchar esr; /* [OUT] encryption service request */
uchar rrqst; /* [OUT] recovery requested */
uchar intfc; /* [OUT] interface changed */
uchar tafc; /* [OUT] tape alert state flag changed */
/* the fields below for LTO only */
uchar rsvd4[6];
uchar overwrite; /* [OUT] overwrite occurs and write mode isn't in append-only
mode */
uchar pcl_p; /* [OUT] potential conflict list present */
uchar rsvd5[6];
uchar hostlogin; /* [OUT] host login has occurred */
uchar sleepmode; /* [OUT] operating in one of the low power states */
char reserved[15];
};

```

**Note:** IBM tape drivers reports the data from the tape drive. The device drivers still return no error on this ioctl when the tape drive returns a good SCSI status even though VHF data is invalid, so the application must determine if the data is valid before to use it. The variable of dinit indicates this VHF data valid or invalid. The variable of wrtp shows a volume is physically write protected only when mprsnt is ON.

An example of the **GET\_VHF\_DEVICE\_STATUS** command is

```
#include <sys/Atape.h>
int rc;
struct get_vhf_device_status qvhf;
bzero(&qvhf,sizeof(struct get_vhf_device_status));
rc = ioctl(fd,GET_VHF_DEVICE_STATUS,&qvhf);
if (rc)
printf("GET_VHF_DEVICE_STATUS fails with rc %d\n",rc);
else
printf ("GET_VHF_DEVICE_STATUS Command Succeeds");
return rc;
```

## Medium changer IOCTL operations

This chapter describes the set of IOCTL commands that provides control and access to the SCSI medium changer functions. These IOCTL operations can be issued to the tape special file (such as **rmt0**), through a separate special file (such as **rmt0.smc**) that was created during the configuration process, or a separate special file (such as **smc0**), to access the medium changer.

When an application opens a **/dev/rmt** special file that is assigned to a drive that has access to a medium changer, these IOCTL operations are also available. The interface to the **/dev/rmt\*.smc** special file provides the application access to a separate medium changer device. When this special file is open, the medium changer is treated as a separate device. While **/dev/rmt\*.smc** is open, access to the IOCTL operations is restricted to **/dev/rmt\*.smc** and any attempt to access them through **/dev/rmt\*** fails.

### Overview

The following IOCTL commands are supported.

#### **SMCIOC\_ELEMENT\_INFO**

Obtain the device element information.

#### **SMCIOC\_MOVE\_MEDIUM**

Move a cartridge from one element to another element.

#### **SMCIOC\_EXCHANGE\_MEDIUM**

Exchange a cartridge in an element with another cartridge.

#### **SMCIOC\_POS\_TO\_ELEM**

Move the robot to an element.

#### **SMCIOC\_INIT\_ELEM\_STAT**

Issue the **SCSI Initialize Element Status** command.

#### **SMCIOC\_INIT\_ELEM\_STAT\_RANGE**

Issue the **SCSI Initialize Element Status with Range** command.

#### **SMCIOC\_INVENTORY**

Return the information about the four element types.

#### **SMCIOC\_LOAD\_MEDIUM**

Load a cartridge from a slot into the drive.

#### **SMCIOC\_UNLOAD\_MEDIUM**

Unload a cartridge from the drive and return it to a slot.

#### **SMCIOC\_PREVENT\_MEDIUM\_REMOVAL**

Prevent medium removal by the operator.

#### **SMCIOC\_ALLOW\_MEDIUM\_REMOVAL**

Allow medium removal by the operator.

## **SMCIOC\_READ\_ELEMENT\_DEVIDS**

Return the device ID element descriptors for drive elements.

## **SMCIOC\_READ\_CARTIDGE\_LOCATION**

Returns the cartridge location information for storage elements in the library.

These IOCTL commands and their associated structures are defined by including the **/usr/include/sys/Atape.h** header file in the C program by using the functions.

## **SMCIOC\_ELEMENT\_INFO**

This IOCTL command obtains the device element information.

The data structure is

```
struct element_info
{
    ushort robot_addr;           /* first robot address */
    ushort robots;               /* number of medium transport elements */
    ushort slot_addr;            /* first medium storage element address */
    ushort slots;                /* number of medium storage elements */
    ushort ie_addr;              /* first import/export element address */
    ushort ie_stations;          /* number of import/export elements */
    ushort drive_addr;            /* first data-transfer element address */
    ushort drives;                /* number of data-transfer elements */
};
```

An example of the **SMCIOC\_ELEMENT\_INFO** command is

```
#include <sys/Atape.h>

struct element_info element_info;

if (!ioctl (smcfd, SMCIOC_ELEMENT_INFO, &element_info))
{
    printf ("The SMCIOC_ELEMENT_INFO ioctl succeeded\n");
    printf ("\nThe element information data is:\n");
    dump_bytes ((uchar *)&element_info, sizeof (struct element_info));
}
else
{
    perror ("The SMCIOC_ELEMENT_INFO ioctl failed");
    smcioc_request_sense();
}
```

## **SMCIOC\_MOVE\_MEDIUM**

This IOCTL command moves a cartridge from one element to another element.

The data structure is

```
struct move_medium
{
    ushort robot;                 /* robot address */
    ushort source;                /* move from location */
    ushort destination;            /* move to location */
    char invert;                  /* invert before placement bit */
};
```

An example of the **SMCIOC\_MOVE\_MEDIUM** command is

```
#include <sys/Atape.h>

struct move_medium move_medium;

move_medium.robot = 0;
move_medium.invert = 0;
move_medium.source = source;
move_medium.destination = dest;

if (!ioctl (smcfd, SMCIOC_MOVE_MEDIUM, &move_medium))
    printf ("The SMCIOC_MOVE_MEDIUM ioctl succeeded\n");
else
```

```

    perror ("The SMCIOC_MOVE_MEDIUM ioctl failed");
    smcioc_request_sense();
}

```

## ***SMCIOC\_EXCHANGE\_MEDIUM***

This IOCTL command exchanges a cartridge in an element with another cartridge. This command is equivalent to two **SCSI Move Medium** commands. The first moves the cartridge from the source element to the **destination1** element. The second moves the cartridge that was previously in the **destination1** element to the **destination2** element. The **destination2** element can be the same as the source element.

The input data structure is

```

struct exchange_medium
{
    ushort robot;           /* robot address */
    ushort source;          /* source address for exchange */
    ushort destination1;   /* first destination address for exchange */
    ushort destination2;   /* second destination address for exchange */
    char invert1;           /* invert before placement into destination1 */
    char invert2;           /* invert before placement into destination2 */
};

```

An example of the **SMCIOC\_EXCHANGE\_MEDIUM** command is

```

#include <sys/Atape.h>

struct exchange_medium exchange_medium;

exchange_medium.robot = 0;
exchange_medium.invert1 = 0;
exchange_medium.invert2 = 0;
exchange_medium.source = 32;      /* slot 32 */
exchange_medium.destination1 = 16; /* drive address 16 */
exchange_medium.destination2 = 35; /* slot 35 */

/* exchange cartridge in drive address 16 with cartridge from slot 32 and */
/* return the cartridge currently in the drive to slot 35 */
if (!ioctl (smcfid, SMCIOC_EXCHANGE_MEDIUM, &exchange_medium))
    printf ("The SMCIOC_EXCHANGE_MEDIUM ioctl succeeded\n");
else
{
    perror ("The SMCIOC_EXCHANGE_MEDIUM ioctl failed");
    smcioc_request_sense();
}

```

## ***SMCIOC\_POS\_TO\_ELEM***

This IOCTL command moves the robot to an element.

The input data structure is

```

struct pos_to_elem
{
    ushort robot;           /* robot address */
    ushort destination;     /* move to location */
    char invert;             /* invert before placement bit */
};

```

An example of the **SMCIOC\_POS\_TO\_ELEM** command is

```

#include <sys/Atape.h>

char buf[10];
struct pos_to_elem pos_to_elem;

pos_to_elem.robot = 0;
pos_to_elem.invert = 0;
pos_to_elem.destination = dest;

if (!ioctl (smcfid, SMCIOC_POS_TO_ELEM, &pos_to_elem))
    printf ("The SMCIOC_POS_TO_ELEM ioctl succeeded\n");
else

```

```

    {
        perror ("The SMCIOC_POS_TO_ELEM ioctl failed");
        smcioc_request_sense();
    }

```

## ***SMCIOC\_INIT\_ELEM\_STAT***

This IOCTL command instructs the medium changer robotic device to issue the **SCSI Initialize Element Status** command.

There is no associated data structure.

An example of the **SMCIOC\_INIT\_ELEM\_STAT** command is

```

#include <sys/Atape.h>

if (!ioctl (smcfd, SMCIOC_INIT_ELEM_STAT, NULL))
    printf ("The SMCIOC_INIT_ELEM_STAT ioctl succeeded\n");
else
{
    perror ("The SMCIOC_INIT_ELEM_STAT ioctl failed");
    smcioc_request_sense();
}

```

## ***SMCIOC\_INIT\_ELEM\_STAT\_RANGE***

This IOCTL command issues the **SCSI Initialize Element Status with Range** command. It is used to audit specific elements in a library by specifying the starting element address and number of elements. Use the **SMCIOC\_INIT\_ELEM\_STAT** IOCTL to audit all elements.

The data structure is

```

struct element_range
{
    ushort element_address;      /* starting element address */
    ushort number_elements;     /* number of elements */
}

```

An example of the **SMCIOC\_INIT\_ELEM\_STAT\_RANGE** command is

```

#include <sys/Atape.h>

struct element_range elements;
/* audit slots 32 to 36 */
elements.element_address = 32;
elements.number_elements = 5;

if (!ioctl (smcfd, SMCIOC_INIT_ELEM_STAT_RANGE, &elements))
    printf ("The SMCIOC_INIT_ELEM_STAT_RANGE ioctl succeeded\n");
else
{
    perror ("The SMCIOC_INIT_ELEM_STAT_RANGE ioctl failed");
    smcioc_request_sense();
}

```

## ***SMCIOC\_INVENTORY***

This IOCTL command returns information about the four element types. The software application processes the input data (the number of elements about which it requires information). Then, it allocates a buffer large enough to hold the output for each element type.

The input data structure is

```

struct element_status
{
    ushort address;           /* element address */
    uint :2,                  /* reserved */
    inenab:1,                 /* media into changer's scope */
    exenab:1,                 /* media out of changer's scope */
    access:1,                  /* robot access allowed */
    except:1,                 /* abnormal element state */
}

```

```

    impexp:1,          /* import/export placed by operator or robot */
    full:1;           /* element contains medium */
uchar resvd1;        /* reserved */
uchar asc;           /* additional sense code */
uchar ascq;          /* additional sense code qualifier */
uint notbus:1,       /* element not on same bus as robot */
:1,                 /* reserved */
idvalid:1,          /* element address valid */
luvalid:1,          /* logical unit valid */
:1,                 /* reserved */
lun:3;              /* logical unit number */
uchar scsi;          /* SCSI bus address */
uchar resvd2;        /* reserved */
uint svalid:1,       /* element address valid */
invert:1,            /* medium inverted */
:6;                 /* reserved */
ushort source;       /* source storage element address */
uchar volume[36];   /* primary volume tag */
uchar resvd3[4];     /* reserved */
};

struct inventory
{
    struct element_status *robot_status; /* medium transport element pages */
    struct element_status *slot_status;  /* medium storage element pages */
    struct element_status *ie_status;    /* import/export element pages */
    struct element_status *drive_status; /* data-transfer element pages */
};

```

An example of the **SMCIOC\_INVENTORY** command is

```

#include <sys/Atape.h>

ushort i;
struct element_status robot_status[1];
struct element_status slot_status[20];
struct element_status ie_status[1];
struct element_status drive_status[1];
struct inventory      inventory;

bzero((caddr_t)robot_status,sizeof(struct element_status));

for (i=0;i<20;i++)
    bzero((caddr_t)(&slot_status[i]),sizeof(struct element_status));

bzero((caddr_t)ie_status,sizeof(struct element_status));
bzero((caddr_t)drive_status,sizeof(struct element_status));

smcioc_element_info();

inventory.robot_status = robot_status;
inventory.slot_status = slot_status;
inventory.ie_status = ie_status;
inventory.drive_status = drive_status;

if (!ioctl (smcfd, SMCIOC_INVENTORY, &inventory))
{
    printf ("\nThe SMCIOC_INVENTORY ioctl succeeded\n");
    printf ("\nThe robot status pages are:\n");

    for (i = 0; i < element_info.robots; i++)
    {
        dump_bytes ((uchar *) (inventory.robot_status+i),
                    sizeof (struct element_status));
        printf ("\n--- more ---");
        getchar();
    }

    printf ("\nThe slot status pages are:\n");
    for (i = 0; i < element_info.slots; i++)
    {
        dump_bytes ((uchar *) (inventory.slot_status+i),
                    sizeof (struct element_status));
        printf ("\n--- more ---");
        getchar();
    }
}

```

```

printf ("\nThe ie status pages are:\n");
for (i = 0; i < element_info.ie_stations; i++)
{
    dump_bytes ((uchar *)(inventory.ie_status+i),
                sizeof (struct element_status));
    printf ("\n--- more ---");
    getchar();
}

printf ("\nThe drive status pages are:\n");
for (i = 0; i < element_info.drives; i++)
{
    dump_bytes ((uchar *)(inventory.drive_status+i),
                sizeof (struct element_status));
    printf ("\n--- more ---");
    getchar();
}
else
{
    perror ("The SMCIOC_INVENTORY ioctl failed");
    smcioc_request_sense();
}

```

### **SMCIOC\_LOAD\_MEDIUM**

This IOCTL command loads a tape from a specific slot into the drive. Or, it loads from the first full slot into the drive if the slot address is specified as zero.

An example of the **SMCIOC\_LOAD\_MEDIUM** command is

```

#include <sys/Atape.h>

/* load cartridge from slot 3 */
if (ioctl (tapefd, SMCIOC_LOAD_MEDIUM,3)<0)
{
    printf ("IOCTL failure. errno=%d\n",errno)
    exit(1);
}

/* load first cartridge from magazine */
if (ioctl (tapefd, SMCIOC_LOAD_MEDIUM,0)<0)
{
    printf ("IOCTL failure. errno=%d\n",errno)
    exit(1);
}

```

### **SMCIOC\_UNLOAD\_MEDIUM**

This IOCTL command moves a tape from the drive and returns it to a specific slot. Or, it moves a tape to the first empty slot in the magazine if the slot address is specified as zero. If the IOCTL is issued to the **/dev/rmt** special file, the tape is automatically rewound and unloaded from the drive first.

An example of the **SMCIOC\_UNLOAD\_MEDIUM** command is

```

#include <sys/Atape.h>

/* unload cartridge to slot 3 */
if (ioctl (tapefd, SMCIOC_UNLOAD_MEDIUM,3)<0)
{
    printf ("IOCTL failure. errno=%d\n",errno)
    exit(1);
}

/* unload cartridge to first empty slot in magazine */
if (ioctl (tapefd, SMCIOC_UNLOAD_MEDIUM,0)<0)
{
    printf ("IOCTL failure. errno=%d\n",errno)
    exit(1);
}

```

## **SMCIOC\_PREVENT\_MEDIUM\_REMOVAL**

This IOCTL command prevents an operator from removing medium from the device until the **SMCIOC\_ALLOW\_MEDIUM\_REMOVAL** command is issued or the device is reset.

There is no associated data structure.

An example of the **SMCIOC\_PREVENT\_MEDIUM\_REMOVAL** command is

```
#include <sys/Atape.h>

if (!ioctl (smcfd, SMCIOC_PREVENT_MEDIUM_REMOVAL, NULL))
    printf ("The SMCIOC_PREVENT_MEDIUM_REMOVAL ioctl succeeded\n");
else
{
    perror ("The SMCIOC_PREVENT_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}
```

## **SMCIOC\_ALLOW\_MEDIUM\_REMOVAL**

This IOCTL command allows an operator to remove medium from the device. This command is used normally after an **SMCIOC\_PREVENT\_MEDIUM\_REMOVAL** command to restore the device to the default state.

There is no associated data structure.

An example of the **SMCIOC\_ALLOW\_MEDIUM\_REMOVAL** command is

```
#include <sys/Atape.h>

if (!ioctl (smcfd, SMCIOC_ALLOW_MEDIUM_REMOVAL, NULL))
    printf ("The SMCIOC_ALLOW_MEDIUM_REMOVAL ioctl succeeded\n");
else
{
    perror ("The SMCIOC_ALLOW_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}
```

## **SMCIOC\_READ\_ELEMENT\_DEVIDS**

This IOCTL command issues the **SCSI Read Element Status** command with the device ID (DVCID) bit set and returns the element descriptors for the data transfer elements. The **element\_address** field specifies the starting address of the first data transfer element. The **number\_elements** field specifies the number of elements to return. The application must allocate a return buffer large enough for the **number\_elements** specified in the input structure.

The input data structure is

```
struct read_element_devids
{
    ushort element_address;           /* starting element address */
    ushort number_elements;          /* number of elements */
    struct element_devid *drive_devid; /* data transfer element pages */
};
```

The output data structure is

```
struct element_devid
{
    ushort address;                 /* element address */
    uint :4,                         /* reserved */
    access:1,                        /* robot access allowed */
    except:1,                        /* abnormal element state */
    :1,                             /* reserved */
    full:1;                         /* element contains medium */
    uchar resvd1;                   /* reserved */
    uchar asc;                       /* additional sense code */
    uchar ascq;                      /* additional sense code qualifier */
    uint notbus:1,                   /* element not on same bus as robot */
    :1,                             /* reserved */
    idvalid:1,                       /* element address valid */
};
```

```

        luvalid:1,           /* logical unit valid */
        :1,                 /* reserved */
        lun:3;              /* logical unit number */
        uchar scsi;          /* scsi bus address */
        uchar resvd2;         /* reserved */
        uint svalid:1,        /* element address valid */
        invert:1,             /* medium inverted */
        :6;                 /* reserved */
        ushort source;        /* source storage element address */
        uint :4,               /* reserved */
        code_set:4;           /* code set X'2' is all ASCII identifier */
        uint :4,               /* reserved */
        ident_type:4;         /* identifier type */
        uchar resvd3;          /* reserved */
        uchar ident_len;        /* identifier length */
        uchar identifier[36];   /* device identification */
    };

```

An example of the **SMCIOC\_READ\_ELEMENT\_DEVIDS** command is

```

#include <sys/Atape.h>

int smcioc_read_element_devids()
{
    int i;
    struct element_devid *elem_devid, *elemp;
    struct read_element_devids devids;
    struct element_info element_info;

    if (ioctl(fd, SMCIOC_ELEMENT_INFO, &element_info))
        return errno;

    if (element_info.drives)
    {
        elem_devid = malloc(element_info.drives * sizeof(struct element_devid));
        if (elem_devid == NULL)
        {
            errno = ENOMEM;
            return errno;
        }
        bzero((caddr_t)elem_devid,element_info.drives * sizeof(struct element_devid));
        devids.drive_devid = elem_devid;
        devids.element_address = element_info.drive_addr;
        devids.number_elements = element_info.drives;

        printf("Reading element device ids...\n");

        if (ioctl (fd, SMCIOC_READ_ELEMENT_DEVIDS, &devids))
        {
            free(elem_devid);
            return errno;
        }

        elemp = elem_devid;
        for (i = 0; i < element_info.drives; i++, elemp++)
        {
            printf("\nDrive Address %d\n",elemp->address);
            if (elemp->except)
                printf(" Drive State ..... Abnormal\n");
            else
                printf(" Drive State ..... Normal\n");
            if (elemp->asc == 0x81 && elemp->ascq ==0x00)
                printf(" ASC/ASCQ ..... %02X%02X (Drive Present)\n",
                       elemp->asc,elemp->ascq);
            else if (elemp->asc == 0x82 && elemp->ascq ==0x00)
                printf(" ASC/ASCQ ..... %02X%02X (Drive Not Present)\n",
                       elemp->asc,elemp->ascq);
            else
                printf(" ASC/ASCQ ..... %02X%02X\n",
                       elemp->asc,elemp->ascq);
            if (elemp->full)
                printf(" Media Present ..... Yes\n");
            else
                printf(" Media Present ..... No\n");
            if (elemp->access)
                printf(" Robot Access Allowed ..... Yes\n");
            else
                printf(" Robot Access Allowed ..... No\n");
            if (elemp->svalid)
                printf(" Source Element Address ..... %d\n",elemp->source);
        }
    }
}

```

```

        else
            printf("  Source Element Address Valid ... No\n");
        if (elemp->invert)
            printf("  Media Inverted ..... Yes\n");
        else
            printf("  Media Inverted ..... No\n");
        if (elemp->notbus)
            printf("  Same Bus as Medium Changer ..... No\n");
        else
            printf("  Same Bus as Medium Changer ..... Yes\n");
        if (elemp->idvalid)
            printf("  SCSI Bus Address ..... %d\n",elemp->scsi);
        else
            printf("  SCSI Bus Address Valid ..... No\n");
        if (elemp->lunvalid)
            printf("  Logical Unit Number ..... %d\n",elemp->lun);
        else
            printf("  Logical Unit Number Valid ..... No\n");
        printf("  Device ID ..... %0.36s\n", elemp->identifier);
    }
}
else
{
    printf("\nNo drives found in element information\n");
}

free(elem_devid);
return errno;
}

```

## **SMCIOC\_READ\_CARTIDGE\_LOCATION**

The **SMCIOC\_READ\_CARTIDGE\_LOCATION** IOCTL is used to return the cartridge location information for storage elements in the library. The **element\_address** field specifies the starting element address to return. The **number\_elements** field specifies how many storage elements are returned. The **data** field is a pointer to the buffer for return data. The buffer must be large enough for the number of elements that are returned. If the storage element contains a cartridge, then the **ASCII identifier** field in return data specifies the location of the cartridge.

**Note:** This IOCTL is supported only on the TS3500 (3584) library.

The data structures that are used with this IOCTL are

```

struct cartridge_location_data
{
    ushort address;           /* element address */
    uint :4;                 /* reserved */
    access:1;                /* robot access allowed */
    except:1;                /* abnormal element state */
    :1;                      /* reserved */
    full:1;                 /* element contains medium */
    uchar resvd1;            /* reserved */
    uchar asc;                /* additional sense code */
    uchar ascq;               /* additional sense code qualifier */
    uchar resvd2[3];          /* reserved */
    uint svalid:1;            /* element address valid */
    invert:1;                /* medium inverted */
    :6;                      /* reserved */
    ushort source;            /* source storage element address */
    uchar volume[36];         /* primary volume tag */
    uint :4;                 /* reserved */
    code_set:4;               /* code set X'2' is all ASCII identifier */
    uint :4;                 /* reserved */
    ident_type:4;             /* identifier type */
    uchar resvd3;              /* reserved */
    uchar ident_len;           /* identifier length */
    uchar identifier[24];      /* slot identification */
};

struct read_cartridge_location
{
    ushort element_address;      /* starting element address */
    ushort number_elements;      /* number of elements */
    struct cartridge_location_data *data; /* storage element pages */
    char reserved[8];            /* reserved */
};

```

### Example of the **SMCIOC\_READ\_CARTRIDGE\_LOCATION** IOCTL

```
#include <sys/Atape.h>

int i;
struct cartridge_location_data *data, *elemp;
struct read_cartridge_location cart_location;
struct element_info element_info;

/* get the number of slots and starting element address */
if (ioctl(fd, SMCIOC_ELEMENT_INFO, &element_info) < 0)
    return errno;

if (element_info.slots == 0)
    return 0;

data = malloc(element_info.slots * sizeof(struct cartridge_location_data));
if (data == NULL)
    return ENOMEM;

/* Read cartridge location for all slots */
bzero(data,element_info.slots * sizeof(struct cartridge_location_data));
cart_location.data = data;
cart_location.element_address = element_info.slot_addr;
cart_location.number_elements = element_info.slots;

if (ioctl (fd, SMCIOC_READ_CARTRIDGE_LOCATION, &cart_location) < 0)
{
    free(data);
    return errno;
}

elemp = data;
for (i = 0; i < element_info.slots; i++, elemp++)
{
    if (elemp->address == 0
        continue;

    printf("Slot Address %d\n", elemp->address);
    if (elemp->except)
        printf(" Slot State ..... Abnormal\n");
    else
        printf(" Slot State ..... Normal\n");
    printf(" ASC/ASCQ ..... %02X%02X\n",
           elemp->asc,elemp->ascq);
    if (elemp->full)
        printf(" Media Present ..... Yes\n");
    else
        printf(" Media Present ..... No\n");
    if (elemp->access)
        printf(" Robot Access Allowed ..... Yes\n");
    else
        printf(" Robot Access Allowed ..... No\n");
    if (elemp->svalid)
        printf(" Source Element Address ..... %d\n",elemp->source);
    else
        printf(" Source Element Address Valid ... No\n");
    if (elemp->invert)
        printf(" Media Inverted ..... Yes\n");
    else
        printf(" Media Inverted ..... No\n");
    printf(" Volume Tag ..... %0.36s\n", elemp->volume);
    printf(" Cartridge Location ..... %0.24s\n", elemp->identifier);

}
free(data);
return 0;
```

## Return codes

This chapter describes the return codes that the device driver generates when an error occurs during an operation. The standard *errno* values are in the AIX **/usr/include/sys/errno.h** header file.

If the return code is input/output error (EIO), the application can issue the **STIOCQRYSENSE** IOCTL command with the LASTERROR option. Or, it can issue the **SIOC\_REQSENSE** IOCTL command to analyze the sense data and determine why the error occurred.

## Codes for all operations

The following codes and their descriptions apply to all operations.

### [EACCES]

Data encryption access denied.

### [EBADF]

A bad file descriptor was passed to the device.

### [EBUSY]

An excessive busy state was encountered in the device.

### [EFAULT]

A memory failure occurred due to an invalid pointer or address.

### [EMEDIA]

An unrecoverable media error was detected in the device.

### [ENOMEM]

Insufficient memory was available for an internal memory operation.

### [ENOTREADY]

The device was not ready for operation, or a tape was not in the drive.

### [ENXIO]

The device was not configured and is not receiving requests.

### [EPERM]

The process does not have permission to complete the wanted function.

### [ETIMEDOUT]

A command that is timed out in the device.

### [ENOCONNECT]

The device did not respond to selection.

### [ECONNREFUSED]

The device driver detected that the device vital product data (VPD) changed. The device must be unconfigured in AIX and reconfigured to correct the condition.

## Open error codes

The following codes and their descriptions apply to **open** operations.

### [EAGAIN]

The device was opened before the **open** operation.

### [EBADF]

A **write** operation was attempted on a device that was opened with the **O\_RDONLY** flag.

### [EBUSY]

The device was reserved by another initiator, or an excessive busy state was encountered.

### [EINVAL]

The operation that is requested has invalid parameters or an invalid combination of parameters, or the device is rejecting open commands.

### [ENOTREADY]

If the device was not opened with the **O\_NONBLOCK** or **O\_NDELAY** flag, then the drive is not ready for operation, or a tape is not in the drive. If a nonblocking flag was used, then the drive is not ready for operation.

### [EWRPROTECT]

An **open** operation with the **O\_RDWR** or **O\_WRONLY** flag was attempted on a write-protected tape.

### [EIO]

An I/O error occurred that indicates a failure to operate the device. Perform the failure analysis.

## [EINPROGRESS]

This *errno* is returned when the extended open flag **SC\_KILL\_OPEN** is used to kill all processes that currently have the device opened.

## Write error codes

The following codes and their descriptions apply to **write** operations.

### [EINVAL]

The operation that is requested has invalid parameters or an invalid combination of parameters.

The number of bytes requested in the **write** operation was not a multiple of the block size for a fixed block transfer.

The number of bytes requested in the **write** operation was greater than the maximum block size allowed by the device for variable block transfers.

### [ENOSPC]

A **write** operation failed because it reached the early warning mark or the programmable early warning zone (PEWZ) while it was in label-processing mode. This return code is returned only once when the early warning or the programmable early warning zone (PEWZ) is reached.

### [ENXIO]

A **write** operation was attempted after the device reached the logical end of the medium.

### [EWRPROTECT]

A **write** operation was attempted on a write-protected tape.

### [EIO]

The physical end of the medium was detected, or a general error occurred that indicates a failure to write to the device. Perform the failure analysis.

## Read error codes

The following codes and their descriptions apply to **read** operations.

### [EBADF]

A **read** operation was attempted on a device opened with the **O\_WRONLY** flag.

### [EINVAL]

The operation that is requested has invalid parameters or an invalid combination of parameters.

The number of bytes requested in the **read** operation was not a multiple of the block size for a fixed block transfer.

The number of bytes requested in the **read** operation was greater than the maximum size allowed by the device for variable block transfers.

### [ENOMEM]

The number of bytes requested in the **read** operation of a variable block record was less than the size of the block. This error is known as an overlength condition.

## Close error codes

The following codes and their descriptions apply to **close** operations.

### [EIO]

An I/O error occurred during the operation. Perform the failure analysis.

### [ENOTREADY]

A command that is issued during **close**, such as a **rewind** command, failed because the device was not ready.

## IOCTL error codes

The following codes and their descriptions apply to IOCTL operations.

### [EINVAL]

The operation that is requested has invalid parameters or an invalid combination of parameters.

This error code also results if the IOCTL is not supported for the device.

### [EWRPROTECT]

An operation that modifies the media was attempted on a write-protected tape or a device opened with the **O\_RDONLY** flag.

### [EIO]

An I/O error occurred during the operation. Perform the failure analysis.

### [ECANCELLED]

The **STIOCTOP** IOCTL with the **st\_op** field that specifies **STERASE\_IMM** was canceled by another process that issued the **STIOC\_CANCEL\_ERASE** IOCTL.

## Linux tape and medium changer device driver

---

IBM supplies a tape drive and medium changer open source device driver for the Linux platform called **lin\_tape**.

## Software interface

### Entry points

Lin\_tape supports the following Linux-defined entry points.

- “[open](#)” on page 87
- “[close](#)” on page 87
- “[read](#)” on page 88
- “[write](#)” on page 88
- “[ioctl](#)” on page 89

#### open

This entry point is driven by the **open** system call.

The programmer can access Lin\_tape devices with one of 3 access modes: **write only**, **read only**, or **read and write**.

Lin\_tape also support the **append open** flag. When the **open** function is called with the **append** flag set to TRUE, Lin\_tape attempts to **rewind and seek** two consecutive filemarks and place the initial tape position between them. **Open append** fails [*errno*: EIO] if no tape is loaded or two consecutive filemarks are not on the loaded tape. **Open append** does not automatically imply write access. Therefore, an access mode must accompany the **append** flag during the **open** operation.

The **open** function issues a SCSI **reserve** command to the target device. If the **reserve** command fails, **open** fails and *errno* EBUSY is returned.

#### close

This entry point is driven explicitly by the **close** system call and implicitly by the operating system at application program termination.

For non-rewinding special files, such as **/dev/IBMtapeOn**, if the last command before the **close** function was a successful **write**, Lin\_tape writes two consecutive filemarks that marks the end of data. It then sets

the tape position between the two consecutive filemarks. If the last command before the **close** function successfully wrote one filemark, then one extra filemark is written that marks the end of data. Then, the tape position is set between the two consecutive filemarks.

For non-rewinding special files, if the last tape command before the close function is **write**, but the write fails with sense key 6 (Unit Attention) and ASC/ASCQ 29/00 (Power On, Reset, or Bus Device Reset Occurred) or sense key 6 and ASC/ASCQ 28/00 (Not Ready to Ready Transition, Medium May Have Changed), Lin\_tape does not write two consecutive tape file marks that mark the end of data during **close** processing. If the last tape command before the close function is **write one file mark** and that command fails with one of the above two errors, Lin\_tape does not write one extra file mark that marks the end of data during close processing.

For rewind devices, such as **/dev/IBMtape0**, if the last command before the **close** function was a successful **write**, Lin\_tape writes two consecutive filemarks that mark the end of data and issues a **rewind** command. If the last command before the **close** function successfully wrote one filemark, one extra filemark is written marking the end of data, and the **rewind** command is issued. If the **write filemark** command fails, no **rewind** command is issued.

The application writers must be aware that a Unit Attention sense data that is presented means that the tape medium might be in an indeterminate condition, and no assumptions can be made about current tape positioning or whether the medium that was previously in the drive is still in the drive. IBM suggests that after a Unit Attention is presented, the tape special file be closed and reopened, label processing/verification be run (to determine that the correct medium is mounted), and explicit commands be run to locate to the wanted location. Extra processing might also be needed for particular applications.

If an **SIOC\_RESERVE ioctl** was issued from an application before **close**, the close function does not release the device; otherwise, it issues the SCSI release command. In both situations, the **close** function attempts to deallocate all resources that are allocated for the device. If, for some reason, Lin\_tape is not able to **close**, an error code is returned.

**Note:** The return code for **close** must always be checked. If **close** is unsuccessful, retry is recommended.

## **read**

This entry point is driven by the **read** system call. The **read** operation can be completed when a tape is loaded in the device.

Lin\_tape supports two modes of **read** operation. If the **read\_past\_filemark** flag is set to TRUE (with the **STIOCSETP** input/output control [IOCTL]), then when a **read** operation encounters a filemark, it returns the number of bytes read before it encounters the filemark and sets the tape position after the filemark. If the **read\_past\_filemark** flag is set to FALSE (by default or with **STIOCSETP** IOCTL), then when a read operation encounters a filemark, if data was read, the **read** function returns the number of bytes read, and positions the tape before the filemark. If no data was read, then **read** returns 0 bytes read and positions the tape after the filemark.

If the **read** function reaches end of the data on the tape, input/output error (EIO) is returned and ASC, ASCQ keys (obtained by request sense **IOCTLs**) indicate the end of data. Lin\_tape also conforms to all SCSI standard **read** operation rules, such as fixed block versus variable block.

## **write**

This entry point is driven by the **write** system call. The **write** operation can be completed when a tape is loaded in the device.

Lin\_tape supports early warning processing. When the **trailer\_labels** flag is set to TRUE (by default or with **STIOCSETP IOCTL** call), Lin\_tape fails with *errno* ENOSPACE only when a **write** operation first encounters the early warning zone for end of tape. After the ENOSPACE error code is returned, Lin\_tape suppresses all warning messages from the device that is generated by subsequent write commands, effectively allowing **write** and **write filemark** commands in the early warning zone. When physical end of tape is reached, error code EIO is returned, and the ASC and ASCQ keys (obtained by the request sense **IOCTL**) confirm the end of physical medium condition. When the **trailer\_labels** flag is set to FALSE (with

**STIOCSETP** IOCTL call), Lin\_tape returns the ENOSPACE *errno* when any **write** command is attempted in the early warning zone.

## ioctl

This entry point provides a set of drive SCSI-specific functions. It allows Linux applications to access and control the features and attributes of the drive device programmatically.

## Medium changer devices

Lin\_tape supports the following Linux entry points for the medium changer devices.

- “[open](#)” on page 89
- “[close](#)” on page 89
- “[ioctl](#)” on page 89

### open

This entry point is driven by the **open** system call. The **open** function attempts a SCSI **reserve** command to the target device. If the **reserve** command fails, **open** fails with *errno* EBUSY.

### close

This entry point is driven explicitly by the **close** system call and implicitly by the operating system at program termination. If an **SIOC\_RESERVE** IOCTL was issued from an application before **close**, the close function does not release the device. Otherwise, it issues the SCSI **release** command. In both situations, the close function attempts to deallocate all resources that are allocated for the device. If, for some reason, Lin\_tape is not able to close, an error code is returned.

## ioctl

This entry point provides a set of medium changer and SCSI-specific functions. It allows Linux applications to access and control the features and attributes of the robotic device programmatically.

## General IOCTL operations

This chapter describes the IOCTL commands that provide access and control to the tape and medium changer devices.

These commands are available for all tape and medium changer devices. They can be issued to any one of the **Lin\_tape** special files.

## Overview

The following IOCTL commands are supported.

### **SIOC\_INQUIRY**

Return the inquiry data.

### **SIOC\_REQSENSE**

Return the sense data.

### **SIOC\_RESERVE**

Reserve the device.

### **SIOC\_RELEASE**

Release the device.

### **SIOC\_TEST\_UNIT\_READY**

Issue the **SCSI Test Unit Ready** command.

**SIOC\_LOG\_SENSE\_PAGE**

Return the log sense data.

**SIOC\_LOG\_SENSE10\_PAGE**

Return the log sense data by using a 10-byte CDB with optional subpage.

**SIOC\_ENH\_LOG\_SENSE**

Return the page data with a requested length from the application if no kernel memory restriction exists.

**SIOC\_MODE\_SENSE\_PAGE**

Return the mode sense data.

**SIOC\_MODE\_SENSE**

Return the mode sense data with optional subpage.

**SIOC\_INQUIRY\_PAGE**

Return the inquiry data for a specific page.

**SIOC\_PASS\_THROUGH**

Pass through custom built SCSI commands.

**SIOC\_QUERY\_PATH**

Return the primary path and information for the first alternate path.

**SIOC\_DEVICE\_PATHS**

Return the primary path and information for all the alternate paths.

**SIOC\_ENABLE\_PATH**

Enable a path from the disabled state.

**SIOC\_DISABLE\_PATH**

Disable a path.

These IOCTL commands and their associated structures are defined in the **IBM\_tape.h** header file, which can be found in **/usr/include/sys** after Lin\_tape is installed. The **IBM\_tape.h** header file must be included in the corresponding C programs that call functions that are provided by Lin\_tape.

All IOCTL commands require a file descriptor of an open file. Use the **open** command to open a device and obtain a valid file descriptor.

**SIOC\_INQUIRY**

This IOCTL command collects the inquiry data from the device.

The data structure is

```
struct inquiry_data {
    uint qual          :3;      /* peripheral qualifier */ */
    uint type          :5;      /* device type */ */
    uint rm            :1;      /* removable medium */ */
    uint mod           :7;      /* device type modifier */ */
    uint iso           :2;      /* ISO version */ */
    uint ecma          :3;      /* EMCA version */ */
    uint ansi          :3;      /* ANSI version */ */
    uint aenc          :1;      /* asynchronous event notification */ */
    uint trmiop        :1;      /* terminate I/O process message */ */
    uint rdf           :2;      /* reserved */ */
    uchar len;          /* response data format */ */
    uchar resvd1;       /* additional length */ */
    uchar resvd2;       /* reserved */ */
    uint             :4;      /* reserved */ */
    uchar mchngr       :1;      /* medium changer mode (SCSI-3 only) */ */
    uchar             :3;      /* reserved */ */
    uint reladr        :1;      /* relative addressing */ */
    wbus32            :1;      /* 32-bit wide data transfers */ */
    wbus16            :1;      /* 16-bit wide data transfers */ */
    sync              :1;      /* synchronous data transfers */ */
    linked            :1;      /* linked commands */ */
    cmdque           :1;      /* reserved */ */
    sftrre            :1;      /* command queueing */ */
    uchar vid[8];       /* soft reset */ */
    uchar pid[16];      /* vendor ID */ */
    uchar             :1;      /* product ID */ */
};
```

```

    uchar revision[4];           /* product revision level      */
    uchar vendor1[20];          /* vendor specific             */
    uchar resvd2[40];           /* reserve                     */
    uchar vendor2[31];          /* vendor specific (padded to 127) */
};

}

```

An example of the **SIOC\_INQUIRY** command is

```

#include <sys/IBM_tape.h>
char vid[9];
char pid[17];
char revision[5];
struct inquiry_data inqdata;
printf("Issuing inquiry...\n");
memset(&inqdata, 0, sizeof(struct inquiry_data));
if (!ioctl(fd, SIOC_INQUIRY, &inqdata)) {
    printf ("The SIOC_INQUIRY ioctl succeeded\n");
    printf ("\nThe inquiry data is:\n");
    /*
     * Just a dump byte won't work because of the compiler
     * bit field mapping
    */
    /* print out structure data field */
    printf("\nInquiry Data:\n");
    printf("Peripheral Qualifer-----0x%02x\n", inqdata.qual);
    printf("Peripheral Device Type-----0x%02x\n", inqdata.type);
    printf("Removal Medium Bit-----%d\n", inqdata.rm);
    printf("Device Type Modifier-----0x%02x\n", inqdata.mod);
    printf("ISO version-----0x%02x\n", inqdata.iso);
    printf("ECMA version-----0x%02x\n", inqdata.ecma);
    printf("ANSI version-----0x%02x\n", inqdata.ansi);
    printf("Asynchronous Event Notification Bit-%d\n", inqdata.aenc);
    printf("Terminate I/O Process Message Bit--%d\n", inqdata.trmio);
    printf("Response Data Format-----0x%02x\n", inqdata.rdf);
    printf("Additional Length-----0x%02x\n", inqdata.len);
    printf("Medium Changer Mode-----0x%02x\n", inqdata.mchngr);
    printf("Relative Addressing Bit-----%d\n", inqdata.reladr);
    printf("32 Bit Wide Data Transfers Bit-----%d\n", inqdata.wbus32);
    printf("16 Bit Wide Data Transfers Bit-----%d\n", inqdata.wbus16);
    printf("Synchronous Data Transfers Bit-----%d\n", inqdata.sync);
    printf("Linked Commands Bit-----%d\n", inqdata.linked);
    printf("Command Queueing Bit-----%d\n", inqdata.cmdque);
    printf("Soft Reset Bit-----%d\n", inqdata.sftre);

    strncpy(vid, inqdata.vid, 8);
    vid[8] = '\0';
    strncpy(pid, inqdata.pid, 16);
    pid[16] = '\0';
    strncpy(revision, inqdata.revision, 4);
    revision[4] = '\0';

    printf("Vendor ID-----%s\n", vid);
    printf("Product ID-----%s\n", pid);
    printf("Product Revision Level-----%s\n", revision);

    dump_bytes(inqdata.vendor1, 20, "vendor1");
    dump_bytes(inqdata.vendor2, 31, "vendor2");
}
else {
    perror ("The SIOC_INQUIRY ioctl failed");
    sioc_request_sense();
}

```

## **SIOC\_REQSENSE**

This IOCTL command returns the device sense data. If the last command resulted in an error, then the sense data is returned for the error. Otherwise, a new sense command is issued to the device.

The data structure is

```

struct request_sense {
    uint valid          :1,      /* sense data is valid      */
    err_code           :7;      /* error code               */
    unchar segnum;        /* segment number           */
    uint fm            :1,      /* filemark detected       */
    eom               :1,      /* end of medium           */
    ili                :1,      /* incorrect length indicator */
    resvd1            :1,      /* reserved                */
};

```

```

    key          :4;      /* sense key                      */
    int   info;           /* information bytes              */
    uchar addlen;         /* additional sense length        */
    uint  cmdinfo;        /* command specific information */
    uchar asc;            /* additional sense code          */
    uchar ascq;           /* additional sense code qualifier */
    uchar fru;            /* field replaceable unit code   */
    uint   sksv     :1,    /* sense key specific valid     */
    cd      :1,    /* control/data                  */
    resvd2  :2,    /* reserved                     */
    bpv     :1,    /* bit pointer valid             */
    sim     :3;    /* system information message    */
    uchar field[2];       /* field pointer                 */
    uchar vendor[109];    /* vendor specific (padded to 127) */
};


```

An example of the **SIOC\_REQSENSE** command is

```

#include <sys/IBM_tape.h>

struct request_sense sense_data;
int rc;
printf("Issuing request sense...\n");
memset(&sense_data, 0, sizeof(struct request_sense));
rc = ioctl(fd, SIOC_REQSENSE, &sense_data);
if (rc == 0)
{
    if(!sense_data.err_code)
        printf("No valid sense data returned.\n");
    else
    {
        /* print out data fields */
        printf("Information Field Valid Bit-----%d\n", sense_data.valid);
        printf("Error Code-----0x%02x\n", sense_data.err_code);
        printf("Segment Number-----0x%02x\n", sense_data.segnum);
        printf("filenamek Detected Bit-----%d\n", sense_data.fm);
        printf("End Of Medium Bit-----%d\n", sense_data.eom);
        printf("Illegal Length Indicator Bit---%d\n", sense_data.il);
        printf("Sense Key-----0x%02x\n", sense_data.key);
        if(sense_data.valid)
            printf("Information Bytes-----0x%02x 0x%02x 0x%02x 0x%02x\n",
                   sense_data.info >> 24, sense_data.info >> 16,
                   sense_data.info >> 8, sense_data.info & 0xFF);
        printf("Additional Sense Length-----0x%02x\n", sense_data.addlen);
        printf("Command Specific Information---0x%02x 0x%02x 0x%02x\n",
               sense_data.cmdinfo >> 24, sense_data.cmdinfo >> 16,
               sense_data.cmdinfo >> 8, sense_data.cmdinfo & 0xFF);
        printf("Additional Sense Code-----0x%02x\n", sense_data.asc);
        printf("Additional Sense Code Qualifier-0x%02x\n", sense_data.ascq);
        printf("Field Replaceable Unit Code----0x%02x\n", sense_data.fru);
        printf("Sense Key Specific Valid Bit---%d\n", sense_data.sksv);
        if(sense_data.sksv)
        {
            printf("Command Data Block Bit--%d\n", sense_data.cd);
            printf("Bit Pointer Valid Bit---%d\n", sense_data.bpv);
            if(sense_data.bpv)
                printf("System Information Message-0x%02x\n", sense_data.sim);
            printf("Field Pointer-----0x%02x%02x\n",
                   sense_data.field[0], sense_data.field[1]);
        }
        dump_bytes(sense_data.vendor, 109, "Vendor");
    }
}
return rc;
}


```

## **SIOC\_RESERVED**

This IOCTL command explicitly reserves the device and prevents it from being released after a close operation.

The device is not released until an **SIOC\_RELEASE** IOCTL command is issued.

The IOCTL command can be used for applications that require multiple open and close processing in a host-sharing environment.

There are no arguments for this IOCTL command.

An example of the **SIOC\_RESERVED** command is

```
#include <sys/IBM_tape.h>
if (!ioctl (fd, SIOC_RESERVED, NULL)) {
    printf ("The SIOC_RESERVED ioctl succeeded\n");
}
else {
    perror ("The SIOC_RESERVED ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_RELEASE**

This IOCTL command explicitly releases the device and allows other hosts to access it. The IOCTL command is used with the **SIOC\_RESERVED** IOCTL command for applications that require multiple open and close processing in a host-sharing environment.

There are no arguments for this IOCTL command.

An example of the **SIOC\_RELEASE** command is

```
#include <sys/IBM_tape.h>
if (!ioctl (fd, SIOC_RELEASE, NULL)) {
    printf ("The SIOC_RELEASE ioctl succeeded\n");
}
else {
    perror ("The SIOC_RELEASE ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_TEST\_UNIT\_READY**

This IOCTL command issues the **SCSI Test Unit Ready** command to the device.

There are no arguments for this IOCTL command.

An example of the **SIOC\_TEST\_UNIT\_READY** command is

```
#include <sys/IBM_tape.h>
if (!ioctl (fd, SIOC_TEST_UNIT_READY, NULL)) {
    printf ("The SIOC_TEST_UNIT_READY ioctl succeeded\n");
}
else {
    perror ("The SIOC_TEST_UNIT_READY ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_LOG\_SENSE\_PAGE, SIOC\_LOG\_SENSE10\_PAGE, and SIOC\_ENH\_LOG\_SENSE**

These IOCTL commands return log sense data from the device. The differences between the three is

- **SIOC\_LOG\_SENSE\_PAGE** allows the user to retrieve a particular log page up to length LOGSENSEPAGE.
- **SIOC\_LOG\_SENSE10\_PAGE** allows for a subpage to be returned up to length LOGSENSEPAGE.
- **SIOC\_ENH\_LOG\_SENSE** returns the page data with a requested length from application if no kernel memory restriction exists.

For both **SIOC\_LOG\_SENSE\_PAGE** and **SIOC\_LOG\_SENSE10\_PAGE**, to obtain the entire log page, the **len** and **parm\_pointer** fields must be set to zero. To obtain the entire log page that starts at a specific parameter code, set the **parm\_pointer** field to the wanted code and the **len** field to zero. To obtain a specific number of parameter bytes, set the **parm\_pointer** field to the wanted code. Then, set the **len** field to the number of parameter bytes plus the size of the log page header (4 bytes). The first 4 bytes of returned data are always the log page header. In the Enhanced log sense page (**SIOC\_ENH\_LOG\_SENSE**), the length cannot be set to zero as it indicates the allocated memory size that **char \*logdatap** is pointing to. The minimum number for this value is 4, as it returns the first 4 bytes that is the log page header. See the appropriate device manual to determine the supported log pages and content.

The data structures are

```

struct log_sense_page {
    uchar page_code;
    unsigned short len;
    unsigned short parm_pointer;
    char data[LOGSENSEPAGE];
};

struct log_sense10_page {
    uchar page_code;
    uchar subpage_code;
    uchar reserved[2];
    unsigned short len;
    unsigned short parm_pointer;
    char data[LOGSENSEPAGE];
};

struct enh_log_sense {
    uchar page_code; /* [IN] Log sense page */
    uchar subpage_code; /* [IN] Log sense sub-page */
    uchar page_control; /* [IN] Page control */
    uchar reserved[5];
    unsigned short len; /* [IN] specific allocation length for logdatap */
    /* by application */
    /* [OUT] the length of return data at */
    /* logdatap from driver */
    unsigned short parm_pointer; /* [IN] specific parameter number at */
    /* which the data begins */
    char *logdatap; /* [IN] the pointer for log sense data allocated*/
    /* by application */
    /* [OUT] log sense data returned from driver */
};

```

The first two IOCTLs are identical, except if a specific subpage is wanted, `log_sense10_page` must be used and `subpage_code` must be assigned by the user application.

An example of the **SIOC\_LOG\_SENSE\_PAGE** command is

```

#include <sys/IBM_tape.h>
struct log_sense_page log_page;
int temp;
/* get log page 0, list of log pages */
log_page.page_code = 0x00;
log_page.len = 0;
log_page.parm_pointer = 0;
if (!ioctl (fd, SIOC_LOG_SENSE_PAGE, &log_page)) {
    printf ("The SIOC_LOG_SENSE_PAGE ioctl succeeded\n");
    dump_bytes(log_page.data, LOGSENSEPAGE);
}
else {
    perror ("The SIOC_LOG_SENSE_PAGE ioctl failed");
    sioc_request_sense();
}
/* get fraction of volume traversed */
log_page.page_code = 0x38;
log_page.len = 0;
log_page.parm_pointer = 0x000F;
if (!ioctl (fd, SIOC_LOG_SENSE_PAGE, &log_page)) {
    temp = log_page.data[sizeof(log_page_header) + 4];
    printf ("The SIOC_LOG_SENSE_PAGE ioctl succeeded\n");
    printf ("Fractional Part of Volume Traversed %x\n",temp);
}
else {
    perror ("The SIOC_LOG_SENSE_PAGE ioctl failed");
    sioc_request_sense();
}

```

An example of the **SIOC\_ENH\_LOG\_SENSE** command is

```

#include <sys/IBM_tape.h>
#define LOG_PAGE_HEADER 4
struct enh_log_sense enh_log_page;
unsigned short length;
memset((char*)&enh_log_page, 0, sizeof(struct enh_log_sense));

```

```

enh_log_page.page_code = 0x17;
enh_log_page.subpage_code = 0x02;
enh_log_page.len = LOG_PAGE_HEADER;
enh_log_page.logdatap = malloc(LOG_PAGE_HEADER);

if(enh_log_page.logdatap == NULL){
    printf ("Unable to malloc LOG_PAGE_HEADER. Closing\n");
    exit(-1);
}

if (!ioctl (fd, SIOC_ENH_LOG_SENSE, &enh_log_page)) {
    printf ("The SIOC_ENH_LOG_SENSE ioctl succeeded\n");
    sprintf(text,"Log enhanced page header 0x%02X subpage 0x%02X",
           enh_log_page.page_code, enh_log_page.subpage_code);
    dump_bytes(enh_log_page.logdatap, enh_log_page.len, text);
}
else {
    perror ("The SIOC_ENH_LOG_SENSE ioctl failed");
}

length = (enh_log_page.logdatap[2] << 8) +enh_log_page.logdatap[3];
free(enh_log_page.logdatap);
enh_log_page.logdatap = NULL;

enh_log_page.len = length;
enh_log_page.logdatap = malloc(length);

if(enh_log_page.logdatap==NULL) {
    printf("Unable to malloc enh_log_page big size %d\n", length);
    if(length > 1024) {
        enh_log_page.logdatap = malloc(1024);
        enh_log_page.len = 1024;
        if(enh_log_page.logdatap == NULL) {
            printf("Unable to malloc enh_log_page 1024 size\n");
            exit(-1);
        }
    }
}

if (!ioctl (fd, SIOC_ENH_LOG_SENSE, &enh_log_page)) {
    printf ("The SIOC_ENH_LOG_SENSE ioctl succeeded\n");
    sprintf(text,"Enhanced Log Sense: page 0x%02X subpage 0x%02X length %d",
           enh_log_page.page_code, enh_log_page.subpage_code);

    dump_bytes(enh_log_page.logdatap, enh_log_page.len, text);
}
else {
    perror ("The SIOC_ENH_LOG_SENSE ioctl failed");
}
free(enh_log_page.logdatap);

```

## ***SIOC\_MODE\_SENSE\_PAGE and SIOC\_MODE\_SENSE***

This IOCTL command returns a mode sense page from the device. The desired page is selected by specifying the **page\_code** in the **mode\_sense\_page** structure. See the appropriate device manual to determine the supported mode pages and content.

The data structures are

```

struct mode_sense_page {
    uchar page_code;
    char data[MAX_MDSNS_LEN];
};

struct mode_sense {
    uchar page_code;
    uchar subpage_code;
    uchar reserved[6];
    uchar cmd_code;
    char data[MAX_MDSNS_LEN];
};

```

The IOCTLs are identical, except that if a specific subpage is desired, **mode\_sense** must be used and **subpage\_code** must be assigned by the user application. Under the current implementation, **cmd\_code** is not assigned by the user and must be left with a value 0.

An example of the **SIOC\_MODE\_SENSE\_PAGE** command is

```
#include <sys/IBM_tape.h>
struct mode_sense_page mode_page;
/* get medium changer mode */
mode_page.page_code = 0x20;
if (!ioctl (fd, SIOC_MODE_SENSE_PAGE, &mode_page)) {
    printf ("The SIOC_MODE_SENSE_PAGE ioctl succeeded\n");
    if (mode_page.data[2] == 0x02)
        printf ("The library is in Random mode.\n");
    else if (mode_page.data[2] == 0x05)
        printf ("The library is in Automatic (Sequential) mode.\n");
}
else {
    perror ("The SIOC_MODE_SENSE_PAGE ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_INQUIRY\_PAGE**

This IOCTL command returns an inquiry page from the device. The desired page is selected by specifying the **page\_code** in the **inquiry\_page** structure. See the appropriate device manual to determine the supported inquiry pages and content.

The data structure is

```
struct inquiry_page {
    char page_code;
    char data[INQUIRYPAGE];
};
```

An example of the **SIOC\_INQUIRY\_PAGE** command is

```
#include <sys/IBM_tape.h>
struct inquiry_page inq_page;
/* get inquiry page x83 */
inq_page.page_code = 0x83;
if (!ioctl (fd, SIOC_INQUIRY_PAGE, &inq_page)) {
    printf ("The SIOC_INQUIRY_PAGE ioctl succeeded\n");
    dump_bytes(inq_page.data, INQUIRYPAGE);
}
else {
    perror ("The SIOC_INQUIRY_PAGE ioctl failed");
    sioc_request_sense();
}
```

## **SCSI\_PASS\_THROUGH**

This IOCTL command passes the built command data block structure with I/O buffer pointers to the lower SCSI layer. Status is returned from the lower SCSI layer to the caller with the **ASC and ASCQ values** and **SenseKey** fields. The ASC and ASCQ and sense key fields are valid only when the **SenseDataValid** field is true.

The data structure is

```
#define SCSI_PASS THROUGH _IOWR('P',0x01,SCSIPassThrough) /* Pass Through */
typedef struct _SCSIPassThrough
{
    uchar   CDB[12];          /* Command Data Block */
    uchar   CommandLength;    /* Command Length */
    uchar * Buffer ;          /* Command Buffer */
    ulong   BufferLength;     /* Buffer Length */
    uchar   DataDirection;    /* Data Transfer Direction */
    ushort  TimeOut;          /* Time Out Value */
    uchar   TargetStatus;     /* Target Status */
    uchar   MessageStatus;    /* Message from host adapter */
    uchar   HostStatus;        /* Host status */
    uchar   DriverStatus;      /* Driver status */
    uchar   SenseDataValid;    /* Sense Data Valid */
    uchar   ASC;               /* ASC key if the SenseDataValid is True */
    uchar   ASCQ;              /* ASCQ key if the SenseDataValid is True */
    uchar   SenseKey;          /* Sense key if the SenseDataValid is True */
} SCSIPassThrough, *PSCSIPassThrough;
#define SCSI_DATA_OUT 1
```

```
#define SCSI_DATA_IN 2
#define SCSI_DATA_NONE 3
```

SCSI\_DATA\_OUT indicates sending data out of the initiator (host bus adapter), also known as write mode.  
 SCSI\_DATA\_IN indicates receiving data into the initiator (host bus adapter), also known as read mode.  
 SCSI\_DATA\_NONE indicates that no data is transferred.

An example of the **SCSI\_PASS\_THROUGH** command is

```
#include <sys/IBM_tape.h>
SCSIPassThrough PassThrough;
memset(&PassThrough, 0, sizeof(SCSIPassThrough));
/* Issue test unit ready command */
PassThrough.CDB[0] = 0x00;
PassThrough.CommandLength = 6;
PassThrough.DataDirection = SCSI_DATA_NONE;
if (!ioctl(fd, SCSI_PASS THROUGH, &PassThrough)) {
    printf ("The SCSI_PASS THROUGH ioctl succeeded\n");
    if((PassThrough.TargetStatus == STATUS_SUCCESS) &&
       (PassThrough.MessageStatus == STATUS_SUCCESS) &&
       (PassThrough.HostStatus == STATUS_SUCCESS) &&
       (PassThrough.DriverStatus == STATUS_SUCCESS))
        printf(" Test Unit Ready returns success\n");

    else {
        printf(" Test Unit Ready failed\n");
        if(PassThrough.SenseDataValid)
            printf("Sense Key %02x, ASC %02x, ASCQ %02x\n",
                   PassThrough.SenseKey, PassThrough.ASC,
                   PassThrough.ASCQ);
    }
}
else {
    perror ("The SIOC SCSI_PASS THROUGH ioctl failed");
    sioc_request_sense();
}
```

## **SIOC\_QUERY\_PATH**

This IOCTL command returns the primary path and the first alternate path information for a physical device.

The data structure is

```
struct scsi_path
{
    char primary_name[30]; /* primary logical device name */
    char primary_parent[30]; /* primary SCSI parent name, "Host" name */
    uchar primary_id; /* primary target address of device, "Id" value*/
    uchar primary_lun; /* primary logical unit of device, "lun" value */
    uchar primary_bus; /* primary SCSI bus for device, "Channel" value*/
    unsigned long long primary_fcp_scsi_id; /* not supported */
    unsigned long long primary_fcp_lun_id; /* not supported */
    unsigned long long primary_fcp_ww_name; /* not supported */
    uchar primary_enabled; /* primary path enabled */
    uchar primary_id_valid; /* primary id/lun/bus fields valid */
    uchar primary_fcp_id_valid; /* not supported */
    uchar alternate_configured; /* alternate path configured */
    char alternate_name[30]; /* alternate logical device name */
    char alternate_parent[30]; /* alternate SCSI parent name */
    uchar alternate_id; /* alternate target address of device */
    uchar alternate_lun; /* alternate logical unit of device */
    uchar alternate_bus; /* alternate SCSI bus for device */
    unsigned long long alternate_fcp_scsi_id; /* not supported */
    unsigned long long alternate_fcp_lun_id; /* not supported */
    unsigned long long alternate_fcp_ww_name; /* not supported */
    uchar alternate_enabled; /* alternate path enabled */
    uchar alternate_id_valid; /* alternate id/lun/bus fields valid */
    uchar alternate_fcp_id_valid; /* not supported */
    uchar primary_drive_port_valid; /* not supported */
    uchar primary_drive_port; /* not supported */
    uchar alternate_drive_port_valid; /* not supported */
    uchar alternate_drive_port; /* not supported */
    uchar primary_fenced; /* primary fenced by disable path ioctl */
    uchar alternate_fenced; /* alternate fenced by disable path ioctl */
    uchar primary_host; /* primary host bus adapter id */
    uchar alternate_host; /* alternate host bus adapter id */
```

```

    char reserved[56];
};


```

An example of the **SIOC\_QUERY\_PATH** command is

```

#include <sys/IBM_tape.h>
struct scsi_path path;
memset(&path, 0, sizeof(struct scsi_path));
printf("Querying SCSI paths...\n");
rc = ioctl(fd, SIOC_QUERY_PATH, &path);
if(rc == 0)
    show_path(&path);


```

## **SIOC\_DEVICE\_PATHS**

This IOCTL command returns the primary path and all of the alternate paths information for a physical device. This IOCTL supports only the 3592 tape drives. The data structure for this IOCTL command is

```

struct device_path_t
{
    char name[30];                                /* logical device name */
    char parent[30];                               /* logical parent name */
    unchar id_valid;                             /* SCSI id/lun/bus fields valid */
    unchar id;                                  /* SCSI target address of device */
    unchar lun;                                 /* SCSI logical unit of device */
    unchar bus;                                /* SCSI bus for device */
    unchar fcp_id_valid;                         /* not supported */
    unsigned long long fcp_scsi_id;             /* not supported */
    unsigned long long fcp_lun_id;               /* not supported */
    unsigned long long fcp_ww_name;              /* not supported */
    unchar enabled;                            /* path enabled */
    unchar drive_port_valid;                    /* not supported */
    unchar drive_port;                          /* not supported */
    unchar fenced;                             /* path fenced by disable path ioctl */
    unchar host;                               /* host bus adapter id */
    char reserved[62];
};

struct device_paths
{
    int number_paths;                           /* number of paths configured */
    struct device_path_t path[MAX_SCSI_PATH];
};


```

An example of this IOCTL command is

```

#include <sys/IBM_tape.h>
struct device_paths device_path;
memset(&device_path, 0, sizeof(struct device_paths));
printf("Querying device paths...\n");
rc = ioctl(fd, SIOC_DEVICE_PATHS, &device_path);
if(rc == 0)
{
    printf("\n");
    for (i=0; i < device_path.number_paths; i++)
    {
        if (i == 0)
            printf("Primary Path Number 1\n");
        else
            printf("Alternate Path Number %d\n", i+1);
        printf(" Logical Device..... %s\n", device_path.path[i].name);
        printf(" Host Bus Adapter..... %s\n", device_path.path[i].parent);

        if (device_path.path[i].id_valid)
        {
            printf(" SCSI Host ID..... %d\n", device_path.path[i].host);
            printf(" SCSI Channel..... %d\n", device_path.path[i].bus);
            printf(" Target ID..... %d\n", device_path.path[i].id);
            printf(" Logical Unit..... %d\n", device_path.path[i].lun);
        }

        if (device_path.path[i].enabled)
            printf(" Path Enabled..... Yes\n");
        else
            printf(" Path Enabled..... No \n");
        if (device_path.path[i].fenced)
            printf(" Path Manually Disabled..... Yes\n");
    }
}


```

```

        else
            printf("  Path Manually Disabled..... No \n");
    }
    printf("Total paths configured..... %d\n",device_path.number_paths);
}

```

### ***SIOC\_ENABLE\_PATH***

This IOCTL enables the path that is specified by the path number. This command supports only the 3592 tape drives.

An example of this IOCTL command is

```

#include <sys/IBM_tape.h>
if (path == PRIMARY_SCSI_PATH)
    printf("Enabling primary SCSI path 1...\n");
else
    printf("Enabling alternate SCSI path %d...\n",path);
rc = ioctl(fd, SIOC_ENABLE_PATH, path);

```

### ***SIOC\_DISABLE\_PATH***

This IOCTL disables the path that is specified by the path number. This command supports only the 3592 tape drives.

An example of this IOCTL command is

```

#include <sys/IBM_tape.h>
if (path == PRIMARY_SCSI_PATH)
    printf("Disabling primary SCSI path 1...\n");
else
    printf("Disabling alternate SCSI path %d...\n",path);
rc = ioctl(fd, SIOC_DISABLE_PATH, path);

```

## **Tape drive IOCTL operations**

The device driver supports the set of tape IOCTL commands that is available with the base Linux operating system. In addition, a set of expanded tape IOCTL commands gives applications access to extra features and functions of the tape drives.

### **Overview**

The following IOCTL commands are supported.

#### **STIOCTOP**

Run the basic tape operations.

#### **STIOCQRYP**

Query the tape device, device driver, and media parameters.

#### **STIOCSETP**

Change the tape device, device driver, and media parameters.

#### **STIOCSYNC**

Synchronize the tape buffers with the tape.

#### **STIOCDM**

Displays and manipulates one or two messages.

#### **STIOCQRYPOS**

Query the tape position and the buffered data.

#### **STIOCSETPOS**

Set the tape position.

**STIOCQRYSENSE**

Query the sense data from the tape device.

**STIOCQRYINQUIRY**

Return the inquiry data.

**STIOC\_LOCATE**

Locate to a certain tape position.

**STIOC\_READ\_POSITION**

Read the current tape position.

**STIOC\_RESET\_DRIVE**

Issue a **SCSI Send Diagnostic** command to reset the tape drive.

**STIOC\_PREVENT\_MEDIUM\_REMOVAL**

Prevent medium removal by an operator.

**STIOC\_ALLOW\_MEDIUM\_REMOVAL**

Allow medium removal by an operator.

**STIOC\_REPORT\_DENSITY\_SUPPORT**

Return supported densities from the tape device.

**MTDEVICE**

Returns the device number that is used for communicating with an Enterprise Tape Library 3494.

**STIOC\_GET\_DENSITY**

Query the current write density format settings on the tape drive. The current density code from the drive **Mode Sense** header, the **Read/Write Control Mode** page default density, and the pending density are returned.

**STIOC\_SET\_DENSITY**

Set a new write density format on the tape drive by using the default and pending density fields. The application can specify a new write density for the currently loaded tape only. Or, it can specify a new write density as a default for all tapes.

**GET\_ENCRYPTION\_STATE**

This IOCTL can be used for application, system, and library-managed encryption. It allows only a query of the encryption status.

**SET\_ENCRYPTION\_STATE**

This IOCTL can be used only for application-managed encryption. It sets the encryption state for application-managed encryption.

**SET\_DATA\_KEY**

This IOCTL can be used only for application-managed encryption. It sets the data key for application-managed encryption.

**STIOC\_QUERY\_PARTITION**

This IOCTL queries for partition information on applicable tapes. It displays maximum number of possible partitions, number of partitions currently on tape, the active partition, the size unit (bytes, kilobytes, and so on), and the sizes of each partition.

**STIOC\_CREATE\_PARTITION**

This IOCTL creates partitions on applicable tapes. The user is allowed to specify the number and type of partitions and the size of each partition.

**STIOC\_SET\_ACTIVE\_PARTITION**

This IOCTL allows the user to set the partition on which to complete tape operations.

**STIOC\_ALLOW\_DATA\_OVERWRITE**

This IOCTL allows tape data to be overwritten when in data safe mode.

**STIOC\_READ\_POSITION\_EX**

This IOCTL reads the tape position and includes support for the long and extended formats.

**STIOC\_LOCATE\_16**

This IOCTL sets the tape position by using a long tape format.

**STIOC\_QUERY\_BLK\_PROTECTION**

This IOCTL queries the current capability and status of Logical Block Protection in the drive.

**STIOC\_SET\_BLK\_PROTECTION**

This IOCTL sets the status of Logical Block Protection in the drive.

**STIOC\_VERIFY\_TAPE\_DATA**

This IOCTL instructs the tape drive to scan the data on its current tape to check for errors.

**STIOC\_QUERY\_RAO**

The IOCTL is used to query the maximum number and size of User Data Segments (UDS) that are supported from tape drive and driver for the wanted **uds\_type**.

**STIOC\_GENERATE\_RAO**

The IOCTL is called to send a **GRAO** list to request that the drive generate a **Recommended Access Order** list.

**STIOC\_RECEIVE\_RAO**

After a **STIOC\_GENERATE\_RAO** IOCTL is completed, the application calls the **STIOC\_RECEIVE\_RAO** IOCTL to receive a recommended access order of UDS from the drive.

**STIOC\_SET\_SPDEV**

This IOCTL is for usage through IBMSpecial open handle only. It sets the drive that processes the command requests, and to do so it needs the serial number of the drive as input.

These IOCTL commands and their associated structures are defined in the **IBM\_tape.h** header file that can be found in the **lin\_tape** source rpm package. This header must be included in the corresponding C program by using the IOCTL commands.

**STIOCTOP**

This IOCTL command runs basic tape operations. The **st\_count** variable is used for many of its operations. Normal error recovery applies to these operations. The device driver can issue several tries to complete them. For all forward movement space operations, the tape position finishes on the end-of-tape side of the record or filemark, and on the beginning-of-tape side of the record or filemark for backward movement.

The input data structure is

```
struct stop {
    short st_op;      /* operations defined below */
    daddr_t st_count; /* how many of them to do (if applicable) */
};
```

The **st\_op** variable is set to one of the following operations.

**STOFFL**

Unload the tape. The **st\_count** parameter does not apply.

**STREW**

Rewind the tape. The **st\_count** parameter does not apply.

**STERASE**

Erase the entire tape. The **st\_count** parameter does not apply.

**STRETEN**

Run the rewind operation. The tape devices run the retension operation automatically when needed.

**STWEOF**

Write **st\_count** number of filemarks.

**STFSF**

Space forward the **st\_count** number of filemarks.

**STRSF**

Space backward the **st\_count** number of filemarks.

**STFSR**

Space forward the **st\_count** number of records.

**STRSR**

Space backward the **st\_count** number of records.

**STTUR**

Issue the **Test Unit Ready** command. The **st\_count** parameter does not apply.

**STLOAD**

Issue the **SCSI Load** command. The **st\_count** parameter does not apply. The operation of the **SCSI Load** command varies depending on the type of device. See the appropriate hardware reference manual.

**STSEOD**

Space forward to the end of the data. The **st\_count** parameter does not apply.

**STEJECT**

Unload the tape. The **st\_count** parameter does not apply.

**STINSRT**

Issue the **SCSI Load** command. The **st\_count** parameter does not apply. The operation of the **SCSI Load** command varies depending on the type of device. See the appropriate hardware reference manual.

**Note:** If zero is used for operations that require the **st\_count** parameter, then the command is not issued to the device. The device driver returns a successful completion.

An example of the **STIOCTOP** command is

```
#include <sys/IBM_tape.h>

struct stop stop;
stop.st_op=STWEOF;
stop.st_count=3;
if (ioctl(tapefd,STIOCTOP,&stop)) {
    printf("ioctl failure. errno=%d",errno);
    exit(errno);
}
```

**STIOCQRYP or STIOCSETP**

The **STIOCQRYP** command allows the program to query the tape device, device driver, and the media parameters. The **STIOCSETP** command allows the program to change the tape device, the device driver, and the media parameters.

Before the **STIOCSETP** command is issued, use the **STIOCQRYP** command to query and fill the fields of the data structure you do not want to change. Then, issue the **STIOCSETP** command to change the selected fields. Changing certain fields, such as **buffered\_mode**, impacts performance. If the **buffered\_mode** field is FALSE, each record that is written to the tape is immediately transferred to the tape. This operation guarantees that each record is on the tape, but it impacts performance.

**Unchangeable parameters**

The following parameters that are returned by the **STIOCQRYP** command cannot be changed by the **STIOCSETP** command.

**hkwrD**

This parameter is accepted but ignored.

**logical\_write\_protect**

This parameter sets the type of logical write protection for the tape that is loaded in the drive.

**write\_protect**

If the currently mounted tape is write protected, this field is set to TRUE. Otherwise, it is set to FALSE.

**min\_blksize**

This parameter is the minimum block size for the device. The driver gets this field by issuing the **SCSI Read Block Limits** command to the device.

**max\_blksize**

This parameter is the maximum block size for the device. The driver gets this field by issuing the **SCSI Read Block Limits** command to the device.

**retain\_reservation**

This parameter is accepted but ignored.

**medium\_type**

This parameter is the media type of the currently loaded tape. Some tape devices support multiple media types and report different values in this field. See the hardware reference guide for the specific tape device to determine the possible values.

**capacity\_scaling**

This parameter sets the capacity or logical length of the current tape. By reducing the capacity of the tape, the tape drive can access data faster. Capacity Scaling is not currently supported in IBMtape.

**density\_code**

This parameter is the density setting for the currently loaded tape. Some tape devices support multiple densities and report the current setting in this field. See the hardware reference guide for the specific tape device to determine the possible values.

**volid**

This field is always set to zero.

**emulate\_autoloader**

This parameter is accepted but ignored.

**record\_space\_mode**

Only **SCSI\_SPACE\_MODE** is supported.

**read\_sili\_bit**

This parameter is accepted but ignored. SILI bit is not supported due to Linux system environment limitations.

## Changeable parameters

The following parameters can be changed by using the **STIOCSETP** IOCTL command.

**trace**

This parameter turns the trace for the tape device On or Off.

**blksize**

This parameter specifies the new effective block size for the tape device. Use 0 for variable block mode.

**compression**

This parameter turns the hardware compression On or Off.

**max\_scsi\_xfer**

This parameter is the maximum transfer size that is allowed per SCSI command. In the IBMtape driver 3.0.3 or lower level, this value is 256 KB (262144 bytes) by default and changeable through the **STIOCSETP** IOCTL. In the IBMtape driver 3.0.5 or above and the open source driver lin\_tape, this parameter is not changeable any more. It is determined by the maximum transfer size of the Host Bus Adapter that the tape drive is attached to.

**trailer\_labels**

If this parameter is set to On, then writing a record past the early warning mark on the tape is allowed. Only the first write operation that detects the early warning mark returns the ENOSPC error code. All subsequent write operations are allowed to continue despite the check conditions that result from writing in the early warning zone (which are suppressed). When the end of the physical volume is reached, EIO is returned.

If this parameter is set to Off, the first write in the early warning zone fails, the ENOSPC error code is returned, and subsequent write operations fail.

#### **rewind\_immediate**

This parameter turns the immediate bit On or Off for subsequent rewind commands. If it is set to On, then the STREW tape operation runs faster. However, the next tape command can take longer to finish because the actual physical rewind operation must complete before the next tape command can start.

#### **logging**

This parameter turns the volume logging for the tape device On or Off.

#### **disable\_sim\_logging**

If this parameter is Off, the SIM/MIM data is automatically retrieved by the IBMtape device driver whenever it is available in the tape device.

#### **disable\_auto\_drive\_dump**

If this parameter is Off, the drive dump is automatically retrieved by the IBMtape device driver whenever a drive dump is in the tape device. It can also be set for all devices at modprobe configuration by adding disable\_auto\_drive\_dump=1.

#### **logical\_write\_protect**

This parameter sets the type of logical write protection for the tape that is loaded in the drive. See the hardware reference guide for the specific device for different types of logical write protect.

#### **capacity\_scaling**

This field can be changed only when the tape is positioned at the beginning of the tape. When a change is accepted, IBMtape rescales the tape capacity by formatting the loaded tape. See the *IBM Enterprise Tape System 3592 SCSI Reference* for the specific device for different types of capacity scaling.

IBM 3592 tape cartridges have two formats available, the 300 GB format and the 60 GB Fast Access format. The format of a cartridge can be queried under program control by issuing the **STIOCQRYP** IOCTL and checking the returned value of capacity\_scaling\_value (in hex).

If the capacity\_scaling\_value is 0x00, your 3592 tape cartridge is in 300 GB format. If the capacity\_scaling\_value is 0x35, your tape cartridge is in 60 GB Fast Access format. If the capacity\_scaling\_value is some other value, your tape cartridge format is undefined. (IBM can later define other supported cartridge formats. If so, they are documented in later versions of the *IBM TotalStorage Enterprise Tape System 3592 SCSI Reference*).

If you want to change your cartridge format, you can use the **STIOCSETP** IOCTL to change the capacity scaling value of your cartridge.



**Warning:** All data on the cartridge is lost when the format is changed.

If you want to set it to the 300 GB format, set capacity\_scaling\_value to 0x00 and capacity\_scaling to SCALE\_VALUE. If you want to set it to the 60 GB Fast Access format, set capacity\_scaling\_value to 0x35 and capacity\_scaling to SCALE\_VALUE. Setting capacity\_scaling to SCALE\_VALUE is required.

**Note:** All data on the tape is deleted and is not recoverable.

#### **read\_past\_file\_mark**

This parameter changes the behavior of the **read** function when a filemark is encountered. If the **read\_past\_filemark** flag is TRUE when a **read** operation encounters a file mark, IBMtape returns the number of bytes read before the filemark is encountered and sets the tape position at the EOT side of the file mark.

If the **read\_past\_filemark** flag is FALSE (by default) when a read operation encounters a filemark, if data was read, the **read** function returns the number of bytes read, and positions the tape at the BOT side of the filemark. If no data was read, the **read** returns 0 bytes and positions the tape at the EOT side of the filemark.

#### **limit\_read\_recov**

If this flag is TRUE, automatic recovery from read errors is limited to 5 seconds. If it is FALSE, the default is restored and the tape drive takes an arbitrary amount of time for read error recovery.

### **limit\_write\_recov**

If this flag is TRUE, automatic recovery from write errors is limited to 5 seconds. If it is FALSE, the default is restored and the tape drive takes an arbitrary amount of time for write error recovery.

### **data\_safe\_mode**

If this flag is TRUE, **data\_safe\_mode** is set in the drive. This action prevents data on the tape from being overwritten to avoid accidental data loss. If the value is FALSE, **data\_safe\_mode** is turned off.

### **pews**

This parameter establishes the programmable early warning zone size. It is a 2-byte numerical value that specifies how many MB before the standard end-of-medium early warning zone to place the programmable early warning indicator. If this value is set to a positive integer, a user application is warned that the tape is running out of space when the tape head reaches the PEW location. If **pews** is set to 0, then there no early warning zone occurs and the user is notified only at the standard early warning location.

The input or output data structure is

```
struct stchgp_s {
    int blksize;                                /* new block size */
    boolean trace;                               /* TRUE = message trace on */
    uint hkwrd;                                 /* trace hook word */
    int sync_count;                             /* obsolete - not used */
    boolean autoload;                            /* on/off autoload feature */
    boolean buffered_mode;                      /* on/off buffered mode */
    boolean compression;                         /* on/off compression */
    boolean trailer_labels;                     /* on/off allow writing after EOM */
    boolean rewind_immediate;                  /* on/off immediate rewinds */
    boolean bus_domination;                    /* obsolete - not used */
    boolean logging;                            /* enable or disable volume logging */
    boolean wwrite_protect;                     /* write_protected media */
    uint min_blksize;                           /* minimum block size */
    uint max_blksize;                           /* maximum block size */
    uint max_scsi_xfer;                         /* maximum scsi tranfer len */
    char volid[16];                            /* volume id */
    uchar acf_mode;                            /* automatic cartridge facility mode*/
#define ACF_NONE      0
#define ACF_MANUAL   1
#define ACF_SYSTEM   2
#define ACF_AUTOMATIC 3
#define ACF_ACCUMULATE 4
#define ACF_RANDOM   5
    uchar record_space_mode;          /* fsr/bsr space mode */
#define SCSI_SPACE_MODE 1
#define AIX_SPACE_MODE 2
    uchar logical_write_protect; /* logical write protect */
#define NO_PROTECT    0
#define ASSOCIATED_PROTECT 1
#define PERSISTENT_PROTECT 2
#define WORM_PROTECT   3
    uchar capacity_scaling;           /* capacity scaling */
#define SCALE_100     0
#define SCALE_75      1
#define SCALE_50      2
#define SCALE_25      3
#define SCALE_VALUE   4
    uchar retain_reservation;        /* retain reservation */
    uchar alt_pathing;              /* alternate pathing active */
    boolean emulate_autoloader;     /* emulate autoloader in random mode*/
    uchar medium_type;              /* tape medium type */
    uchar density_code;             /* tape density code */
    boolean disable_sim_logging;    /* disable sim/mim error logging */
    boolean read_sili_bit;           /* SILI bit setting for read commands*/
    uchar read_past_filemark;       /* fixed block read pass the filemark*/
    boolean disable_auto_drive_dump; /* disable auto drive dump logging*/
    uchar capacity_scaling_value;   /* hex value of capacity scaling */
    boolean wfm_immediate;          /* buffer write file mark */
    boolean limit_read_recov;        /* limit read recovery to 5 seconds */
    boolean limit_write_recov;       /* limit write recovery to 5 seconds */
    boolean data_safe_mode;          /* turn data safe mode on/off */
    uchar pews[2];                 /* programmable early warn zone size*/
    uchar reserve_type;             /* if set persistent reserve will be used */
    uchar reserved[12];
};
```

An example of the **STIOCQRYP** and **STIOCSETP** commands is

```
#include <sys/IBM_tape.h>
struct stchgp_s stchgp;
/* get current parameters */
if (ioctl(tapefd,STIOCQRYP,&stchgp)) {
    printf("ioctl failure. errno=%d",errno);
    exit(errno);
}
/* set new parameters */
stchgp.rewind_immediate=1;
stchgp.trailer_labels=1;
if (ioctl(tapefd,STIOCSETP,&stchgp)) {
    printf("IOCTL failure. errno=%d",errno);
    exit(errno);
}
```

## STIOCSYNC

This IOCTL command immediately flushes the tape buffers to the tape. There are no arguments for this IOCTL command.

An example of the **STIOCSYNC** command is

```
#include <sys/IBM_tape.h>
if (ioctl(tapefd,STIOCSYNC,NULL)) {
    printf("ioctl failure. errno=%d",errno);
    exit(errno);
}
```

## STIOCDM

This IOCTL command shows and manipulates one or two messages on the message display. The message that is sent with this call does not always remain on the display. It depends on the current state of the tape device. Refer to the IBM 3590 manuals for a description of the message display functions.

The input data structure is

```
#define MAXMSGLEN 8
struct stdm_s
{
    char dm_function;           /* function code */
    /* function selection */
#define DMSTATUSMSG   0x00      /* general status message */
#define DMDVMSG      0x20      /* demount verify message */
#define DMMIMMED     0x40      /* mount with immediate action indicator */
#define DMDEMIMMED   0xE0      /* demount/mount with immediate action */
    /* message control */
#define DMMSG0        0x00      /* display message 0 */
#define DMMSG1        0x04      /* display message 1 */
#define DMFLASHMSG0   0x08      /* flash message 0 */
#define DMFLASHMSG1   0x0C      /* flash message 1 */
#define DMAALTERNATE  0x10      /* alternate message 0 and message 1 */
    char dm_msg0[MAXMSGLEN];   /* message 0 */
    char dm_msg1[MAXMSGLEN];   /* message 1 */
};
```

An example of the **STIOCDM** command is

```
#include <sys/IBM_tape.h>
struct stdm_s stdm;
memset(&stdm, 0, sizeof(struct stdm_s));
stdm.dm_func = DMSTATUSMSG|DMMSG0;
bcopy("SSG", stdm.dm_msg0, 8);
if(ioctl(tapefd, STIOCDM, &stdm)<0)
{
    printf("IOCTL failure, errno = %d", errno);
    exit(errno);
}
```

## **STIOCQRYPOS**

This command queries the tape position. Tape position is defined as the location where the next read or write operation occurs. The query function can be used independently of, or with, the **STIOCSETPOS** IOCTL command.

A write filemark of count 0 is always issued to the drive, which flushes all data from the buffers to the tape media. After the write filemark finishes, the query is issued.

After a **query** operation, the **curpos** field is set to an unsigned integer that represents the current position.

The **eot** field is set to TRUE if the tape is positioned between the early warning and the physical end of the tape. Otherwise, it is set to FALSE.

The **lbot** field is valid only if the last command was a **write** command. If a query is issued and the last command was not a write, **lbot** contains the value **LBOT\_UNKNOWN**.

**Note:** **lbot** indicates the last block of data that is transferred to the tape.

The number of blocks and number of bytes currently in the tape device buffers is returned in the **num\_blocks** and **num\_bytes** fields.

The **bot** field is set to TRUE if the tape position is at the beginning of the tape. Otherwise, it is set to FALSE.

The returned **partition\_number** field is the current partition of the loaded tape.

The block ID of the next block of data to be transferred to or from the physical tape is returned in the **tapepos** field.

The position data structure is

```
typedef unsigned int blockid_t;
struct stpos_s {
    char      block_type;          /* Format of block ID information */
#define QP_LOGICAL 0               /* SCSI logical block ID format */
#define QP_PHYSICAL 1              /* Vendor-specific block ID format */
    boolean   eot;                /* Position is after early warning, */
                                /* before physical end of tape. */
    blockid_t curpos;             /* For query pos, current position.*/
                                /* For set pos, position to go to. */
    blockid_t lbot;               /* Last block written to tape. */
#define LBOT_NONE 0xFFFFFFFF       /* No blocks written to tape.*/
#define LBOT_UNKNOWN 0xFFFFFFFFFF /* Unable to determine info. */
    uint      num_blocks;          /* Number of blocks in buffer. */
    uint      num_bytes;           /* Number of bytes in buffer. */
    boolean   bot;                /* Position is at beginning of tape*/
    uchar    partition_number;     /* Current partition number on tape*/
    uchar    reserved1[2];
    blockid_t tapepos;            /* Next block to be transferred. */
    uchar    reserved2[48];
};
```

An example of the **STIOCQRYPOS** command is

```
#include <sys/IBM_tape.h>
struct stpos_s stpos;
stpos.block_type=QP_PHYSICAL;
if (ioctl(tapefd,STIOCQRYPOS,&stpos)) {
    printf("ioctl failure.  errno=%d",errno);
    exit(errno);
}
oldposition=stpos.curpos;
```

## **STIOCSETPOS**

This IOCTL command issues a high speed **locate** operation to the position specified on the tape. It uses the same position data structure that is described for **STIOCQRYPOS**, however, only the **block\_type** and **curpos** fields are used during a **set** operation. **STIOCSETPOS** can be used independently of or with **STIOCQRYPOS**.

The **block\_type** must be set to either **QP\_PHYSICAL** or **QP\_LOGICAL**. However, there is no difference in how IBMtape processes the request.

An example of the **STIOCQRYPOS** and **STIOCSETPOS** commands is

```
#include <sys/IBM_tape.h>
struct stpos_s stpos;
stpos.block_type=QP_LOGICAL;
if (ioctl(tapefd,STIOCQRYPOS,&stpos)) {
    printf("ioctl failure. errno=%d",errno);
    exit(errno);
}
oldposition=stpos.curpos;

stpos.curpos=oldposition;
stpos.block_type=QP_LOGICAL;
if (ioctl(tapefd,STIOCSETPOS,&stpos)) {
    printf("ioctl failure. errno=%d",errno);
    exit(errno);
}
```

## STIOCQRYSENSE

This IOCTL command returns the last sense data that is collected from the tape device. Or, it issues a new **Request Sense** command and returns the collected data. If **sense\_type** equals **LASTERROR**, then the sense data is valid only if the last tape operation had an error that caused a sense command to be issued to the device. If the sense data is valid, then the IOCTL command finishes successfully, and the **len** field is set to a value greater than zero. The **residual\_count** field contains the residual count from the last operation.

The input or output data structure is

```
#define MAXSENSE 255
struct stsense_s {
    /* input */
    char sense_type;          /* fresh (new sense) or sense from last error */
    #define FRESH 1            /* Initiate a new sense command */
    #define LASTERROR 2         /* Return sense gathered from */
                                /* the last SCSI sense command. */
    /* output */
    uchar sense[MAXSENSE];    /* actual sense data */
    int len;                  /* length of valid sense data returned */
    int residual_count;       /* residual count from last operation */
    uchar reserved[60];
};
```

An example of the **STIOCQRYSENSE** command is

```
#include <sys/IBM_tape.h>
struct stsense_s stsense;
stsense.sense_type=LASTERROR;
#define MEDIUM_ERROR 0x03
if (ioctl(tapefd,STIOCQRYSENSE,&stsense)) {
    printf("ioctl failure. errno=%d",errno);
    exit(errno);
}
if ((stsense.sense[2]&0x0F)==MEDIUM_ERROR) {
    printf("We're in trouble now!");
    exit(SENSE_KEY(&stsense.sense));
}
```

## STIOCQRYINQUIRY

This IOCTL command returns the inquiry data from the device. The data is divided into standard and vendor-specific portions.

The output data structure is

```
/*inquiry data info */
struct inq_data_s {
    BYTE b0;
    /*macros for accessing fields of byte 1 */
    #define PERIPHERAL_QUALIFIER(x) ((x->b0 &0xE0)>>5)
```

```

#define PERIPHERAL_CONNECTED 0x00
#define PERIPHERAL_NOT_CONNECTED 0x01
#define LUN_NOT_SUPPORTED 0x03
#define PERIPHERAL_DEVICE_TYPE(x) ((x->b0 &0x1F)
#define DIRECT_ACCESS 0x00
#define SEQUENTIAL_DEVICE 0x01
#define PRINTER_DEVICE 0x02
#define PROCESSOR_DEVICE 0x03
#define CD_ROM_DEVICE 0x05
#define OPTICAL_MEMORY_DEVICE 0x07
#define MEDIUM_CHANGER_DEVICE 0x08
#define UNKNOWN 0x1F
BYTE b1;
/*macros for accessing fields of byte 2 */
#define RMB(x) (((x->b1 &0x80)>>7) /*removable media bit */
#define FIXED 0
#define REMOVABLE 1
#define device_type_qualifier(x) (x->b1 &0x7F) /*vendor specific */
BYTE b2;
/*macros for accessing fields of byte 3 */
#define ISO_Version(x) (((x->b2 &0xC0)>>6)
#define ECMA_Version(x) (((x->b2 &0x38)>>3)
#define ANSI_Version(x) (x->b2 &0x07)
#define NONSTANDARD 0
#define SCSI1 1
#define SCSI2 2
#define SCSI3 3
BYTE b3;
/*macros for accessing fields of byte 4 */
/* asynchronous event notification */
#define AENC(x) (((x->b3 &0x80)>>7)
/* support terminate I/O process message? */
#define TrmIOP(x) (((x->b3 &0x40)>>6)
#define Response_Data_Format(x) (x->b3 &0xF)
#define SCSI1INQ 0 /* SCSI-1 standard inquiry data format */
#define CCSINQ 1 /* CCS standard inquiry data format */
#define SCSI2INQ 2 /* SCSI-2 standard inquiry data format */
BYTE additional_length; /* bytes following this field minus 4 */
BYTE res5;
BYTE b6;
#define MChngr(x) (((x->b6 & 0x08)>>3)
BYTE b7;
/*macros for accessing fields of byte 7 */
#define RelAdr(x) (((x->b7 &0x80)>>7)
/* the following fields are true or false */
#define WBus32(x) (((x->b7 &0x40)>>6)
#define WBus16(x) (((x->b7 &0x20)>>5)
#define Sync(x) (((x->b7 &0x10)>>4)
#define Linked(x) (((x->b7 &0x08)>>3)
#define CmdQue(x) (((x->b7 &0x02)>>1)
#define SftRe(x) (x->b7 &0x01)
char vendor_identification [8 ];
char product_identification [16 ];
char product_revision_level [4 ];
};

struct st_inquiry
{
    struct inq_data_s standard;
    BYTE vendor_specific [255-sizeof(struct inq_data_s)];
};

```

An example of the **STIOCQRYINQUIRY** command is

```

struct st_inquiry inqd;
if (ioctl(tapefd,STIOCQRYINQUIRY,&inqd)) {
    printf("ioctl failure.  errno=%d\n",errno);
    exit(errno);
}
if (ANSI_Version(((struct inq_data_s *)&(inqd.standard)))==SCSI2)
printf("Hey!  We have a SCSI-2 device\n");

```

## **STIOC\_LOCATE**

This IOCTL command causes the tape to be positioned at the specified block ID. The block ID used for the command must be obtained by using the **STIOC\_READ\_POSITION** command.

An example of the **STIOC\_LOCATE** command is

```
#include <sys/IBM_tape.h>
unsigned int current_blockid;

/* read current tape position */
if (ioctl(tapefd,STIOC_READ_POSITION,&current_blockid)) {
    printf("ioctl failure. errno=%d\n",errno);
    exit(1);
}

/* restore current tape position */
if (ioctl(tapefd,STIOC_LOCATE,current_blockid)) {
    printf("ioctl failure. errno=%d\n",errno);
    exit(1);
}
```

### ***STIOC\_READ\_POSITION***

This IOCTL command returns the block ID of the current position of the tape. The block ID returned from this command can be used with the **STIOC\_LOCATE** command to set the position of the tape.

An example of the **STIOC\_READ\_POSITION** command is

```
#include <sys/IBM_tape.h>
unsigned int current_blockid;
/* read current tape position */
if (ioctl(tapefd,STIOC_READ_POSITION,&current_blockid)) {
    printf("ioctl failure. errno=%d\n",errno);
    exit(1);
}
/* restore current tape position */
if (ioctl(tapefd,STIOC_LOCATE,current_blockid)) {
    printf("ioctl failure. errno=%d\n",errno);
    exit(1);
}
```

### ***STIOC\_RESET\_DRIVE***

This IOCTL command issues a **SCSI Send Diagnostic** command to reset the tape drive. There are no arguments for this IOCTL command.

An example of the **STIOC\_RESET\_DRIVE** command is

```
/* reset the tape drive */
if (ioctl(tapefd,STIOC_RESET_DRIVE,NULL)) {
    printf("ioctl failure. errno=%d\n",errno);
    exit(errno);
}
```

### ***STIOC\_PREVENT\_MEDIUM\_REMOVAL***

This IOCTL command prevents an operator from removing media from the device until the **STIOC\_ALLOW\_MEDIUM\_REMOVAL** command is issued or the device is reset.

There is no associated data structure.

An example of the **STIOC\_PREVENT\_MEDIUM\_REMOVAL** command is

```
#include <sys/IBM_tape.h>
if (!ioctl (tapefd, STIOC_PREVENT_MEDIUM_REMOVAL, NULL))
    printf ("The STIOC_PREVENT_MEDIUM_REMOVAL ioctl succeeded\n");
else {
    perror ("The STIOC_PREVENT_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}
```

### ***STIOC\_ALLOW\_MEDIUM\_REMOVAL***

This IOCTL command allows an operator to remove media from the device. This command is normally used after the **STIOC\_PREVENT\_MEDIUM\_REMOVAL** command to restore the device to the default state.

There is no associated data structure.

An example of the **STIOC\_ALLOW\_MEDIUM\_REMOVAL** command is

```
#include <sys/IBM_tape.h>
if (!ioctl (tapefd, STIOC_ALLOW_MEDIUM_REMOVAL, NULL))
    printf ("The STIOC_ALLOW_MEDIUM_REMOVAL ioctl succeeded\n");
else {
    perror ("The STIOC_ALLOW_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}
```

## **STIOC\_REPORT\_DENSITY\_SUPPORT**

This IOCTL command issues the **SCSI Report Density Support** command to the tape device. It returns either ALL supported densities or only supported densities for the currently mounted media. The **media** field specifies which type of report is requested. The **number\_reports** field is returned by the device driver and indicates how many density reports in the **reports** array field were returned.

The data structures that are used with this IOCTL is

```
struct density_report {
    uchar primary_density_code; /* primary density code */
    uchar secondary_density_code; /* secondary density code */
    uint wrtok :1, /* write ok, device can write this format */
        dup :1, /* zero if density only reported once */
        deflt :1, /* current density is default format */
        :5; /* reserved */
    char reserved[2]; /* reserved */
    uint bits_per_mm :24; /* bits per mm */
    ushort media_width; /* media width in millimeters */
    ushort tracks; /* tracks */
    uint capacity; /* capacity in megabytes */
    char assigning_org[8]; /* assigning organization in ASCII */
    char density_name[8]; /* density name in ASCII */
    char description[20]; /* description in ASCII */
};

struct report_density_support {
    uchar media; /* report all or current media as defined above */
    ushort number_reports; /* number of density reports returned in array */
    struct density_report reports[MAX_DENSITY_REPORTS];
};
```

Examples of the **STIOC\_REPORT\_DENSITY\_SUPPORT** command are

```
#include <sys/IBM_tape.h>
int stioc_report_density_support(void)
{
    int i;
    struct report_density_support density;
    printf("Issuing Report Density Support for ALL supported media...\n");
    density.media = ALL_MEDIA_DENSITY;
    if (ioctl(fd, STIOC_REPORT_DENSITY_SUPPORT, &density) != 0)
        return errno;
    printf("Total number of densities reported:
        %d\n", density.number_reports);
    for (i = 0; i<density.number_reports; i++) {
        printf("\n");
        printf(" Density Name..... %0.8s\n",
            density.reports[i].density_name);
        printf(" Assigning Organization..... %0.8s\n",
            density.reports[i].assigning_org);
        printf(" Density Name..... %0.8s\n",
            density.reports[i].density_name);
        printf(" Description..... %0.20s\n",
            density.reports[i].description);
        printf(" Primary Density Code..... %02X\n",
            density.reports[i].primary_density_code);
        printf(" Secondary Density Code..... %02X\n",
            density.reports[i].secondary_density_code);
        if (density.reports[i].wrtok)
            printf(" Write OK..... Yes\n");
        else
            printf(" Write OK..... No\n");
        if (density.reports[i].dup)
```

```

        printf(" Duplicate..... Yes\n");
    else
        printf(" Duplicate..... No\n");
    if (density.reports[i].deflt)
        printf(" Default..... Yes\n");
    else
        printf(" Default..... No\n");
    printf(" Bits per MM..... %d\n",
           density.reports[i].bits_per_mm);
    printf(" Media Width (millimeters).... %d\n",
           density.reports[i].media_width);
    printf(" Tracks..... %d\n",
           density.reports[i].tracks);
    printf(" Capacity (megabytes)..... %d\n",
           density.reports[i].capacity);
    if (opcode) {
        printf ("\nHit enter> to continue?");
        getchar();
    }
}
printf("\nIssuing Report Density Support for CURRENT media...\n");
density.media = CURRENT_MEDIA_DENSITY;
if (ioctl(fd, STIOC_REPORT_DENSITY_SUPPORT, &density) != 0)
    return errno;
for (i = 0; i<density.number_reports; i++) {
    printf("\n");
    printf(" Density Name..... %0.8s\n",
           density.reports[i].density_name);
    printf(" Assigning Organization..... %0.8s\n",
           density.reports[i].assigning_org);
    printf(" Description..... %0.20s\n",
           density.reports[i].description);
    printf(" Primary Density Code..... %02X\n",
           density.reports[i].primary_density_code);
    printf(" Secondary Density Code..... %02X\n",
           density.reports[i].secondary_density_code);
    if (density.reports[i].wrtok)
        printf(" Write OK..... Yes\n");
    else
        printf(" Write OK..... No\n");
    if (density.reports[i].dup)
        printf(" Duplicate..... Yes\n");
    else
        printf(" Duplicate..... No\n");
    if (density.reports[i].deflt)
        printf(" Default..... Yes\n");
    else
        printf(" Default..... No\n");
    printf(" Bits per MM..... %d\n",
           density.reports[i].bits_per_mm);
    printf(" Media Width (millimeters).... %d\n",
           density.reports[i].media_width);
    printf(" Tracks..... %d\n",
           density.reports[i].tracks);
    printf(" Capacity (megabytes)..... %d\n",
           density.reports[i].capacity);
}
return errno;
}

```

## **MTDEVICE (Obtain Device Number)**

This IOCTL command obtains the device number that is used for communicating with a 3494 library.

An example of the **MTDEVICE** command is

```

int device;
if(ioctl(tapefd, MTDEVICE, &device)<0)
{
    printf("IOCTL failure, errno = %d\n", errno);
    exit(errno);
}
printf("Device number is %X\n", device);

```

## **STIOC\_GET\_DENSITY and STIOC\_SET\_DENSITY**

The **STIOC\_GET\_DENSITY** IOCTL is used to query the current write density format settings on the tape drive. The current density code from the drive **Mode Sense** header, the **Read/Write Control Mode** page default density and pending density are returned.

The **STIOC\_SET\_DENSITY** IOCTL is used to set a new write density format on the tape drive by using the default and pending density fields. The density code field is not used and ignored on this IOCTL. The application can specify a new write density for the current loaded tape only or as a default for all tapes. Refer to the examples below.

The application must get the current density settings first before the current settings are modified. If the application specifies a new density for the current loaded tape only, then the application must issue another set density IOCTL after the current tape is unloaded and the next tape is loaded to either the default maximum density or a new density to ensure the tape drive uses the correct density. If the application specifies a new default density for all tapes, the setting remains in effect until changed by another set density IOCTL or the tape drive is closed by the application.

The structure for the **STIOC\_GET\_DENSITY** and **STIOC\_SET\_DENSITY** IOCTLs is

```
struct density_data_t
{
    char density_code;           /* mode sense header density code */
    char default_density;        /* default write density */
    char pending_density;        /* pending write density */
    char reserved[9];
};
```

### **Note:**

1. These IOCTLs are supported only on tape drives that can write multiple density formats. Refer to the Hardware Reference for the specific tape drive to determine whether multiple write densities are supported. If the tape drive does not support these IOCTLs, *errno* EINVAL is returned.
2. The device driver always sets the default maximum write density for the tape drive on every open system call. Any previous **STIOC\_SET\_DENSITY** IOCTL values from the last open are not used.
3. If the tape drive detects an invalid density code or cannot run the operation on the **STIOC\_SET\_DENSITY** IOCTL, the *errno* is returned and the current drive density settings before the IOCTL are restored.
4. The **struct density\_data\_t** defined in the header file is used for both IOCTLs. The **density\_code** field is not used and ignored on the **STIOC\_SET\_DENSITY** IOCTL.

## **Examples**

```
struct density_data_t data;

/* open the tape drive */
/* get current density settings */
rc = ioctl(fd, STIOC_GET_DENSITY, %data);

/* set 3592 J1A density format for current loaded tape only */
data.default_density = 0x7F;
data.pending_density = 0x51;
rc = ioctl(fd, STIOC_SET_DENSITY, %data);

/* unload tape */
/* load next tape */
/* set 3592 E05 density format for current loaded tape only */
data.default_density = 0x7F;
data.pending_density = 0x52;
rc = ioctl(fd, STIOC_SET_DENSITY, %data);

/* unload tape */
/* load next tape */
/* set default maximum density for current loaded tape */
data.default_density = 0;
```

```

data.pending_density = 0;
rc = ioctl(fd, STIOC_SET_DENSITY, %data);

/* close the tape drive          */
/* open the tape drive          */
/* set 3592 J1A density format for current loaded tape and all subsequent tapes */
data.default_density = 0x51;
data.pending_density = 0x51;

rc = ioctl(fd, STIOC_SET_DENSITY, %data);

```

## ***GET\_ENCRYPTION\_STATE***

This IOCTL command queries the drive's encryption method and state. The data structure that is used for this IOCTL is as follows on all of the supported operating systems

```

struct encryption_status
{
    uchar encryption_capable;      /* (1)Set this field as a boolean based on the
                                    capability of the drive */

    uchar encryption_method;        /* (2)Set this field to one of the following */
    #define METHOD_NONE            0 /* Only used in GET_ENCRYPTION_STATE */
    #define METHOD_LIBRARY          1 /* Only used in GET_ENCRYPTION_STATE */
    #define METHOD_SYSTEM           2 /* Only used in GET_ENCRYPTION_STATE */
    #define METHOD_APPLICATION      3 /* Only used in GET_ENCRYPTION_STATE */
    #define METHOD_CUSTOM            4 /* Only used in GET_ENCRYPTION_STATE */
    #define METHOD_UNKNOWN           5 /* Only used in GET_ENCRYPTION_STATE */

    uchar encryption_state;         /* (3) Set this field to one of the following */
    #define STATE_OFF               0 /* Used in GET/SET_ENCRYPTION_STATE */
    #define STATE_ON                1 /* Used in GET/SET_ENCRYPTION_STATE */
    #define STATE_NA                2 /* Only used in GET_ENCRYPTION_STATE*/
    uchar[13] reserved;
};

```

An example of the **GET\_ENCRYPTION\_STATE** command is

```

int qry_encrytion_state (void)
{
    int rc = 0;
    struct encryption_status encryption_status_t;

    printf("issuing query encryption status...\n");
    memset(&encryption_status_t, 0, sizeof(struct encryption_status));
    rc = ioctl(fd, GET_ENCRYPTION_STATE, );&encryption_status_t
    if(rc == 0)
    {
        if(encryption_status_t.encryption_capable)
            printf("encryption capable.....Yes\n");
        else
            printf("encryption capable.....No\n");
        switch(encryption_status_t.encryption_method)
        {
            case METHOD_NONE:
                printf("encryption method.....METHOD_NONE\n");
                break;
            case METHOD_LIBRARY:
                printf("encryption method.....METHOD_LIBRARY\n");
                break;
            case METHOD_SYSTEM:
                printf("encryption method.....METHOD_SYSTEM\n");
                break;
            case METHOD_APPLICATION:
                printf("encryption method.....METHOD_APPLICATION\n");
                break;
            case METHOD_CUSTOM:
                printf("encrypiton method.....METHOD_CUSTOM\n");
                break;
            case METHOD_UNKNOWN:
                printf("encryption method.....METHOD_UNKNOWN\n");
                break;
            default:
                printf("encription method.....Error\n");
        }
        switch(encryption_status_t.encryption_state)
        {

```

```

        case STATE_OFF:
            printf("encryption state.....OFF\n");
            break;
        case STATE_ON:
            printf("encryption state.....ON\n");
            break;
        case STATE_NA:
            printf("encryption state.....NA\n");
            break;
        default:
            printf("encryption state.....Error\n");
    }
}

return rc;
}

```

## **SET\_ENCRYPTION\_STATE**

This IOCTL command allows setting the encryption state only for application-managed encryption. On unload, some drive settings might be reset to default. To set the encryption state, the application must issue this IOCTL after a tape is loaded and at BOP.

The data structure that is used for this IOCTL is the same as the one for **GET\_ENCRYPTION\_STATE**. An example of the **SET\_ENCRYPTION\_STATE** command is

```

int set_encryption_state(int option)
{
    int rc = 0;
    struct encryption_status encryption_status_t;

    printf("issuing query encryption status...\n");
    memset(&encryption_status_t, 0, sizeof(struct encryption_status));
    rc = ioctl(fd, GET_ENCRYPTION_STATE, );&encryption_status_t
    if(rc < 0) return rc;

    if(option == 0)
        encryption_status_t.encryption_state = STATE_OFF;
    else if(option == 1)
        encryption_status_t.encryption_state = STATE_ON;
    else
    {
        printf("Invalid parameter.\n");
        return -EINVAL;
    }

    printf("Issuing set encryption state.....\n");
    rc = ioctl(fd, SET_ENCRYPTION_STATE, &encryption_status_t);

    return rc;
}

```

## **SET\_DATA\_KEY**

This IOCTL command allows the data key to be set only for application-managed encryption. The data structure that is used for this IOCTL is as follows on all of the supported operating systems.

```

struct data_key
{
    uchar[12] data_key_index;
    uchar data_key_index_length;
    uchar[15] reserved1;
    uchar[32] data_key;
    uchar[48] reserved2;
};

```

An example of the **SET\_DATA\_KEY** command is

```

int set_datakey(void)
{
    int rc = 0;
    struct data_key encryption_data_key_t;

    printf("Issuing set encryption data key.....\n");

```

```

        memset(&encryption_data_key_t, 0, sizeof(struct data_key));
        /* fill in your data key here, then issue the following ioctl*/
        rc = ioctl(fd, SET_DATA_KEY, &encryption_data_key_t);
        return rc;
    }
}

```

## **STIOC\_QUERY\_PARTITION**

This IOCTL queries and displays information for tapes that support partitioning. The data structure that is used for this IOCTL is

```

#define MAX_PARTITIONS 255
struct query_partition {
    uchar max_partitions;
    uchar active_partition;
    uchar number_of_partitions;
    uchar size_unit;
    ushort size[MAX_PARTITIONS];
    char reserved[32];
};

```

- **max\_partitions** is the maximum number of partitions that the tape allows.
- **active\_partition** is the current partition to which tape operations apply.
- **number\_of\_partitions** is the number of partitions currently on the tape.
- **size\_unit** describes the units for the size of the tape, which is given as a logarithm to the base 10.

For example, 0 refers to  $10^0 = 1$ , the most basic unit, which is bytes. All sizes that are reported are in bytes. 3 refers to  $10^3$ , or kilobytes. Size is an array of the size of the partitions on tape, one array element per partition, in size\_units.

An example of the **STIOC\_QUERY\_PARTITION** IOCTL is

```

int stioc_query_partition()
{
    struct query_partition qry;
    int rc = 0, i = 0;

    memset(&qry, '\0', sizeof(struct query_partition));
    printf("Issuing IOCTL...\n");
    rc = ioctl(fd, STIOC_QUERY_PARTITION, &qry);

    if(rc) {
        printf("Query partition failed: %d\n", rc);
        goto EXIT_LABEL;
    } /* if */

    printf("\nmax possible partitions: %d\n", qry.max_partitions);
    printf("number currently on tape: %d\n", qry.number_of_partitions);
    printf("active: %d\n", qry.active_partition);
    printf("unit: %d\n", qry.size_unit);

    for(i = 0; i < qry.number_of_partitions; i++)
        printf("size[%d]: %d\n", i, qry.size[i]);

    EXIT_LABEL:

    return rc;
} /* stioc_query_partition() */

```

## **STIOC\_CREATE\_PARTITION**

This IOCTL creates partitions on tapes that support partitioning. The data structure that is used for this IOCTL is

```

#define IDP_PARTITION      (1)
#define SDP_PARTITION      (2)
#define FDP_PARTITION      (3)
struct tape_partition {
    uchar type;
    uchar number_of_partitions;
    uchar size_unit;
    ushort size[MAX_PARTITIONS];
};

```

```

    char reserved[32];
};
```

**Type** is the type of partition, whether **IDP\_PARTITION** (initiator defined partition), **SDP\_PARTITION** (select data partition), or **FDP\_PARTITION** (fixed data partition). The behavior of these options is described in the SCSI reference for your tape drive.

- **number\_of\_partitions** is the number of partitions the user wants to create.
- **size\_unit** is as defined in the **STIOC\_QUERY\_PARTITION** section.
- **size** is an array of requested sizes, in size\_units, one array element per partition.

An example of the **STIOC\_CREATE\_PARTITION** IOCTL is

```

int stioc_create_partition()
{
    int rc = 0, i = 0, char_cap = 0, short_cap = 0;
    struct tape_partition crt;
    char* input = NULL;

    char_cap = pow(2, sizeof(char) * BITS_PER_BYTE) - 1;
    short_cap = pow(2, sizeof(short) * BITS_PER_BYTE) - 1;

    input = malloc(DEF_BUF_SIZE / 16);
    if(!input) {
        rc = ENOMEM;
        goto EXIT_LABEL;
    } /* if */
    memset(input, '\0', DEF_BUF_SIZE / 16);

    memset(&crt, '\0', sizeof(struct tape_partition));

    while(atoi(input) < IDP_PARTITION || atoi(input) > FDP_PARTITION + 1) {
        printf("%d) IDP_PARTITION\n", IDP_PARTITION);
        printf("%d) SDP_PARTITION\n", SDP_PARTITION);
        printf("%d) FDP_PARTITION\n", FDP_PARTITION);
        printf("%d) Cancel\n", FDP_PARTITION + 1);
        printf("\nPlease select: ");

        fgets(input, DEF_BUF_SIZE / 16, stdin);
        if(atoi(input) == FDP_PARTITION + 1) {
            rc = 0;
            goto EXIT_LABEL;
        } /* if */
    } /* while */

    crt.type = atoi(input);

    memset(input, '\0', DEF_BUF_SIZE / 16);
    while(input[0] < '1' || input[0] > '9') {
        printf("Enter desired number of partitions (0 to cancel): ");
        fgets(input, DEF_BUF_SIZE / 16, stdin);
        if(input[0] == '0') {
            rc = 0;
            goto EXIT_LABEL;
        } /* if */

        if(atoi(input) > MAX_PARTITIONS) {
            printf("Please select number <= %d\n", MAX_PARTITIONS);
            input[0] = '\0';
        } /* if */
    } /* while */

    crt.number_of_partitions = atoi(input);

    if(crt.type == IDP_PARTITION && crt.number_of_partitions > 1) {
        memset(input, '\0', DEF_BUF_SIZE / 16);
        while(input[0] < '0' || input[0] > '9') {
            printf("Enter size unit (0 to cancel): ");
            fgets(input, DEF_BUF_SIZE / 16, stdin);
            if(input[0] == '0') {
                rc = 0;
                goto EXIT_LABEL;
            } /* if */
            if(atoi(input) > char_cap) {
                printf("Please select number <= %d\n", char_cap);
                input[0] = '\0';
            } /* if */
        } /* while */
    }
```

```

crt.size_unit = atoi(input);

for(i = 0; i < crt.number_of_partitions; i++) {
    memset(input, '\0', DEF_BUF_SIZE / 16);
    while(input[0] != '-' &&
          (input[0] < '0' || input[0] > '9')) {
        printf("Enter size[%d] (0 to cancel, < 0 for "\n
               "remaining space on cartridge): ", i);
        fgets(input, DEF_BUF_SIZE / 16, stdin);
        if(input[0] == '0') {
            rc = 0;
            goto EXIT_LABEL;
        } /* if */

        if(atoi(input) > short_cap) {
            printf("Please select number <= %d\n",
                   short_cap);
            input[0] = '\0';
        } /* if */
    } /* while */
    if(input[0] == '-' && atoi(&input[1]) > 0)
        crt.size[i] = 0xFFFF;
    else crt.size[i] = atoi(input);
} /* for */
} /* if */

printf("Issuing IOCTL...\n");
rc = ioctl(fd, STIOC_CREATE_PARTITION, &crt);

if(rc) {
    printf("Create partition failed: %d\n", rc);
    goto EXIT_LABEL;
} /* if */

EXIT_LABEL:

if(input) free(input);
return rc;
} /* stioc_create_partition() */

```

## **STIOC\_SET\_ACTIVE\_PARTITION**

This IOCTL allows the user to specify the partition on which to run subsequent tape operations. The data structure that is used for this IOCTL is

```

struct set_active_partition {
    uchar partition_number;
    unsigned long long logical_block_id;
    char reserved[32];
};

```

- **partition\_number** is the number of the requested active partition.
- **logical\_block\_id** is the requested block position within the new active partition.

An example of the **STIOC\_SET\_ACTIVE\_PARTITION** IOCTL is

```

int stioc_set_partition()
{
    int rc = 0;
    struct set_active_partition set;
    char* input = NULL;

    input = malloc(DEF_BUF_SIZE / 16);
    if(!input) {
        rc = ENOMEM;
        goto EXIT_LABEL;
    } /* if */
    memset(input, '\0', DEF_BUF_SIZE / 16);

    memset(&set, '\0', sizeof(struct set_active_partition));
    while(input[0] < '0' || input[0] > '9') {
        printf("Select partition (< 0 to cancel): ");
        fgets(input, DEF_BUF_SIZE / 16, stdin);

        if(input[0] == '-' && atoi(&input[1]) > 0) {
            rc = 0;
            goto EXIT_LABEL;
        } /* if */
    } /* while */
}

```

```

} /* if */

    if(atoi(input) > MAX_PARTITIONS) {
        printf("Please select number < %d\n", MAX_PARTITIONS);
        input[0] = '\0';
    } /* if */
} /* while */
set.partition_number = atoi(input);

printf("Issuing IOCTL...\n");
rc = ioctl(fd, STIOC_SET_ACTIVE_PARTITION, &set);
if(rc) {
    printf("Set partition failed: %d\n", rc);
    goto EXIT_LABEL;
} /* if */

EXIT_LABEL:

    if(input) free(input);
    return rc;
} /* stioc_set_partition() */

```

### **STIOC\_ALLOW\_DATA\_OVERWRITE**

This IOCTL allows data on the tape to be overwritten when in data safe mode. The data structure that is used for this IOCTL is

```

struct allow_data_overwrite {
    uchar partition_number;
    unsigned long long logical_block_id;
    uchar allow_format_overwrite;
    char reserved[32];
};

```

- **partition\_number** is the number of the drive partition on which to allow the overwrite.
- **logical\_block\_id** is the block that you want to overwrite.
- **allow\_format\_overwrite**, if set to TRUE, instructs the tape drive to allow a format of the tape and accept the **CREATE\_PARTITION** ioctl.

If **allow\_format\_overwrite** is TRUE, **partition\_number** and **logical\_block\_id** are ignored.

An example of the use of the **STIOC\_ALLOW\_DATA\_OVERWRITE** IOCTL is

```

int stioc_allow_overwrite()
{
    int rc = 0, i = 0, brk = FALSE;
    struct allow_data_overwrite ado;
    char* input = NULL;

    memset(&ado, '\0', sizeof(struct allow_data_overwrite));
    input = malloc(DEF_BUF_SIZE / 4);
    if(!input) {
        rc = ENOMEM;
        goto EXIT_LABEL;
    } /* if */
    memset(input, '\0', DEF_BUF_SIZE / 4);

    while(input[0] < '0' || input[0] > '1') {
        printf("0. Write Data 1. Create Partition (< 0 to cancel): ");
        fgets(input, DEF_BUF_SIZE / 4, stdin);

        if(input[0] == '-' && atoi(&input[1]) > 0) {
            rc = 0;
            goto EXIT_LABEL;
        } /* if */
    } /* while */

    ado.allow_format_overwrite = atoi(&input[0]);
    switch(ado.allow_format_overwrite) {
    case 0:
        memset(input, '\0', DEF_BUF_SIZE / 4);
        while((input[0] < '0' || input[0] > '9') && (input[0] < 'a' || input[0] > 'f')) {
            brk = FALSE;
            printf("Enter partition in hex (< 0 to cancel): 0x");
            fgets(input, DEF_BUF_SIZE / 4, stdin);
    }
}

```

```

        if(input[0] == '-' && atoi(&input[1]) > 0) {
            rc = 0;
            goto EXIT_LABEL;
        } /* if */

        while(strlen(input) &&
              isspace(input[strlen(input) - 1]))
            input[strlen(input) - 1] = '\0';
        if(!strlen(input)) continue;

        for(i = 0; i < strlen(input); i++) {
            if(input[i] >= 'A' && input[i] <= 'F')
                input[i] = input[i] - 'A' + 'a';
            else if(((input[i] < '0' || input[i] > '9') &&
                     (input[i] < 'a' || input[i] > 'f')) || i >= sizeof(unchar) * 2) {
                printf("Input must be from 0 to 0xFF\n");
                memset(input, '\0', DEF_BUF_SIZE / 4);
                brk = TRUE;
                break;
            } /* else if */
        } /* for */
        if(brk) continue;

    } /* while */

ado.partition_number = char_to_hex(input);

memset(input, '\0', DEF_BUF_SIZE / 4);
while((input[0] < '0' || input[0] > '9') && (input[0] < 'a' || input[0] > 'f')) {
    brk = FALSE;
    printf("Enter block ID in hex (< 0 to cancel): 0x");
    fgets(input, DEF_BUF_SIZE / 4, stdin);

    if(input[0] == '-' && atoi(&input[1]) > 0) {
        rc = 0;
        goto EXIT_LABEL;
    } /* if */

    while(strlen(input) &&
          isspace(input[strlen(input) - 1]))
        input[strlen(input) - 1] = '\0';
    if(!strlen(input)) continue;

    for(i = 0; i < strlen(input); i++) {
        if(input[i] >= 'A' && input[i] <= 'F')
            input[i] = input[i] - 'A' + 'a';
        else if(((input[i] < '0' || input[i] > '9') && (input[i] < 'a' || input[i] > 'f')) || i >= sizeof(unsigned long long) * 2) {
            printf("Input out of range\n");
            memset(input, '\0', DEF_BUF_SIZE / 4);
            brk = TRUE;
            break;
        } /* else if */
    } /* for */
    if(brk) continue;
}

} /* while */

ado.logical_block_id = char_to_hex(input);
break;
case 1:
    break;
default:
    assert(!"Unreachable.");
} /* switch */

printf("Issuing IOCTL...\n");
rc = ioctl(fd, STIOC_ALLOW_DATA_OVERWRITE, &ado);

if(rc) {
    printf("Allow data overwrite failed: %d\n", rc);
    goto EXIT_LABEL;
} /* if */

EXIT_LABEL:

if(input) free(input);

```

```
        return rc;
} /* stioc_allow_overwrite() */
```

## STIOC\_READ\_POSITION\_EX

This IOCTL returns tape position with support for the short, long, and extended formats. The definitions and data structures that are used for this IOCTL follow. See the **READ\_POSITION** section of your tape drive's SCSI documentation for details on the **short\_data\_format**, **long\_data\_format**, and **extended\_data\_format** structures.

```
#define RP_SHORT_FORM (0x00)
#define RP_LONG_FORM (0x06)
#define RP_EXTENDED_FORM (0x08)

struct short_data_format {
#if defined __LITTLE_ENDIAN
    uchar bpew : 1;
    uchar perr : 1;
    uchar lolu : 1;
    uchar rsvd : 1;
    uchar bycu : 1;
    uchar locu : 1;
    uchar eop : 1;
    uchar bop : 1;
#elif defined __BIG_ENDIAN
    uchar bop : 1;
    uchar eop : 1;
    uchar locu : 1;
    uchar bycu : 1; uchar rsvd : 1;
    uchar lolu : 1;
    uchar perr : 1;
    uchar bpew : 1;
#else
    error
#endif
    uchar active_partition;
    char reserved[2];
    uchar first_logical_obj_position[4];
    uchar last_logical_obj_position[4];
    uchar num_buffer_logical_obj[4];
    uchar num_buffer_bytes[4];
    char reserved1;
};

struct long_data_format {
#if defined __LITTLE_ENDIAN
    uchar bpew : 1;
    uchar rsvd2 : 1;
    uchar lonu : 1;
    uchar mpu : 1;
    uchar rsvd1 : 2;
    uchar eop : 1;
    uchar bop : 1;
#elif defined __BIG_ENDIAN
    uchar bop : 1;
    uchar eop : 1;
    uchar rsvd1 : 2;
    uchar mpu : 1;
    uchar lonu : 1;
    uchar rsvd2 : 1;
    uchar bpew : 1;
#else
    error
#endif
    char reserved[6];
    uchar active_partition;
    uchar logical_obj_number[8];
    uchar logical_file_id[8];
    uchar obsolete[8];
};

struct extended_data_format {
#if defined __LITTLE_ENDIAN
    uchar bpew : 1;
    uchar perr : 1;
    uchar lolu : 1;
    uchar rsvd : 1;
    uchar bycu : 1;
```

```

        uchar locu : 1;
        uchar eop : 1;
        uchar bop : 1;
#if defined __BIG_ENDIAN
        uchar bop : 1;
        uchar eop : 1;
        uchar locu : 1;
        uchar bycu : 1;
        uchar rsvd : 1;
        uchar lolu : 1;
        uchar perr : 1;
        uchar bpew : 1;
#else
        error
#endif
        uchar active_partition;
        uchar additional_length[2];
        uchar num_buffer_logical_obj[4];
        uchar first_logical_obj_position[8];
        uchar last_logical_obj_position[8];
        uchar num_buffer_bytes[8];
        uchar reserved;
};

struct read_tape_position {
    uchar data_format;
    union {
        struct short_data_format rp_short;
        struct long_data_format rp_long;
        struct extended_data_format rp_extended;
    } rp_data;
};

```

**data\_format** is the format in which you want to receive your data, as defined here. It can take the value RP\_SHORT\_FORM, RP\_LONG\_FORM, or RP\_EXTENDED\_FORM. When the IOCTL finishes, data is returned to the corresponding structure within the **rp\_data** union.

An example of the use of the **STIOC\_READ\_POSITION\_EX** IOCTL is

```

int stioc_read_position_ex(void)
{
    int rc = 0;
    char* input = NULL;
    struct read_tape_position rp = {0};

    printf("Note: only supported on LTO 5 and higher drives\n");
    input = malloc(DEF_BUF_SIZE / 16);
    if(!input) {
        rc = ENOMEM;
        goto EXIT_LABEL;
    } /* if */
    memset(input, '\0', DEF_BUF_SIZE / 16);

    while(input[0] == '\0' || atoi(input) < 0 || atoi(input) > 3) {
        printf("0) Cancel\n");
        printf("1) Short Form\n");
        printf("2) Long Form\n");
        printf("3) Extended Form\n");

        printf("\nPlease select: ");

        fgets(input, DEF_BUF_SIZE / 16, stdin);
        if(!atoi(input)) {
            rc = 0;
            goto EXIT_LABEL;
        } /* if */
    } /* while */

    memset(&rp, '\0', sizeof(struct read_tape_position));

    switch(atoi(input)) {
    case 1:
        rp.data_format = RP_SHORT_FORM;
        break;
    case 2:
        rp.data_format = RP_LONG_FORM;
        break;
    case 3:
        rp.data_format = RP_EXTENDED_FORM;
        break;
    }
}

```

```

default:
    rc = EINVAL;
    goto EXIT_LABEL;
} /* switch */

rc = ioctl(fd, STIOC_READ_POSITION_EX, &rp);

if(rc) {
    printf("Cannot get position: %d\n", rc = errno);
    goto EXIT_LABEL;
} /* if */

print_read_position_ex(&rp);

EXIT_LABEL:
    if(input) free(input);
    return rc;
} /* stioc_read_position_ex() */

```

## ***STIOC\_LOCATE\_16***

This IOCTL sets the tape position by using the long tape format. The definitions and structure that are used for this IOCTL are

```

#define LOGICAL_ID_BLOCK_TYPE (0x00)
#define LOGICAL_ID_FILE_TYPE (0x01)

struct set_tape_position {
    uchar logical_id_type;
    unsigned long long logical_id;
    char reserved[32];
};

```

*logical\_id\_type* can take the values *LOGICAL\_ID\_BLOCK\_TYPE* or *LOGICAL\_ID\_FILE\_TYPE*. The values specify whether the tape head is located to the block with the specified *logical\_id* or to the file with the specified *logical\_id*. An example on how to use the ***STIOC\_LOCATE\_16*** IOCTL follows. The snippet assumes the declaration of global variables *filetype* and *blockid*.

```

int stioc_locate_16(void)
{
    int rc = 0;
    struct set_tape_position pos;

    memset(&pos, '\0', sizeof(struct set_tape_position));
    printf("\nLocating to %s ID %u (0x%08X)...\\n",
           filetype ? "File" : "Block", blockid, blockid);

    pos.logical_id_type = filetype;
    pos.logical_id = (long long) blockid;

    rc = ioctl(fd, STIOC_LOCATE_16, &pos);
    return rc;
} /* stioc_locate_16() */

```

## ***STIOC\_QUERY\_BLK\_PROTECTION***

This IOCTL queries capability and status of the drive's Logical Block Protection. The structures and defines are

```

#define LBP_DISABLE          (0x00)
#define REED_SOLOMON_CRC     (0x01)

struct logical_block_protection {
    uchar lbp_capable;
    uchar lbp_method;
    uchar lbp_info_length;
    uchar lbp_w;
    uchar lbp_r;
    uchar rbdp;
    uchar reserved[26];
};

```

The *lbp\_capable* is set to True if the drive supports logical block protection, or False otherwise.

A lbp\_method method of LBP\_DISABLE indicates that the logical block protection feature is turned off. A value of REED\_SOLOMON\_CRC indicates that logical block protection is used, with a Reed-Solomon cyclical redundancy check algorithm to run the block protection.

The lbp\_w indicates that logical block protection is run for write commands. The lbp\_x indicates that logical block protection is run for read commands. The rbdp indicates that logical block protection is run for recover buffer data. To use this IOCTL, issue the following call.

```
rc = ioctl(fd, STIOC_QUERY_BLK_PROTECTION, &lbp);
```

## **STIOC\_SET\_BLK\_PROTECTION**

This IOCTL sets status of the drive's Logical Block Protection. All fields are configurable except **lbp\_capable** and **reserved**. The structures and defines are the same as for **STIOC\_QUERY\_BLK\_PROTECTION**. To use this IOCTL, issue the following call.

```
rc = ioctl(fd, STIOC_SET_BLK_PROTECTION, &lbp);
```

## **STIOC\_VERIFY\_TAPE\_DATA**

This IOCTL instructs the tape drive to scan the data on its current tape to check for errors. The structure is defined as follows.

```
struct verify_data {  
    #if defined __LITTLE_ENDIAN  
        uchar fixed : 1;  
        uchar bytcmp : 1;  
        uchar immed : 1;  
        uchar vbf : 1;  
        uchar vlbpm : 1;  
        uchar vte : 1;  
        uchar reserved1 : 2;  
    #elif defined __BIG_ENDIAN  
        uchar reserved1 : 2;  
        uchar vte : 1;  
        uchar vlbpm : 1;  
        uchar vbf : 1;  
        uchar immed : 1;  
        uchar bytcmp : 1;  
        uchar fixed : 1;  
    #else  
        error  
    #endif  
    uchar verify_length[3];  
    uchar reserved2[15];  
};
```

vte instructs the drive to verify from the current tape head position to end of data.

vlbpm instructs the drive to verify that the logical block protection method that is specified in the **Control Data Protection** mode page is used for each block.

If vbf is set, then the **verify\_length** field contains the number of filemarks to be traversed, rather than the number of blocks or bytes.

immed specifies that status is to be returned immediately after the command descriptor block is validated. Otherwise, the command does not return status until the entire operation finishes.

bytcmp is set to 0.

fixed indicates a fixed-block length, and that **verify\_length** is interpreted as blocks rather than bytes.

**verify\_length** specifies the length to verify in files, blocks or bytes, depending on the values of the **vbf** and **fixed** fields. If **vte** is set to 1, **verify\_length** is ignored.

An example of the use of **STIOC\_VERIFY\_TAPE\_DATA** is as follows.

```
int stioc_verify()
{
    int rc = 0, i = 0, cont = TRUE, len = 0;
    char* input = NULL;
    struct verify_data* vfy = NULL;

    struct {
        char* desc;
        int idx;
    } table[] = {
        {"Verify to EOD", VFY_VTE},
        {"Verify Logical Block Protection", VFY_VLBPM},
        {"Verify by Filemarks", VFY_VBF},
        {"Return immediately", VFY_IMMED},
        {"Fixed", VFY_FIXED},
        {NULL, 0}
    };

    input = malloc(DEF_BUF_SIZE / 16);
    if(!input) {
        rc = ENOMEM;
        goto EXIT_LABEL;
    } /* if */
    memset(input, '\0', DEF_BUF_SIZE / 16);

    vfy = malloc(sizeof(struct verify_data));
    if(!vfy) {
        rc = ENOMEM;
        goto EXIT_LABEL;
    } /* if */
    memset(vfy, '\0', sizeof(struct verify_data));

    printf("\n");
    for(i = 0; table[i].desc; i++) {
        while(tolower(input[0]) != 'y' && tolower(input[0]) != 'n') {
            printf("%s (y/n/c to cancel)? ", table[i].desc);
            fgets(input, DEF_BUF_SIZE / 16, stdin);
            if(tolower(input[0]) == 'c') {
                rc = 0;
                goto EXIT_LABEL;
            } /* if */
        } /* while */

        if(tolower(input[0]) == 'y') {
            switch(table[i].idx) {
                case VFY_VTE: vfy->vte = 1; break;
                case VFY_VLBPM: vfy->vlbpm = 1; break;
                case VFY_VBF: vfy->vbf = 1; break;
                case VFY_IMMED: vfy->immed = 1; break;
                default: break;
            } /* switch */
        } /* if */
        memset(input, '\0', DEF_BUF_SIZE / 16);
    } /* for */

    if(!vfy->vte) {
        while(cont) {
            cont = FALSE;

            printf("Verify length in decimal (c to cancel): ");
            fgets(input, DEF_BUF_SIZE / 16, stdin);

            while(strlen(input) && isspace(input[strlen(input)-1]))
                input[strlen(input) - 1] = '\0';

            if(!strlen(input)) {
                cont = TRUE;
                continue;
            } /* if */

            if(tolower(input[0]) == 'c') {
                rc = 0;
                goto EXIT_LABEL;
            } /* if */

            for(i = 0; i < strlen(input); i++) {
                if(!isdigit(input[i])) {
                    memset(input, '\0', DEF_BUF_SIZE / 16);
                    cont = TRUE;
                }
            }
        }
    }
}
```

```

        } /* if */
    } /* for */

    } /* while */

    len = atoi(input);
    vfy->verify_length[0] = (len >> 16) & 0xFF;
    vfy->verify_length[1] = (len >> 8) & 0xFF;
    vfy->verify_length[2] = len & 0xFF;
} /* if */

rc = ioctl(fd, STIOC_VERIFY_TAPE_DATA, &vfy);
printf("VERIFY_TAPE_DATA returned %d\n", rc);
if(rc) printf("errno: %d\n", errno);

EXIT_LABEL:

if(input) free(input);
if(vfy) free(vfy);
return rc;
} /* stioc_verify() */

```

## **STIOC\_QUERY\_RAO**

The IOCTL is used to query the maximum number and size of User Data Segments (UDS) that are supported from tape drive and driver for the wanted **uds\_type**.

The application calls this IOCTL before the **STIOC\_GENERATE\_RAO** and **STIOC\_RECEIVE\_RAO** IOCTLs are issued. The application uses the return data to limit the number of UDS requested in the **GENERATE\_RAO** IOCTL.

The structure that is defined for this IOCTL is

```

struct query_rao_info{
    char uds_type;           /* [IN]   0: UDS_WITHOUT_GEOMETRY   */
                           /*          1: UDS_WITH_GEOMETRY      */
    char reserved[7];
    ushort max_uds_number;  /* [OUT] Max UDS number supported from drive */
    ushort max_uds_size;    /* [OUT] Max single UDS size supported from   */
                           /*          drive in byte          */
    ushort max_host_uds_number; /* [OUT] Max UDS number supported from driver */
};

```

An example of the **QUERY\_RAO\_INFO** command is

```

#include <sys/IBMtape.h>
int rc;
struct query_rao_info stQueryRao;

bzero( (void *) &stQueryRao, sizeof(struct query_rao_info));
stQueryRao.udt_type = uds_type;
rc = ioctl(fd, STIOC_QUERY_RAO, &stQueryRao);

if(rc)
    printf("STIOC_QUERY_RAO fails with rc: %d\n", rc);
else{
    max_host_uds_num = stQueryRao.max_host_uds_number;
    max_uds_size = stQueryRao.max_uds_size;
}
return rc;

```

## **STIOC\_GENERATE\_RAO**

The IOCTL is called to send a **GRAO** list to request that the drive generate a **Recommended Access Order** list. The process method is either 1 or 2 to create a **RAO** list, and the type of UDS is either with or without the geometry. The **uds\_number** must be not larger than **max\_host\_uds\_number** in the **STIOC\_QUERY\_RAO** IOCTL. The application allocates a block of memory with **graol\_list\_leng** (**uds\_number**\***sizeof(struct grao\_uds\_desc)**+8) for the pointer of **graol\_list**.

The structure for the **STIOC\_GENERATE\_RAO** IOCTL is

```
struct generate_rao {
    char process;          /* [IN] Requested process to generate RAO list */
    /* 0: no reorder UDS and no calculate */
    /* locate time(not currently supported */
    /* by the drive) */
    /* 1: no reorder UDS but calculate locate */
    /* time */
    /* 2: reorder UDS and calculate locate time*/
    char uds_type;         /* [IN] 0: UDS_WITHOUT_GEOMETRY */
    /* 1: UDS_WITH_GEOMETRY */
    char reserved1[2];
    uint grao_list_leng;   /* [IN] The data length is
allocated for GRAO */                                */
    list
    char *grao_list;       /* [IN] the pointer is allocated to the size */
    /* of grao_list_leng (uds_number */
    /* * sizeof(struct grao_uds_desc) */
    /* + sizeof(struct grao_list_header)) */
    /* and contains the data of GRAO */
    /* parameter list. The uds number is */
    /* less than max_host_uds_number in */
    /* QUERY_RAO ioctl. */
    char reserved2[8];
};
```

The **grao** list header and UDS segments make up the parameter data and are to be put in the following order.

```
-- List Header
-- UDS Segment Descriptor (first)
.....
-- UDS Segment Descriptor (last)
```

The device driver does not supply the header or UDS segment Descriptor structures. That structure is to be supplied by the application.

Examples of the data structures are

```
struct grao_list_header{
    uchar reserved[4];
    char addl_data[4]; /* additional data */           */
} ;

struct grao_uds_desc{
    uchar desc_leng[2]; /* descriptor length */        */
    char reserved[3];
    char uds_name[10]; /* uds name given by application */
    uchar partition; /* Partition number 0-n to overwrite */
    uchar beginning_loi[8]; /* Beginning logical object ID */
    uchar ending_loi[8]; /* Ending logical object ID */
};
```

A sample of **STIOC\_GENERATE\_RAO** is

```
#include<sys/IBM_tape.h>

int rc;
struct generate_rao grao;
bzero(&grao,sizeof(struct generate_rao));
grao.process=2;
grao.udt_type=uds_type;
grao.grao_list_leng=max_host_uds_num * sizeof(struct grao_uds_desc)
+ sizeof(struct grao_list_header);
if(!(grao.grao_list=malloc(grao.grao_list_leng)))
{
    perror("Failure allocating memory");
    return (errno);
}

memset(grao.grao_list, 0, grao.grao_list_leng);
rc=ioctl(fd,GENERATE_RAO,&grao);

if (rc)
```

```

    printf("GENERATE_RAO fails with rc%d\n",rc);
else
    printf("GENERATE_RAO succeeds\n");

free(rrao.rrao_list);

return rc;

```

## ***STIOC\_RECEIVE\_RAO***

After a **STIOC\_GENERATE\_RAO** IOCTL is completed, the application calls the **STIOC\_RECEIVE\_RAO** IOCTL to receive a recommended access order of UDS from the drive. To avoid a system crash, it is important that the application allocates a large enough block of memory for the **\*rrao\_list** pointer and notifies the driver of the allocated size. It is done by indicating the size of the buffer in bytes to the *rrao\_list\_leng* variable as an input to the **receive\_rao\_list** structure.

The structure for the **STIOC\_RECEIVE\_RAO** IOCTL is

```

struct receive_rao_list {
    uint     rrao_list_offset; /* [IN] The offset of receive RAO list to */
                           /* begin returning data */
    uint     rrao_list_leng;  /* [IN/OUT] number byte of data length */
                           /* [IN] The data length is allocated for RRAO */
                           /* list by application the length is */
                           /* (max_uds_size * uds_number + */
                           /* sizeof(struct rrao_list_header) */
                           /* max_uds_size is reported in */
                           /* sizeof(struct rrao_list_header) */
                           /* uds_number is the total UDS number */
                           /* requested from application in */
                           /* GENERATE_RAO ioctl */
                           /* [OUT] the data length is actual returned */
                           /* in RRAO list from the driver */
                           /* [IN/OUT] the data pointer of RRAO list */
    char    *rrao_list;
    char    reserved[8];
};

```

The sample code is

```

#include <sys/IBMtape.h>

int rc;
struct receive_rao_list rrao;

bzero(&rrao,sizeof(struct receive_rao_list));
rrao.rrao_list_offset=0;
rrao.rrao_list_leng= max_host_uds_num * max_uds_size + 8;
/* 8 is the header of rrao list */

if (!(rrao.rrao_list=malloc(rrao.rrao_list_leng)))
{
    perror("Failure allocating memory");
    return (errno);
}

memset(rrao.rrao_list, 0, rrao.rrao_list_leng);

rc=ioctl(fd,STIOC_RECEIVE_RAO,&rrao);

if (rc)
    printf("STIOC_RECEIVE_RAO fails with rc %d\n",rc);
else
    printf("rrao_list_leng %d\n",rrao.rrao_list_leng);

free(rrao.rrao_list);

return rc;

```

## ***STIOC\_SET\_SPDEV***

With the latest lin\_tape versions, the IBMSpecial device is created. It allows the use of ioctls for preemption purposes. Applications must use it cautiously and manage persistent reservation properly.

This ioctl is for usage through IBMSpecial open handle only. It sets the drive that processes the command requests, and to do so it needs the serial number of the drive as input. If /dev/IBMSpecial is not created, it is not supported.

The data structure is

```
#define DD_MAX_DEVICE_SERIAL      36
struct sp_dev{
    char device_serial[DD_MAX_DEVICE_SERIAL];
};
```

An example of the **STIOC\_SET\_SPDEV** command is

```
#include <sys/IBM_tape.h>
struct sp_dev spd;

setDriveSN(spd.device_serial);
if (!ioctl (fd, STIOC_SET_SPDEV, &spd)) {
printf ("The STIOC_SET_SPDEV ioctl succeeded\n");
}
else {
perror ("The STIOC_SET_SPDEV ioctl failed the drive to work with was not set");
}
```

When the **STIOC\_SET\_SPDEV** ioctl succeeds, it is possible to send any of these ioctls to the drive previously set and identified by serial number: **STIOC\_READ\_RESERVEKEYS**, **STIOC\_READ\_RESERVATIONS**, **STIOC\_REGISTER\_KEY**, **STIOC\_REMOVE\_REGISTRATION**, **STIOC\_CLEAR\_ALL\_REGSITRATION**.

## Tape drive compatibility IOCTL operations

The following IOCTL commands help provide compatibility for previously compiled programs. Where practical, such programs must be recompiled to use the preferred IOCTL commands in the Lin\_tape device driver.

### MTIOCTOP

This IOCTL command is similar in function to the **st MTIOCTOP** command. It is provided as a convenience for precompiled programs that call that IOCTL command. Refer to **/usr/include/sys/mtio.h** or **/usr/include/linux/mtio.h** for information on the **MTIOCTOP** command.

### MTIOCGET

This IOCTL command is similar in function to the **st MTIOCGET** command. It is provided as a convenience for precompiled programs that call that IOCTL command. Refer to **/usr/include/sys/mtio.h** or **/usr/include/linux/mtio.h** for information on the **MTIOCGET** command.

### MTIOCPOS

This IOCTL command is similar in function to the **st MTIOCPOS** command. It is provided as a convenience for precompiled programs that call that IOCTL command. Refer to **/usr/include/sys/mtio.h** or **/usr/include/linux/mtio.h** for information on the **MTIOCPOS** command.

## Medium changer IOCTL operations

This chapter describes the IOCTL commands that provide access and control of the SCSI medium changer functions. These IOCTL operations can be issued to the medium changer special file, such as **IBMchanger0**.

The following IOCTL commands are supported.

**SMCIOC\_ELEMENT\_INFO**

Obtain the device element information.

**SMCIOC\_MOVE\_MEDIUM**

Move a cartridge from one element to another element.

**SMCIOC\_EXCHANGE\_MEDIUM**

Exchange a cartridge in an element with another cartridge.

**SMCIOC\_POS\_TO\_ELEM**

Move the robot to an element.

**SMCIOC\_INIT\_ELEM\_STAT**

Issue the **SCSI Initialize Element Status** command.

**SMCIOC\_INIT\_ELEM\_STAT\_RANGE**

Issue the **SCSI Initialize Element Status with Range** command.

**SMCIOC\_INVENTORY**

Return the information about the four element types.

**SMCIOC\_LOAD\_MEDIUM**

Load a cartridge from a slot into the drive.

**SMCIOC\_UNLOAD\_MEDIUM**

Unload a cartridge from the drive and return it to a slot.

**SMCIOC\_PREVENT\_MEDIUM\_REMOVAL**

Prevent medium removal by the operator.

**SMCIOC\_ALLOW\_MEDIUM\_REMOVAL**

Allow medium removal by the operator.

**SMCIOC\_READ\_ELEMENT\_DEVIDS**

Return the device id element descriptors for drive elements.

## SCSI IOCTL commands

These IOCTL commands and their associated structures are defined in the **IBM\_tape.h** header file, which can be found in /usr/include/sys after Lin\_tape is installed. The **IBM\_tape.h** header file is included in the corresponding C program by using the functions.

**SMCIOC\_ELEMENT\_INFO**

This IOCTL command obtains the device element information.

The data structure is

```
struct element_info {  
    ushort robot_addr; /* first robot address */  
    ushort robots; /* number of medium transport elements */  
    ushort slot_addr; /* first medium storage element address */  
    ushort slots; /* number of medium storage elements */  
    ushort ie_addr; /* first import/export element address */  
    ushort ie_stations; /* number of import/export elements */  
    ushort drive_addr; /* first data-transfer element address */  
    ushort drives; /* number of data-transfer elements */  
};
```

An example of the **SMCIOC\_ELEMENT\_INFO** command is

```
#include <sys/IBM_tape.h>  
struct element_info element_info;  
if (!ioctl (smcfid, SMCIOC_ELEMENT_INFO, &element_info)) {  
    printf ("The SMCIOC_ELEMENT_INFO ioctl succeeded\n");  
    printf ("\nThe element information data is:\n");  
    dump_bytes ((unchar *) &element_info, sizeof (struct element_info));  
}  
else {  
    perror ("The SMCIOC_ELEMENT_INFO ioctl failed");  
}
```

```
    smcioc_request_sense();  
}
```

## **SMCIOC\_MOVE\_MEDIUM**

This IOCTL command moves a cartridge from one element to another element.

The data structure is

```
struct move_medium {  
    ushort robot;      /* robot address */  
    ushort source;    /* move from location */  
    ushort destination; /* move to location */  
    char invert;      /* invert before placement bit */  
};
```

An example of the **SMCIOC\_MOVE\_MEDIUM** command is

```
#include <sys/IBM_tape.h>  
struct move_medium move_medium;  
move_medium.robot = 0;  
move_medium.invert = 0;  
move_medium.source = source;  
move_medium.destination = dest;  
if (!ioctl (smcfd, SMCIOC_MOVE_MEDIUM, &move_medium))  
    printf ("The SMCIOC_MOVE_MEDIUM ioctl succeeded\n");  
else {  
    perror ("The SMCIOC_MOVE_MEDIUM ioctl failed");  
    smcioc_request_sense();  
}
```

## **SMCIOC\_EXCHANGE\_MEDIUM**

This IOCTL command exchanges a cartridge in an element with another cartridge. This command is equivalent to two **SCSI Move Medium** commands. The first moves the cartridge from the source element to the **destination1** element. The second moves the cartridge that was previously in the **destination1** element to the **destination 2** element. This function is available only in the IBM 3584 UltraScalable tape library. The **destination2** element can be the same as the source element.

The input data structure is

```
struct exchange_medium {  
    ushort robot;      /* robot address */  
    ushort source;    /* source address for exchange */  
    ushort destination1; /* first destination address for exchange */  
    ushort destination2; /* second destination address for exchange */  
    char invert1;      /* invert before placement into destination1 */  
    char invert2;      /* invert before placement into destination2 */  
};
```

An example of the **SMCIOC\_EXCHANGE\_MEDIUM** command is

```
#include <sys/IBM_tape.h>  
struct exchange_medium exchange_medium;  
exchange_medium.robot = 0;  
exchange_medium.invert1 = 0;  
exchange_medium.invert2 = 0;  
exchange_medium.source = 32;      /* slot 32 */  
exchange_medium.destination1 = 16; /* drive address 16 */  
exchange_medium.destination2 = 35; /* slot 35 */  
  
/* exchange cartridge in drive address 16 with cartridge from */  
/* slot 32 and return the cartridge currently in the drive to */  
/* slot 35 */  
if (!ioctl (smcfd, SMCIOC_EXCHANGE_MEDIUM, &exchange_medium))  
    printf("The SMCIOC_EXCHANGE_MEDIUM ioctl succeeded\n");  
else {  
    perror ("The SMCIOC_EXCHANGE_MEDIUM ioctl failed");  
    smcioc_request_sense();  
}
```

## **SMCIOC\_POS\_TO\_ELEM**

This IOCTL command moves the robot to an element.

The input data structure is

```
struct pos_to_elem {  
    ushort robot; /* robot address */  
    ushort destination; /* move to location */  
    char invert; /* invert before placement bit */  
};
```

An example of the **SMCIOC\_POS\_TO\_ELEM** command is

```
#include <sys/IBM_tape.h>  
struct pos_to_elem pos_to_elem;  
pos_to_elem.robot = 0;  
pos_to_elem.invert = 0;  
pos_to_elem.destination = dest;  
if (!ioctl (smcfd, SMCIOC_POS_TO_ELEM, &pos_to_elem))  
    printf ("The SMCIOC_POS_TO_ELEM ioctl succeeded\n");  
else {  
    perror ("The SMCIOC_POS_TO_ELEM ioctl failed");  
    smcioc_request_sense();  
}
```

## **SMCIOC\_INIT\_ELEM\_STAT**

This IOCTL command instructs the medium changer robotic device to issue the **SCSI Initialize Element Status** command.

There is no associated data structure.

An example of the **SMCIOC\_INIT\_ELEM\_STAT** command is

```
#include <sys/IBM_tape.h>  
if (!ioctl (smcfd, SMCIOC_INIT_ELEM_STAT, NULL))  
    printf ("The SMCIOC_INIT_ELEM_STAT ioctl succeeded\n");  
else {  
    perror ("The SMCIOC_INIT_ELEM_STAT ioctl failed");  
    smcioc_request_sense();  
}
```

## **SMCIOC\_INIT\_ELEM\_STAT\_RANGE**

This IOCTL command issues the **SCSI Initialize Element Status with Range** command and audits specific elements in a library by specifying the starting element address and number of elements. Use the **SMCIOC\_INIT\_ELEM\_STAT** IOCTL to audit all elements.

The data structure is

```
struct element_range {  
    ushort element_address; /* starting element address */  
    ushort number_elements; /* number of elements */  
};
```

An example of the **SMCIOC\_INIT\_ELEM\_STAT\_RANGE** command is

```
#include <sys/IBM_tape.h>  
struct element_range elements;  
/* audit slots 32 to 36 */  
elements.element_address = 32;  
elements.number_elements = 5;  
if (!ioctl (smcfd, SMCIOC_INIT_ELEM_STAT_RANGE, &elements))  
    printf ("The SMCIOC_INIT_ELEM_STAT_RANGE ioctl succeeded\n");  
else {  
    perror ("The SMCIOC_INIT_ELEM_STAT_RANGE ioctl failed");  
    smcioc_request_sense();  
}
```

**Note:** Use the **SMCIOG\_INVENTORY** IOCTL command to obtain the current version after this IOCTL command is issued.

## **SMCIOC\_INVENTORY**

This IOCTL command returns the information about the four element types. The software application processes the input data (the number of elements about which it requires information). Then, it allocates a buffer large enough to hold the output for each element type.

The input data structure is

```
struct element_status {
    ushort address; /* element address */
    uint reserved1 :2; /* reserved */
    inenab :1; /* media into changer's scope */
    exenab :1; /* media out of changer's scope */
    access :1; /* robot access allowed */
    except :1; /* abnormal element state */
    impexp :1; /* import/export placed by operator or robot */
    full :1; /* element contains medium */
    uchar resvd1; /* reserved */
    uchar asc; /* additional sense code */
    uchar ascq; /* additional sense code qualifier */
    uint notbus :1; /* element not on same bus as robot */
    :1; /* reserved */
    idvalid :1; /* element address valid */
    luvalid :1; /* logical unit valid */
    :1; /* reserved */
    lun :3; /* logical unit number */
    uchar scsi; /* SCSI bus address */
    uchar resvd2; /* reserved */
    uint svalid :1; /* element address valid */
    invert :1; /* medium inverted */
    :6; /* reserved */
    ushort source; /* source storage element address */
    uchar volume[36]; /* primary volume tag */
    uchar resvd3[4]; /* reserved */
};

struct inventory {
    struct element_status *robot_status; /* medium transport elem pgs */
    struct element_status *slot_status; /* medium storage elem pgs */
    struct element_status *ie_status; /* import/export elem pgs */
    struct element_status *drive_status; /* data-transfer elem pgs */
};
```

An example of the **SMCIOC\_INVENTORY** command is

```
#include <sys/IBM_tape.h>
ushort i;
struct element_info element_info;
struct element_status robot_status[1];
struct element_status slot_status[20];
struct element_status ie_status[1];
struct element_status drive_status[1];
struct inventory inventory;
bzero((caddr_t)robot_status,sizeof(struct element_status));
for (i=0;i<20;i++)
    bzero((caddr_t)&slot_status[i],sizeof(struct element_status));
bzero((caddr_t)&ie_status[sizeof(struct element_status)]);
bzero((caddr_t)&drive_status[sizeof(struct element_status)]);
smcioc_element_info(&element_info);
inventory.robot_status = robot_status;
inventory.slot_status = slot_status;
inventory.ie_status = ie_status;
inventory.drive_status = drive_status;
if (!ioctl (smcfid, SMCIOC_INVENTORY, &inventory)) {
    printf ("\nThe SMCIOC_INVENTORY ioctl succeeded\n");
    printf ("\nThe robot status pages are:\n");
    for (i = 0; i<element_info.robots; i++) {
        dump_bytes ((unchar *)(&robot_status[i]), sizeof (struct
element_status));
        printf ("\n--- more ---");
        getchar();
    }
    printf ("\nThe slot status pages are:\n");
    for (i = 0; i<element_info.slots; i++) {
        dump_bytes ((unchar *)(&slot_status[i]), sizeof (struct
```

```

        element_status));
        printf ("\n--- more ---");
        getchar();
    }
    printf ("\nThe ie status pages are:\n");
    for (i = 0; i<element_info.ie_stations; i++) {
        dump_bytes ((unchar *)(&ie_status[i]), sizeof (struct
element_status));
        printf ("\n--- more ---");
        getchar();
    }
    printf ("\nThe drive status pages are:\n");
    for (i = 0; i<element_info.drives; i++) {
        dump_bytes ((unchar *)(&drive_status[i]), sizeof (struct element_status));
        printf ("\n--- more ---");
        getchar();
    }
}
else {
    perror ("The SMCIOC_INVENTORY ioctl failed");
    smcioc_request_sense();
}

```

### **SMCIOC\_LOAD\_MEDIUM**

This IOCTL command loads a tape from a specific slot into the drive. Or, it loads from the first full slot into the drive if the slot address is specified as zero.

An example of the **SMCIOC\_LOAD\_MEDIUM** command is

```

#include <sys/IBM_tape.h>
/* load cartridge from slot 3 */
if (ioctl (tapefd, SMCIOC_LOAD_MEDIUM,3)) {
    printf ("IOCTL failure. errno=%d\n",errno);
    exit(1);
}
/* load first cartridge from magazine */
if (ioctl (tapefd, SMCIOC_LOAD_MEDIUM,0)) {
    printf ("IOCTL failure. errno=%d\n",errno);
    exit(1);
}

```

### **SMCIOC\_UNLOAD\_MEDIUM**

This IOCTL command moves a tape from the drive and returns it to a specific slot. Or, it moves a tape to the first empty slot in the magazine if the slot address is specified as zero. An **unload/offline** command must be sent to the tape first, otherwise, this IOCTL command fails with *errno* EIO.

An example of the **SMCIOC\_UNLOAD\_MEDIUM** command is

```

#include <sys/IBM_tape.h>
/* unload cartridge to slot 3 */
if (ioctl (tapefd, SMCIOC_UNLOAD_MEDIUM,3)) {
    printf ("IOCTL failure. errno=%d\n",errno);
    exit(1);
}
/* unload cartridge to first empty slot in magazine */
if (ioctl (tapefd, SMCIOC_UNLOAD_MEDIUM,0)) {
    printf ("IOCTL failure. errno=%d\n",errno);
    exit(1);
}

```

### **SMCIOC\_PREVENT\_MEDIUM\_REMOVAL**

This IOCTL command prevents an operator from removing medium from the device until the **SMCIOC\_ALLOW\_MEDIUM\_REMOVAL** command is issued or the device is reset. There is no associated data structure.

An example of the **SMCIOC\_PREVENT\_MEDIUM\_REMOVAL** command is

```

#include <sys/IBM_tape.h>
if (!ioctl (smcfid, SMCIOC_PREVENT_MEDIUM_REMOVAL, NULL))
    printf ("The SMCIOC_PREVENT_MEDIUM_REMOVAL ioctl succeeded\n");

```

```

else {
    perror ("The SMCIOC_PREVENT_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}

```

## **SMCIOC\_ALLOW\_MEDIUM\_REMOVAL**

This IOCTL command allows an operator to remove medium from the device. This command is normally used after an **SMCIOC\_PREVENT\_MEDIUM\_REMOVAL** command to restore the device to the default state. There is no associated data structure.

An example of the **SMCIOC\_ALLOW\_MEDIUM\_REMOVAL** command is

```

#include <sys/IBM_tape.h>
if (!ioctl (smcfd, SMCIOC_ALLOW_MEDIUM_REMOVAL, NULL))
    printf ("The SMCIOC_ALLOW_MEDIUM_REMOVAL ioctl succeeded\n");
else {
    perror ("The SMCIOC_ALLOW_MEDIUM_REMOVAL ioctl failed");
    smcioc_request_sense();
}

```

## **SMCIOC\_READ\_ELEMENT\_DEVIDS**

This IOCTL command issues the **SCSI Read Element Status** command with the device ID(DVCID) bit set and returns the element descriptors for the data transfer elements. The **element\_address** field specifies the starting address of the first data transfer element. The **number\_elements** field specifies the number of elements to return. The application must allocate a return buffer large enough for the number of elements that are specified in the input structure.

The input data structure is

```

struct read_element_devids {
    ushort element_address;           /* starting element address */
    ushort number_elements;          /* number of elements */
    struct element_devid *drive_devid; /* data transfer element pages */
};

```

The output data structure is

```

struct element_devid {
    ushort address;                  /* element address */ */
    uint      :4,                   /* reserved */ */
    access   :1,                   /* robot access allowed */ */
    except   :1,                   /* abnormal element state */ */
    :1,                      /* reserved */ */
    full    :1;                   /* element contains medium */ */
    uchar resvd1;                 /* reserved */ */
    uchar asc;                     /* additional sense code */ */
    uchar ascq;                    /* additional sense code qualifier */ */
    uint notbus :1,                /* element not on same bus as robot */ */
    :1,                      /* reserved */ */
    idvalid :1,                   /* element address valid */ */
    luvalid :1,                   /* logical unit valid */ */
    :1,                      /* reserved */ */
    lun     :3;                   /* logical unit number */ */
    uchar scsi;                   /* scsi bus address */ */
    uchar resvd2;                 /* reserved */ */
    uint svalid :1,                /* element address valid */ */
    invert  :1,                   /* medium inverted */ */
    :6;                      /* reserved */ */
    ushort source;                /* source storage element address */ */
    uint      :4,                   /* reserved */ */
    code_set :4;                  /* code set X'2' is all ASCII identifier*/ */
    uint      :4,                   /* reserved */ */
    ident_type :4;                /* identifier type */ */
    uchar resvd3;                 /* reserved */ */
    uchar ident_len;              /* identifier length */ */
    uchar identifier[36];          /* device identification */ */
};

```

An example of the **SMCIOC\_READ\_ELEMENT\_DEVIDS** command is

```
#include <sys/IBM_tape.h>
int smcioc_read_element_devids() {
int i;
struct element_devid *elem_devid, *elemp;
struct read_element_devids devids;
struct element_info element_info;
if (ioctl(fd, SMCIOC_ELEMENT_INFO, &element_info)) return errno;
if (element_info.drives) {
    elem_devid = malloc(element_info.drives
        * sizeof(struct element_devid));
    if (elem_devid == NULL) {
        errno = ENOMEM;
        return errno;
    }
    bzero((caddr_t)elem_devid,element_info.drives
        * sizeof(struct element_devid));
    devids.drive_devid = elem_devid;
    devids.element_address = element_info.drive_addr;
    devids.number_elements = element_info.drives;
    printf("Reading element device ids?\n");
    if (ioctl (fd, SMCIOC_READ_ELEMENT_DEVIDS, &devids)) {
        free(elem_devid);
        return errno;
    }
    elemp = elem_devid;
    for (i = 0; i<element_info.drives; i++, elemp++) {
        printf("\nDrive Address %d\n",elemp->address);
        if (elemp->except)
            printf(" Drive State ..... Abnormal\n");
        else
            printf(" Drive State ..... Normal\n");
        if (elemp->asc == 0x81 && elemp->ascq ==0x00)
            printf(" ASC/ASCQ ..... %02X%02X (Drive Present)\n",
                elemp->asc,elemp->ascq);
        else if (elemp->asc == 0x82 && elemp->ascq ==0x00)
            printf(" ASC/ASCQ ..... %02X%02X (Drive Not Present)\n",
                elemp->asc,elemp->ascq);
        else
            printf(" ASC/ASCQ ..... %02X%02X\n",
                elemp->asc,elemp->ascq);
        if (elemp->full)
            printf(" Media Present ..... Yes\n");
        else
            printf(" Media Present ..... No\n");
        if (elemp->access)
            printf(" Robot Access Allowed ..... Yes\n");
        else
            printf(" Robot Access Allowed ..... No\n");
        if (elemp->svalid)
            printf(" Source Element Address ..... %d\n",
                elemp->source);
        else
            printf(" Source Element Address Valid ..... No\n");
        if (elemp->invert)
            printf(" Media Inverted ..... Yes\n");
        else
            printf(" Media Inverted ..... No\n");
        if (elemp->notbus)
            printf(" Same Bus as Medium Changer ..... No\n");
        else
            printf(" Same Bus as Medium Changer ..... Yes\n");
        if (elemp->idvalid)
            printf(" SCSI Bus Address ..... %d\n",
                elemp->scsi);
        else
            printf(" SCSI Bus Address Valid ..... No\n");
        if (elemp->lunvalid)
            printf(" Logical Unit Number ..... %d\n",
                elemp->lun);
        else
            printf(" Logical Unit Number Valid ..... No\n");
        printf(" Device ID ..... %0.36s\n",
            elemp->identifier);
    }
    else {
        printf("\nNo drives found in element information\n");
    }
    free(elem_devid);
    return errno;
}
```

## Return codes

This chapter describes error codes that are generated by Lin\_tape when an error occurs during an operation. On error, the operation returns negative one (-1), and the external variable *errno* is set to one of the listed error codes. *Errno* values are defined in **/usr/include/errno.h** (and other files that it includes). Application programs must include **errno.h** to interpret the return codes.

**Note:** For error code EIO, an application can retrieve more information from the device itself. Issue the **STIOCQRYSENSE** IOCTL command when the **sense\_type** equals **LASTERROR**, or the **SIOC\_REQSENSE** IOCTL command, to retrieve sense data. Then, analyze the sense data by using the appropriate hardware or SCSI reference for that device.

### General error codes

The following codes apply to all operations.

[EBUSY]	An excessively busy state was encountered in the device.
[EFAULT]	A memory failure occurred due to an invalid pointer or address.
[EIO]	An error due to one of the following conditions: <ul style="list-style-type: none"><li>• An unrecoverable media error was detected by the device.</li><li>• The device was not ready for operation or a tape was not in the drive.</li><li>• The device did not respond to SCSI selection.</li><li>• A bad file descriptor was passed to the device.</li></ul>
[ENOMEM]	Insufficient memory was available for an internal memory operation.
[ENXIO]	The device was not configured and is not receiving requests.
[EPERM]	The process does not have permission to run the desired function.
[ETIMEDOUT]	A command timed out in the device.

### Open error codes

The following codes apply to **open** operations.

[EACCES]	The <b>open</b> requires write access when the cartridge loaded in the drive is physically write-protected.
[EAGAIN]	The device was already open when an <b>open</b> was attempted.
[EBUSY]	The device was reserved by another initiator or an excessively busy state was encountered.
[EINVAL]	The operation that is requested has invalid parameters or an invalid combination of parameters, or the device is rejecting <b>open</b> commands.
[EIO]	An I/O error occurred that indicates a failure to operate the device. Run failure analysis.
[ENOMEM]	Insufficient memory was available for an internal memory operation.
[EPERM]	One of the following situations occurred: <ul style="list-style-type: none"><li>• An <b>open</b> operation with the <b>O_RDWR</b> or <b>O_WRONLY</b> flag was attempted on a write-protected tape.</li><li>• A write operation was attempted on a device that was opened with the <b>O_RDONLY</b> flag.</li></ul>

## Close error codes

The following codes apply to **close** operations.

[EBUSY]	The SCSI subsystem was busy.
[EFAULT]	Memory reallocation failed.
[EIO]	A command that is issued during <b>close</b> , such as a <b>rewind</b> command, failed because the device was not ready. An I/O error occurred during the operation. Run failure analysis.

## Read error codes

The following codes apply to **read** operations.

[EFAULT]	Failure copying from user to kernel space or vice versa.
[EINVAL]	One of the following situations occurred: <ul style="list-style-type: none"><li>• The operation that is requested has invalid parameters or an invalid combination of parameters.</li><li>• The number of bytes requested in the <b>read</b> operation was not a multiple of the block size for a fixed block transfer.</li><li>• The number of bytes requested in the <b>read</b> operation was greater than the maximum size allowed by the device for variable block transfers.</li><li>• A read for multiple fixed odd-byte-count blocks was issued.</li></ul>
[ENOMEM]	One of the following situations occurred: <ul style="list-style-type: none"><li>• The number of bytes requested in the <b>read</b> operation of a variable block record was less than the size of the block. This error is known as an overlength condition.</li><li>• Insufficient memory was available for an internal memory operation.</li></ul>
[EPERM]	A <b>read</b> operation was attempted on a device that was opened with the <b>O_WRONLY</b> flag.

## Write error codes

The following codes apply to **write** operations.

[EFAULT]	Failure copying from user to kernel space or vice versa.
[EINVAL]	One of the following conditions occurred: <ul style="list-style-type: none"><li>• The operation that is requested has invalid parameters or an invalid combination of parameters.</li><li>• The number of bytes requested in the <b>write</b> operation was not a multiple of the block size for a fixed block transfer.</li><li>• The number of bytes requested in the <b>write</b> operation was greater than the maximum block size allowed by the device for variable block transfers.</li></ul>
[EIO]	The physical end of the medium was detected, or it is a general error that indicates a failure to write to the device. Perform failure analysis.
[ENOMEM]	Insufficient memory was available for an internal memory operation.

[ENOSPC] A **write** operation failed because it reached the early warning mark. This error code is returned only one time when the early warning is reached and **trailer\_labels** is set to true. A **write** operation was attempted after the device reached the logical end of the medium and **trailer\_labels** were set to false.

[EPERM] A **write** operation was attempted on a write protected tape.

## IOCTL error codes

The following codes apply to IOCTL operations.

[EBUSY] SCSI subsystem was busy.

[EFAULT] Failure copying from user to kernel space or vice versa.

[EINVAL] The operation that is requested has invalid parameters or an invalid combination of parameters. This error code also results if the IOCTL command is not supported by the device. For example, if you are attempting to issue tape drive IOCTL commands to a SCSI medium changer. An invalid or nonexistent IOCTL command was specified.

[EIO] An I/O error occurred during the operation. Run failure analysis.

[ENOMEM] Insufficient memory was available for an internal memory operation.

[ENOSYS] The underlying function for this IOCTL command does not exist on this device. (Other devices might support the function.)

[EPERM] An operation that modifies the media was attempted on a write-protected tape or a device that was opened with the **O\_RDONLY** flag.

## Windows tape and medium changer device drivers

---

### Windows programming interface

The programming interface conforms to the standard Microsoft Windows Server tape device drivers interface. It is detailed in the Microsoft Developer Network (MSDN) Software Development Kit (SDK) and Windows Driver Kit (WDK). Common documentation for these similar devices are indicated by 200x.

Windows IBMTape is conformed by two sets of device drivers,

- **ibmtpxxx.sys**, which supports the IBM Tape Storage tape drives, where
  - **ibmtp.sys**, **ibmtpbs.sys**, **ibmtpft.sys**.
- **ibmcgxxx.sys**, which supports the IBM Tape Storage medium changer, where
  - **ibmcg.sys**, **ibmcgbs.sys**, **ibmcgft.sys**.

The programming interface conforms to the standard Microsoft Windows 200x tape device driver interface. It is detailed in the Microsoft Developer Network (MSDN) Software Development Kit (SDK), and Windows Development Kit (WDK).

### User-callable entry points

The following user-callable tape driver entry points are supported under **ibmtpxxx.sys**.

- CreateFile
- CloseHandle
- DeviceIoControl

- EraseTape
- GetTapeParameters
- GetTapePosition
- GetTapeStatus
- PrepareTape
- ReadFile
- SetTapeParameters
- SetTapePosition
- WriteFile
- WriteTapemark

## Tape Media Changer driver entry points

If the Removable Storage Manager is stopped, then the following user-callable tape media changer driver entry points are supported under **ibmcgxxx.sys**:

- CreateFile
- CloseHandle
- DeviceIoControl

Users who want to write application programs to issue commands to IBM Tape Storage device drivers must obtain a license to the MSDN and the Microsoft Visual C++ Compiler. Users also need access to IBM hardware reference manuals for IBM Tape Storage devices.

Programs that access the IBM Tape Storage device driver must complete the following steps:

1. Include the following files in the application.

```
#include <ntddscsi.h>
#include <ntddchgr.h>
#include <ntddtape.h> /* Modified as indicated below */
```

2. Add the following lines to **ntddtape.h**.

```
#define LB_ACCESS FILE_READ_ACCESS | FILE_WRITE_ACCESS
#define M_MTI(x) CTL_CODE(IOCTL_BASE+2,x,METHOD_BUFFERED, LB_ACCESS)
#define IOCTL_TAPE_OBTAIN_SENSE CTL_CODE(IOCTL_TAPE_BASE, 0x0819,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_OBTAIN_VERSION CTL_CODE(IOCTL_TAPE_BASE, 0x081a,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_LOG_SELECT CTL_CODE(IOCTL_TAPE_BASE, 0x081c,
METHOD_BUFFERED, FILE_READ_ACCESS | FILE_WRITE_ACCESS)
#define IOCTL_TAPE_LOG_SENSE CTL_CODE(IOCTL_TAPE_BASE, 0x081d,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_LOG_SENSE10 CTL_CODE(IOCTL_TAPE_BASE, 0x0833,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_ENH_TAPE_LOG_SENSE10 CTL_CODE(IOCTL_TAPE_BASE, 0x0835, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define IOCTL_TAPE_REPORT_MEDIA_DENSITY CTL_CODE(IOCTL_TAPE_BASE, 0x081e,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_OBTAIN_MTDEVICE (M_MTI(16))
#define IOCTL_CREATE_PARTITION CTL_CODE(IOCTL_TAPE_BASE, 0x0826, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_QUERY_PARTITION CTL_CODE(IOCTL_TAPE_BASE, 0x0825, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_SET_ACTIVE_PARTITION CTL_CODE(IOCTL_TAPE_BASE, 0x0827, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_QUERY_DATA_SAFE_MODE CTL_CODE(IOCTL_TAPE_BASE, 0x0823, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_SET_DATA_SAFE_MODE CTL_CODE(IOCTL_TAPE_BASE, 0x0824, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_ALLOW_DATA_OVERWRITE CTL_CODE(IOCTL_TAPE_BASE, 0x0828, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_SET_PEW_SIZE
CTL_CODE(IOCTL_TAPE_BASE, 0x082c, METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_QUERY_PEW_SIZE
CTL_CODE(IOCTL_TAPE_BASE, 0x082b, METHOD_BUFFERED, FILE_READ_ACCESS )
```

```

#define IOCTL_VERIFY_TAPE_DATA
CTL_CODE(IOCTL_TAPE_BASE, 0x082A, METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_QUERY_RAO_INFO    CTL_CODE(IOCTL_TAPE_BASE, 0x082E, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define IOCTL_GENERATE_RAO      CTL_CODE(IOCTL_TAPE_BASE, 0x082F, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define IOCTL_RECEIVE_RAO       CTL_CODE(IOCTL_TAPE_BASE, 0x0834, METHOD_BUFFERED,
FILE_READ_ACCESS )

```

## CreateFile

The **CreateFile** entry point is called to make the driver and device ready for input/output (I/O). Installing the driver in non-exclusive mode allows several handles to the same Tape Storage device. If the driver was installed in exclusive mode, by default only one active handle is allowed to a given Tape Storage device. However, if more than one handle is needed for the same device, the **dwCreationDisposition** parameter can be set to OPEN\_ALWAYS in CreateFile() function to create an extra handle. By using this flag the resulting handle has limited functions. The driver allows the following IOCTLs only:

```

IOCTL_STORAGE_PERSISTENT_RESERVE_IN
IOCTL_STORAGE_PERSISTENT_RESERVE_OUT
IOCTL_SCSI_PASS_THROUGH_DIRECT with Cdb[0] set to INQUIRY (0x12)

```

The following code fragment illustrates a call to the **CreateFile** routine:

```

HANDLE ddHandle0, ddHandle1; // file handle for LUN0 and LUN1

/*
** Open for reading/writing on LUN0,
** where the device special file name is in the form of tapex and
** x is the logical device 0 to n - can be determined from Registry
*/
/*
** Open for media mover operations on LUN1,
** where the device special file name is in the form of
** changerx and x is the logical device 0 to n - can be determined from Registry
*/

ddHandle0 = CreateFile(
    "\\\.\tape0",
    DWORD dwDesiredAccess,
    DWORD dwShareMode,
    LPSECURITY_ATTRIBUTES lpSecurityAttributes,
    DWORD dwCreationDisposition,
    DWORD dwFlagsAndAttributes,
    HANDLE hTemplateFile
);

ddHandle1 = CreateFile(
    "\\\.\changer0",
    DWORD dwDesiredAccess,
    DWORD dwShareMode,
    LPSECURITY_ATTRIBUTES lpSecurityAttributes,
    DWORD dwCreationDisposition,
    DWORD dwFlagsAndAttributes,
    HANDLE hTemplateFile
);

/* Print msg if open failed for handle 0 or 1 */
if(ddHandle1 == INVALID_HANDLE_VALUE)
{
    printf("open failed for LUNn\n");
    printf("System Error = %d\n",GetLastError());
    exit (-1);
}

```

The **CloseHandle** entry point is called to stop I/O to the driver and device. The following code fragment illustrates a call to the **CloseHandle** routine:

```

BOOL rc;

rc = CloseHandle(
    ddHandle0
);

if (!rc)
{
    printf("close failed\n");
    printf("System Error = %d\n",GetLastError());
    exit (-1);
}

```

```
}
```

where ddHandle0 is the open file handle returned by the CreateFile call.

## ReadFile

The **ReadFile** entry point is called to read data from tape. The caller provides a buffer address and length, and the driver returns data from the tape to the buffer. The amount of data that is returned never exceeds the length parameter.

See “[Variable and fixed block read/write processing](#)” on page 166 for a full discussion of the read/write processing feature.

The following code fragment illustrates a **ReadFile** call to the driver:

```
BOOL rc;  
rc = ReadFile(  
    hFile,  
    lpBuffer,  
    nBufferSize,  
    lpBytesRead,  
    lpOverlapped  
);  
if(rc)  
{  
    if (*lpBytesRead > 0)  
        printf("Read %d bytes\n", *lpBytesRead);  
    else  
        printf("Read found file mark\n");  
}  
else  
{  
    printf("Error on read\n");  
    printf("System Error = %d\n", GetLastError());  
    exit (-1);  
}
```

Where **hFile** is the open file handle, **lpBuffer** is the address of a buffer in which to place the data, **nBufferSize** is the number of bytes to be read, and **lpBytesRead** is the number of bytes read.

If the function succeeds, the return value **rc** is nonzero.

## WriteFile

The **WriteFile** entry point is called to write data to the tape. The caller provides the address and length of the buffer to be written to tape. The physical limitations of the drive can cause the write to fail. One example is attempting to write past the physical end of the tape.

See “[Variable and fixed block read/write processing](#)” on page 166 for a full discussion of the read/write processing feature.

The following code fragment illustrates a call to the **WriteFile** routine:

```
BOOL rc;  
rc = WriteFile(  
    hFile,  
    lpBuffer,  
    nBufferSize,  
    lpNumberOfBytesWritten,  
    lpOverlapped  
);  
if (!rc)  
{  
    printf("Error on write\n");  
    printf("System Error = %d\n", GetLastError());  
    exit (-1);  
}
```

Where **hFile** is the open file handle, **lpBuffer** is the buffer address, and **nBufferSize** is the size of the buffer in bytes.

If the function succeeds, the return value rc is nonzero. The application also verifies that all the requested data was written by examining the **lplpNumberofBytesWritten** parameter. See “[Write Tapemark](#)” on page 143 for details on committing data on the media.

## **Write Tapemark**

Application writers who are using the **WriteFile** entry point to write data to tape must understand that the tape device buffers data in its memory and writes that data to the media as those device buffers fill. Thus, a **WriteFile** call might return a successful return code, but the data might not be on the media yet. Calling the **WriteTapemark** entry point and receiving a good return code, however, ensures that data is committed to tape media properly if all previous **WriteFile** calls were successful. However, applications that write large amounts of data to tape might not want to wait until writing a tapemark to know whether previous data was written to the media properly. For example:

```
WriteTapemark(
    HANDLE hDevice,
    DWORD dwTapemarkType,
    DWORD dwTapemarkCount,
    BOOL bImmediate
);
```

**dwTapemarkType** is the type of operation requested.

The only type that is supported is

**TAPE\_FILEMARKS**

The **WriteTapemark** entry point might also be called with the **dwTapemarkCount** parameter set to 0 and the **bImmediate** parameter that is set to FALSE. This action commits any uncommitted data that is written by previous **WriteFile** calls (since the last call to **WriteTapemark**) to the media. If no error is returned by the **WriteFile** calls and the **WriteTapemark** call, the application can assume that all data is committed to the media successfully.

## **SetTapePosition**

The **SetTapePosition** entry point is called to seek to a particular block of media data. For example:

```
SetTapePosition(
    HANDLE hDevice,
    DWORD dwPositionMethod,
    DWORD dwPartition,
    DWORD dwOffsetLow,
    DWORD dwOffsetHigh,
    BOOL bImmediate
);
```

**dwPositionMethod** is the type of positioning.

For Tape Storage devices, the following types of tapemarks and immediate values are supported.

<b>TAPE_ABSOLUTE_BLOCK</b>	bImmediate TRUE or FALSE
<b>TAPE_LOGICAL_BLOCK</b>	bImmediate TRUE or FALSE

For Tape Storage devices, there is no difference between the absolute and logical block addresses.

<b>TAPE_REWIND</b>	bImmediate TRUE or FALSE
<b>TAPE_SPACE_END_OF_DATA</b>	bImmediate FALSE
<b>TAPE_SPACE_FILEMARKS</b>	bImmediate FALSE
<b>TAPE_SPACE_RELATIVE_BLOCKS</b>	bImmediate FALSE
<b>TAPE_SPACE_SEQUENTIAL_FMKS</b>	

## **GetTapePosition**

The **GetTapePosition** entry point is called to retrieve the current tape position. For example:

```
GetTapePosition(
    HANDLE hDevice,
    DWORD dwPositionType,
    LPDWORD lpdwPartition,
    LPDWORD lpdwOffsetLow,
    LPDWORD lpdwOffsetHigh
);
```

**dwPositionType** is the type of positioning.

**TAPE\_ABSOLUTE\_POSITION** or **TAPE\_LOGICAL\_POSITION** might be specified but only the absolute position is returned.

## **SetTapeParameters**

The **SetTapeParameters** entry point is called to either specify the block size of a tape or set tape device data compression. The data structures are

```
struct{ // structure used by operation SET_TAPE_MEDIA_INFORMATION
    ULONG BlockSize;
}TAPE_SET_MEDIA_PARAMETERS;

struct{ // structure used by operation SET_TAPE_DRIVE_INFORMATION
    BOOLEAN ECC;                      // Not Supported
    BOOLEAN Compression;              // Only compression can be set
    BOOLEAN DataPadding;              // Not Supported
    BOOLEAN ReportSetmarks;           // Not Supported
    ULONG EOTWarningZoneSize;          // Not Supported
}TAPE_SET_DRIVE_PARAMETERS;

SetTapeParameters(
    HANDLE hDevice,
    DWORD dwOperation,
    LPVOID lpParameters
);
```

**dwOperation** is the type of information to set (**SET\_TAPE\_MEDIA\_INFORMATION** or **SET\_TAPE\_DRIVE\_INFORMATION**). For **SET\_TAPE\_DRIVE\_INFORMATION**, only compression is changeable.

**lpParameters** is the address of either a **TAPE\_SET\_MEDIA\_PARAMETERS** or a **TAPE\_SET\_DRIVE\_PARAMETERS** data structure that contains the parameters.

## **GetTapeParameters**

The **GetTapeParameters** entry point is called to get information that describes the tape or the tape drive.

The data structures are

```
struct{ // structure used by GET_TAPE_MEDIA_INFORMATION
    LARGE_INTEGER Capacity; /* invalid for Magstar */
    LARGE_INTEGER Remaining; /* invalid for Magstar */
    DWORD BlockSize;
    DWORD PartitionCount;
    BOOLEAN WriteProtected;
}TAPE_GET_MEDIA_PARAMETERS;

struct{ // structure used by GET_TAPE_DRIVE_INFORMATION
    BOOLEAN ECC;
    BOOLEAN Compression;
    BOOLEAN DataPadding;
    BOOLEAN ReportSetmarks;
    ULONG DefaultBlockSize;
    ULONG MaximumBlockSize;
    ULONG MinimumBlockSize;
    ULONG MaximumPartitionCount;
    ULONG FeaturesLow;
    ULONG FeaturesHigh;
}
```

```
    ULONG   EOTWarningZoneSize;
}TAPE_GET_DRIVE_PARAMETERS;
```

The following code fragment illustrates a call to the **GetTapeParameters** routine.

```
DWORD rc;  
  
rc = GetTapeParameters(  
    hDevice,  
    dwOperation,  
    lpdwSize,  
    lpParameters  
);  
  
if (rc)  
{  
    printf("Error on GetTapeParameters\n");  
    printf("System Error = %d\n", GetLastError());  
    exit (-1);  
}
```

Where **hDevice** is the open file handle, **dwOperation** is the type of information requested (**GET\_TAPE\_MEDIA\_INFORMATION** or **GET\_TAPE\_DRIVE\_INFORMATION**), and **lpParameters** is the address of the returned data parameter structure.

If the function succeeds, the return value **rc** is **ERROR\_SUCCESS**.

### **PrepareTape**

The **PrepareTape** entry point is called to prepare the tape for access or removal. For example,

```
PrepareTape(  
    hDevice,  
    dwOperation,  
    bImmediate  
)
```

**dwOperation** is the type of operation requested.

The following types of operations and immediate values are supported:

<b>TAPE_LOAD</b>	bImmediate TRUE or FALSE
<b>TAPE_LOCK</b>	bImmediate FALSE
<b>TAPE_UNLOAD</b>	bImmediate TRUE or FALSE
<b>TAPE_UNLOCK</b>	bImmediate FALSE

### **EraseTape**

The **EraseTape** entry point is called to erase all or a part of a tape. The erase is completed from the current location. For example:

```
EraseTape(  
    hDevice,  
    dwEraseType,  
    bImmediate  
)
```

**dwEraseType** is the type of operation requested.

The following types of operations and immediate values are supported.

<b>TAPE_ERASE_LONG</b>	bImmediate TRUE or FALSE
------------------------	--------------------------

## **GetTapeStatus**

The **GetTapeStatus** entry point is called to determine whether the tape device is ready to process tape commands. For example,

```
GetTapeStatus(  
    HANDLE hDevice  
) ;
```

`hDevice` is the handle to the device for which to get the device status.

## **DeviceIoControl**

The **DeviceIoControl** function is described in the *Microsoft Developer Network (MSDN) Software Developer Kit*(SDK) and *Windows Driver Kit* (WDK).

The **DeviceIoControl** function sends a control code directly to a specified device driver, causing the corresponding device to complete the specified operation.

```
BOOL DeviceIoControl(  
    HANDLE hDevice,           // handle to device of interest  
    DWORD dwIoControlCode,    // control code of operation to perform  
    LPVOID lpInBuffer,        // pointer to buffer to supply input data  
    DWORD nInBufferSize,      // size of input buffer  
    LPVOID lpOutBuffer,       // pointer to buffer to receive output data  
    DWORD nOutBufferSize,     // size of output buffer  
    LPDWORD lpBytesReturned,  // pointer to variable to receive output byte count  
    LPOVERLAPPED lpOverlapped // pointer to overlapped structure for \  
                           // asynchronous operation  
) ;
```

Following is a list of the supported `dwIoControlCode` codes that are described in the MSDN WDK and used through the **DeviceIoControl** API.

### **IOCTL\_SCSI\_PASS\_THROUGH**

Tape and medium changer.

### **IOCTL\_SCSI\_PASS\_THROUGH\_DIRECT**

Tape and medium changer.

### **IOCTL\_STORAGE\_RESERVE**

Tape and medium changer.

### **IOCTL\_STORAGE\_RELEASE**

Tape and medium changer.

### **IOCTL\_STORAGE\_PERSISTENT\_RESERVE\_IN**

Tape and medium changer.

### **IOCTL\_STORAGE\_PERSISTENT\_RESERVE\_OUT**

Tape and medium changer.

### **IOCTL\_CHANGER\_EXCHANGE\_MEDIUM**

Medium changer not all changers.

### **IOCTL\_CHANGER\_GET\_ELEMENT\_STATUS**

Medium changer if bar code Reader then VolTags supported.

### **IOCTL\_CHANGER\_GET\_PARAMETERS**

Medium changer.

### **IOCTL\_CHANGER\_GET\_PRODUCT\_DATA**

Medium changer.

### **IOCTL\_CHANGER\_GET\_STATUS**

Medium changer.

### **IOCTL\_CHANGER\_INITIALIZE\_ELEMENT\_STATUS**

Medium changer with range not supported by all changers.

### **IOCTL\_CHANGER\_MOVE\_MEDIUM**

Medium changer.

### **IOCTL\_CHANGER\_SET\_ACCESS**

Medium changer for IE Port only and not for all changers.

### **IOCTL\_CHANGER\_SET\_POSITION**

Medium changer only some devices support the transport object.

An example of the use of **SCSI Pass Through** is contained in the sample code **SPTI.C** in the WDK.

The function call **DeviceIoControl** is described in the SDK and examples of its use are shown in the WDK.

## **Medium Changer IOCTLs**

### ***IOCTL commands***

Not all source or destination addresses, exchanges, moves, or operations are allowed for a particular IBM Medium Changer. The user must issue an **IOCTL\_CHANGER\_GET\_PARAMETER** to determine the type of operations that are allowed by a specific changer device. Further information on allowable commands for a particular changer can be found in the IBM hardware reference for that device. It is recommended that the user have a copy of the hardware reference before any applications for the changer device are constructed.

### **IOCTL\_CHANGER\_EXCHANGE\_MEDIUM**

The media from the source element is moved to the first destination element. The medium that previously occupied the first destination element is moved to the second destination element (the second destination element might be the same as the source) by sending an **ExchangeMedium (0xA6)** SCSI command to the device. The input data is a structure of **CHANGER\_EXCHANGE\_MEDIUM**. This command is not supported by all devices.

### **IOCTL\_CHANGER\_GET\_ELEMENT\_STATUS**

Returns the status of all elements or of a specified number of elements of a particular type by sending a **ReadElementStatus (0xB8)** SCSI command to the device. The input and output data is a structure of **CHANGER\_ELEMENT\_STATUS**.

### **IOCTL\_CHANGER\_GET\_PARAMETERS**

Returns the capabilities of the changer. The output data is in a structure of **GET\_CHANGER\_PARAMETERS**.

### **IOCTL\_CHANGER\_GET\_PRODUCT\_DATA**

Returns the product data for the changer. The output data is in a structure of **CHANGER\_PRODUCT\_DATA**.

### **IOCTL\_CHANGER\_GET\_STATUS**

Returns the status of the changer by sending a **TestUnitReady (0x00)** SCSI command to the device.

### **IOCTL\_CHANGER\_INITIALIZE\_ELEMENT\_STATUS**

Initializes the status of all elements or a range of a particular element by sending an **InitializeElementStatus (0x07)** or **InitializeElementStatusWithRange (0xE7)** SCSI command to the device. The input data is a structure of **CHANGER\_INITIALIZE\_ELEMENT\_STATUS**.

### **IOCTL\_CHANGER\_MOVE\_MEDIUM**

Moves a piece of media from a source to a destination by sending a **MoveMedia (0xA5)** SCSI command to the device. The input data is a structure of **CHANGER\_MOVE\_MEDIUM**.

## **IOCTL\_CHANGER\_REINITIALIZE\_TRANSPORT**

Physically recalibrates a transport element by sending a **RezeroUnit (0x01)** SCSI command to the device. The input data is a structure of **CHANGER\_ELEMENT**. This command is not supported by all devices.

## **IOCTL\_CHANGER\_SET\_ACCESS**

Sets the access state of the changers IE port by sending a **PrevenAllowMediumRemoval (0x1E)** SCSI command to the device. The input data is a structure of **CHANGER\_SET\_ACCESS**.

## **IOCTL\_CHANGER\_SET\_POSITION**

Sets the changers robotic transport to a specified address by sending a **PositionToElement (0x2B)** SCSI command to the device. The input data is a structure of **CHANGER\_SET\_POSITION**.

### **Preempt reservation**

A reservation can be preempted by issuing the appropriate IOCTL. The current reservation key is needed to successfully preempt the reservation. The current reservation key can be queried by issuing an **IOCTL\_STORAGE\_PERSISTENT\_RESERVE\_IN** IOCTL:

```
PERSISTENT_RESERVE_COMMAND prcmd = { 0 };
PPRI_RESERVATION_LIST prsl = NULL;
ULONG AdditionalLength = 0;
DWORD BytesReturned = 0;
INT iStatus = 0, i, j;

UCHAR *bufDataRead = (UCHAR *) malloc(SENSE_BUFFER_SIZE * 2);
ZeroMemory(bufDataRead, SENSE_BUFFER_SIZE * 2);

prcmd.Size = sizeof(PERSISTENT_RESERVE_COMMAND);

prcmd.PR_IN.ServiceAction = RESERVATION_ACTION_READ_KEYS;
prcmd.PR_IN.AllocationLength = sizeof(PRI_REGISTRATION_LIST);

for (i = 0; i < 2; i++) {
    if (0 == i)
        AdditionalLength = sizeof(PRI_RESERVATION_LIST);

    iStatus = DeviceIoControl(tape,
                               IOCTL_STORAGE_PERSISTENT_RESERVE_IN,
                               &prcmd,
                               prcmd.Size,
                               bufDataRead,
                               AdditionalLength,
                               &BytesReturned,
                               NULL);

    if (0 == iStatus) {
        free(bufDataRead);
        return FALSE;
    }
    prsl = (PPRI_RESERVATION_LIST)bufDataRead;

    if (0 == i) {
        AdditionalLength = (ULONG)((prsl->AdditionalLength[0] & 0xff) << 12);
        AdditionalLength |= (ULONG)((prsl->AdditionalLength[1] & 0xff) << 8);
        AdditionalLength |= (ULONG)((prsl->AdditionalLength[2] & 0xff) << 4);
        AdditionalLength |= (ULONG) (prsl->AdditionalLength[3] & 0xff);

        AdditionalLength += sizeof(PRI_RESERVATION_LIST);

        prcmd.PR_IN.AllocationLength = AdditionalLength;
    } else if (1 == i) {
        for (j = 0; (j * sizeof(PRI_RESERVATION_DESCRIPTOR)
                    + sizeof(PRI_RESERVATION_LIST)) <= AdditionalLength; j++) {
            printf("\nReservation 0x%08x%08x being examined at
descriptor index %d.\n",
                   ((LARGE_INTEGER*)(prsl->Reservations[j].ReservationKey))
                   ->HighPart,
                   ((LARGE_INTEGER*)(prsl->Reservations[j].ReservationKey)));
    }
}
```

```

        ->LowPart,
        j);
}
}
}

```

When the reservation key is known, it can be preempted, meaning a new host can be the reservation holder, thus being able to interact with the target as needed.

```

PRI_RESERVATION_DESCRIPTOR reservation = { 0 };
PERSISTENT_RESERVE_COMMAND prcmd = { 0 };
PRO_PARAMETER_LIST prolist = { 0 };
DWORD BytesReturned = 0;
INT iStatus = 0;

UCHAR bufDataRead[sizeof(PERSISTENT_RESERVE_COMMAND)
+ sizeof(PRO_PARAMETER_LIST)] = { 0 };

prcmd.Size = sizeof(PERSISTENT_RESERVE_COMMAND) + sizeof(PRO_PARAMETER_LIST);

prcmd.PR_OUT.ParameterList[1] = 0x18;
prcmd.PR_OUT.ServiceAction = 0x3;
prcmd.PR_OUT.Type = 0x3;

query_reserve(tape, &reservation);

RtlCopyMemory(prolist.ReservationKey, reservation.ReservationKey, 8);
RtlCopyMemory(prolist.ServiceActionReservationKey,
reservation.ReservationKey, 8);

RtlCopyMemory(bufDataRead, &prcmd, sizeof(PERSISTENT_RESERVE_COMMAND));
RtlCopyMemory(bufDataRead + sizeof(PERSISTENT_RESERVE_COMMAND),
&prolist, sizeof(PRO_PARAMETER_LIST));

iStatus = DeviceIoControl(tape,
IOCTL_STORAGE_PERSISTENT_RESERVE_OUT,
bufDataRead,
sizeof(bufDataRead),
bufDataRead,
sizeof(bufDataRead),
&BytesReturned,
NULL);

```

The query\_reserve function can be implemented as explained earlier. Finally, if the reservation needs to be preempted on a different system than the original reservation holder, the OPEN\_ALWAYS flag comes in handy. It allows the user to query the target's serial number, then queries the reservation key, and preempts the reservation.

Caution is advised when preempting reservations due to inherent risk of data loss if done incorrectly. Applications must make sure that they are clearing or preempting the appropriate reservation.

## Vendor-specific (IBM) device IOCTLs for DeviceIoControl

The following descriptions are of the IBM vendor-specific ioctl requests for tape and changer.

```

/*
This macro is defined in ntddk.h and devioctl.h
#define CTL_CODE(DeviceType, Function, Method, Access) \
    (((DeviceType) << 16) | ((Access) << 14) | ((Function) << 2) | (Method))
*/

```

The following **ioctl** commands are supported by the **ibmtp.sys** driver through DeviceIoControl.

```

#define LB_ACCESS FILE_READ_ACCESS | FILE_WRITE_ACCESS
#define M_MTI(x) CTL_CODE(IOCTL_BASE+2,x,METHOD_BUFFERED, LB_ACCESS)
#define IOCTL_TAPE_OBTAIN_SENSE CTL_CODE(IOCTL_TAPE_BASE, 0x0819,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_OBTAIN_VERSION CTL_CODE(IOCTL_TAPE_BASE, 0x081a,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_LOG_SELECT CTL_CODE(IOCTL_TAPE_BASE, 0x081c,
METHOD_BUFFERED, FILE_READ_ACCESS | FILE_WRITE_ACCESS)
#define IOCTL_TAPE_LOG_SENSE CTL_CODE(IOCTL_TAPE_BASE, 0x081d,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_LOG_SENSE10 CTL_CODE(IOCTL_TAPE_BASE, 0x0833,

```

```

METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_ENH_TAPE_LOG_SENSE10 CTL_CODE(IOCTL_TAPE_BASE, 0x0835, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define IOCTL_TAPE_REPORT_MEDIA_DENSITY CTL_CODE(IOCTL_TAPE_BASE, 0x081e,
METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_TAPE_OBTAIN_MTDEVICE (M_MTI(16))
#define IOCTL_CREATE_PARTITION CTL_CODE(IOCTL_TAPE_BASE, 0x0826, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_QUERY_PARTITION CTL_CODE(IOCTL_TAPE_BASE, 0x0825, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_SET_ACTIVE_PARTITION CTL_CODE(IOCTL_TAPE_BASE, 0x0827, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_QUERY_DATA_SAFE_MODE CTL_CODE(IOCTL_TAPE_BASE, 0x0823, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_SET_DATA_SAFE_MODE CTL_CODE(IOCTL_TAPE_BASE, 0x0824, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_ALLOW_DATA_OVERWRITE CTL_CODE(IOCTL_TAPE_BASE, 0x0828, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define IOCTL_SET_PEW_SIZE
CTL_CODE(IOCTL_TAPE_BASE, 0x082c, METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_QUERY_PEW_SIZE
CTL_CODE(IOCTL_TAPE_BASE, 0x082b, METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_VERIFY_TAPE_DATA
CTL_CODE(IOCTL_TAPE_BASE, 0x082a, METHOD_BUFFERED, FILE_READ_ACCESS )
#define IOCTL_QUERY_RAO_INFO CTL_CODE(IOCTL_TAPE_BASE, 0x082e, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define IOCTL_GENERATE_RAO CTL_CODE(IOCTL_TAPE_BASE, 0x082f, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define IOCTL_RECEIVE_RAO CTL_CODE(IOCTL_TAPE_BASE, 0x0834, METHOD_BUFFERED,
FILE_READ_ACCESS )

```

## ***IOCTL\_TAPE\_OBTAIN\_SENSE***

Issue this command after an error occurs to obtain sense information that is associated with the most recent error. To guarantee that the application can obtain sense information that is associated with an error, the application must issue this command before other commands to the device are issued. Subsequent operations (other than **IOCTL\_TAPE\_OBTAIN\_SENSE**) reset the sense data field before the operation is run.

This IOCTL is only available for the tape path.

The following output structure is completed by the **IOCTL\_TAPE\_OBTAIN\_SENSE** command that is passed by the caller.

```

#define IOCTL_TAPE_OBTAIN_SENSE CTL_CODE(IOCTL_TAPE_BASE, 0x0819, METHOD_BUFFERED,
FILE_READ_ACCESS)
#define SENSE_BUFFER_SIZE 96 /* Default request sense buffer size for \
Windows 200x */

typedef struct _TAPE_OBTAIN_SENSE {
    ULONG SenseDataLength;
    // The number of bytes of valid sense data.
    // Will be zero if no error with sense data has occurred.
    // The only sense data available is that of the last error.
    CHAR SenseData[SENSE_BUFFER_SIZE];
} TAPE_OBTAIN_SENSE, *PTAPE_OBTAIN_SENSE;

```

An example of the **IOCTL\_TAPE\_OBTAIN\_SENSE** command is

```

DWORD cb;
TAPE_OBTAIN_SENSE sense_data;
DeviceIoControl(hDevice,
    IOCTL_TAPE_OBTAIN_SENSE,
    NULL,
    0,
    &sense_data,
    (long)sizeof(TAPE_OBTAIN_SENSE),
    &cb,
    (LPOVERLAPPED) NULL);

```

## ***IOCTL\_TAPE\_OBTAIN\_VERSION***

Issue this command to obtain the version of the device driver. It is in the form of a null terminated string.

This IOCTL is only for the tape path.

The following output structure is completed by the **IOCTL\_TAPE\_OBTAIN\_VERSION** command.

```
#define IOCTL_TAPE_OBTAIN_VERSION      CTL_CODE(IOCTL_TAPE_BASE, 0x081a, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define MAX_DRIVER_VERSIONID_LENGTH 12

typedef struct _TAPE_OBTAIN_VERSION {
    CHAR VersionId[MAX_DRIVER_VERSIONID_LENGTH];
} TAPE_OBTAIN_VERSION, *PTAPE_OBTAIN_VERSION;
```

An example of the **IOCTL\_TAPE\_OBTAIN\_VERSION** command is

```
DWORD cb;
TAPE_OBTAIN_VERSION code_version;
DeviceIoControl(hDevice,
                 IOCTL_TAPE_OBTAIN_VERSION,
                 NULL,
                 0,
                 &code_version,
                 (long)sizeof(TAPE_OBTAIN_VERSION),
                 &cb,
                 (LPOVERLAPPED) NULL);
```

## **IOCTL\_TAPE\_LOG\_SELECT**

This command resets all log pages that can be reset on the device to their default values. This IOCTL is only for the tape path.

The IOCTL code for **IOCTL\_TAPE\_LOG\_SELECT** is defined as follows.

```
#define IOCTL_TAPE_LOG_SELECT      CTL_CODE(IOCTL_TAPE_BASE, 0x081c, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS)
```

An example of this command to reset all log pages follows.

```
DWORD cb;
DeviceIoControl(hDevice,
                 IOCTL_TAPE_LOG_SELECT,
                 NULL,
                 0,
                 NULL,
                 0,
                 &cb,
                 (LPOVERLAPPED) NULL);
```

## **IOCTL\_TAPE\_LOG\_SENSE**

Issue this command to obtain the log data of the requested log page from IBM Tape Storage device. The data that is returned is formatted according to the IBM Tape Storage hardware reference.

This IOCTL is only for the tape path.

The following input/output structure is used by the **IOCTL\_TAPE\_LOG\_SENSE** command.

```
#define IOCTL_TAPE_LOG_SENSE      CTL_CODE(IOCTL_TAPE_BASE, 0x081d, METHOD_BUFFERED,
FILE_READ_ACCESS )
#define MAX_LOG_SENSE 1024 // Maximum number of bytes the command will return

typedef struct _TAPE_LOG_SENSE_PARAMETERS{
    UCHAR PageCode; // The requested log page code
    UCHAR PC; // PC = 0 for maximum values, 1 for current value, 3 for power-on values
    UCHAR PageLength[2]; /* Length of returned data, filled in by the command */
    UCHAR LogData[MAX_LOG_SENSE]; /* Log data, filled in by the command */
} TAPE_LOG_SENSE_PARAMETERS, *PTAPE_LOG_SENSE_PARAMETERS;
```

An example of the **IOCTL\_TAPE\_LOG\_SENSE** command is

```
DWORD cb;
TAPE_LOG_SENSE_PARAMETERS logsense;
logsense.PageCode=0;
```

```

logSense.PC = 1;

DeviceIoControl(hDevice,
                 IOCTL_TAPE_LOG_SENSE,
                 &logSense,
                 (long)sizeof(TAPE_LOG_SENSE_PARAMETERS),
                 &logSense,
                 (long)sizeof(TAPE_LOG_SENSE_PARAMETERS),
                 &cb,
                 (LPOVERLAPPED) NULL);

```

## ***IOCTL\_TAPE\_LOG\_SENSE10***

Issue this command to obtain the log data of the requested log page/subpage from IBM Tape Storage device. The data returned is formatted according to the IBM Tape Storage hardware reference. This IOCTL is only for the tape path.

The following input/output structure is used by the **IOCTL\_TAPE\_LOG\_SENSE10** command.

```

#define IOCTL_TAPE_LOG_SENSE10           CTL_CODE(IOCTL_TAPE_BASE, 0x0833, METHOD_BUFFERED,
FILE_READ_ACCESS)
#define MAX_LOG_SENSE 1024 // Maximum number of bytes the command will return

typedef struct _TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE{
    UCHAR      PageCode;          /* [IN] Log sense page */
    UCHAR      SubPageCode;       /* [IN] Log sense subpage */
    UCHAR      PC;               /* [IN] PC bit to be consistent with
                                previous Log Sense IOCTL */
    UCHAR[2]   reserved[2];       /* unused */
    ULONG     PageLength;        /* [OUT] number of valid bytes in data
                                (log_page_header_size+page_length)*/
    ULONG     parm_pointer;      /* [IN] specific parameter number at which
                                the data begins */
    CHAR      LogData[MAX_LOG_SENSE_DATA]; /* [OUT] log sense data */
} TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE, *PTAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE;

```

An example of the **IOCTL\_TAPE\_LOG\_SENSE10** command is

```

DWORD cb;
TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE logSense;
logSense.PageCode=0x10;
logSense.SubPageCode=0x01;
logSense.PC = 1;

DeviceIoControl(hDevice,
                 IOCTL_TAPE_LOG_SENSE10,
                 &logSense, (long)sizeof(TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE),
                 &logSense, (long)sizeof(TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE),
                 &cb, (LPOVERLAPPED) NULL);

```

## ***IOCTL\_ENH\_TAPE\_LOG\_SENSE10***

Issue this command to obtain the log data of the requested log page/subpage from IBM Tape Storage device. The data that is returned is formatted according to the IBM Tape Storage hardware reference. This IOCTL is only for the tape path and is enhanced so the application can set the page length and provide the buffer enough to get the data back. The following input/output structure is used by the **IOCTL\_ENH\_TAPE\_LOG\_SENSE10** command.

```

#define IOCTL_ENH_TAPE_LOG_SENSE10           CTL_CODE(IOCTL_TAPE_BASE, 0x0835, METHOD_BUFFERED,
FILE_READ_ACCESS)

typedef struct _ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE{
    UCHAR      PageCode;          /* [IN] Log sense page */
    UCHAR      SubPageCode;       /* [IN] Log sense subpage */
    UCHAR      PC;               /* [IN] PC bit */
    UCHAR[5]   reserved[5];       /* unused */
    ULONG     Length;            /* [IN][OUT] number of valid bytes in data */
                                /* (log_page_header_size+page_length) */
    ULONG     parm_pointer;      /* [IN] specific parameter number at which
                                the data begins */
    CHAR      LogData[1];         /* [IN] log sense buffer allocated by
                                application */
} ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE, *PENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE;

```

```
        /* [OUT] log sense data */  
} ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE, *PENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE;
```

An example of the **IOCTL\_ENH\_TAPE\_LOG\_SENSE10** command is

```
DWORD cb;  
char *logSense;  
int pageLength = 256;  
long lsize = sizeof(ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE) - sizeof  
(CHAR) /*LogData[1]*/ + pageLength  
  
logSense = malloc (lsize);  
(ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE)logSense->PageCode=0x10;  
(ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE)logSense->SubPageCode=0x01;  
(ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE)logSense->PC = 1;  
(ENH_TAPE_LOG_SENSE_PARAMETERS_WITH_SUBPAGE)logSense->Length = pageLength;  
DeviceIoControl(hDevice,  
                IOCTL_ENH_TAPE_LOG_SENSE10,  
                &logSense, (long)lsize,  
                &logSense, (long)lsize,  
                &cb, (LPOVERLAPPED) NULL);
```

## **IOCTL\_TAPE\_REPORT\_MEDIA\_DENSITY**

Issue this command to obtain the media density information on the loaded media in the drive. If there is no media load, the command fails. This IOCTL is only for the tape path.

The following output structure is completed by the **IOCTL\_TAPE\_REPORT\_MEDIA\_DENSITY** command.

```
#define IOCTL_TAPE_REPORT_MEDIA_DENSITY CTL_CODE(IOCTL_TAPE_BASE, 0x081e, METHOD_BUFFERED,  
FILE_READ_ACCESS )  
  
typedef struct_TAPE_REPORT_DENSITY{  
    ULONG PrimaryDensityCode;      /* Primary Density Code */  
    ULONG SecondaryDensityCode;    /* Secondary Density Code */  
    BOOLEAN WriteOk;              /* 0 = does not support writing in this format */  
                                /* 1 = support writing in this format */  
    ULONG BitsPerMM;             /* Bits Per mm */  
    ULONG MediaWidth;            /* Media Width */  
    ULONG Tracks;                /* Tracks */  
    ULONG Capacity;              /* Capacity in MegaBytes */  
} TAPE_REPORT_DENSITY, *PTAPE_REPORT_DENSITY;
```

An example of the **IOCTL\_TAPE\_REPORT\_MEDIA\_DENSITY** command is

```
DWORD cb;  
TAPE_REPORT_DENSITY tape_reportden;  
  
DeviceIoControl (hDevice,  
                IOCTL_TAPE_REPORT_MEDIA_DENSITY,  
                NULL,  
                0,  
                &tape_reportden,  
                (long)sizeof(TAPE_REPORT_DENSITY),  
                &cb,  
                (LPOVERLAPPED) NULL);
```

## **IOCTL\_TAPE\_OBTAIN\_MTDEVICE**

Issue this command to obtain the device number. An error is returned if it is issued against a 3570 drive.

The following output structure is filled in by the **IOCTL\_TAPE\_OBTAIN\_MTDEVICE** command.

```
typedef ULONG TAPE_OBTAIN_MTDEVICE, *PTAPE_OBTAIN_MTDEVICE;
```

An example of the **IOCTL\_TAPE\_OBTAIN\_MTDEVICE** command is

```
BOOL result;  
DWORD cb;  
TAPE_OBTAIN_MTDEVICE mt_device;  
  
result = DeviceIoControl(gp->ddHandle0,  
                        IOCTL_TAPE_OBTAIN_MTDEVICE,  
                        NULL,
```

```

    0,
    &mt_device,
    (long)sizeof(TAPE_OBTAIN_MTDEVICE),
    &cb,
    (LPOVERLAPPED) NULL);

if(result)
    printf(fp, "\nMTDevice Info : %x\n\n", mt_device);
else
    /* Error handling code */

```

## ***IOCTL\_TAPE\_GET\_DENSITY***

The IOCTL code for **IOCTL\_TAPE\_GET\_DENSITY** is defined as follows.

```
#define IOCTL_TAPE_GET_DENSITY \
CTL_CODE(IOCTL_TAPE_BASE, 0x000c, METHOD_BUFFERED, \
FILE_READ_ACCESS | FILE_WRITE_ACCESS).
```

The IOCTL reports density for supported devices by using the following structure.

```
typedef struct _TAPE_DENSITY
{
    UCHAR ucDensityCode;
    UCHAR ucDefaultDensity;
    UCHAR ucPendingDensity;
} TAPE_DENSITY, *PTAPE_DENSITY;
```

An example of the **IOCTL\_TAPE\_GET\_DENSITY** command is

```
TAPE_DENSITY tape_density = {0};

rc = DeviceIoControl(hDevice,
IOCTL_TAPE_GET_DENSITY,
NULL,
0,
&tape_density,
sizeof(TAPE_DENSITY),
&cb,
(LPOVERLAPPED) NULL);
```

## ***IOCTL\_TAPE\_SET\_DENSITY***

The IOCTL code for **IOCTL\_TAPE\_SET\_DENSITY** is defined as follows.

```
#define IOCTL_TAPE_SET_DENSITY \
CTL_CODE(IOCTL_TAPE_BASE, 0x000d, METHOD_BUFFERED, \
FILE_READ_ACCESS | FILE_WRITE_ACCESS)
```

The IOCTL sets density for supported devices by using the following structure.

```
typedef struct _TAPE_DENSITY
{
    UCHAR ucDensityCode;
    UCHAR ucDefaultDensity;
    UCHAR ucPendingDensity;
} TAPE_DENSITY, *PTAPE_DENSITY;
```

ucDensityCode is ignored. ucDefaultDensity and ucPendingDensity are set by using the tape drive's mode page 0x25. Caution must be taken when this IOCTL is issued. An incorrect tape density might lead to data corruption.

An example of the **IOCTL\_TAPE\_SET\_DENSITY** command is

```
TAPE_DENSITY tape_density;

// Modify fields of tape_density. For details, see the SCSI specification
// for your hardware.

rc = DeviceIoControl(hDevice,
IOCTL_TAPE_SET_DENSITY,
&tape_density,
sizeof(TAPE_DENSITY),
```

```
NULL,  
0,  
&cb,  
(LPOVERLAPPED) NULL);
```

## ***IOCTL\_TAPE\_GET\_ENCRYPTION\_STATE***

This IOCTL command queries the drive's encryption method and state.

The IOCTL code for **IOCTL\_TAPE\_GET\_ENCRYPTION\_STATE** is defined as follows.

```
#define IOCTL_TAPE_GET_ENCRYPTION_STATE CTL_CODE(IOCTL_TAPE_BASE, 0x0820,  
METHOD_BUFFERED, FILE_READ_ACCESS )
```

The IOCTL gets encryption states for supported devices by using the following structure.

```
typedef struct _ENCRYPTION_STATUS  
{  
    UCHAR ucEncryptionCapable; /* (1)Set this field as a boolean based on  
                                the capability of the drive */  
    UCHAR ucEncryptionMethod; /* (2)Set this field to one of the  
                            defines METHOD_* below */  
    UCHAR ucEncryptionState; /* (3)Set this field to one of the  
                            #defines STATE_* below */  
    UCHAR aucReserved[13];  
} ENCRYPTION_STATUS, *PENCRYPTION_STATUS;
```

#defines for METHOD.

```
#define ENCRYPTION_METHOD_NONE          0 /* Only used in  
                                         GET_ENCRYPTION_STATE */  
#define ENCRYPTION_METHOD_LIBRARY        1 /* Only used in  
                                         GET_ENCRYPTION_STATE */  
#define ENCRYPTION_METHOD_SYSTEM         2 /* Only used in  
                                         GET_ENCRYPTION_STATE */  
#define ENCRYPTION_METHOD_APPLICATION    3 /* Only used in  
                                         GET_ENCRYPTION_STATE */  
#define ENCRYPTION_METHOD_CUSTOM         4 /* Only used in  
                                         GET_ENCRYPTION_STATE */  
#define ENCRYPTION_METHOD_UNKNOWN        5 /* Only used in  
                                         GET_ENCRYPTION_STATE */
```

#defines for STATE.

```
#define ENCRYPTION_STATE_OFF 0 /* Used in GET/SET_ENCRYPTION_STATE */  
#define ENCRYPTION_STATE_ON   1 /* Used in GET/SET_ENCRYPTION_STATE */  
#define ENCRYPTION_STATE_NA  2 /* Only used in GET_ENCRYPTION_STATE*/
```

An example of the **IOCTL\_TAPE\_GET\_ENCRYPTION\_STATE** command is

```
ENCRYPTION_STATUS scEncryptStat;  
DeviceIoControl(hDevice,  
                IOCTL_TAPE_GET_ENCRYPTION_STATE,  
                &scEncryptStat,  
                sizeof(ENCRYPTION_STATUS),  
                &scEncryptStat,  
                sizeof(ENCRYPTION_STATUS),  
                &cb  
                (LPOVERLAPPED) NULL);
```

## ***IOCTL\_TAPE\_SET\_ENCRYPTION\_STATE***

This IOCTL command allows only set encryption state for application-managed encryption.

**Note:** On unload, some drive settings might be reset to default. To set the encryption state, the application must issue this IOCTL after a tape is loaded and at BOP.

The data structure that is used for this IOCTL is the same as for **IOCTL\_GET\_ENCRYPTION\_STATE**.

```
#define IOCTL_TAPE_SET_ENCRYPTION_STATE CTL_CODE(IOCTL_TAPE_BASE, 0x0821,  
METHOD_BUFFERED,  
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
```

An example of the **IOCTL\_TAPE\_SET\_ENCRYPTION\_STATE** command is

```
ENCRYPTION_STATUS scEncryptStat;
DeviceIoControl(hDevice,
                IOCTL_TAPE_SET_ENCRYPTION_STATE,
                &scEncryptStat,
                sizeof(ENCRYPTION_STATUS),
                &scEncryptStat
                sizeof(ENCRYPTION_STATUS),
                &cb,
                (LPOVERLAPPED) NULL);
```

## **IOCTL\_TAPE\_SET\_DATA\_KEY**

This IOCTL command is used to set the data key only for application-managed encryption.

The IOCTL sets data keys for supported devices by using the following structure.

```
#define IOCTL_TAPE_SET_DATA_KEY CTL_CODE(IOCTL_TAPE_BASE, 0x0822,
                                         METHOD_BUFFERED,
                                         FILE_READ_ACCESS | FILE_WRITE_ACCESS )

#define DATA_KEY_INDEX_LENGTH      12
#define DATA_KEY_RESERVED1_LENGTH 15
#define DATA_KEY_LENGTH            32
#define DATA_KEY_RESERVED2_LENGTH 48
typedef struct _DATA_KEY
{
    UCHAR aucDataKeyIndex[DATA_KEY_INDEX_LENGTH];
    UCHAR ucDataKeyIndexLength;
    UCHAR aucReserved1[DATA_KEY_RESERVED1_LENGTH];
    UCHAR aucDataKey[DATA_KEY_LENGTH];
    UCHAR aucReserved2[DATA_KEY_RESERVED2_LENGTH];
} DATA_KEY, *PDATA_KEY;
```

An example of the **IOCTL\_TAPE\_SET\_DATA\_KEY** command is

```
DATA_KEY scDataKey;
/* fill in your data key and data key length, then issue DeviceIoControl */
DeviceIoControl(hDevice,
                IOCTL_TAPE_SET_DATA_KEY,
                &scDataKey,
                sizeof(DATA_KEY),
                &scDataKey,
                sizeof(DATA_KEY),
                &cb,
                (LPOVERLAPPED) NULL);
```

## **IOCTL\_CREATE\_PARTITION**

This command is used to create one or more partitions on the tape. The tape must be at BOT (partition 0 logical block id 0) before the command is issued or it fails. The application must either issue this **IOCTL\_CREATE\_PARTITION** after a tape is initially loaded or issue the **IOCTL\_SET\_ACTIVE\_PARTITION** with the **partition\_number** and **logical\_clock\_id** fields that are set to 0 first.

The structure that is used to create partitions is

```
#define IOCTL_CREATE_PARTITION           CTL_CODE(IOCTL_TAPE_BASE, 0x0826, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
typedef struct _TAPE_PARTITION{
    UCHAR type;                      /* Type of tape partition to create      */
    UCHAR number_of_partitions;        /* Number of partitions to create       */
    UCHAR size_unit;                 /* IDP size unit of partition sizes below */
    USHORT size[MAX_PARTITIONS];     /* Array of partition sizes in size units */
                                    /* for each partition, 0 to (number - 1) */
                                    /* Size can not be 0 and one partition */
                                    /* size must be 0xFFFF to use the      */
                                    /* remaining capacity on the tape.    */
    UCHAR partition_method;          /* partitioning type                   */
    UCHAR reserved [31];
} TAPE_PARTITION, *PTAPE_PARTITION;
```

An example of the **IOCTL\_CREATE\_PARTITION** command is

```
DWORD cb;
TAPE_PARTITION tape_partition
...
DeviceIoControl(gp->ddHandle0,
                IOCTL_CREATE_PARTITION,
                &tape_partition,
                (long)sizeof(TAPE_PARTITION),
                NULL,
                0,
                &cb,
                (LPOVERLAPPED) NULL);
```

## **IOCTL\_QUERY\_PARTITION**

This command returns partition information for the current loaded tape.

The following output structure is completed by the **IOCTL\_QUERY\_PARTITION** command.

```
#define IOCTL_QUERY_PARTITION      CTL_CODE(IOCTL_TAPE_BASE, 0x0825, METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define MAX_PARTITIONS 255

typedef struct _QUERY_PARTITION{
    UCHAR max_partitions;          /* Max number of supported partitions */
    UCHAR active_partition;        /* current active partition on tape */
    UCHAR number_of_partitions;    /* Number of partitions from 1 to max */
    UCHAR size_unit;              /* Size unit of partition sizes below */
    USHORT size[MAX_PARTITIONS];   /* Array of partition sizes in size units */
                                    /* for each partition, 0 to (number - 1) */
    UCHAR partition_method;        /* partitioning type */
    char reserved [31];
} QUERY_PARTITION, *PQUERY_PARTITION;
```

An example of the **IOCTL\_QUERY\_PARTITION** command is

```
DWORD cb;
QUERY_PARTITION tape_query_partition;
DeviceIoControl(gp->ddHandle0,
                IOCTL_QUERY_PARTITION,
                NULL,
                0,
                &tape_query_partition,
                (long)sizeof(QUERY_PARTITION),
                &cb,
                (LPOVERLAPPED) NULL);
```

## **IOCTL\_SET\_ACTIVE\_PARTITION**

This command is used to set the current active partition that is used on tape and locate to a specific logical block id within the partition. If the logical block id is 0, the tape is positioned at BOP. If the partition number specified is 0 along with a logical block id 0, the tape is positioned at both BOP and BOT.

The structure for **IOCTL\_SET\_ACTIVE\_PARTITION** command is

```
#define IOCTL_SET_ACTIVE_PARTITION      CTL_CODE(IOCTL_TAPE_BASE, 0x0827,
METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
typedef struct _SET_ACTIVE_PARTITION{
    UCHAR partition_number;        /* Partition number 0-n to change to */
    ULONGLONG logical_block_id;   /* Blockid to locate to within partition */
    char reserved[32];
} SET_ACTIVE_PARTITION, *PSET_ACTIVE_PARTITION;
```

An example of the **IOCTL\_SET\_ACTIVE\_PARTITION** command is

```
DWORD cb;
SET_ACTIVE_PARTITION set_partition;
...
DeviceIoControl(gp->ddHandle0,
                IOCTL_SET_ACTIVE_PARTITION,
                &set_partition,
                (long)sizeof(SET_ACTIVE_PARTITION),
```

```

NULL,
0,
&cb,
(LPOVERLAPPED) NULL);

```

### ***IOCTL\_QUERY\_DATA\_SAFE\_MODE***

This command reports if the Data Safe Mode is enabled or disabled.

The following output structure is completed by the **IOCTL\_QUERY\_DATA\_SAFE\_MODE** command.

```

#define IOCTL_QUERY_DATA_SAFE_MODE      CTL_CODE(IOCTL_TAPE_BASE, 0x0823,
METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
typedef struct _DATA_SAFE_MODE{
    ULONG value;
} DATA_SAFE_MODE, *PDATA_SAFE_MODE;

```

An example of the **IOCTL\_QUERY\_DATA\_SAFE\_MODE** command is

```

DWORD cb;
DATA_SAFE_MODE tapeDataSafeMode;
DeviceIoControl(gp->ddHandle0,
                IOCTL_QUERY_DATA_SAFE_MODE,
                NULL,
                0,
                &tapeDataSafeMode,
                (long)sizeof(DATA_SAFE_MODE),
                &cb,
                (LPOVERLAPPED) NULL);

```

### ***IOCTL\_SET\_DATA\_SAFE\_MODE***

This command enables or disables Data Safe Mode.

The structure that is used to enable or disable Data Safe Mode is the same as **IOCTL\_QUERY\_DATA\_SAFE\_MODE**.

An example of the **IOCTL\_SET\_DATA\_SAFE\_MODE** command is

```

#define IOCTL_SET_DATA_SAFE_MODE      CTL_CODE(IOCTL_TAPE_BASE, 0x0824,
METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
DATA_SAFE_MODE tapeDataSafeMode;
...
DeviceIoControl(gp->ddHandle0,
                IOCTL_SET_DATA_SAFE_MODE,
                &tapeDataSafeMode,
                (long)sizeof(DATA_SAFE_MODE),
                NULL,
                0,
                &cb,
                (LPOVERLAPPED) NULL);

```

### ***IOCTL\_ALLOW\_DATA\_OVERWRITE***

This command allows previously written data on the tape to be overwritten. This action happens when append only mode is enabled on the drive with either a write type command or a format command is allowed on the **IOCTL\_CREATE\_PARTITION**. Before this IOCTL is issued, the application must locate to the partition number and logical block id within the partition where the data overwrite or format occurs.

The data structure that is used for **IOCTL\_ALLOW\_DATA\_OVERWRITE** to enable or disable is

```

#define IOCTL_ALLOW_DATA_OVERWRITE     CTL_CODE(IOCTL_TAPE_BASE, 0x0828,
METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
typedef struct ALLOW_DATA_OVERWRITE{
    UCHAR partition_number;           /* Partition number 0-n to overwrite */
    ULONGULONG logical_block_id;     /* Blockid to overwrite to within partition */
    UCHAR allow_format_overwrite;    /* allow format if in data safe mode */
    UCHAR reserved[32];
} ALLOW_DATA_OVERWRITE, *PALLOW_DATA_OVERWRITE;

```

An example of the **IOCTL\_ALLOW\_DATA\_OVERWRITE** command is

```
ALLOW_DATA_OVERWRITE tapeAllowDataOverwrite;
...
DeviceIoControl(gp->ddHandle0,
                IOCTL_ALLOW_DATA_OVERWRITE,
                &tapeAllowDataOverwrite,
                (long)sizeof(ALLOW_DATA_OVERWRITE),
                NULL,
                0,
                &cbs,
                (LPOVERLAPPED) NULL);
```

## **IOCTL\_READ\_TAPE\_POSITION**

This command returns Position data in either the short, long, or extended form. The type of data to return is specified by setting the **data\_format** field to either **RP\_SHORT\_FORM**, **RP\_LONG\_FORM**, or **RP\_EXTENDED\_FORM**.

The data structures that are used with this IOCTL are

```
#define IOCTL_READ_TAPE_POSITION           CTL_CODE(IOCTL_TAPE_BASE, 0x0829,
METHOD_BUFFERED, FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define RP_SHORT_FORM          0x00
#define RP_LONG_FORM           0x06
#define RP_EXTENDED_FORM       0x08

typedef struct _SHORT_DATA_FORMAT {
    UCHAR bop:1,      /* beginning of partition */
    eop:1,      /* end of partition */
    locu:1,     /* 1 means num_buffer_logical_obj field is unknown */
    bycu:1,     /* 1 means the num_buffer_bytes field is unknown */
    rsvd :1,
    lolu:1,     /* 1 means the first and last logical obj position fields
                   are unknown */
    perr: 1,    /* 1 means the position fields have overflowed and
                   cannot be reported */
    bpew :1;    /* beyond programmable early warning */
    UCHAR active_partition; /* current active partition */
    UCHAR reserved[2];
    UCHAR first_logical_obj_position[4]; /* current logical object position */
    UCHAR last_logical_obj_position[4]; /* next logical object to be
                                         transferred to tape */
    UCHAR num_buffer_logical_obj[4]; /* number of logical objects in buffer */
    UCHAR num_buffer_bytes[4]; /* number of bytes in buffer */
    UCHAR reserved1; /* instead of the commented reserved1 */
} SHORT_DATA_FORMAT, *PSHORT_DATA_FORMAT;

typedef struct _LONG_DATA_FORMAT {
    UCHAR bop:1, /* beginning of partition */
    eop:1, /* end of partition */
    rsvd1:2,
    mpu:1, /* 1 means the logical file id field is unknown */
    lonu:1,/* 1 means either the partition number or logical obj number field
             are unknown */
    rsvd2:1,
    bpew :1; /* beyond programmable early warning */
    CHAR reserved[6];
    UCHAR active_partition; /* current active partition */
    UCHAR logical_obj_number[8]; /* current logical object position */
    UCHAR logical_file_id[8]; /* number of filemarks from bop and
                               current logical position */
    UCHAR obsolete[8];
} LONG_DATA_FORMAT, *PLONG_DATA_FORMAT;

typedef struct _EXTENDED_DATA_FORMAT {
    UCHAR bop:1, /* beginning of partition */
    eop:1, /* end of partition */
    locu:1, /* 1 means num_buffer_logical_obj field is unknown */
    bycu:1, /* 1 means the num_buffer_bytes field is unknown */
    rsvd :1,
    lolu:1, /* 1 means the first and last logical obj position fields
              are unknown */
    perr: 1,/* 1 means the position fields have overflowed and
              can not be reported */
    bpew :1; /* beyond programmable early warning */
    UCHAR active_partition; /* current active partition */
    UCHAR additional_length[2];
}
```

```

UCHAR num_buffer_logical_obj[4]; /* number of logical objects in buffer */
UCHAR first_logical_obj_position[8]; /* current logical object position */
UCHAR last_logical_obj_position[8]; /* next logical object to be
                                         transferred to tape */
UCHAR num_buffer_bytes[8]; /* number of bytes in buffer */
UCHAR reserved;
} EXTENDED_DATA_FORMAT, *PEXTENDED_DATA_FORMAT;

typedef struct READ_TAPE_POSITION{
    UCHAR data_format; /* Specifies the return data format either
                           short, long or extended*/
    union
    {
        SHORT_DATA_FORMAT rp_short;
        LONG_DATA_FORMAT rp_long;
        EXTENDED_DATA_FORMAT rp_extended;
        UCHAR reserved[64];
    } rp_data;
} READ_TAPE_POSITION, *PREAD_TAPE_POSITION;

```

An example of the **READ\_TAPE\_POSITION** command is

```

DWORD cb;
READ_TAPE_POSITION tapePosition;
*rc_ptr = DeviceIoControl(gp->ddHandle0,
                         IOCTL_READ_TAPE_POSITION,
                         &tapePosition,
                         (long)sizeof(READ_TAPE_POSITION),
                         &tapePosition,
                         (long)sizeof(READ_TAPE_POSITION),
                         &cb,
                         (LPOVERLAPPED) NULL);

```

### ***IOCTL\_SET\_TAPE\_POSITION***

This command is used to position the tape in the current active partition to either a logical block id or logical filemark. The **logical\_id\_type** field in the IOCTL structure specifies either a logical block or logical filemark.

The data structure that is used with this IOCTL is

```

#define IOCTL_SET_TAPE_POSITION_LOCATE16 CTL_CODE(IOCTL_TAPE_BASE, 0x0830,
METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
#define LOGICAL_ID_BLOCK_TYPE 0x00
#define LOGICAL_ID_FILE_TYPE 0x01

typedef struct _SET_TAPE_POSITION{
    UCHAR logical_id_type; /* Block or file as defined above */
    ULONGLONG logical_id; /* logical object or logical file to position to */
    UCHAR reserved[32];
} SET_TAPE_POSITION, *PSET_TAPE_POSITION;

```

An example of the **SET\_TAPE\_POSITION** command is

```

DWORD cb;
SET_TAPE_POSITION tapePosition;

*rc_ptr = DeviceIoControl(gp->ddHandle0,
                         IOCTL_SET_TAPE_POSITION_LOCATE16,
                         &tapePosition,
                         (long)sizeof(SET_TAPE_POSITION)
                         NULL,
                         0,
                         &cb,
                         (LPOVERLAPPED) NULL);

```

### ***IOCTL\_QUERY\_LBP***

This command returns logical block protection information. The following output structure is completed by the **IOCTL\_QUERY\_LBP** command.

```

#define IOCTL_QUERY_LBP     CTL_CODE(IOCTL_TAPE_BASE, 0x0831,
METHOD_BUFFERED,

```

```

FILE_READ_ACCESS | FILE_WRITE_ACCESS )
typedef struct _LOGICAL_BLOCK_PROTECTION {
    UCHAR lbp_capable; /* [OUTPUT] the capability of lbp for QUERY ioctl only */
    UCHAR lbp_method; /* lbp method used for QUERY [OUTPUT] and SET [INPUT] ioctls */
#define LBP_DISABLE 0x00
#define REED_SOLOMON_CRC 0x01
    UCHAR lbp_info_length; /* lbp info length for QUERY [OUTPUT] and SET [INPUT]
                           ioctls */
    UCHAR lbp_w;          /* protection info included in write data */
                           /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    UCHAR lbp_r;          /* protection info included in read data */
                           /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    UCHAR rbdp;           /* protection info included in recover buffer data */
                           /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    UCHAR reserved[26];
}LOGICAL_BLOCK_PROTECTION, *PLOGICAL_BLOCK_PROTECTION;

```

An example of the **IOCTL\_QUERY\_LBP** command is

```

*rc_ptr = DeviceIoControl(gp->ddHandle0,
                          IOCTL_QUERY_LBP,
                          NULL,
                          0,
                          &tape_query_LBP,
                          (long)sizeof(LOGICAL_BLOCK_PROTECTION),
                          &cb,
                          (LPOVERLAPPED) NULL);

```

## **IOCTL\_SET\_LBP**

This command sets logical block protection information. The following input structure is sent to the **IOCTL\_SET\_LBP** command.

```

#define IOCTL_SET_LBP      CTL_CODE(IOCTL_TAPE_BASE, 0x0832,
METHOD_BUFFERED,
FILE_READ_ACCESS | FILE_WRITE_ACCESS )
typedef struct _LOGICAL_BLOCK_PROTECTION {
    UCHAR lbp_capable; /* [OUTPUT] the capability of lbp for QUERY ioctl only */
    UCHAR lbp_method; /* lbp method used for QUERY [OUTPUT] and SET [INPUT]
                      ioctls */
#define LBP_DISABLE 0x00
#define REED_SOLOMON_CRC 0x01
    UCHAR lbp_info_length; /* lbp info length for QUERY [OUTPUT] and SET [INPUT]
                           ioctls */
    UCHAR lbp_w;          /* protection info included in write data */
                           /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    UCHAR lbp_r;          /* protection info included in read data */
                           /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    UCHAR rbdp;           /* protection info included in recover buffer data */
                           /* a boolean for QUERY [OUTPUT] and SET [INPUT] ioctls */
    UCHAR reserved[26];
}LOGICAL_BLOCK_PROTECTION, *PLOGICAL_BLOCK_PROTECTION;

```

An example of the **IOCTL\_SET\_LBP** command is

```

*rc_ptr = DeviceIoControl(gp->ddHandle0,
                          IOCTL_SET_LBP,
                          &tape_set_LBP,
                          (long)sizeof(LOGICAL_BLOCK_PROTECTION),
                          NULL,
                          0,
                          &cb,
                          (LPOVERLAPPED) NULL);

```

## **IOCTL\_SET\_PEW\_SIZE**

This command is used to set Programmable Early Warning size.

```

#define IOCTL_SET_PEW_SIZE
CTL_CODE(IOCTL_TAPE_BASE, 0x082C, METHOD_BUFFERED, FILE_READ_ACCESS )

```

The structure that is used to set PEW size is

```

typedef struct _PEW_SIZE{
    USHORT value;
} PEW_SIZE, *PPEW_SIZE;

```

An example of the **IOCTL\_SET\_PEW\_SIZE** command is

```

DWORD cb;
PEW_SIZE pew_size;
...
DeviceIoControl(gp->ddHandle0,
IOCTL_SET_PEW_SIZE,
&pew_size, (long)sizeof(PEW_SIZE),
NULL,
0,
&cb,
(LPOVERLAPPED) NULL);

```

### ***IOCTL\_QUERY\_PEW\_SIZE***

This command is used to query Programmable Early Warning size.

```

#define IOCTL_QUERY_PEW_SIZE
    CTL_CODE(IOCTL_TAPE_BASE, 0x082B, METHOD_BUFFERED, FILE_READ_ACCESS )

```

The structure that is used to query PEW size is

```

typedef struct _PEW_SIZE{
    USHORT value;
} PEW_SIZE, *PPEW_SIZE;

```

An example of the **IOCTL\_QUERY\_PEW\_SIZE** command is

```

DWORD cb;
PEW_SIZE pew_size;
...
DeviceIoControl(gp->ddHandle0,
IOCTL_QUERY_PEW_SIZE,
NULL,
0,
&pew_size,
(long)sizeof(PEW_SIZE),
&cb,
(LPOVERLAPPED) NULL);

```

### ***IOCTL\_VERIFY\_TAPE\_DATA***

This command is used to verify tape data. It uses the drive's error detection and correction hardware to determine whether it can be recovered from the tape. It also checks whether the protection information is present and validates correctly on logical block on the medium. It returns a failure or a success.

```

#define IOCTL_VERIFY_TAPE_DATA
    CTL_CODE(IOCTL_TAPE_BASE, 0x082A, METHOD_BUFFERED, FILE_READ_ACCESS )

```

The structure that is used to verify tape data is

```

typedef struct _VERIFY_DATA {
    UCHAR reserved : 2; /* Reserved */ */
    UCHAR vte: 1; /* [IN] verify to end-of-data */ */
    UCHAR vlpbm: 1; /* [IN] verify logical block protection information */ */
    UCHAR vbf: 1; /* [IN] verify by filemarks */ */
    UCHAR immed: 1; /* [IN] return SCSI status immediately */ */
    UCHAR bytcmp: 1; /* [IN] No use currently */ */
    UCHAR fixed: 1; /* [IN] set Fixed bit to verify the length of */ */
                     /* each logical block */ */
    UCHAR reseved[15];
    ULONG verify_length; /* [IN] amount of data to be verified */ */
} VERIFY_DATA, *PVERIFY_DATA;

```

An example of the **IOCTL\_VERIFY\_DATA** command is

```
DWORD cb;
VERIFY_DATA verify_data;
...
DeviceIoControl(gp->ddHandle0,
IOCTL_VERIFY_TAPE_DATA,
&verify_data,
sizeof(VERIFY_DATA),
NULL,
0,
&cb,
(LPOVERLAPPED) NULL);
```

### **IOCTL\_QUERY\_RAO\_INFO**

This command is used to query the maximum UDS number and UDS size before the **Recommended Access Order** list is generated and received (TS1140 or later, and LTO 9 FH). The structure for the **IOCTL\_QUERY\_RAO\_INFO** command is

```
#define UDS_WITHOUT_GEOMETRY 0
#define UDS_WITH_GEOMETRY 1

typedef struct _QUERY_RAO_INFO{
    CHAR uds_type; /*[IN] 0: UDS_WITHOUT_GEOMETRY
                    1: UDS_WITH_GEOMETRY */
    CHAR reserved[7]; USHORT max_uds_number; /* [OUT] Max UDS number supported
                                                from drive */
    USHORT max_uds_size; /* [OUT] Max single UDS size supported
                           from drive in bytes */
    USHORT max_host_uds_number; /* [OUT] Max UDS number supported from
                                 driver */
} QUERY_RAO_INFO, *PQUERY_RAO_INFO;
```

An example of the **IOCTL\_QUERY\_RAO\_INFO** command is

```
QUERY_RAO_INFO qRAO;
...
qRAO.udt_type = udstype; //UDS_WITHOUT_GEOMETRY or UDS_WITH_GEOMETRY
*rc_ptr = DeviceIoControl(hDevice,
                          IOCTL_QUERY_RAO_INFO,&
                          qRAO,
                          sizeof(QUERY_RAO_INFO),
                          &qRAO,
                          sizeof(QUERY_RAO_INFO),
                          &cb,
                          (LPOVERLAPPED) NULL);
```

### **IOCTL\_GENERATE\_RAO**

This command is used to generate a **Recommended Access Order** list (TS1140 or later, and LTO 9 FH). The UDS list is required as input. Use **USD\_DESCRIPTOR** to build this list. The structure for the **IOCTL\_GENERATE\_RAO** command is

```
#define UDS_WITHOUT_GEOMETRY 0
#define UDS_WITH_GEOMETRY 1

typedef struct _GENERATE_RAO{
    CHAR process; /* [IN] Requested process to generate RAO list */
    /* 0: no reorder UDS and no calculate locate time */
    /* (not currently supported by the drive) */
    /* 1: no reorder UDS but calculate locate time */
    /* 2: reorder UDS and calculate locate time */
    CHAR uds_type; /* [IN] 0: UDS_WITHOUT_GEOMETRY */
    /* 1: UDS_WITH_GEOMETRY */
    CHAR reserved1[2];
    ULONG grao_list_length; /* [IN] The data length is allocated for GRAO list. */
    CHAR reserved2[8];
    CHAR grao_list[1]; /* [IN] the pointer is allocated to the size
                        /* of grao_list_leng
                        /* (uds_number * sizeof(struct grao_uds_desc))
                        /* +sizeof(struct grao_list_header))
                        /* and contains the data of GRAO parameter list. */
                        /* The uds number isn't larger */
} GENERATE_RAO;
```

```

        /*      than max_host_uds_number in QUERY_RAO ioctl. */
    } GENERATE_RAO, *PGENERATE_RAO;

typedef struct _UDS_DESCRIPTOR{
    CHAR descLength[2];
    CHAR reserved[3];
    CHAR UDSName[10];
    CHAR PartNum;
    CHAR beginningLOI[8];
    CHAR endingLOI[8];
}UDS_DESCRIPTOR, *PUDS_DESCRIPTOR;

```

An example of the **IOCTL\_GENERATE\_RAO** command is

```

#define UDS_SIZE 32
#define HEADER_SIZE 8

char *pGRAO;
...
long lGRAOsize = sizeof(GENERATE_RAO)-sizeof(CHAR)/*grao_list[1]
/* + udsamount*UDS_SIZE + HEADER_SIZE;
pGRAO = malloc (lGRAOsize);
...
((PGRAO)pGRAO)->process = process;
((PGRAO)pGRAO)->uds_type = udstype;
((PGRAO)pGRAO)->grao_list_length = udsamount*UDS_SIZE + HEADER_SIZE;
...
((PGRAO_LIST_HEADER)((PGRAO)pGRAO)->grao_list)->addl_data = ((PGRAO)pGRAO)->
grao_list_length - HEADER_SIZE;
PopulateUDS ((PGRAO_LIST_HEADER)((PGRAO)pGRAO)->grao_list)+HEADER_SIZE, udsamount);

*rc_ptr = DeviceIoControl(hDevice,
                           IOCTL_GENERATE_RAO,
                           pGRAO,
                           lGRAOsize,
                           NULL,
                           0,
                           &cb,
                           LPOVERLAPPED) NULL);

```

## **IOCTL\_RECEIVE\_RAO**

This command is used to receive the generated **Recommended Access Order** list (TS1140 or later, and LTO 9 FH).

The structure for **IOCTL\_RECEIVE\_RAO** command is

```

typedef struct _RECEIVE_RAO_LIST {
    ULONG rrao_list_offset; /* [IN] The offset of receive RAO list to
                           /* begin returning data */
    ULONG rrao_list_length; /* [IN/OUT] number byte of data length */
                           /* [IN] The data length is allocated for RRAO list
                           /* by application */
                           /* the length is (max_uds_size * uds_number + sizeof */
                           /* (struct rrao_list_header) */
                           /* max_uds_size is reported in QUERY_RAO_INFO ioctl */
                           /* uds_number is the total UDS number requested */
                           /* from application */
                           /* in GENERATE_RAO ioctl */
                           /* [OUT] the data length is actual returned in RRAO list */
                           /* from the driver */
    CHAR reserved[8];
    CHAR rrao_list[1]; /* [IN/OUT] the data pointer of RRAO list */
}RECEIVE_RAO_LIST, *PRECEIVE_RAO_LIST;

```

An example of the **IOCTL\_RECEIVE\_RAO** command is

```

#define HEADER_SIZE 8
...
char pRRAO;
...
long lRRAOsize=sizeof(RECEIVE_RAO_LIST)-sizeof(CHAR)/*rrao_list[1]
/* + udsamount*udssize+UDS_HEADER;
...
pRRAO = malloc (lRRAOsize);
...

```

```
((PRECEIVE_RAO_LIST) pRRAO)->rراو_list_offset = offset;
((PRECEIVE_RAO_LIST) pRRAO)->rراو_list_length = udsamount*udssize+UDS_HEADER;
...
*rc_ptr = DeviceIoControl(hDevice,
                           IOCTL_RECEIVE_RAO,
                           pRRAO,
                           lRRAOsize,
                           pRRAO,
                           lRRAOsize,
                           &cb,
                           (LPOVERLAPPED) NULL);
```

### **IOCTL\_CHANGE\_OBTAIN\_SENSE**

Issue this command after an error occurs to obtain sense information that is associated with the most recent error. To guarantee that the application can obtain sense information that is associated with an error, the application must issue this command before other commands are issued to the device. Subsequent operations (other than **IOCTL\_CHANGE\_OBTAIN\_SENSE**) reset the sense data field before the operation is run.

This IOCTL is only available for the changer path.

```
#define IOCTL_CHANGE_BASE           FILE_DEVICE_CHANGE
#define IOCTL_CHANGE_OBTAIN_SENSE   CTL_CODE(IOCTL_CHANGE_BASE, 0x0819, METHOD_BUFFERED, FILE_READ_ACCESS)
```

The following output structure is completed by the **IOCTL\_CHANGE\_OBTAIN\_SENSE** command that is passed by the caller.

```
#define MAG_SENSE_BUFFER_SIZE 96 /* Default request sense buffer size for \
Windows 200x */
typedef struct _CHANGE_OBTAIN_SENSE {
    ULONG SenseDataLength;      // The number of bytes of valid sense data.
                                // Will be zero if no error with sense data has occurred.
                                // The only sense data available is that of the last error.
    CHAR SenseData[MAG_SENSE_BUFFER_SIZE];
} CHANGE_OBTAIN_SENSE, *PCHANGE_OBTAIN_SENSE;
```

An example of the **IOCTL\_CHANGE\_OBTAIN\_SENSE** command is

```
DWORD cb;
CHANGE_OBTAIN_SENSE sense_data;
DeviceIoControl(hDevice,
IOCTL_CHANGE_OBTAIN_SENSE,
NULL,
0,
&sense_data,
(long)sizeof(CHANGE_OBTAIN_SENSE),
&cb,
(LPOVERLAPPED) NULL);
```

### **IOCTL\_MODE\_SENSE**

This command is used to get the Mode Sense Page/Subpage.

```
***** GENERIC SCSI IOCTL *****

#define IOCTL_IBM_BASE ((('IBM' << 8) | FILE_DEVICE_SCSI)

#define DEFINE_IBM_IOCTL(x) CTL_CODE(IOCTL_IBM_BASE, x, METHOD_BUFFERED, \
FILE_READ_ACCESS | FILE_WRITE_ACCESS)
#define IOCTL_MODE_SENSE DEFINE_IBM_IOCTL(0x003)
```

The structure that is used for this IOCTL is

```
typedef struct _MODE_SENSE_PARAMETERS
{
    UCHAR page_code;          /* [IN] mode sense page code */
    UCHAR subpage_code;       /* [IN] mode sense subpage code */
    UCHAR reserved[6];
    UCHAR cmd_code;           /* [OUT] SCSI Command Code: this field is set with */
                                /* SCSI command code which the */
                                /* device responded. */
```

```

/*      x'5A' = Mode Sense (10)          */
/*      x'1A' = Mode Sense (6)           */
CHAR data[MAX_MODESENSEPAGE]; /* [OUT] whole mode sense data include header,
block descriptor and page */
} MODE_SENSE_PARAMETERS, *PMODE_SENSE_PARAMETERS;

```

An example of the **IOCTL\_MODE\_SENSE** command is

```

DWORD cb;
MODE_SENSE_PARAMETERS mode_sense;
...
DeviceIoControl(gp->ddHandle0,
IOCTL_MODE_SENSE,
&mode_sense,
sizeof(MODE_SENSE_PARAMETERS),
NULL,
0,
&cb,
(LPOVERLAPPED) NULL);

```

## Variable and fixed block read/write processing

In Windows , tape APIs can be configured to manipulate tapes that use either fixed block size or variable block size.

If variable block size is wanted, the block size must be set to zero. The **SetTapeParameters** function must be called specifying the **SET\_TAPE\_MEDIA\_INFORMATION** operation. The function requires the use of a **TAPE\_SET\_MEDIA\_PARAMETERS** structure. The BlockSize member of the structure must be set to the wanted block size. Any block size other than 0 sets the media parameters to fixed block size. The size of the block is equal to the BlockSize member.

In fixed block mode, the size of all data buffers used for reading and writing must be a multiple of the block size. To determine the fixed block size, the **GetTapeParameters** function must be used. Specifying the **GET\_TAPE\_MEDIA\_INFORMATION** operation yields a **TAPE\_GET\_MEDIA\_PARAMETERS** structure. The BlockSize member of this structure reports the block size of the tape. The size of buffers that are used in read and write operations must be a multiple of the block size. This mode allows multiple blocks to be transferred in a single operation. In fixed block mode, transfer of odd block sizes (for example, 999 bytes) is not supported.

When reading or writing variable sized blocks, the operation cannot exceed the maximum transfer length of the Host Bus Adapter. This length is the length of each transfer page (typically 4 K) times the number of transfer pages (if set to 16, the maximum transfer length for variable sized transfers is 64 K). This number can be modified by changing the *scatter-gather* variable in the system registry, but this action is not recommended because it uses up scarce system resources.

Reading a tape that contains variable sized blocks can be accomplished even without knowing what size the blocks are. If a buffer is large enough to read the data in a block, then the data is read without any errors. If the buffer is larger than a block, then only data in a single block is read and the tape is advanced to the next block.

The size of the block is returned by the read operation in the **\*pBytesRead** parameter. If a data buffer is too small to contain all of the data in a block, then a couple of things occur. First, the data buffer contains data from the tape, but the read operation fails and **GetLastError** returns **ERROR\_MORE\_DATA**. This error value indicates that more data is in the block to be read. Second, the tape is advanced to the next block. To reread the previous block, the tape must be repositioned to the wanted block and a larger buffer must be specified. It is best to specify as large a buffer as possible so that this issue does not occur.

If a tape contains fixed size blocks, but the tape media parameters are set to variable block size, then no assumptions are made regarding the size of the blocks on the tape. Each read operation behaves as described. The sizes of the blocks on the tape are treated as variable, but happen to be the same size. If a tape has variable size blocks, but the tape media parameters are set to fixed block size, then the size of all blocks on the tape are expected to be the same fixed size. Reading a block of a tape in this situation fails and **GetLastError** returns **ERROR\_INVALID\_BLOCK\_LENGTH**. The only exception is if the block size in the media parameters is the same as the size of the variable block and the size of the read buffer happens to be a multiple of the size of the variable block.

If **ReadFile** encounters a tapemark, the data up to the tapemark is read and the function fails. (The **GetLastError** function returns an error code that indicates that a tapemark was encountered.) The tape is positioned past the tapemark, and an application can call **ReadFile** again to continue reading.

## Event log

The **ibmtpxxx** and **ibmcgxxx** device drivers log certain data to the Event Log when exceptions are encountered.

To interpret this event data, the user must be familiar with the following components:

- Microsoft Event Viewer
- The SDK and WDK components from the Microsoft Development Network (MSDN)
- Tape Storage hardware terminology
- SCSI terminology

Several bytes of "Event Detail" data are logged under Source = **ibmtpxxx** or **ibmcgxxx**.

The following description texts are expected:

- The description for Event ID (0) in Source (**ibmtpxxx**) was not found. It contains the following insertion strings: \Device\Tapex.

The user must view the event data in Word format to properly decode the data.

Table 5 on page 167 and Table 6 on page 168 indicate the hexadecimal offsets, names, and definitions for Tape Storage **ibmtpxxx** and **ibmcgxxx** event data.

Table 5. Tape Storage <i>ibmtpxxx</i> , and <i>ibmcgxxx</i> event data		
Offset	Name	Definition
0x00-0x01	DumpDataSize	Indicates the size in bytes required for any DumpData the driver places in the packet.
0x02	RetryCount	Indicates how many times the driver tried the operation and encountered this error.
0x03	MajorFunctionCode	Indicates the IRP_MJ_XXX from the driver's I/O stack location in the current IRP (from NTDDK.H).
0x0C-0x0F	ErrorCode	For the Tape Storage device driver, it is 0.
0x10-0x13	UniqueErrorValue	Reserved
0x14-0x17	FinalStatus	Indicates the value that is set in the I/O status block of the IRP when it was completed or the STATUS_XXX returned by a support routine the driver called (from NTSTATUS.H).
0x1C-0x1F	IoControlCode	For the Tape Storage device driver, it indicates the I/O control code from the driver's I/O stack location in the current IRP if the MajorFunctionCode is IRP_MJ_DEVICE_CONTROL. Otherwise, this value is 0.
0x28	Beginning of Dump Data	The following items are variable in length. See the DDK and SCSI documentation for details.
0x38	Beginning of SRB structure	The SCSI Request Block (from NTDDK.H).
0X68	Beginning of CDB structure	The Command Descriptor Block (from SCSI.H).

*Table 5. Tape Storage ibmtpxxx, and ibmcgxxx event data (continued)*

Offset	Name	Definition
0x78	Beginning of SCSI Sense Data	(from SCSI.H). If the first word in this field is 0x00DF0000 (SCSI error marker) or 0x00EF0000 (Non-SCSI error marker), no valid sense information was available for this error.

For example, **ibmcgxxx** logs the following error when a move medium is attempted and the destination element is full. Explanations of selected fields follow.

```

0000: 006c000f 00c40001 00000000 c004000b
0010: bcde7f48 c0000284 00000000 00000000
0020: 00000000 00000000 00000000 000052f4
0030: 00000000 00000000 004000c4 02000003
0040: 600c00ff 00000028 00000000 00000258
0050: 00000000 814dac28 00000000 bcde7f48
0060: 81841000 00000000 a5600000 00200010
0070: 00000000 00000000 70000500 00000058
0080: 00000000 3b0dff02 00790000 0000093e
0090: 00000000

```

*Table 6. Tape Storage ibmtpxxx, and ibmcgxxx event data*

Field	Value	Definition
DumpDataSize	0x006C	6C hex (108 dec) bytes of dump data, beginning at byte 28 hex.
RetryCount	0x00	The first time that the operation is attempted (no retries).
MajorFunctionCode	0x0F	IRP_MJ_INTERNAL_DEVICE_CONTROL
FinalStatus	0xC0000284	STATUS_DESTINATION_ELEMENT_FULL
IoControlCode	0x00000000	-
SRB	0x004000C4...	From NTDDK.H, the first word of the SRB indicates the length of the SRB (40 hex bytes, 64 dec bytes), the function code (0x00), and the SrbStatus (from SRB.H, 0xC4 = SRB_STATUS_AUTOSENSE_VALID, SRB_STATUS_QUEUE_FROZEN, SRB_STATUS_ERROR).
CDB	0xA5...	From SCSI.H, the first byte of the CDB is the operation code. 0xA5 = SCSIOP_MOVE_MEDIUM.
Sense Data	0x70000500...	From SCSI.H, the first word of the sense data indicates the error code (0x70), the segment number (0x00), and the sense key (0x05, corresponding to an illegal SCSI request).

## Notices

References in this publication to IBM products, programs, or services do not imply that IBM intends to make these available in all countries (or regions) in which IBM operates.

Any references to an IBM program or other IBM product in this publication is not intended to state or imply that only IBM's program or other product may be used. Any functionally equivalent program that does not infringe any of IBM's intellectual property rights may be used instead of the IBM product. Evaluation and verification of operation in conjunction with other products, except those expressly designed by IBM, is the user's responsibility.

IBM may have patents or pending patent applications covering subject matter in this document. The furnishing of this document does not give you any license to these patents. You may send license inquiries, in writing, to:

IBM Director of Licensing  
IBM Corporation  
North Castle Drive  
Armonk, NY 10504-1785  
U.S.A.

For license inquiries regarding double-byte character set (DBCS) information, contact the IBM Intellectual Property Department in your country or send inquiries, in writing, to:

Intellectual Property Licensing  
Legal and Intellectual Property Law  
IBM Japan, Ltd  
19-21, Nihonbashi-Hakozakicho, Chuo-ku  
Tokyo 103-8510, Japan

The following paragraph does not apply to the United Kingdom or any other country (or region) where such provisions are inconsistent with local law:

INTERNATIONAL BUSINESS MACHINES CORPORATION PROVIDES THIS PUBLICATION "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY, OR FITNESS FOR A PARTICULAR PURPOSE. Some states (or regions) do not allow disclaimer of express or implied warranties in certain transactions, therefore, this statement cannot apply to you.

This information could include technical inaccuracies or typographical errors. Changes are periodically made to the information herein; these changes are incorporated in new editions of the publication. IBM may make improvements and/or changes in the products and/or programs described in this publication at any time without notice.

IBM may use or distribute any of the information you supply in any way it believes appropriate without incurring any obligation to you.

The ITDT-SE and ITDT-GE software uses Henry Spencer's regular expression library that is subject to the following copyright notice:

"Copyright 1992, 1993, 1994, 1997 Henry Spencer. All rights reserved. This software is not subject to any license of the American Telephone and Telegraph Company or of the Regents of the University of California.

Permission is granted to anyone to use this software for any purpose on any computer system, and to alter it and redistribute it, subject to the following restrictions:

1. The author is not responsible for the consequences of use of this software, no matter how awful, even if they arise from flaws in it.
2. The origin of this software must not be misrepresented, either by explicit claim or by omission. Since few users ever read sources, credits must appear in the documentation.
3. Altered versions must be plainly marked as such, and must not be misrepresented as being the original software. Since few users ever read sources, credits must appear in the documentation.
4. This notice cannot be removed or altered.

## Trademarks

The following terms are trademarks of International Business Machines Corporation in the United States, other countries (or regions), or both:

AIX

IBMLink

RS/6000

System z®

AIX 5L

Magstar

S/390®

Tivoli®

FICON®	Micro Channel	StorageSmart	TotalStorage
HyperFactor	Netfinity	System i	Virtualization Engine
i5/OS	POWER5	System p	xSeries
iSeries	ProtecTIER®	System Storage	z9
IBM	pSeries	System x	zSeries

Adobe and Acrobat are either registered trademarks or trademarks of Adobe Systems Incorporated in the United States, and/or other countries.

Intel, Itanium, and Pentium are trademarks of Intel Corporation in the United States, other countries (or regions), or both.

Java™ and all Java-based trademarks are trademarks of Oracle, Inc. in the United States, other countries, or both.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and Windows 2000 are trademarks of Microsoft Corporation in the United States, other countries (or regions), or both.

UNIX is a registered trademark of The Open Group in the United States and other countries (or regions).

Other company, product, and service names may be trademarks or service marks of others.

## How to send your comments

---

Your feedback is important in helping to provide the most accurate and highest quality information.

To submit any comments about this publication or any other IBM storage product documentation:

- Send your comments by email to [ibmkc@us.ibm.com](mailto:ibmkc@us.ibm.com). Be sure to include the following information:
  - Exact publication title and version
  - Page, table, or illustration numbers that you are commenting on
  - A detailed description of any information that should be changed

## Terms and conditions for IBM Documentation

---

Permissions for the use of these publications are granted subject to the following terms and conditions.

### IBM Privacy Policy

We intend to protect your personal information and to maintain its integrity. IBM implements reasonable physical, administrative, and technical safeguards to help us protect your personal information from unauthorized access, use, and disclosure. For example, the products only provide IBM data about the asset usage and configuration and do not reflect private use of the asset. When diagnostics are required to be sent to IBM, and a problem is submitted, that data is routed directly to a secured infrastructure. Only individuals with a need to know are given access while working to resolve your problem. When appropriate, we also require that our suppliers protect such information from unauthorized access, use, and disclosure.

Visit the IBM Privacy Policy for additional information on this topic at <https://www.ibm.com/privacy/details/us/en/>.

### Applicability

These terms and conditions are in addition to any terms of use for the IBM website.

## **Personal use**

You can reproduce these publications for your personal, noncommercial use provided that all proprietary notices are preserved. You cannot distribute, display, or make derivative work of these publications, or any portion thereof, without the express consent of IBM.

## **Commercial use**

You can reproduce, distribute, and display these publications solely within your enterprise provided that all proprietary notices are preserved. You cannot make derivative works of these publications, or reproduce, distribute, or display these publications or any portion thereof outside your enterprise, without the express consent of IBM.

## **Rights**

Except as expressly granted in this permission, no other permissions, licenses, or rights are granted, either express or implied, to the Publications or any information, data, software or other intellectual property contained therein.

IBM reserves the right to withdraw the permissions that are granted herein whenever, in its discretion, the use of the publications is detrimental to its interest or as determined by IBM, the above instructions are not being properly followed.

You cannot download, export, or reexport this information except in full compliance with all applicable laws and regulations, including all United States export laws and regulations.

IBM MAKES NO GUARANTEE ABOUT THE CONTENT OF THESE PUBLICATIONS. THE PUBLICATIONS ARE PROVIDED "AS-IS" AND WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT, AND FITNESS FOR A PARTICULAR PURPOSE.

## **IBM trademarks**

IBM, the IBM logo, and ibm.com<sup>®</sup> are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available at <http://www.ibm.com/legal/copytrade.shtml>.

# Index

## A

AIX Device Driver (Atape) [8–11, 13–15, 23, 37, 75, 84](#)

## C

Close error codes [138](#)  
Closing the special file [14](#)  
common functions [3](#)  
CreateFile [141](#)

## D

Device and volume information logging [14, 15](#)  
DeviceIoControl [146](#)

## E

EraseTape [145](#)

## F

features [3](#)  
fixed block read/write processing [166](#)

## G

General error codes [137](#)  
General IOCTL operations [23, 89](#)  
GetTapeParameters [144](#)  
GetTapePosition [144](#)  
GetTapeStatus [146](#)

## I

Introduction [8, 9](#)  
IOCTL commands [147](#)  
IOCTL error codes [139](#)

## L

legal  
    terms and conditions [170](#)  
lin\_tape [87](#)  
Lin\_tape [87, 89, 129](#)  
Linux device driver (IBMtape) [87, 89, 99, 129, 130, 137–139](#)  
Linux-defined entry points [87](#)  
Log file [15](#)

## M

media partitioning [4](#)  
Medium changer devices [89](#)  
Medium changer IOCTL operations [75, 129, 130](#)  
Medium Changer IOCTLs [147](#)

## O

Open error codes [137](#)  
Opening the special file for I/O [11](#)  
overview [89, 99](#)  
Overview [23, 37, 75](#)

## P

Persistent reservation support [16](#)  
PrepareTape [145](#)  
programming interface [139–147, 149, 166](#)

## R

Read error codes [138](#)  
ReadFile [142](#)  
Reading from the special file [13](#)  
Reading with the TAPE\_READ\_REVERSE extended parameter [13](#)  
Reading with the TAPE\_SHORT\_READ extended parameter [13](#)  
Return codes [84, 137–139](#)

## S

SCSI IOCTL commands [130](#)  
SetTapeParameters [144](#)  
SetTapePosition [143](#)  
software interface [87, 89](#)  
Software interface for medium changer devices [9](#)  
Software interface for tape devices [8](#)  
Special files [9–11, 13, 14](#)  
Special files for 3490E, 3590, Magstar MP or 7332 tape devices [9](#)  
Special files for 3575, 7331, 7334, 7336, or 7337 medium changer devices [10](#)

## T

tape device driver [16](#)  
Tape drive compatibility IOCTL operations [129](#)  
Tape drive IOCTL operations [99](#)  
Tape IOCTL operations [37](#)  
tape Media Changer driver entry points [140–146](#)  
The extended open operation [11](#)

## V

variable block read/write processing [166](#)  
Vendor-specific device IOCTLs for DeviceIoControl [149](#)

## W

Windows 200x  
    event log [167](#)

Windows NT device driver

event log [167](#)

Write error codes [138](#)

Write Tapemark [143](#)

WriteFile [142](#)

Writing to the special file [13](#)

**IBM.**<sup>®</sup>